

# Ensemble Methods (Part II)

## Gradient Boosting

Schwartz

September 12, 2016

# The most powerful prediction algorithm on the planet

In “The BellKor Solution to the Netflix Grand Prize” Yehuda Koren reported on the use of gradient boosted decision trees (GBDT) in the top performing algorithm.

The Netflix Prize was an open competition for the best collaborative filtering algorithm to predict user ratings for films. The competition was open to anyone not connected with Netflix or a resident of Cuba, Iran, Syria, North Korea, Myanmar or Sudan. On 21 September 2009, the grand prize of \$1,000,000 was given to the Pragmatic Chaos team which bested Netflix’s own prediction algorithm by 10.06%.

The screenshot shows the Netflix Prize website's leaderboard. A large red stamp on the right side of the header reads "COMPLETED". The page title is "Leaderboard". Below the title, a table lists the top 8 teams, their scores, improvement percentages, and submission times. The winning team, "BellKor's Pragmatic Chaos", is highlighted in a blue box with the text "Grand Prize - RMSE = 0.8567 - Winning Team: BellKor's Pragmatic Chaos".

Rank	Team Name	Best Test Score	% Improvement	Best Submit Time
1	<a href="#">BellKor's Pragmatic Chaos</a>	0.8567	10.06	2009-07-26 18:18:28
2	<a href="#">The Ensemble</a>	0.8567	10.06	2009-07-26 18:38:22
3	<a href="#">Grand Prize Team</a>	0.8582	9.90	2009-07-10 21:24:40
4	<a href="#">Opera Solutions and Vandelay United</a>	0.8588	9.84	2009-07-10 01:12:31
5	<a href="#">Vandelay Industries !</a>	0.8591	9.81	2009-07-10 00:32:20
6	<a href="#">PragmaticTheory</a>	0.8594	9.77	2009-06-24 12:06:56
7	<a href="#">BellKor in BigChaos</a>	0.8601	9.70	2009-05-13 08:14:09
8	<a href="#">Dace</a>	0.8612	9.59	2009-07-24 17:18:43

# Objectives

- ▶ Understand gradient boosting

# Objectives

- ▶ Understand gradient boosting
  - ▶ as an ensemble of sequential “weak learners”

# Objectives

- ▶ Understand gradient boosting
  - ▶ as an ensemble of sequential “weak learners”
  - ▶ as a general algorithm

# Objectives

- ▶ Understand gradient boosting
  - ▶ as an ensemble of sequential “weak learners”
  - ▶ as a general algorithm
  - ▶ as a series of regression on sequential residuals

# Objectives

- ▶ Understand gradient boosting
  - ▶ as an ensemble of sequential “weak learners”
  - ▶ as a general algorithm
  - ▶ as a series of regression on sequential residuals
  - ▶ as a general gradient decent algorithm

# Objectives

- ▶ Understand gradient boosting
  - ▶ as an ensemble of sequential “weak learners”
  - ▶ as a general algorithm
  - ▶ as a series of regression on sequential residuals
  - ▶ as a general gradient decent algorithm
  - ▶ in its earliest/latest incarnation AdaBoost/XGBoost

# Objectives

- ▶ Understand gradient boosting
  - ▶ as an ensemble of sequential “weak learners”
  - ▶ as a general algorithm
  - ▶ as a series of regression on sequential residuals
  - ▶ as a general gradient decent algorithm
  - ▶ in its earliest/latest incarnation AdaBoost/XGBoost
- ▶ Understand the tuning parameters of tree-based boosting

# Objectives

- ▶ Understand gradient boosting
  - ▶ as an ensemble of sequential “weak learners”
  - ▶ as a general algorithm
  - ▶ as a series of regression on sequential residuals
  - ▶ as a general gradient decent algorithm
  - ▶ in its earliest/latest incarnation AdaBoost/XGBoost
- ▶ Understand the tuning parameters of tree-based boosting
  - ▶ Tree parameters versus boosting parameters

# Objectives

- ▶ Understand gradient boosting
  - ▶ as an ensemble of sequential “weak learners”
  - ▶ as a general algorithm
  - ▶ as a series of regression on sequential residuals
  - ▶ as a general gradient decent algorithm
  - ▶ in its earliest/latest incarnation AdaBoost/XGBoost
- ▶ Understand the tuning parameters of tree-based boosting
  - ▶ Tree parameters versus boosting parameters
  - ▶ Boosting implementation/tuning via grid search

# Objectives

- ▶ Understand gradient boosting
  - ▶ as an ensemble of sequential “weak learners”
  - ▶ as a general algorithm
  - ▶ as a series of regression on sequential residuals
  - ▶ as a general gradient decent algorithm
  - ▶ in its earliest/latest incarnation AdaBoost/XGBoost
- ▶ Understand the tuning parameters of tree-based boosting
  - ▶ Tree parameters versus boosting parameters
  - ▶ Boosting implementation/tuning via grid search
- ▶ Understand partial dependency plots

# Objectives

- ▶ Understand gradient boosting
  - ▶ as an ensemble of sequential “weak learners”
  - ▶ as a general algorithm
  - ▶ as a series of regression on sequential residuals
  - ▶ as a general gradient decent algorithm
  - ▶ in its earliest/latest incarnation AdaBoost/XGBoost
- ▶ Understand the tuning parameters of tree-based boosting
  - ▶ Tree parameters versus boosting parameters
  - ▶ Boosting implementation/tuning via grid search
- ▶ Understand partial dependency plots
- ▶ *Critique and Sell* gradient boosting

# Sequential “weak learners”

1. Student 1 takes exam

## Sequential “weak learners”

1. Student 1 takes exam
2. Student 2 takes exam *knowing how student 1 did*

## Sequential “weak learners”

1. Student 1 takes exam
2. Student 2 takes exam *knowing how student 1 did*  
⋮
3. Student  $k$  takes exam *knowing how student  $k - 1$  did*

## Sequential “weak learners”

1. Student 1 takes exam
2. Student 2 takes exam *knowing how student 1 did*  
⋮
3. Student  $k$  takes exam *knowing how student  $k - 1$  did*  
⋮
4. Student  $n$  takes exam *knowing how student  $n - 1$  did*

# Thought Experiment

- ▶ *Sequential learners* focus effort on outcomes predicted poorly by previous learners as opposed to those predicted accurately

## Thought Experiment

- ▶ *Sequential learners* focus effort on outcomes predicted poorly by previous learners as opposed to those predicted accurately
- ▶ *Weak learners* consistently predict outcomes but with (high bias, low variance) accuracy only  $\epsilon$  slightly better than chance

# Thought Experiment

- ▶ *Sequential learners* focus effort on outcomes predicted poorly by previous learners as opposed to those predicted accurately
- ▶ *Weak learners* consistently predict outcomes but with (high bias, low variance) accuracy only  $\epsilon$  slightly better than chance

Can many *sequential weak learners*  
be used to make a powerful prediction tool?

# Thought Experiment

- ▶ *Sequential learners* focus effort on outcomes predicted poorly by previous learners as opposed to those predicted accurately
- ▶ *Weak learners* consistently predict outcomes but with (high bias, low variance) accuracy only  $\epsilon$  slightly better than chance

Can many *sequential weak learners*  
be used to make a powerful prediction tool?

Do you think such an approach could suffer from overfitting?

## Sequential Estimators Notation

- ▶ Let  $W$  be an untrained weak learner and  $\phi_k(\mathbf{x})$  be an estimator that predicts outcome  $Y$  based on features  $\mathbf{x}$

Define a new estimator

$$\phi_{k+1}(\mathbf{x}) = \phi_k(\mathbf{x}) + \alpha_k W(\mathbf{x}, \gamma_k)$$

where  $W$  is trained on  $\mathbf{x}$  “in the presence” of  $\phi_k(\mathbf{x})$

## Sequential Estimators Notation

- ▶ Let  $W$  be an untrained weak learner and  $\phi_k(\mathbf{x})$  be an estimator that predicts outcome  $Y$  based on features  $\mathbf{x}$

Define a new estimator

$$\phi_{k+1}(\mathbf{x}) = \phi_k(\mathbf{x}) + \alpha_k W(\mathbf{x}, \gamma_k)$$

where  $W$  is trained on  $\mathbf{x}$  “in the presence” of  $\phi_k(\mathbf{x})$

- ▶ E.g., let  $W$  be a tree and  $\gamma_k = \text{number of splits} = 1$

## Sequential Estimators Notation

- ▶ Let  $W$  be an untrained weak learner and  $\phi_k(\mathbf{x})$  be an estimator that predicts outcome  $Y$  based on features  $\mathbf{x}$

Define a new estimator

$$\phi_{k+1}(\mathbf{x}) = \phi_k(\mathbf{x}) + \alpha_k W(\mathbf{x}, \gamma_k)$$

where  $W$  is trained on  $\mathbf{x}$  “in the presence” of  $\phi_k(\mathbf{x})$

- ▶ E.g., let  $W$  be a tree and  $\gamma_k = \text{number of splits} = 1$   
So  $W$  is a stump – a weak learner\*

## Sequential Estimators Notation

- ▶ Let  $W$  be an untrained weak learner and  $\phi_k(\mathbf{x})$  be an estimator that predicts outcome  $Y$  based on features  $\mathbf{x}$

Define a new estimator

$$\phi_{k+1}(\mathbf{x}) = \phi_k(\mathbf{x}) + \alpha_k W(\mathbf{x}, \gamma_k)$$

where  $W$  is trained on  $\mathbf{x}$  “in the presence” of  $\phi_k(\mathbf{x})$

- ▶ E.g., let  $W$  be a tree and  $\gamma_k = \text{number of splits} = 1$   
So  $W$  is a stump – a weak learner\*
- ▶  $W$  is fit to *improve* the available prediction of  $Y$  from  $\phi_k(\mathbf{x})$

## Sequential Estimators Notation

- ▶ Let  $W$  be an untrained weak learner and  $\phi_k(\mathbf{x})$  be an estimator that predicts outcome  $Y$  based on features  $\mathbf{x}$

Define a new estimator

$$\phi_{k+1}(\mathbf{x}) = \phi_k(\mathbf{x}) + \alpha_k W(\mathbf{x}, \gamma_k)$$

where  $W$  is trained on  $\mathbf{x}$  “in the presence” of  $\phi_k(\mathbf{x})$

- ▶ E.g., let  $W$  be a tree and  $\gamma_k = \text{number of splits} = 1$   
So  $W$  is a stump – a weak learner\*
- ▶  $W$  is fit to *improve* the available prediction of  $Y$  from  $\phi_k(\mathbf{x})$
- ▶ But the negligible impact of  $W^*$  will be further muted by  $\alpha_k$
- ▶  $\alpha_k$  **controls the learning rate of the boosting algorithm**

## Compare and Contrast

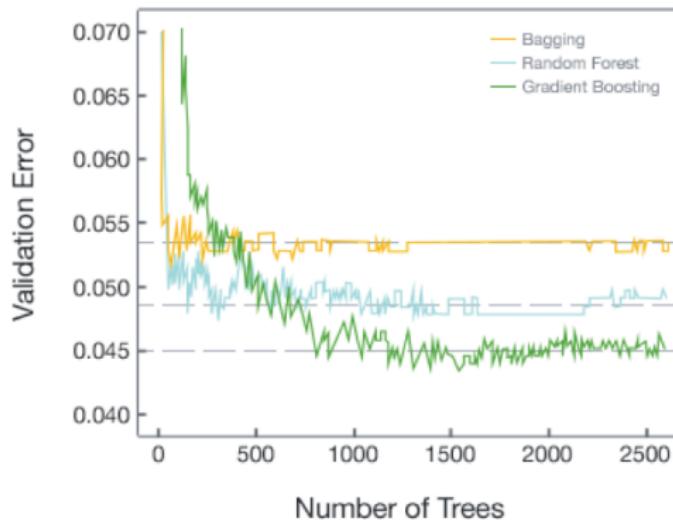
- (A) Tree
- (B) Bagging
- (C) Random Forests
- (D) Sequential Weak Learners

## Compare and Contrast

- (A) Tree  $\hat{f}^{(1)}(\mathbf{x})$
- (B) Bagging  $\frac{1}{m} \sum \hat{f}^{(k)}(\mathbf{x})$
- (C) Random Forests  $\frac{1}{m} \sum \hat{f}^{(k)}(\mathbf{x})$ , with small  $\text{Cor}[f^{(k)}, f^{(k')}]$
- (D) Sequential Weak Learners  $\sum \alpha_k \hat{W}(\mathbf{x}, \gamma_k)$

## Compare and Contrast

- (A) Tree  $\hat{f}^{(1)}(\mathbf{x})$
- (B) Bagging  $\frac{1}{m} \sum \hat{f}^{(k)}(\mathbf{x})$
- (C) Random Forests  $\frac{1}{m} \sum \hat{f}^{(k)}(\mathbf{x})$ , with small  $\text{Cor}[f^{(k)}, f^{(k')}]$
- (D) Sequential Weak Learners  $\sum \alpha_k \hat{W}(\mathbf{x}, \gamma_k)$




# Fitting regressions on sequential residuals

- if we use squared error loss

$$L(Y_i, \phi_k(\mathbf{x}_i)) = \frac{1}{2}(Y_i - \phi_k(\mathbf{x}_i))^2$$

then the direction of the gradient is vector  $\epsilon$  of residuals




# Fitting regressions on sequential residuals

- if we use squared error loss

$$L(Y_i, \phi_k(\mathbf{x}_i)) = \frac{1}{2}(Y_i - \phi_k(\mathbf{x}_i))^2$$

then the direction of the gradient is vector  $\epsilon$  of residuals




- Moving from  $\phi_k(\mathbf{x}_i)$  to  $\phi_k(\mathbf{x}_i) + \epsilon$  would be zero loss

# Fitting regressions on sequential residuals

- if we use squared error loss

$$L(Y_i, \phi_k(\mathbf{x}_i)) = \frac{1}{2}(Y_i - \phi_k(\mathbf{x}_i))^2$$

then the direction of the gradient is vector  $\epsilon$  of residuals




- Moving from  $\phi_k(\mathbf{x}_i)$  to  $\phi_k(\mathbf{x}_i) + \epsilon$  would be zero loss
- But we **don't** move exactly to  $\phi_k(\mathbf{x}_i) + \epsilon$  because

# Fitting regressions on sequential residuals

- ▶ if we use squared error loss

$$L(Y_i, \phi_k(\mathbf{x}_i)) = \frac{1}{2}(Y_i - \phi_k(\mathbf{x}_i))^2$$

then the direction of the gradient is vector  $\epsilon$  of residuals




- ▶ Moving from  $\phi_k(\mathbf{x}_i)$  to  $\phi_k(\mathbf{x}_i) + \epsilon$  would be zero loss
- ▶ But we **don't** move exactly to  $\phi_k(\mathbf{x}_i) + \epsilon$  because
  - ▶  $W(\mathbf{x}, \gamma_k)$  is a weak learner and probably can't, and

# Fitting regressions on sequential residuals

- if we use squared error loss

$$L(Y_i, \phi_k(\mathbf{x}_i)) = \frac{1}{2}(Y_i - \phi_k(\mathbf{x}_i))^2$$

then the direction of the gradient is vector  $\epsilon$  of residuals



- Moving from  $\phi_k(\mathbf{x}_i)$  to  $\phi_k(\mathbf{x}_i) + \epsilon$  would be zero loss
- But we **don't** move exactly to  $\phi_k(\mathbf{x}_i) + \epsilon$  because
  - $W(\mathbf{x}, \gamma_k)$  is a weak learner and probably can't, and
  - we only move to  $\phi_k(\mathbf{x}) + \alpha_k \hat{W}(\mathbf{x}, \gamma_k)$  at learning rate  $\alpha_k$

## Fitting regressions on sequential residuals

1. Set  $\hat{f}^{(0)}(\mathbf{x}_i) = 0$  and  $\hat{\epsilon}_i^{(1)} = Y_i$

## Fitting regressions on sequential residuals

1. Set  $\hat{f}^{(0)}(\mathbf{x}_i) = 0$  and  $\hat{\epsilon}_i^{(1)} = Y_i$
2. For  $k = 1, \dots, m$

## Fitting regressions on sequential residuals

1. Set  $\hat{f}^{(0)}(\mathbf{x}_i) = 0$  and  $\hat{\epsilon}_i^{(1)} = Y_i$
2. For  $k = 1, \dots, m$ 
  - 2.1 Fit a tree  $\hat{f}^{(k)}$  with  $T$  leaves to  $\hat{\epsilon}$  using features  $\mathbf{x}$

## Fitting regressions on sequential residuals

1. Set  $\hat{f}^{(0)}(\mathbf{x}_i) = 0$  and  $\hat{\epsilon}_i^{(1)} = Y_i$
2. For  $k = 1, \dots, m$ 
  - 2.1 Fit a tree  $\hat{f}^{(k)}$  with  $T$  leaves to  $\hat{\epsilon}$  using features  $\mathbf{x}$
  - 2.2 Update the estimator  $\hat{f}^{(k+1)} = f^{(k)} + \alpha_k f^{(k-1)}$

## Fitting regressions on sequential residuals


1. Set  $\hat{f}^{(0)}(\mathbf{x}_i) = 0$  and  $\hat{\epsilon}_i^{(1)} = Y_i$
2. For  $k = 1, \dots, m$ 
  - 2.1 Fit a tree  $\hat{f}^{(k)}$  with  $T$  leaves to  $\hat{\epsilon}$  using features  $\mathbf{x}$
  - 2.2 Update the estimator  $\hat{f}^{(k+1)} = f^{(k)} + \alpha_k f^{(k-1)}$
  - 2.3 Update the residuals  $\hat{\epsilon}_i^{(k+1)} = \hat{\epsilon}_i^{(k)} - \alpha_k f^{(k)}(\mathbf{x}_i)$

## Fitting regressions on sequential residuals

1. Set  $\hat{f}^{(0)}(\mathbf{x}_i) = 0$  and  $\hat{\epsilon}_i^{(1)} = Y_i$
2. For  $k = 1, \dots, m$ 
  - 2.1 Fit a tree  $\hat{f}^{(k)}$  with  $T$  leaves to  $\hat{\epsilon}$  using features  $\mathbf{x}$
  - 2.2 Update the estimator  $\hat{f}^{(k+1)} = f^{(k)} + \alpha_k f^{(k-1)}$
  - 2.3 Update the residuals  $\hat{\epsilon}_i^{(k+1)} = \hat{\epsilon}_i^{(k)} - \alpha_k f^{(k)}(\mathbf{x}_i)$
3. Return the boosted model  $\hat{f}(\mathbf{x}_i) = \sum_{k=1}^m \alpha_k f^{(k)}(\mathbf{x}_i)$


# Fitting regressions on sequential residuals

1. Set  $\hat{f}^{(0)}(\mathbf{x}_i) = 0$  and  $\hat{\epsilon}_i^{(1)} = Y_i$
2. For  $k = 1, \dots, m$ 
  - 2.1 Fit a tree  $\hat{f}^{(k)}$  with  $T$  leaves to  $\hat{\epsilon}$  using features  $\mathbf{x}$
  - 2.2 Update the estimator  $\hat{f}^{(k+1)} = f^{(k)} + \alpha_k f^{(k-1)}$
  - 2.3 Update the residuals  $\hat{\epsilon}_i^{(k+1)} = \hat{\epsilon}_i^{(k)} - \alpha_k f^{(k)}(\mathbf{x}_i)$
3. Return the boosted model  $\hat{f}(\mathbf{x}_i) = \sum_{k=1}^m \alpha_k f^{(k)}(\mathbf{x}_i)$




# Fitting regressions on sequential residuals

1. Set  $\hat{f}^{(0)}(\mathbf{x}_i) = 0$  and  $\hat{\epsilon}_i^{(1)} = Y_i$
2. For  $k = 1, \dots, m$ 
  - 2.1 Fit a tree  $\hat{f}^{(k)}$  with  $T$  leaves to  $\hat{\epsilon}$  using features  $\mathbf{x}$
  - 2.2 Update the estimator  $\hat{f}^{(k+1)} = f^{(k)} + \alpha_k f^{(k-1)}$
  - 2.3 Update the residuals  $\hat{\epsilon}_i^{(k+1)} = \hat{\epsilon}_i^{(k)} - \alpha_k f^{(k)}(\mathbf{x}_i)$
3. Return the boosted model  $\hat{f}(\mathbf{x}_i) = \sum_{k=1}^m \alpha_k f^{(k)}(\mathbf{x}_i)$



# Fitting regressions on sequential residuals

1. Set  $\hat{f}^{(0)}(\mathbf{x}_i) = 0$  and  $\hat{\epsilon}_i^{(1)} = Y_i$
2. For  $k = 1, \dots, m$ 
  - 2.1 Fit a tree  $\hat{f}^{(k)}$  with  $T$  leaves to  $\hat{\epsilon}$  using features  $\mathbf{x}$
  - 2.2 Update the estimator  $\hat{f}^{(k+1)} = f^{(k)} + \alpha_k f^{(k-1)}$
  - 2.3 Update the residuals  $\hat{\epsilon}_i^{(k+1)} = \hat{\epsilon}_i^{(k)} - \alpha_k f^{(k)}(\mathbf{x}_i)$
3. Return the boosted model  $\hat{f}(\mathbf{x}_i) = \sum_{k=1}^m \alpha_k f^{(k)}(\mathbf{x}_i)$



<https://www.r-bloggers.com/an-attempt-to-understand-boosting-algorithms/>

## Loss Functions

- ▶ The *loss function*  $L(Y_i, \hat{Y}_i)$  specifies prediction accuracy of  $\hat{Y}_i$

## Loss Functions

- ▶ The *loss function*  $L(Y_i, \hat{Y}_i)$  specifies prediction accuracy of  $\hat{Y}_i$
- ▶ The prediction  $\hat{Y}_i$  can be changed

## Loss Functions

- ▶ The *loss function*  $L(Y_i, \hat{Y}_i)$  specifies prediction accuracy of  $\hat{Y}_i$
- ▶ The prediction  $\hat{Y}_i$  can be changed
- ▶ The gradient of the loss function with respect to  $\hat{Y}_i$

$$\frac{\partial L(Y_i, \hat{Y}_i)}{\partial \hat{Y}_i}$$

gives the instantaneous  $\Delta$  in  $L(Y_i, \hat{Y}_i)$  at  $\hat{Y}_i \rightarrow \hat{Y}_i + \epsilon$

## Loss Functions

- ▶ The *loss function*  $L(Y_i, \hat{Y}_i)$  specifies prediction accuracy of  $\hat{Y}_i$
- ▶ The prediction  $\hat{Y}_i$  can be changed
- ▶ The gradient of the loss function with respect to  $\hat{Y}_i$

$$\frac{\partial L(Y_i, \hat{Y}_i)}{\partial \hat{Y}_i}$$

gives the instantaneous  $\Delta$  in  $L(Y_i, \hat{Y}_i)$  at  $\hat{Y}_i \rightarrow \hat{Y}_i + \epsilon$

**I.e., it tells how the loss function changes as  $\hat{Y}_i$  increases**  
(this is a very simple idea with somewhat complex notation)

# Gradient Decent

- ▶ But if we have partial derivatives we can use gradient decent

$$\nabla_{\hat{\mathbf{Y}}} L(\mathbf{Y}, \hat{\mathbf{Y}}) = \begin{bmatrix} \frac{\partial L(Y_1, \hat{Y}_1)}{\partial \hat{Y}_1} \\ \frac{\partial L(Y_2, \hat{Y}_2)}{\partial \hat{Y}_2} \\ \vdots \\ \frac{\partial L(Y_i, \hat{Y}_i)}{\partial \hat{Y}_i} \\ \vdots \\ \frac{\partial L(Y_n, \hat{Y}_n)}{\partial \hat{Y}_n} \end{bmatrix}$$


# Gradient Decent

- ▶ But if we have partial derivatives we can use gradient decent

$$\nabla_{\hat{\mathbf{Y}}} L(\mathbf{Y}, \hat{\mathbf{Y}}) = \begin{bmatrix} \frac{\partial L(Y_1, \hat{Y}_1)}{\partial \hat{Y}_1} \\ \frac{\partial L(Y_2, \hat{Y}_2)}{\partial \hat{Y}_2} \\ \vdots \\ \frac{\partial L(Y_i, \hat{Y}_i)}{\partial \hat{Y}_i} \\ \vdots \\ \frac{\partial L(Y_n, \hat{Y}_n)}{\partial \hat{Y}_n} \end{bmatrix}$$

to move in the direction of greatest decrease in loss

## Gradient Decent



The direction of maximal increase is the the gradient of a function

*(The negative of the gradient is the direction of maximal decrease)*

# Gradient Boosting

- If we take  $\hat{Y}_i$  to be our current prediction  $\phi_k(\mathbf{x}_i)$

## Gradient Boosting

- If we take  $\hat{Y}_i$  to be our current prediction  $\phi_k(\mathbf{x}_i)$ , the negative partial derivative of the loss function with respect to  $\phi_k(\mathbf{x}_i)$

$$\Delta_i^k = -\frac{\partial L(Y_i, \phi_k(\mathbf{x}_i))}{\partial \phi_k(\mathbf{x}_i)}$$

is the direction of greatest decrease in cost for  $\hat{Y}_i$

## Gradient Boosting

- If we take  $\hat{Y}_i$  to be our current prediction  $\phi_k(\mathbf{x}_i)$ , the negative partial derivative of the loss function with respect to  $\phi_k(\mathbf{x}_i)$

$$\Delta_i^k = -\frac{\partial L(Y_i, \phi_k(\mathbf{x}_i))}{\partial \phi_k(\mathbf{x}_i)}$$

is the direction of greatest decrease in cost for  $\hat{Y}_i$

⇒ We should update  $\phi_k(\mathbf{x}_i) \rightarrow \phi_{k+1}(\mathbf{x}_i)$  in that direction

## Gradient Boosting

- If we take  $\hat{Y}_i$  to be our current prediction  $\phi_k(\mathbf{x}_i)$ , the negative partial derivative of the loss function with respect to  $\phi_k(\mathbf{x}_i)$

$$\Delta_i^k = -\frac{\partial L(Y_i, \phi_k(\mathbf{x}_i))}{\partial \phi_k(\mathbf{x}_i)}$$

is the direction of greatest decrease in cost for  $\hat{Y}_i$

⇒ We should update  $\phi_k(\mathbf{x}_i) \rightarrow \phi_{k+1}(\mathbf{x}_i)$  in that direction

- But  $\hat{\mathbf{Y}} = \{\hat{Y}_1, \hat{Y}_2, \dots, \hat{Y}_n\}$  so we compromise between all  $\Delta_i^k$

$$\hat{W}(\cdot, \gamma_k) = \operatorname{argmin}_{W(\cdot, \gamma_k)} \sum_{i=1}^n D(\Delta_i^k, W(\mathbf{x}_i, \gamma_k))$$

and set

$$\phi_{k+1}(\mathbf{x}) = \phi_k(\mathbf{x}) + \alpha_k \hat{W}(\mathbf{x}, \gamma_k)$$

where  $\hat{W}(\cdot, \gamma_k)$  is a *weak learner* boosted at *learning rate*  $\alpha_k$

# Gradient Boosting *Tuning*

Gradient Boosting is typically applied using trees

# Gradient Boosting *Tuning*


Gradient Boosting is typically applied using trees

- ▶ Shallow trees (stumps?) provide weak learners/reduce tuning

# Gradient Boosting Tuning

Gradient Boosting is typically applied using trees

- ▶ Shallow trees (stumps?) provide weak learners/reduce tuning
- ▶ Nonetheless, number of leaves, depth, minimum split size, minimum leaf size, and restricted features\* can be tuned...



# Gradient Boosting Tuning

Gradient Boosting is typically applied using trees

- ▶ Shallow trees (stumps?) provide weak learners/reduce tuning
- ▶ Nonetheless, number of leaves, depth, minimum split size, minimum leaf size, and restricted features\* can be tuned...
- ▶ The learning rate/number of trees  $\alpha/m$  needs to be tuned



```
from sklearn.ensemble import GradientBoostingRegressor
from sklearn.grid_search import GridSearchCV
param_grid = {'learning_rate': [0.001, 0.01, 0.1],
              'max_depth': [4, 6],
              'min_samples_leaf': [4, 8, 16],
              'max_features': [0.75, 0.9, 1]}
model = GradientBoostingRegressor(n_estimators=3000)
gs_cv = GridSearchCV(model, param_grid).fit(X, y)

gs_cv.best_params_, gs_cv.best_score_, gs_cv.best_estimator_
```

# Gradient Boosting Tuning

Gradient Boosting is typically applied using trees



- ▶ Shallow trees (stumps?) provide weak learners/reduce tuning
- ▶ Nonetheless, number of leaves, depth, minimum split size, minimum leaf, size, and restricted features\* can be tuned...
- ▶ The learning rate/number of trees  $\alpha/m$  needs to be tuned
- ▶ Stochastic Gradient Boosting using subsamples also available



# Gradient Boosting Tuning

Gradient Boosting is typically applied using trees

- ▶ Shallow trees (stumps?) provide weak learners/reduce tuning
- ▶ Nonetheless, number of leaves, depth, minimum split size, minimum leaf, size, and restricted features\* can be tuned...
- ▶ The learning rate/number of trees  $\alpha/m$  needs to be tuned
- ▶ Stochastic Gradient Boosting using subsamples also available
- ▶ Careful parameter tuning is necessary for good performance



# Buy or Sell Gradient Boosting?

- ▶ Pros
- ▶ Cons

# Buy or Sell Gradient Boosting?

- ▶ Pros
  - ▶ Netflix Prize
- ▶ Cons

# Buy or Sell Gradient Boosting?

- ▶ Pros
  - ▶ Netflix Prize
- ▶ Cons
  - ▶ Extensive tuning required

# Buy or Sell Gradient Boosting?

- ▶ Pros
  - ▶ Netflix Prize
- ▶ Cons
  - ▶ Extensive tuning required
  - Unlike Bagging/Random Forests, Boosting *CAN* overfit

# Buy or Sell Gradient Boosting?

- ▶ Pros
  - ▶ Netflix Prize
- ▶ Cons
  - ▶ Extensive tuning required
  - ▶ Unlike Bagging/Random Forests, Boosting *CAN* overfit
  - ▶ Inherently non-parallelizable

# Buy or Sell Gradient Boosting?

- ▶ Pros
  - ▶ Netflix Prize
- ▶ Cons
  - ▶ Extensive tuning required
  - ▶ Unlike Bagging/Random Forests, Boosting *CAN* overfit
  - ▶ Inherently non-parallelizable

Why does this work better than Random Forests, etc.?

# More loss functions

## Regression

Squared Loss

$$\frac{1}{2}(y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2$$

Absolute Loss

$$|Y_i - \mathbf{x}_i^T \boldsymbol{\beta}|$$

## Classification

Log Loss



$$1 - \frac{1}{1+e^{-\mathbf{x}_i^T \boldsymbol{\beta}}}$$

(Bernoulli Deviance/Logistic Reg.)

Exponential Loss

$$\exp(-y_i \cdot \mathbf{x}_i^T \boldsymbol{\beta})$$

(Ada Boost)



# AdaBoost (Adaptive Boosting)

$$Y \in \{-1, 1\}$$

1. Set weights  $w_i^{(1)} = \frac{1}{n}$

# AdaBoost (Adaptive Boosting)

$$Y \in \{-1, 1\}$$

1. Set weights  $w_i^{(1)} = \frac{1}{n}$
2. For  $k = 1, \dots, m$

# AdaBoost (Adaptive Boosting)

$$Y \in \{-1, 1\}$$

1. Set weights  $w_i^{(1)} = \frac{1}{n}$

2. For  $k = 1, \dots, m$

2.1 Fit  $W(\cdot, \gamma_k)$  to the weighted data using weights  $w_i^{(k)}$

# AdaBoost (Adaptive Boosting)

$$Y \in \{-1, 1\}$$

1. Set weights  $w_i^{(1)} = \frac{1}{n}$

2. For  $k = 1, \dots, m$

2.1 Fit  $W(\cdot, \gamma_k)$  to the weighted data using weights  $w_i^{(k)}$

2.2 Compute error  $\hat{\epsilon}_k = \frac{\sum_{Y_i \neq W(x_i, \gamma_k)} w_i^{(k)}}{\sum w_i^{(k)}}$

# AdaBoost (Adaptive Boosting)

$$Y \in \{-1, 1\}$$

1. Set weights  $w_i^{(1)} = \frac{1}{n}$

2. For  $k = 1, \dots, m$

    2.1 Fit  $W(\cdot, \gamma_k)$  to the weighted data using weights  $w_i^{(k)}$

    2.2 Compute error  $\hat{\epsilon}_k = \frac{\sum_{Y_i \neq W(x_i, \gamma_k)} w_i^{(k)}}{\sum w_i^{(k)}}$

    2.3 Compute  $\alpha_k = \log \left( \frac{1}{\hat{\epsilon}_k} - 1 \right)$

# AdaBoost (Adaptive Boosting)

$$Y \in \{-1, 1\}$$

1. Set weights  $w_i^{(1)} = \frac{1}{n}$

2. For  $k = 1, \dots, m$

    2.1 Fit  $W(\cdot, \gamma_k)$  to the weighted data using weights  $w_i^{(k)}$

    2.2 Compute error  $\hat{\epsilon}_k = \frac{\sum_{Y_i \neq W(x_i, \gamma_k)} w_i^{(k)}}{\sum w_i^{(k)}}$

    2.3 Compute  $\alpha_k = \log \left( \frac{1}{\hat{\epsilon}_k} - 1 \right)$

    2.4 Set  $w_i^{(k+1)} = w_i^{(k)} e^{\alpha_k 1_{[Y_i \neq W(x_i, \gamma_k)]}}$

# AdaBoost (Adaptive Boosting)

$$Y \in \{-1, 1\}$$

1. Set weights  $w_i^{(1)} = \frac{1}{n}$

2. For  $k = 1, \dots, m$

    2.1 Fit  $W(\cdot, \gamma_k)$  to the weighted data using weights  $w_i^{(k)}$

    2.2 Compute error  $\hat{\epsilon}_k = \frac{\sum_{Y_i \neq W(x_i, \gamma_k)} w_i^{(k)}}{\sum w_i^{(k)}}$

    2.3 Compute  $\alpha_k = \log \left( \frac{1}{\hat{\epsilon}_k} - 1 \right)$


    2.4 Set  $w_i^{(k+1)} = w_i^{(k)} e^{\alpha_k 1_{[Y_i \neq W(x_i, \gamma_k)]}}$

3. Return  $G(\mathbf{x}_i) = \text{sign} (\sum_{k=1}^m \alpha_k W(\mathbf{x}_i, \gamma_k))$

# AdaBoost (Adaptive Boosting)

$$Y \in \{-1, 1\}$$


1. Set weights  $w_i^{(1)} = \frac{1}{n}$
2. For  $k = 1, \dots, m$ 
  - 2.1 Fit  $W(\cdot, \gamma_k)$  to the weighted data using weights  $w_i^{(k)}$
  - 2.2 Compute error  $\hat{\epsilon}_k = \frac{\sum_{Y_i \neq W(x_i, \gamma_k)} w_i^{(k)}}{\sum w_i^{(k)}}$
  - 2.3 Compute  $\alpha_k = \log \left( \frac{1}{\hat{\epsilon}_k} - 1 \right)$
  - 2.4 Set  $w_i^{(k+1)} = w_i^{(k)} e^{\alpha_k 1_{[Y_i \neq W(x_i, \gamma_k)]}}$
3. Return  $G(\mathbf{x}_i) = \text{sign} (\sum_{k=1}^m \alpha_k W(\mathbf{x}_i, \gamma_k))$



# AdaBoost (Adaptive Boosting)

$$Y \in \{-1, 1\}$$

1. Set weights  $w_i^{(1)} = \frac{1}{n}$
2. For  $k = 1, \dots, m$ 
  - 2.1 Fit  $W(\cdot, \gamma_k)$  to the weighted data using weights  $w_i^{(k)}$
  - 2.2 Compute error  $\hat{\epsilon}_k = \frac{\sum_{Y_i \neq W(x_i, \gamma_k)} w_i^{(k)}}{\sum w_i^{(k)}}$
  - 2.3 Compute  $\alpha_k = \log \left( \frac{1}{\hat{\epsilon}_k} - 1 \right)$
  - 2.4 Set  $w_i^{(k+1)} = w_i^{(k)} e^{\alpha_k \mathbf{1}_{[Y_i \neq W(x_i, \gamma_k)]}}$
3. Return  $G(\mathbf{x}_i) = \text{sign} (\sum_{k=1}^m \alpha_k W(\mathbf{x}_i, \gamma_k))$



\*Sean Sall's slides have an appendix showing that AdaBoost is exponential loss

# XGBoost (eXtreme Gradient Boosting)

- ▶ Faster/More Efficient
- ▶ Features are binned into percentiles for splitting (smaller search space)
- ▶ Handles missing data
- ▶ Handles mixed data types
- ▶ Allows for out-of-core computation
- ▶ Builds in smart memory management
- ▶ Install using pip – it has an sklearn interface