

# Clustering

Jack Bennetto

March 8, 2016

# Objectives

Today's objectives:

- Explain the difference between **supervised** and **unsupervised** learning
- Implement a **k-means** algorithm for clustering
- Choose the best  $k$  using the **elbow method** or **silhouette scores**
- Implement and interpret **hierarchical clustering**
- Discuss how **curse of dimensionality** affects clustering

# Agenda

## Morning:

- Supervised/unsupervised learning
- Clustering
- k-means algorithm

## Afternoon:

- Curse of dimensionality
- How to choose k
- Hierarchical and other clustering methods

# Supervised learning

Most of what you've learned so far

- Linear & logistic regression with lasso or ridge regularization
- Decision trees, bagging, random forest, boosting
- SVM
- kNN

Label == target == endogenous variable == dependent variable ==  $y$

# Unsupervised learning

No labels. No target.

Why use it?

# Unsupervised learning

No labels. No target.

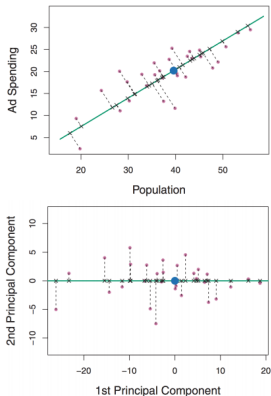
Used for:

- EDA
- Discovering latent variables
- Feature engineering
- Preprocessing

# Unsupervised learning

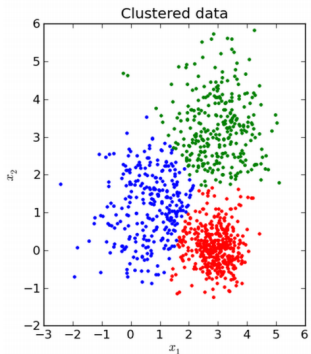
## PCA

Low-dim representation of data that explains good fraction of variance



## Clustering

Find homogenous subgroups among data



# Clustering Problem

Divide data into **distinct subgroups** such that observations within each group are similar.

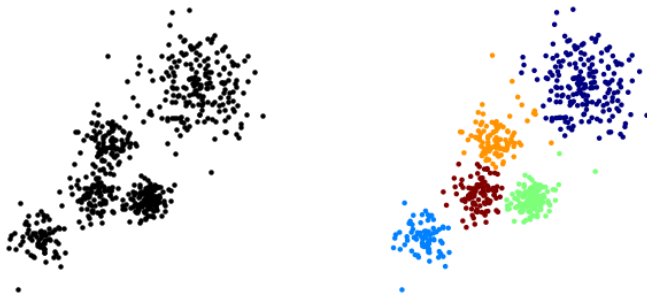


Figure 1:



# Various Algorithms

There are several approaches to clustering, each with variations.

- k-means clustering
- Hierarchical clustering
- Density-based clustering (DBSCAN)
- Distribution-based clustering
- ...

How do we measure how good the clustering is?

# Within-Cluster Sum of Squares

Measures the goodness of a clustering

$$W(C) = \sum_{k=1}^K \frac{1}{n_k} \sum_{C(i)=k} \sum_{C(j)=k} \|x_i - x_j\|^2$$

How long will it take to optimize this?

# Within-Cluster Sum of Squares

Measures the goodness of a clustering

$$W(C) = \sum_{k=1}^K \frac{1}{n_k} \sum_{C(i)=k} \sum_{C(j)=k} \|x_i - x_j\|^2$$

How long will it take to optimize this?

Do you need to normalize?

# k-means Algorithm

## The k-means algorithm

- Choose a number of clusters  $k$
- Randomly assign each point to a cluster
- Repeat:
  - ▶ a. For each of  $k$  clusters, compute cluster *centroid* by taking mean vector of points in the cluster
  - ▶ b. Assign each data point to cluster for which centroid is closest (Euclidean)

... until clusters stop changing

# k-means Algorithm

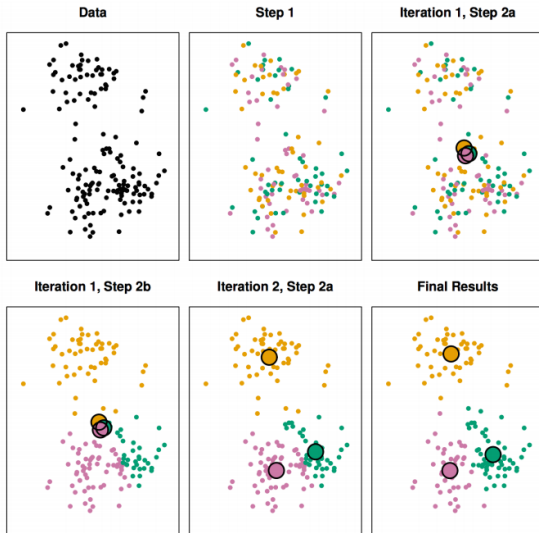


Figure 2: The k-means algorithm

k-means finds a local minimum, and sometimes a bad one.

One alternative: use random points as cluster center.

k-means++ is the same algorithm but with a different start.

- Choose one point for first center.
- Repeat:
  - ▶ Calculate distance from each point to the nearest center  $d_i$
  - ▶ Choose a point to be the next center, randomly, using a weighed probability  $d_i^2$

... until k centers have been choosen.

# The Curse of Dimensionality

Random variation in extra dimensions can many hide significant differences between clusters.

The more dimensions there are, the worse the problem.

More than 10 dimensions: consider PCA first.

# How Many Clusters?

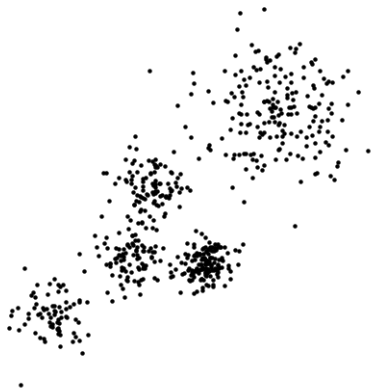


Figure 3:  
Clustering



# How Many Clusters?

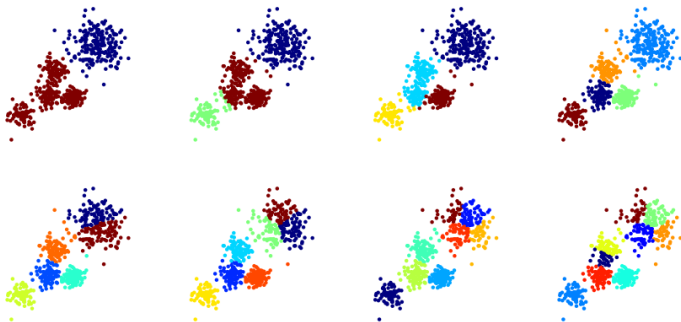


Figure 4:

# Choosing K

Can we use within-cluster sum of squares (WCSS) to choose  $k$ ?

# Choosing K

More clusters  $\implies$  lower WCSS

Several measures for the “best” K - no easy answer

- The Elbow Method
- Silhouette Score
- GAP Statistic

# Choosing K – The Elbow Method

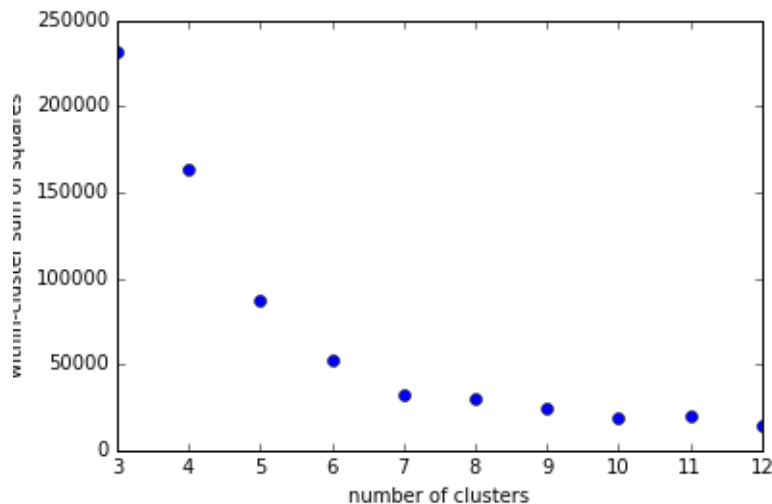


Figure 5:  
Clustering

# Choosing K – Silhouette Score

For each point  $x_i$ :

- $a(i)$  average dissimilarity of  $x_i$  with points in the same cluster
- $b(i)$  average dissimilarity of  $x_i$  with points in the nearest cluster
  - ▶ “nearest” means cluster with the smallest  $b(i)$

$$\text{silhouette}(i) = \frac{b(i) - a(i)}{\max(a(i), b(i))}$$

What's the range of silhouette scores?

# Choosing K – Silhouette Score

Silhouette score is between 1 and -1

- near 1: very small tight cluster.
- 0: at the edge of two clusters; could be in either.
- $< 0$ : oops.

The higher the the average silhouette score, the tighter and more separated the clusters.

# Choosing K – Silhouette Score

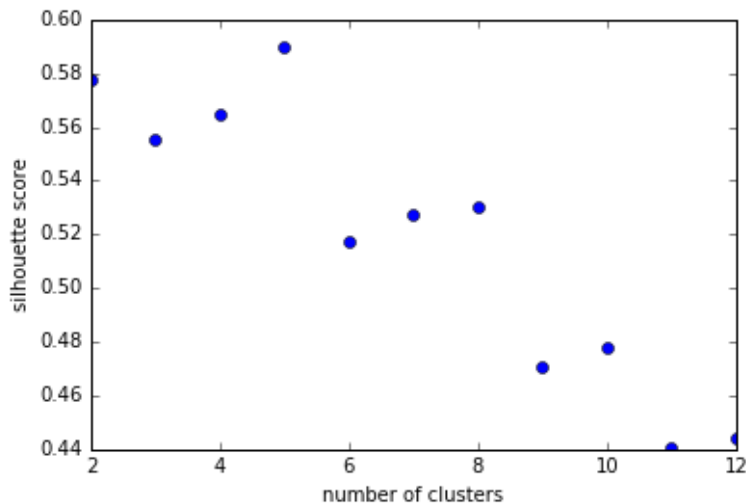


Figure 6:  
Clustering

# Silhouette Graph

(see notebook)



# Choosing K – GAP Statistic

For each  $k$ , compare  $W(k)$  (within-cluster sum of squares) with that of randomly generated “reference distributions”

Generate  $B$  distributions

$$Gap(k) = \frac{1}{B} \sum_{b=1}^B \log W_{kb} - \log W_k$$

Choose smallest  $k$  such that  $Gap(k) \geq Gap(k+1) - s_{k+1}$

where  $s_k$  is the standard error of  $Gap(k)$

# Hierarchical Clustering

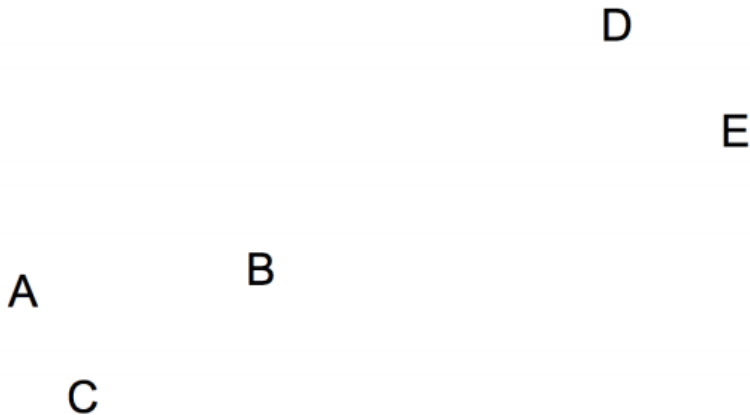
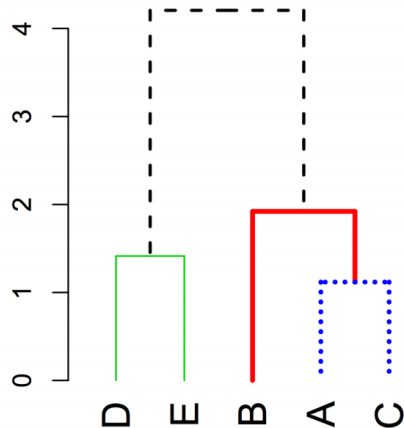
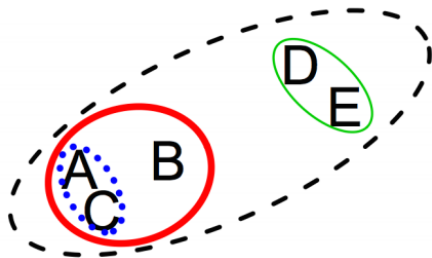


Figure 7:

# Hierarchical Clustering



# Hierarchical Clustering

## Algorithm

- Assign each point to its own cluster
- Repeat:
- Compute distances between clusters
- Merge closest clusters . . . until all are merged

How do we define dissimilarity between clusters?

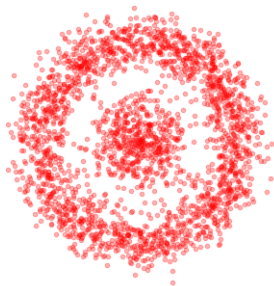
# Hierarchical Clustering – Linkage

How do we define dissimilarity between clusters?

- **Complete:** Maximum pairwise dissimilarity between points in clusters – good
- **Average:** Average of pairwise dissimilarity between points in clusters – also good
- **Single:** Minimum pairwise dissimilarity between points in clusters – not as good; can lead to long narrow clusters
- **Centroid:** Dissimilarity between centroids – used in genomics; risk of inversions

# Problems with k-means

k-means has limitations.



Two parameters (number of clusters not specified)

- $\epsilon$ : distance between points for them to be connected
- minPts: number of connected points for a point to be a “core” point

A cluster is all connected core points, plus others within  $\epsilon$  of one of those. Other points are noise.

# Distribution-based clustering

Assume clusters follow some (generally gaussian) distribution

Find distributions with the **maximum likelihood** to produce this result

...except you don't know which point is part of which cluster, so you need to add some hidden variables and follow an **expectation-maximization** (EM) algorithm.