

Sampling and Estimation

Brian J. Mann

Estimation

Objectives

This morning we'll talk about estimating statistical distributions from observed data

- ▶ Recall what the expected value and variance of a random variable are
- ▶ Use Maximum Likelihood Estimation to estimate a parametric distribution from observed data
- ▶ Use the Method of Moments to estimate a parametric distribution from observed data
- ▶ Understand how Kernel Density Estimation estimates a non-parametric distribution from observed data

Expected Value

Recall that the *expected value* of a discrete random variable is the weighted sum:

$$E[X] = P(X = x_1) * x_1 + P(X = x_2) * x_2 + \cdots + P(X = x_n) * x_n$$

For a continuous random variable with density function f :

$$E[X] = \int xf(x) dx$$

Variance (1/2)

The *variance* of a random variable X is the expected value of the square difference from the mean:

$$\begin{aligned}\text{Var}(X) &= E[(X - E[X])^2] \\ &= E[X^2 - 2XE[X] + E[X]^2] \\ &= E[X^2] - E[X]^2\end{aligned}$$

Variance (2/2)

For a discrete random variable:

$$\text{Var}(X) = \sum_i P(X = x_i) * (x_i - E[X])^2$$

For a continuous random variable with density f :

$$\text{Var}(X) = \int (x - E[X])^2 f(x) dx$$

Inference

Parametric

- ▶ Assumes the data is drawn from a class of distributions determined by numeric parameters
- ▶ For example $Norm(\mu, \sigma)$, $Poisson(\lambda)$, or $Binom(n, p)$
- ▶ Determine which parameters are the best fit for the data

Non-Parametric

- ▶ Make no assumption about the family of distribution the data is drawn from
- ▶ More flexible
- ▶ Less interpretable, often hard to compute anything about the inferred distribution

Maximum Likelihood Estimation (MLE) (Parametric)

Assume each data point is drawn independently from the same distribution with density $f(x|\theta)$. Since the draws are independent the joint density function is

$$f(x_1, x_2, \dots, x_n|\theta) = f(x_1|\theta) * f(x_2|\theta) * \dots * f(x_n|\theta)$$

If we have a formula for f in terms of the parameters θ , we can find the values of theta which maximizes the *likelihood*

$$\mathcal{L}(\theta|x_1, x_2, \dots, x_n) = f(x_1, x_2, \dots, x_n|\theta) = \prod f(x_i|\theta)$$

or equivalently the *log-likelihood*

$$\ell(\theta|x_1, x_2, \dots, x_n) = \sum \log f(x_i|\theta)$$

$$\hat{\theta} = \operatorname{argmax}_{\theta} \log f(x_1, x_2, \dots, x_n|\theta)$$

Example - MLE (1/3)

Suppose we flip a coin N times and get H heads. We want an estimate for how biased the coin is. Each flip is a Bernoulli trial with parameter p . The joint distribution is $\text{Binom}(N, p)$, so we need to find p which minimizes

$$\log p^H (1 - p)^{N-H}$$

Example - MLE (2/3)

Suppose that $x_1, \dots, x_n \sim N(\mu, \sigma)$. Ignoring some constants, the likelihood function is

$$\begin{aligned}\mathcal{L}(\mu, \sigma) &= \prod \frac{1}{\sigma} \exp\left(-\frac{1}{2\sigma^2}(x_i - \mu)^2\right) \\ &= \frac{1}{\sigma^n} \exp\left(-\frac{1}{2\sigma^2} \sum (x_i - \mu)^2\right) \\ &= \frac{1}{\sigma^n} \exp\left(-\frac{nS^2}{2\sigma^2}\right) \exp\left(-\frac{n(\bar{x} - \mu)^2}{2\sigma^2}\right)\end{aligned}$$

where \bar{x} is the sample mean and $S^2 = \frac{1}{n} \sum (x_i - \bar{x})^2$ is a biased estimator for the variance.

Example = MLE (3/3)

The log-likelihood is

$$\ell(\mu, \sigma) = -n \log \sigma - \frac{nS^2}{2\sigma^2} - \frac{n(\bar{x} - \mu)^2}{2\sigma^2}$$

Maximizing this gives estimates $\bar{\mu} = \bar{x}$ and $\bar{\sigma} = S$.

Maximum A Posteriori (MAP) (Parametric)

Generalization of MLE where we assume some prior distribution on the parameters θ

$$\mathcal{L}(\theta|x) = \frac{f(x|\theta)g(\theta)}{\int_{\Theta} f(x|t)g(t)dt}$$

To find the optimal θ we find

$$\hat{\theta} = \operatorname{argmax}_{\theta} \frac{f(x|\theta)g(\theta)}{\int_{\Theta} f(x|t)g(t)dt} = \operatorname{argmax}_{\theta} f(x|\theta)g(\theta)$$

To get MLE, assume a uniform prior on θ so that the function g disappears from the *argmax* above

Method of Moments (MOM) (Parametric)

Older method, generally MLE is preferred. But good to know anyway.

- ▶ A *moment* of a distribution is $E[X]$, $E[X^2]$, $E[X^3]$, \dots
- ▶ $E[X]$ is the first moment, $E[X^2]$ is the second moment, etc. \dots
- ▶ Use the moments to derive as many equations as parameters, and then solve

Example - MOM (1/2)

Suppose we flip a coin N times again, and get H heads. Let's use MOM this time to estimate p , the probability of flipping a head. Since the number of heads of N flips is modeled by a Binomial distribution we can compute the first moment

$$E[X] = Np$$

Since we have a single unknown, we stop at the first moment. We compute the sample first moment $\bar{x} = H$ and set this equal to theoretical first moment

$$H = Np$$

So we estimate

$$\hat{p} = H/N$$

Example - MOM (2/2)

Suppose we have data sampled from a symmetric uniform distribution with unknown bounds $X \sim \text{Unif}(-b, b)$. The first moment is

$$E[X] = 0$$

so that doesn't help. The second moment is

$$E[X^2] = \text{Var}(X) + E[X]^2 = \text{Var}(X) = b^2/3$$

Computing the sample variance s^2

$$s^2 = b^2/3$$

so that

$$\hat{b} = \sqrt{3s^2}$$

Kernel Density Estimation (KDE) (Non-Parametric)

A *kernel* is another word for a density function of a distribution with mean 0.

Kernel Density Estimation estimates a distribution empirically given data by summing kernels centered at each point. The density function of the kernel density estimate is:

$$\hat{f}_h(x) = \frac{1}{nh} \sum_i K\left(\frac{x - x_i}{h}\right)$$

K is a kernel. The parameter h is called the *bandwidth*, and it's analogous to the width of bins in a histogram.

Example - KDE




Figure 1: KDE for $x_1 = -2.1$, $x_2 = -1.3$, $x_3 = -0.4$, $x_4 = 1.9$, $x_5 = 5.1$, $x_6 = 6.2$

Sampling

Objectives

This afternoon you will

- ▶ Discover the Central Limit Theorem
- ▶ Apply the Central Limit Theorem to construct confidence intervals for the mean of a population
- ▶ Use bootstrapping to construct confidence intervals for any population statistic

Statistical Discovery in General

1. Ask a question
2. Design an experiment
3. Collect data (Sampling)
4. Analyze data (Estimation/Inference)
5. Repeat




Figure 2: Statistical Inference

Make sure you have good data!

Your results are only as good as your data. Garbage in, garbage out.

- ▶ Your data should be representative of the population
- ▶ Important to make sure there is no bias when designing your experiment or “randomly” sampling
- ▶ For example, if you want to estimate the average height of a person in the US, but all the people you measure are in the 90th percentile for weight, something is wrong!

Simple Random Sampling

The most common way to sample from a population is called *simple random sampling*

- ▶ Each subject has an equal chance of being selected from the population
- ▶ If your population is x_1, \dots, x_n , to sample choose a number uniformly at random from $1, \dots, n$ and select that observation

Central Limit Theorem

One of the most important results in classical statistical inference is the *Central Limit Theorem* which says that if X_1, X_2, \dots, X_n are i.i.d. random variables with mean μ and variance σ^2 then their mean

$$\bar{X} = \frac{X_1 + \dots + X_n}{n}$$

is approximately normally distributed with mean μ and variance $\frac{\sigma^2}{n}$

$$\bar{X} \sim N\left(\mu, \frac{\sigma}{\sqrt{n}}\right)$$

Example - CLT

Recall that $\text{Binom}(n, p)$ is the sum of n independent Bernoulli trials with parameter p . This means that

$$\text{Binom}(n, p) \sim N\left(np, \sqrt{np(1-p)}\right)$$

Why??

Confidence Intervals (1/2)

A *confidence interval* is an interval estimate of the true parameter of your population

- ▶ An α confidence interval is an interval centered around estimated parameter which contains the true value of that parameter with *confidence* α (α is usually 99%, 95%, 90%, or 80%)
- ▶ In other words, if you resample or rerun the experiment many times, α percent of the time the true value will be in the computed confidence interval
- ▶ It is *not* a statement that the true value of the parameter is contained in the interval with a certain probability

Confidence Intervals (2/2)

For $n \geq 30$ a 95% confidence interval for the mean is

$$\left(\bar{x} - 1.96 \frac{\sigma}{\sqrt{n}}, \bar{x} + 1.96 \frac{\sigma}{\sqrt{n}}\right)$$

- ▶ Why??
- ▶ Since we don't know σ , use the sample standard deviation s instead
- ▶ If n is small, the central limit theorem does not guarantee normality. We need a t -distribution instead

$$\bar{x} \pm t_{(\alpha/2, n-1)} \frac{s}{\sqrt{n}}$$

Check for Mastery

Using Python, sample 100 times from a normal distribution. Compute the sample mean and a 95% confidence interval. Is the true mean in your interval?!? Rerun your code several time and see if you find an interval which doesn't contain the true mean.

Bootstrapping (1/2)

Another way to generate confidence intervals for a population parameter is through a process called bootstrapping

- ▶ Simple idea: sample from your observed data *with replacement* B times
- ▶ With these B samples, compute the statistic (i.e. mean, median, variance, etc. . .) of interest and then estimate the sample variance
- ▶ Computationally expensive

Bootstrapping (2/2)

How to bootstrap:

Start with n i.i.d. samples X_1, \dots, X_n .

For i from 1 to B :

1. Sample X_1^*, \dots, X_n^* with replacement from your data
2. Compute your sample statistic $\theta_i^* = g(X_1^*, \dots, X_n^*)$

Then compute

$$v_{boot} = \frac{1}{B} \sum_{b=1}^B \left(\theta_b^* - \frac{1}{B} \sum_{r=1}^B \theta_r^* \right)^2$$

which is the sample variance of your statistic

Bootstrap Confidence Intervals (The Normal Interval)

There are a few different ways to build bootstrap confidence intervals that rely on differing assumptions. The first is the *normal interval*

- ▶ If your parameter is approximately normally distributed (like the mean of a sample with $n > 30$) your interval will be

$$\theta_n \pm z_{\alpha/2} \hat{se}_{boot}$$

where $\theta_n = g(X_1, \dots, X_n)$ is your estimate of the parameter, z is standard normal (e.g. for 95% it is 1.96), and $\hat{se}_{boot} = \sqrt{v_{boot}}$ is the bootstrap estimated standard error of your parameter

Bootstrap Confidence Intervals (Percentile Method)

Let θ_β^* be the β sample quantile of your bootstrap sample statistics $(\theta_1^*, \dots, \theta_B^*)$. Then an α bootstrap percentile interval is

$$C_n = (\theta_{1-\alpha/2}^*, \theta_{\alpha/2}^*)$$

Why Bootstrap?

Why would we use bootstrapping over standard confidence intervals?

- ▶ Small sample size
- ▶ The distribution of the statistic is complicated or hard to compute