

Optimization in Data Science

Brian J. Mann

Feb 29, 2016

Objectives

- ▶ Understand How Gradient Descent Works
- ▶ Use Gradient Descent to Optimize the Cost Function For Logistic Regression
- ▶ Stochastic Gradient Descent and Newton's Method

Agenda

- ▶ Morning

1. What is Gradient Descent and why do we need it?
2. An example of gradient descent
3. Using Gradient Descent to Solve Logistic Regression

- ▶ Afternoon

1. Stochastic Gradient Descent and Examples
2. Newton's Method and Examples

Cost Functions

- ▶ Machine learning often involves fitting a model to test data
- ▶ The best fit is often determined using a *cost function* or *likelihood function*

- ▶ Linear Regression:

$$\sum (y_i - \beta^T x_i)^2$$

- ▶ Logistic Regression:

$$\sum y_i \log g(\beta^T x_i) + (1 - y_i) \log(1 - g(\beta^T x_i))$$

$$\left(g(z) = \frac{1}{1 + e^{-z}} \right)$$

Linear Regression

- ▶ The cost function $\sum (y_i - \beta^T x_i)^2$ can be represented in matrix format:

$$\|y - X\beta\|^2$$

- ▶ Has a closed-form solution for the minimum

$$\beta = (X^T X)^{-1} X^T y$$

- ▶ $(X^T X)^{-1}$ very hard to compute if you have many features (say $> 10,000$)
- ▶ Is there a quicker way?

Logistic Regression

- ▶ The log-likelihood function

$$\sum y_i \log g(\beta^T x_i) + (1 - y_i) \log(1 - g(\beta^T x_i))$$

has no such closed form for its maximum.

- ▶ How will you find the maximum?

Ideas?

Gradient Descent

- ▶ Algorithm for finding the minimum of a function
- ▶ Question: Can be used to find maxima by _____?

Recall

- ▶ The *gradient* of a multivariate function $f(x_1, \dots, x_n)$ is

$$\nabla f(a) = \left(\frac{\partial f}{\partial x_1}(a), \dots, \frac{\partial f}{\partial x_n}(a) \right)$$

- ▶ $\nabla f(a)$ points in the direction of greatest increase of f at a

Gradient Descent

- ▶ Minimize f
- ▶ Choose:
 - ▶ a starting point x
 - ▶ *learning rate* α
 - ▶ threshold ϵ
- ▶ Move in the direction of $-\nabla f(x)$:
 - ▶ Set $y = x - \alpha \nabla f(x)$
- ▶ If $\frac{|f(x) - f(y)|}{|f(x)|} < \epsilon$, return $f(y)$ as the min, and y as the argmin

Gradient Descent

- ▶ α is called the *step-size* or *learning rate*
 - ▶ If ' α ' is too small, convergence takes a long time
 - ▶ If ' α ' is too big, can overshoot the minimum

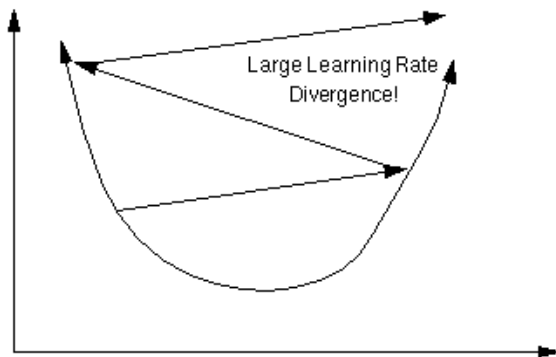


Figure 1: α too large

Choosing Alpha

- ▶ If the value of

$$\frac{|\nabla f(x) - \nabla f(y)|}{|x - y|}$$

is bounded above by some number $L(\nabla f)$ then

$$\alpha \leq \frac{1}{L(\nabla f)}$$

will converge.

- ▶ For example:
 - ▶ $f(x) = x^2$
 - ▶ $L(\nabla f) = 2$
 - ▶ $\alpha = 1/2$ will be the best value

Adaptive Step Size

- ▶ Change α at each iteration
- ▶ (Barzilai and Borwein 1998)
 - ▶ Suppose x_i is the value of x at the iteration i
 - ▶ $\Delta x = x_i - x_{i-1}$
 - ▶ $\Delta g(x) = \nabla f(x_i) - \nabla f(x_{i-1})$
 - ▶ At each step

$$\alpha = \frac{\Delta g(x)^T \Delta x}{\|\Delta g(x)\|^2}$$

is a good choice of α

Convergence Criteria

Choices:

- ▶ $\frac{|f(x)-f(y)|}{|f(x)|} < \epsilon$
- ▶ Max number of iterations
- ▶ Magnitude of gradient $|\nabla f| < \epsilon$

Gradient Ascent

- ▶ To maximize f , we can minimize $-f$
- ▶ Still use almost the same algorithm
 - ▶ Just replace

$$y = x - \alpha \nabla f(x)$$

with

$$y = x + \alpha \nabla f(x)$$

Some Examples

- ▶ Examples

What Can Go Wrong

- ▶ Where do you think gradient descent fails?

Example

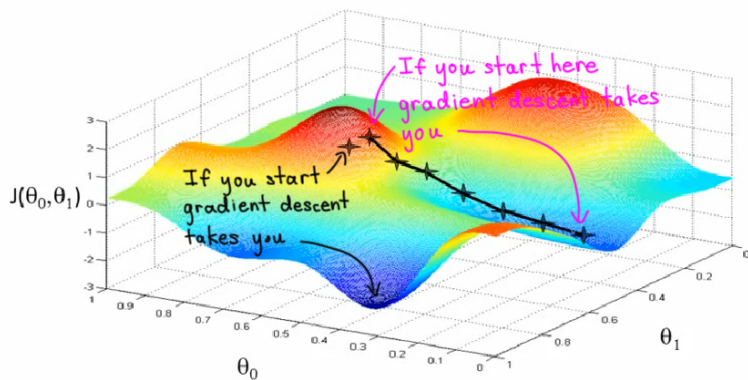


Figure 2: Non-convex function

More Bad Things

- ▶ Need differentiable and convex cost/likelihood function
- ▶ Only finds local extrema
- ▶ Poor performance without feature scaling

Back to Logistic Regression

- ▶ Trying to maximize the log-likelihood function

$$\ell(\beta) = \sum y_i \log g(\beta^T x_i) + (1 - y_i) \log(1 - g(\beta^T x_i))$$

- ▶ To use gradient ascent: need to compute $\nabla \ell(\beta)$

More Logistic Regression

First, let's compute the derivative of the sigmoid function g :

$$\begin{aligned}\frac{d}{dz}g(z) &= \frac{d}{dz} \frac{1}{1 + e^{-z}} \\&= \frac{d}{dz} (1 + e^{-z})^{-1} \\&= - (1 + e^{-z})^{-2} (-e^{-z}) \\&= \frac{1}{1 + e^{-z}} \left(\frac{e^{-z}}{1 + e^{-z}} \right) \\&= g(z) \left(\frac{1 + e^{-z} - 1}{1 + e^{-z}} \right) \\&= g(z) \left(\frac{1 + e^{-z}}{1 + e^{-z}} - \frac{1}{1 + e^{-z}} \right) \\&= g(z) (1 - g(z))\end{aligned}$$

Figure 3: $g'(z)$

More Logistic Regression

- Using this and the chain rule, compute $\frac{\partial \ell}{\partial \beta_i}$

$$\begin{aligned}\frac{\partial}{\partial \beta_j} \ell(\beta) &= \frac{\partial}{\partial \beta_j} \left(\sum_{i=1}^n y_i \log(g(\beta^T x_i)) + (1 - y_i) \log(1 - g(\beta^T x_i)) \right) \\&= \sum_{i=1}^n \left(\frac{y_i}{g(\beta^T x_i)} - \frac{1 - y_i}{1 - g(\beta^T x_i)} \right) \frac{\partial}{\partial \beta_j} (g(\beta^T x_i)) \\&= \sum_{i=1}^n \left(\frac{y_i}{g(\beta^T x_i)} - \frac{1 - y_i}{1 - g(\beta^T x_i)} \right) g(\beta^T x_i) (1 - g(\beta^T x_i)) \frac{\partial}{\partial \beta_j} (\beta^T x_i) \\&= \sum_{i=1}^n \left(\frac{y_i}{g(\beta^T x_i)} - \frac{1 - y_i}{1 - g(\beta^T x_i)} \right) g(\beta^T x_i) (1 - g(\beta^T x_i)) x_{ij}\end{aligned}$$

Figure 4: Computing $\nabla \ell$

More Logistic Regression

- Simplifying:

$$\begin{aligned}\frac{\partial}{\partial \beta_j} \ell(\beta) &= \sum_{i=1}^n \left(y_i (1 - g(\beta^T x_i)) - (1 - y_i) g(\beta^T x_i) \right) x_{ij} \\ &= \sum_{i=1}^n \left(y_i - y_i g(\beta^T x_i) - g(\beta^T x_i) + y_i g(\beta^T x_i) \right) x_{ij} \\ &= \sum_{i=1}^n \left(y_i - g(\beta^T x_i) \right) x_{ij} \\ &= \sum_{i=1}^n \left(y_i - f(x_i) \right) x_{ij}\end{aligned}$$

Figure 5: Computing $\nabla \ell$

More Logistic Regression

- ▶ This is what you'll use to update the value of β in each iteration of gradient descent

Stochastic Gradient Descent

Why Not Regular Gradient Descent?

- ▶ Can you think of some problems with gradient descent as we learned it this morning?

Problems with Gradient Descent

- ▶ Memory constrained
 - ▶ Need to store all data in memory
- ▶ CPU constrained
 - ▶ Cost function is a function of *all* data
- ▶ What if you are getting new data continuously?

Solution

- ▶ Only use a single data point, or a small subset of your data!

Algorithm

- ▶ Same as gradient descent except **at each step compute the cost function by using just one observation**
- ▶ For example in linear regression, instead of computing the gradient of

$$\sum_i (y_i - \beta^T x_i)^2$$

randomly select some x_i, y_i and compute the gradient of

$$y_i - \beta^T x_i$$

Properties

- ▶ Faster than *batch* Gradient Descent on average
- ▶ Prone to oscillation around an optimum
- ▶ Only requires one observation in memory at once

Variants

- ▶ Can use a small subset of your data instead of a single observations
 - ▶ “Minibatch” Stochastic GD
- ▶ “Online” Stochastic GD updates the model by performing a gradient descent step each time a new observation is collected

Newton's Method

What Is It?

- ▶ Optimization technique similar to gradient descent
- ▶ Uses a root-finding method applied to $f'(x)$

Algorithm in One Dimension

- ▶ Choose initial x_0
- ▶ While $f'(x) > \epsilon$:

$$x_{i+1} = x_i - \frac{f'(x_i)}{f''(x_i)}$$

Higher Dimensions

► $y_{i+1} = y_i - H(y_i)^{-1} \nabla f(y_i)$

($H(a) = \left[\frac{\partial^2 f}{\partial x_i \partial x_j}(a) \right]$ is the *Hessian* matrix, the matrix of second partial derivatives at a)

Problems

- ▶ Hessian might be singular, or computation can be slow
- ▶ Can diverge with a bad starting guess