

Gradient Descent and such

Schwartz

July 22, 2017

A brief history of Optimization

- 300 bc Euclid considers minimal distance from point to a line & proves square is the greatest area rectangle
- 1615 Kepler optimizes dimensions of wine barrel & formulates an early version of the (classical) secretary problem while looking for a new wife
- 1636 Fermat shows derivatives vanish at extremes & light travels between two points in minimal time
- 1660s Newton & Leibniz create the mathematical basis of calculus and hence optimization calculus
- 1696 Johann & Jacob Bernoulli study Brachistochrone problem – calculus optimization is born
- 1712 König shows that the shape of honeycomb is optimal. The French Academy of Sciences declares the phenomenon as divine guidance
- 1740 Euler's publication begins the research on a general theory of calculus optimization
- 1754 Lagrange makes his first of many findings regarding calculus optimization at age 19
- 1900's The first optimization algorithms are presented by Weierstrass, Steiner, Hamilton and Jacobi
- 1806 Legendre presents the least square method, which also Gauss claims to have invented
- 1815 "The Law of Diminishing Returns" (introduced simultaneously by Malthus, Torrens, West, and Ricardo) uses a (quasi) concave function
- 1826 Fourier formulates linear programming (LP) for solving mechanics and probability problems
- 1847 Cauchy presents the gradient method
- 1857 Gibbs shows chemical equilibrium is minimum energy
- 1870s The marginalist revolution in economics shifts the focus of economists to maximizing individuals utility
- 1880s Convexity theory created – Jensen introduces convex functions in 1905 – Minkowski convex sets in 1911
- 1917 Hancock publishes the first text book on optimization: "Theory of Minima and Maxima"
- 1917 Thompson's "On Growth and Form" applies optimization to analyze the forms of living organisms
- 1928 Ramsey studies optimal economic growth which becomes optimal growth theory in the 1950's
- 1932 Menger generalizes the traveling salesman problem
- 1939 Kantorovich presents LP-model & solution algorithm and receives Nobel prize in 1975 with Koopmans
- 1944 Neuman and Morgenstern, and Wald (1947) solve sequential problems w/ dynamic programming (DP)
- 1947 Dantzig (USAF) presents the Simplex method for LP-problems, Neumann establishes duality theory
- 1950's Electronic calculation initiates algorithmic research
- 1951 Markowitz presents portfolio theory using quadratic programing (QP) and receives the 1990 Nobel prize
- 1954 Ford & Fulkerson introduce combinatorial optimization for network research problems
- 1960's Space race sparks optimal control theory research
- 1970s Complexity analysis influences optimization theory
- 1980's Heuristic global optimization algorithms for large scale problems gain popularity as computers improve

Objectives

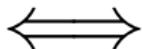
- ▶ Loss Function
 - ▶ Cost Function
 - ▶ Objective Function
-
- ▶ Gradient Descent
 - ▶ Stochastic Gradient Descent
 - ▶ Newton's Method

Exercise of the day

Find parameters that maximizes
the fit of model to data

Exercise of the day

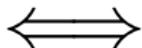
Find parameters that maximizes
the fit of model to data



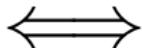
minimize distance between
predicted and observed values

Exercise of the day

Find parameters that maximizes
the fit of model to data



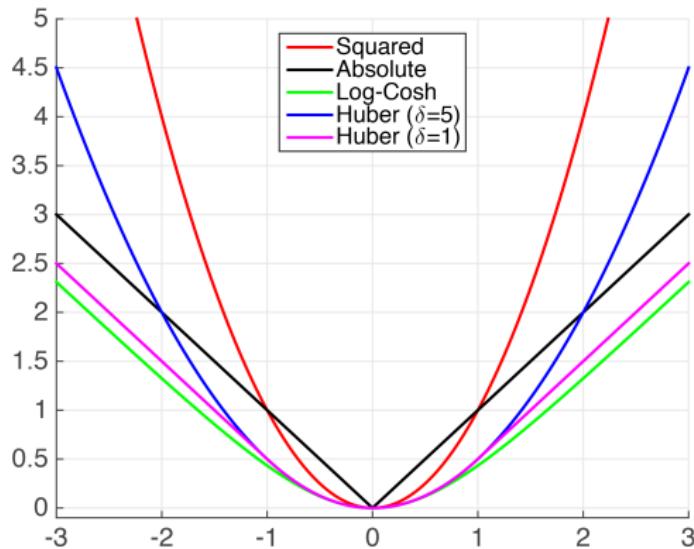
minimize distance between
predicted and observed values



$$\hat{Y}_i \approx Y_i$$

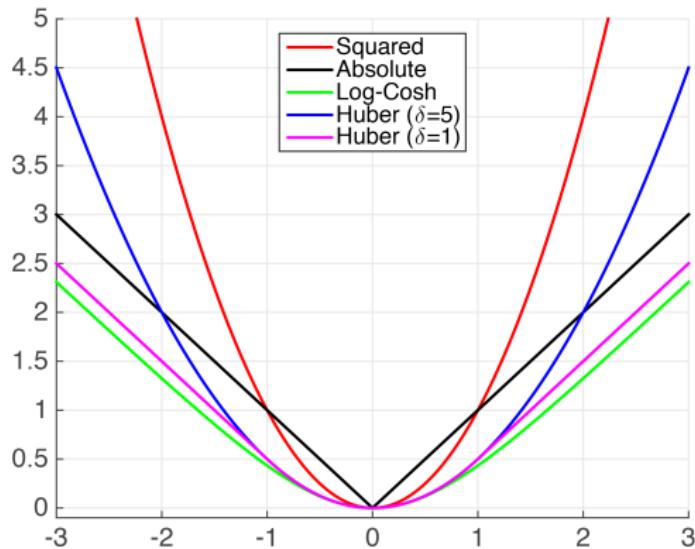
Regression Loss Functions

$$Y_i - \hat{Y}_i$$



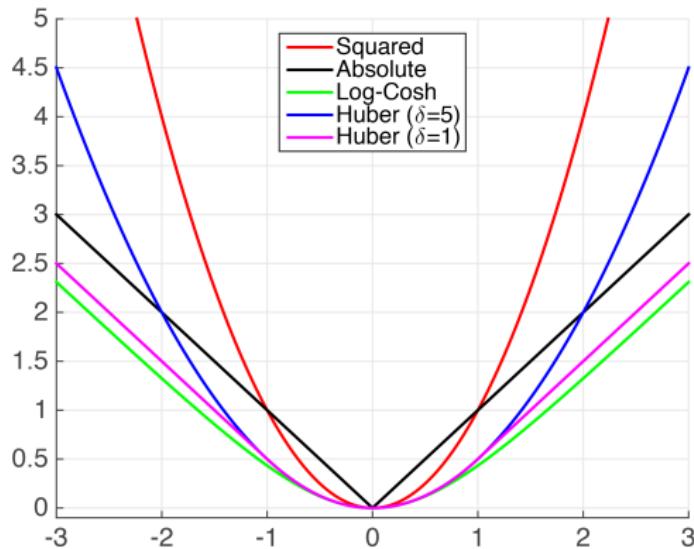
Regression Loss Functions

$$(Y_i - \hat{Y}_i)^2$$



Regression Loss Functions

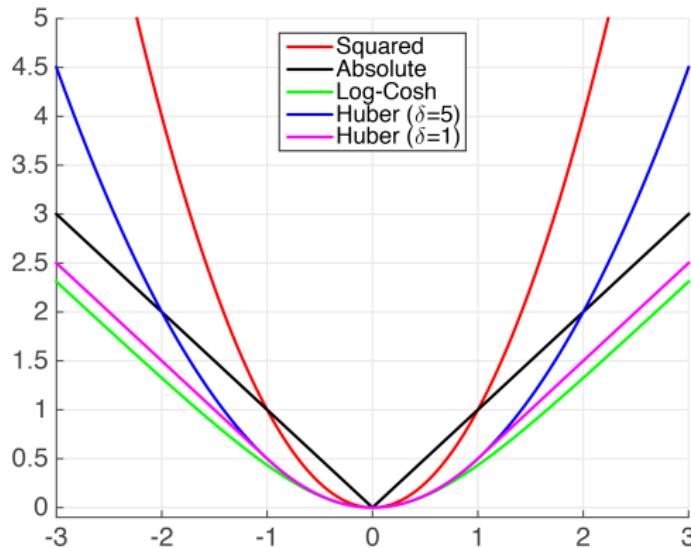
$$|Y_i - \hat{Y}_i|$$



Regression Loss Functions

$$\ln(\cosh(Y_i - \hat{Y}_i))$$

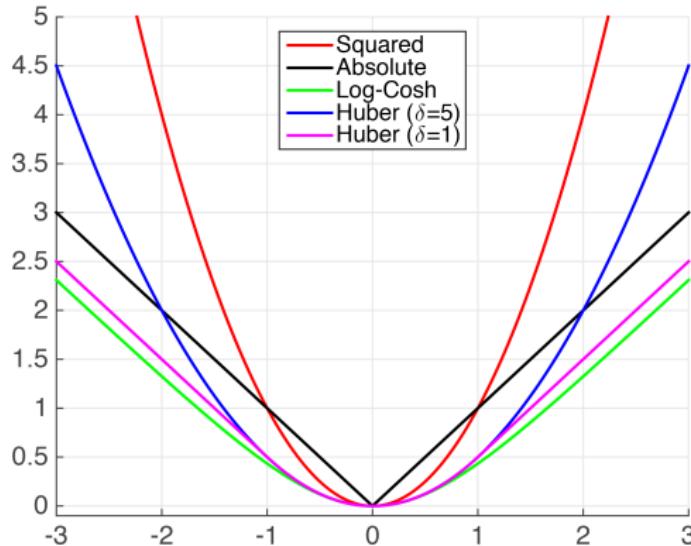
the good ol' *hyperbolic cosine* function



Regression Loss Functions

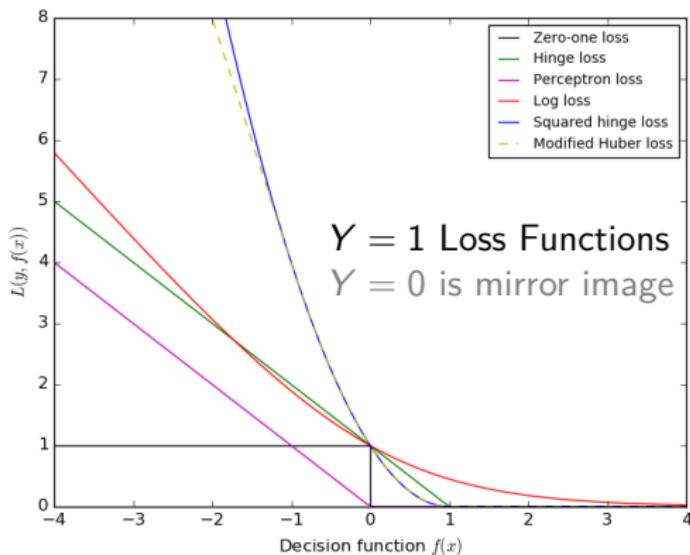
$$L_\delta(Y_i - \hat{Y}_i)$$

$$L_\delta(a) = \begin{cases} \frac{1}{2}a^2 & : |a| < \delta \\ \delta(|a| - \frac{1}{2}\delta) & : o.w. \end{cases}$$



Classification Loss Functions

$$Y \in \{0, 1\} : \min_p -Y \log p - (1 - Y) \log(1 - p) \quad p = \frac{1}{1 + e^{-x^T \beta}}$$

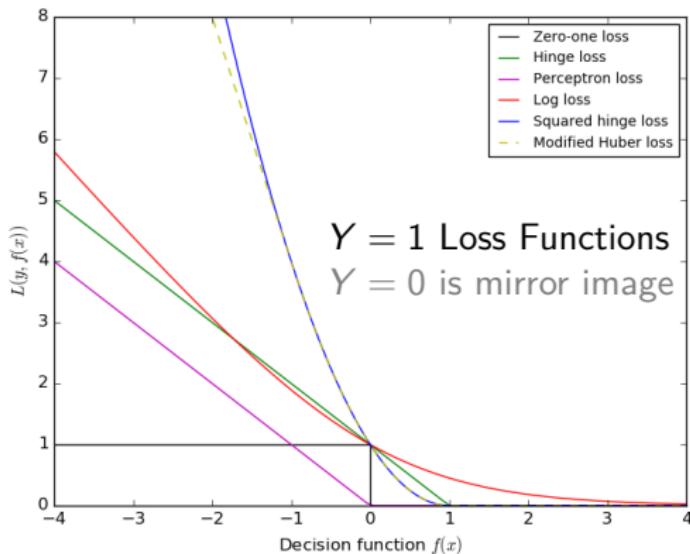


Classification Loss Functions

$$Y \in \{0, 1\} : \min_p -Y \log p - (1 - Y) \log(1 - p) \quad p = \frac{1}{1 + e^{-x^T \beta}}$$

\iff

$$Y \in \{-1, 1\} : \min_p \frac{1}{\ln 2} \ln \left(1 + e^{-Y \ln \left(\frac{p}{1-p} \right)} \right) \quad \ln \left(\frac{p}{1-p} \right) = x^T \beta$$



Cost function

- ▶ $\sum(Y_i - x_i^T \beta)^2$

- ▶ $-\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$

$$g^{-1}(z) = \frac{1}{1 + e^{-z}}$$

Cost function

- ▶ $\sum(Y_i - x_i^T \beta)^2 = (\mathbf{Y} - \mathbf{x}\beta)^T(\mathbf{Y} - \mathbf{x}\beta) = \|\mathbf{Y} - \mathbf{x}\beta\|^2$
- ▶ $-\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$

$$g^{-1}(z) = \frac{1}{1 + e^{-z}}$$

Cost function

- ▶ $\sum(Y_i - x_i^T \beta)^2 = (\mathbf{Y} - \mathbf{x}\beta)^T(\mathbf{Y} - \mathbf{x}\beta) = \|\mathbf{Y} - \mathbf{x}\beta\|^2$

This cost function is minimized at $(\mathbf{x}^T \mathbf{x})^{-1} \mathbf{x}^T \mathbf{Y}$

- ▶ $-\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$

$$g^{-1}(z) = \frac{1}{1 + e^{-z}}$$

Cost function

- ▶ $\sum(Y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2 = (\mathbf{Y} - \mathbf{x}\boldsymbol{\beta})^T(\mathbf{Y} - \mathbf{x}\boldsymbol{\beta}) = \|\mathbf{Y} - \mathbf{x}\boldsymbol{\beta}\|^2$

This cost function is minimized at $(\mathbf{x}^T \mathbf{x})^{-1} \mathbf{x}^T \mathbf{Y}$

- ▶ $-\sum Y_i \log g^{-1}(\mathbf{x}_i^T \boldsymbol{\beta}) + (1 - Y_i) \log (1 - g^{-1}(\mathbf{x}_i^T \boldsymbol{\beta}))$

$$g^{-1}(z) = \frac{1}{1 + e^{-z}}$$

This cost function is minimized at... oh, no analytical solution...

Cost function

- ▶ $\sum(Y_i - \mathbf{x}_i^T \boldsymbol{\beta})^2 = (\mathbf{Y} - \mathbf{x}\boldsymbol{\beta})^T(\mathbf{Y} - \mathbf{x}\boldsymbol{\beta}) = \|\mathbf{Y} - \mathbf{x}\boldsymbol{\beta}\|^2 + \lambda\|\boldsymbol{\beta}\|$
This cost function is minimized at $(\mathbf{x}^T \mathbf{x})^{-1} \mathbf{x}^T \mathbf{Y}$

- ▶ $-\sum Y_i \log g^{-1}(\mathbf{x}_i^T \boldsymbol{\beta}) + (1 - Y_i) \log (1 - g^{-1}(\mathbf{x}_i^T \boldsymbol{\beta})) + \lambda\|\boldsymbol{\beta}\|$

$$g^{-1}(z) = \frac{1}{1 + e^{-z}}$$

This cost function is minimized at... oh, no analytical solution...

Regularized cost functions also don't have closed form solutions

Objective function

- ▶ An objective function is a target of an optimization procedure
- ▶ A cost function *IS* an objective function (which we minimize)

Objective function

- ▶ An objective function is a target of an optimization procedure
- ▶ A cost function *IS* an objective function (which we minimize)
- ▶ But there are obviously *other* objective functions
- ▶ E.g., a likelihood is an objective function (which we maximize)

Objective function

- ▶ An objective function is a target of an optimization procedure
- ▶ A cost function *IS* an objective function (which we minimize)
- ▶ But there are obviously *other* objective functions
- ▶ E.g., a likelihood is an objective function (which we maximize)
- ▶ MLE maximizes the (log) likelihood, e.g.,

$$\underset{\beta}{\operatorname{argmax}} \ (2\pi\sigma^2)^{-\frac{n}{2}} e^{-\frac{1}{2\sigma^2}(\mathbf{Y}-\mathbf{x}\beta)^T(\mathbf{Y}-\mathbf{x}\beta)}$$

or

$$\underset{\beta}{\operatorname{argmax}} \ \prod \left(\frac{1}{1 + e^{-\mathbf{x}_i^T \beta}} \right)^{Y_i} \left(\frac{1}{1 + e^{\mathbf{x}_i^T \beta}} \right)^{1-Y_i}$$

Objective function

- ▶ An objective function is a target of an optimization procedure
- ▶ A cost function *IS* an objective function (which we minimize)
- ▶ But there are obviously *other* objective functions
- ▶ E.g., a likelihood is an objective function (which we maximize)
- ▶ MLE maximizes the (log) likelihood, e.g.,

$$\underset{\beta}{\operatorname{argmax}} \frac{1}{(2\pi\sigma^2)^{\frac{n}{2}}} e^{-\frac{1}{2\sigma^2}(\mathbf{Y}-\mathbf{x}\beta)^T(\mathbf{Y}-\mathbf{x}\beta)}$$

or

$$\underset{\beta}{\operatorname{argmax}} \prod \left(\frac{1}{1 + e^{-\mathbf{x}_i^T \beta}} \right)^{Y_i} \left(\frac{1}{1 + e^{\mathbf{x}_i^T \beta}} \right)^{1-Y_i}$$

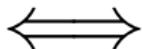
- ▶ But in Machine Learning
the standard orientation and nomenclature is “minimize cost”

Exercise of the day

Find parameters that maximizes
the fit of model to data

Exercise of the day

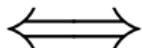
Find parameters that maximizes
the fit of model to data



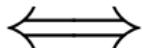
minimize distance between
predicted and observed values

Exercise of the day

Find parameters that maximizes
the fit of model to data



minimize distance between
predicted and observed values



$$\hat{Y}_i \approx Y_i$$

Gradients

For some function $f(\mathbf{x})$, the gradient

$$\nabla_{\mathbf{x}} f \left(\mathbf{x}^{(0)} \right) = \left(\frac{\partial f}{\partial x_1} \left(\mathbf{x}^{(0)} \right), \frac{\partial f}{\partial x_2} \left(\mathbf{x}^{(0)} \right), \dots, \frac{\partial f}{\partial x_p} \left(\mathbf{x}^{(0)} \right) \right)$$

collects the instantaneous slopes (derivatives) with respect to each variable x_j of \mathbf{x} and then evaluates each one at point $\mathbf{x}^{(0)}$

Gradients

For some function $f(\mathbf{x})$, the gradient

$$\nabla_{\mathbf{x}} f(\mathbf{x}^{(0)}) = \left(\frac{\partial f}{\partial x_1}(\mathbf{x}^{(0)}), \frac{\partial f}{\partial x_2}(\mathbf{x}^{(0)}), \dots, \frac{\partial f}{\partial x_p}(\mathbf{x}^{(0)}) \right)$$

collects the instantaneous slopes (derivatives) with respect to each variable x_j of \mathbf{x} and then evaluates each one at point $\mathbf{x}^{(0)}$

⇒ $\nabla f(\mathbf{x}^{(0)})$ points in the direction of greatest increase in f at $\mathbf{x}^{(0)}$

Gradients

For some function $f(\mathbf{x})$, the gradient

$$\nabla_{\mathbf{x}} f(\mathbf{x}^{(0)}) = \left(\frac{\partial f}{\partial x_1}(\mathbf{x}^{(0)}), \frac{\partial f}{\partial x_2}(\mathbf{x}^{(0)}), \dots, \frac{\partial f}{\partial x_p}(\mathbf{x}^{(0)}) \right)$$

collects the instantaneous slopes (derivatives) with respect to each variable x_j of \mathbf{x} and then evaluates each one at point $\mathbf{x}^{(0)}$

- ⇒ $\nabla f(\mathbf{x}^{(0)})$ points in the direction of greatest increase in f at $\mathbf{x}^{(0)}$
- ⇒ The magnitude $|\nabla f(\mathbf{x}^{(0)})|$ is proportional to f 's steepness at $\mathbf{x}^{(0)}$

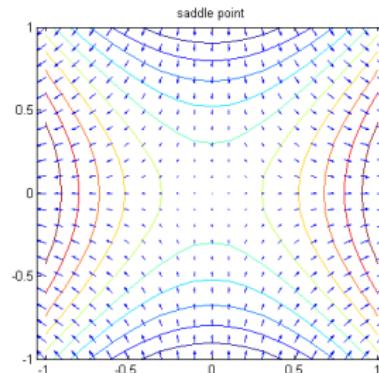
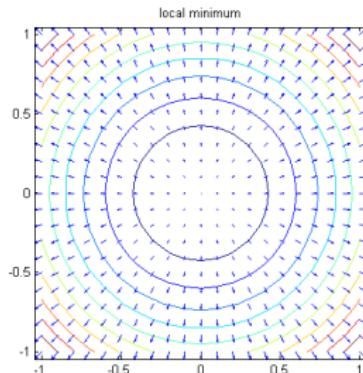
Gradients

For some function $f(\mathbf{x})$, the gradient

$$\nabla_{\mathbf{x}} f(\mathbf{x}^{(0)}) = \left(\frac{\partial f}{\partial x_1}(\mathbf{x}^{(0)}), \frac{\partial f}{\partial x_2}(\mathbf{x}^{(0)}), \dots, \frac{\partial f}{\partial x_p}(\mathbf{x}^{(0)}) \right)$$

collects the instantaneous slopes (derivatives) with respect to each variable x_j of \mathbf{x} and then evaluates each one at point $\mathbf{x}^{(0)}$

- ⇒ $\nabla f(\mathbf{x}^{(0)})$ points in the direction of greatest increase in f at $\mathbf{x}^{(0)}$
- ⇒ The magnitude $|\nabla f(\mathbf{x}^{(0)})|$ is proportional to f 's steepness at $\mathbf{x}^{(0)}$



Gradient Descent

<http://vis.supstat.com/2013/03/gradient-descent-algorithm-with-r/>

0. Choose step size α and precision threshold ϵ
1. Select starting point $\mathbf{x}^{(0)}$, set $i = 1$
2. Update $\mathbf{x}^{(t)} = \mathbf{x}^{(t-1)} - \alpha \nabla f(\mathbf{x}^{(t-1)})$
3. If $\frac{|f(\mathbf{x}^{(t-1)})| - |f(\mathbf{x}^{(t)})|}{|f(\mathbf{x}^{(t-1)})|} < \epsilon$, return $\min |f(\mathbf{x}^{(t)})|$ & $\operatorname{argmin} \mathbf{x}^{(t)}$
4. else, return to step 2.

Gradient Descent

<http://vis.supstat.com/2013/03/gradient-descent-algorithm-with-r/>

0. Choose step size α and precision threshold ϵ
1. Select starting point $\mathbf{x}^{(0)}$, set $i = 1$
2. Update $\mathbf{x}^{(t)} = \mathbf{x}^{(t-1)} - \alpha \nabla f(\mathbf{x}^{(t-1)})$
3. If $\frac{|f(\mathbf{x}^{(t-1)})| - |f(\mathbf{x}^{(t)})|}{|f(\mathbf{x}^{(t-1)})|} < \epsilon$, return $\min |f(\mathbf{x}^{(t)})|$ & $\operatorname{argmin} \mathbf{x}^{(t)}$
4. else, return to step 2. [how do we choose α and ϵ ?]

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$
- ▶ E.g., if $f(x) = x^2$, then $\nabla f(x) = f'(x) = 2x$

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$
- ▶ E.g., if $f(x) = x^2$, then $\nabla f(x) = f'(x) = 2x$
- ▶ So $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} = \frac{2(\mathbf{a} - \mathbf{a}')}{(\mathbf{a} - \mathbf{a}')} = 2$, and $\alpha = 1/2$ is optimal

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$
- ▶ E.g., if $f(x) = x^2$, then $\nabla f(x) = f'(x) = 2x$
- ▶ So $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} = \frac{2(\mathbf{a} - \mathbf{a}')}{(\mathbf{a} - \mathbf{a}')} = 2$, and $\alpha = 1/2$ is optimal

And an (adaptive) way to choose α :

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$
- ▶ E.g., if $f(x) = x^2$, then $\nabla f(x) = f'(x) = 2x$
- ▶ So $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} = \frac{2(\mathbf{a} - \mathbf{a}')}{(\mathbf{a} - \mathbf{a}')} = 2$, and $\alpha = 1/2$ is optimal

And an (adaptive) way to choose α :

- ▶ Let $\Delta \mathbf{x}^{(t)} = \mathbf{x}^{(t)} - \mathbf{x}^{(t-1)}$

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$
- ▶ E.g., if $f(x) = x^2$, then $\nabla f(x) = f'(x) = 2x$
- ▶ So $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} = \frac{2(\mathbf{a} - \mathbf{a}')}{(\mathbf{a} - \mathbf{a}')} = 2$, and $\alpha = 1/2$ is optimal

And an (adaptive) way to choose α :

- ▶ Let $\Delta \mathbf{x}^{(t)} = \mathbf{x}^{(t)} - \mathbf{x}^{(t-1)}$
and let $\Delta \nabla f(\mathbf{x}^{(t)}) = \nabla f(\mathbf{x}^{(t)}) - \nabla f(\mathbf{x}^{(t-1)})$

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$
- ▶ E.g., if $f(x) = x^2$, then $\nabla f(x) = f'(x) = 2x$
- ▶ So $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} = \frac{2(\mathbf{a} - \mathbf{a}')}{(\mathbf{a} - \mathbf{a}')} = 2$, and $\alpha = 1/2$ is optimal

And an (adaptive) way to choose α (relying on "smoothness" of f):

- ▶ Let $\Delta \mathbf{x}^{(t)} = \mathbf{x}^{(t)} - \mathbf{x}^{(t-1)}$
and let $\Delta \nabla f(\mathbf{x}^{(t)}) = \nabla f(\mathbf{x}^{(t)}) - \nabla f(\mathbf{x}^{(t-1)})$
- ▶ For step $t + 1$, let $\alpha = \frac{\Delta \nabla f(\mathbf{x}^{(t)})^T \Delta \mathbf{x}^{(t)}}{\|\Delta \nabla f(\mathbf{x}^{(t)})\|^2} \rightarrow$
 - large if the change in gradient is the same direction as just moved
 - large if the gradient is changing

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$
- ▶ E.g., if $f(x) = x^2$, then $\nabla f(x) = f'(x) = 2x$
- ▶ So $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} = \frac{2(\mathbf{a} - \mathbf{a}')}{(\mathbf{a} - \mathbf{a}')} = 2$, and $\alpha = 1/2$ is optimal

And an (adaptive) way to choose α (relying on "smoothness" of f):

- ▶ Let $\Delta \mathbf{x}^{(t)} = \mathbf{x}^{(t)} - \mathbf{x}^{(t-1)}$
and let $\Delta \nabla f(\mathbf{x}^{(t)}) = \nabla f(\mathbf{x}^{(t)}) - \nabla f(\mathbf{x}^{(t-1)})$
- ▶ For step $t + 1$, let $\alpha = \frac{\Delta \nabla f(\mathbf{x}^{(t)})^T \Delta \mathbf{x}^{(t)}}{\|\Delta \nabla f(\mathbf{x}^{(t)})\|^2} \rightarrow$
 - large if the change in gradient is the same direction as just moved
 - large if the gradient is changing

And here's some potential stopping criterion:

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$
- ▶ E.g., if $f(x) = x^2$, then $\nabla f(x) = f'(x) = 2x$
- ▶ So $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} = \frac{2(\mathbf{a} - \mathbf{a}')}{(\mathbf{a} - \mathbf{a}')} = 2$, and $\alpha = 1/2$ is optimal

And an (adaptive) way to choose α (relying on "smoothness" of f):

- ▶ Let $\Delta \mathbf{x}^{(t)} = \mathbf{x}^{(t)} - \mathbf{x}^{(t-1)}$
and let $\Delta \nabla f(\mathbf{x}^{(t)}) = \nabla f(\mathbf{x}^{(t)}) - \nabla f(\mathbf{x}^{(t-1)})$
- ▶ For step $t + 1$, let $\alpha = \frac{\Delta \nabla f(\mathbf{x}^{(t)})^T \Delta \mathbf{x}^{(t)}}{\|\Delta \nabla f(\mathbf{x}^{(t)})\|^2} \rightarrow$
 - large if the change in gradient is the same direction as just moved
 - large if the gradient is changing

And here's some potential stopping criterion:

a. $\frac{|f(\mathbf{x}^{(t-1)})| - |f(\mathbf{x}^{(t)})|}{|f(\mathbf{x}^{(t-1)})|} < \epsilon$

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$
- ▶ E.g., if $f(x) = x^2$, then $\nabla f(x) = f'(x) = 2x$
- ▶ So $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} = \frac{2(\mathbf{a} - \mathbf{a}')}{(\mathbf{a} - \mathbf{a}')} = 2$, and $\alpha = 1/2$ is optimal

And an (adaptive) way to choose α (relying on "smoothness" of f):

- ▶ Let $\Delta \mathbf{x}^{(t)} = \mathbf{x}^{(t)} - \mathbf{x}^{(t-1)}$
and let $\Delta \nabla f(\mathbf{x}^{(t)}) = \nabla f(\mathbf{x}^{(t)}) - \nabla f(\mathbf{x}^{(t-1)})$
- ▶ For step $t + 1$, let $\alpha = \frac{\Delta \nabla f(\mathbf{x}^{(t)})^T \Delta \mathbf{x}^{(t)}}{\|\Delta \nabla f(\mathbf{x}^{(t)})\|^2} \rightarrow$
 - large if the change in gradient is the same direction as just moved
 - large if the gradient is changing

And here's some potential stopping criterion:

- $\frac{|f(\mathbf{x}^{(t-1)})| - |f(\mathbf{x}^{(t)})|}{|f(\mathbf{x}^{(t-1)})|} < \epsilon$
- Max number of iterations

Step Size α and Stopping Criterion ϵ

Here's one way to choose α :

- ▶ If $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} < c \in \mathbb{R}$, choose $\alpha \leq 1/c$
- ▶ E.g., if $f(x) = x^2$, then $\nabla f(x) = f'(x) = 2x$
- ▶ So $\frac{|\nabla f(\mathbf{a}) - \nabla f(\mathbf{a}')|}{|\mathbf{a} - \mathbf{a}'|} = \frac{2(\mathbf{a} - \mathbf{a}')}{(\mathbf{a} - \mathbf{a}')} = 2$, and $\alpha = 1/2$ is optimal

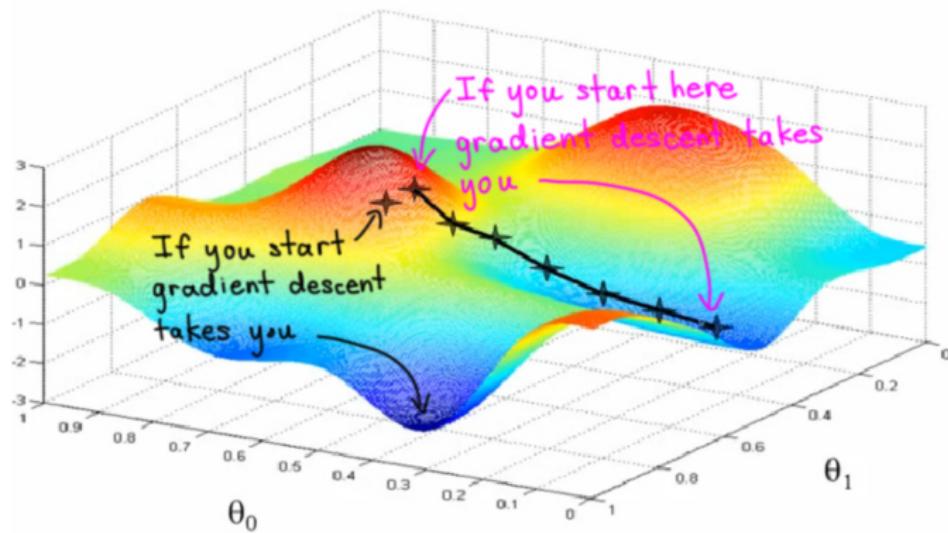
And an (adaptive) way to choose α (relying on "smoothness" of f):

- ▶ Let $\Delta \mathbf{x}^{(t)} = \mathbf{x}^{(t)} - \mathbf{x}^{(t-1)}$
and let $\Delta \nabla f(\mathbf{x}^{(t)}) = \nabla f(\mathbf{x}^{(t)}) - \nabla f(\mathbf{x}^{(t-1)})$
- ▶ For step $t + 1$, let $\alpha = \frac{\Delta \nabla f(\mathbf{x}^{(t)})^T \Delta \mathbf{x}^{(t)}}{\|\Delta \nabla f(\mathbf{x}^{(t)})\|^2} \rightarrow$
 - large if the change in gradient is the same direction as just moved
 - large if the gradient is changing

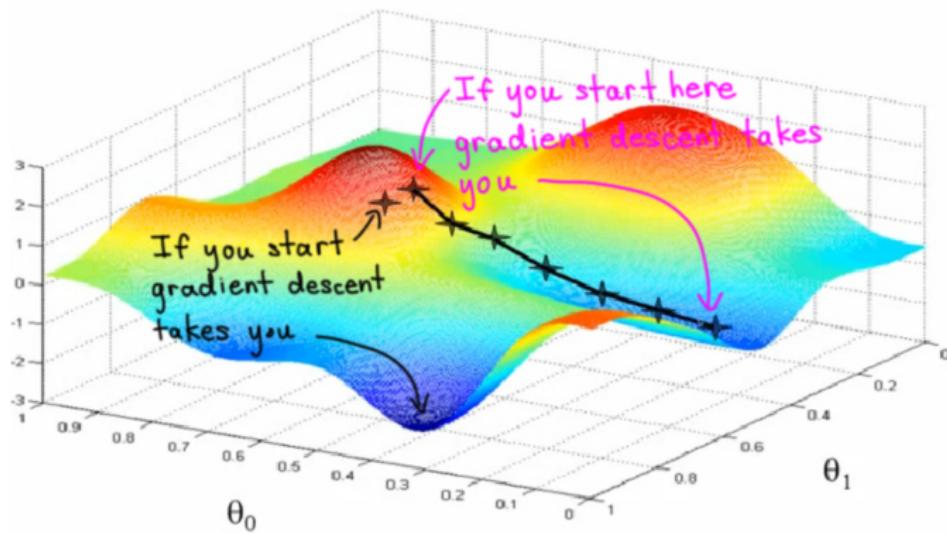
And here's some potential stopping criterion:

- $\frac{|f(\mathbf{x}^{(t-1)})| - |f(\mathbf{x}^{(t)})|}{|f(\mathbf{x}^{(t-1)})|} < \epsilon$
- Max number of iterations
- $|\nabla f| < \epsilon$

Pitfalls (literally)

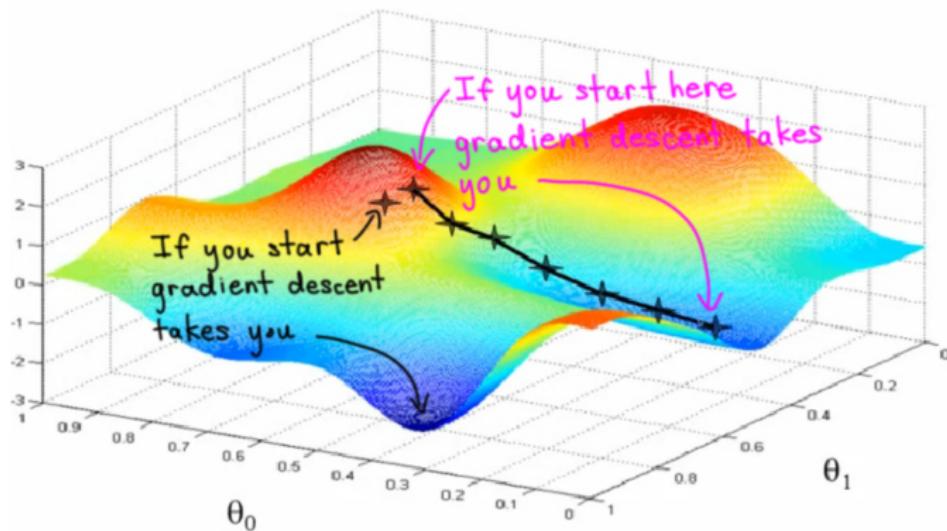


Pitfalls (literally)



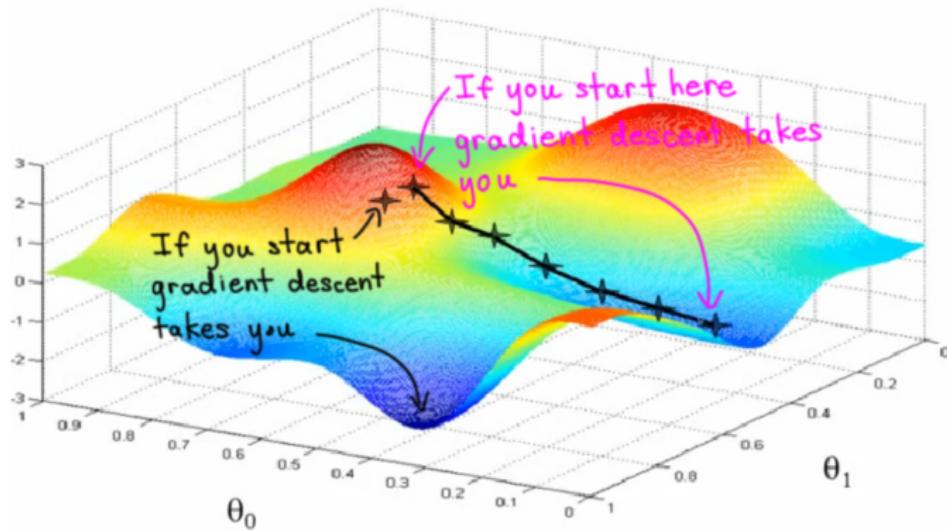
- ▶ Convexity required to guarantee *global minimum*

Pitfalls (literally)



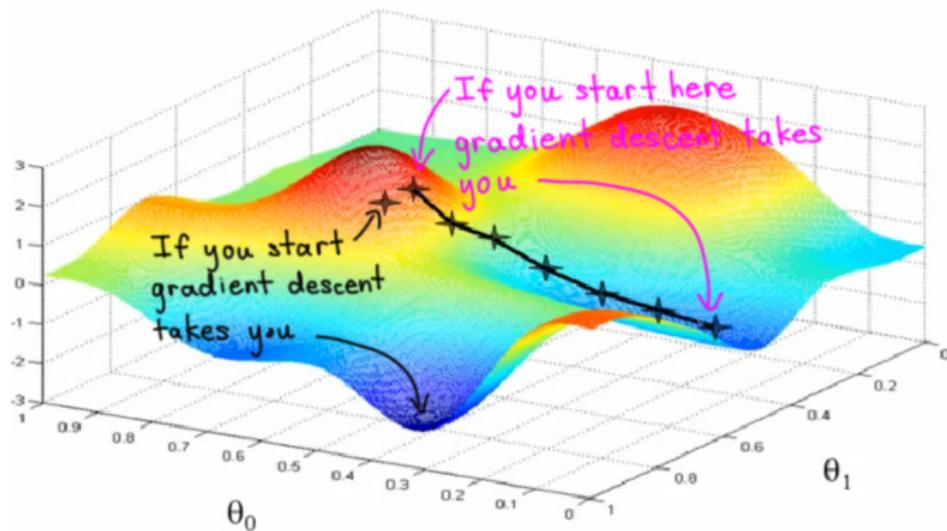
- ▶ Convexity required to guarantee *global minimum*
- ▶ Differentiable cost function required for *gradient descent*

Pitfalls (literally)



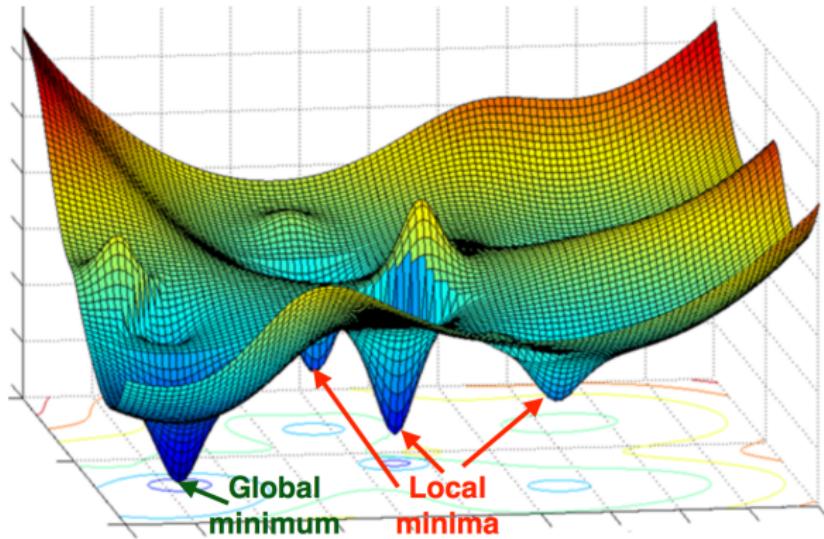
- ▶ Convexity required to guarantee *global minimum*
- ▶ Differentiable cost function required for *gradient descent*
- ▶ Convergence is asymptotic...

Pitfalls (literally)



- ▶ Convexity required to guarantee *global minimum*
- ▶ Differentiable cost function required for *gradient descent*
- ▶ Convergence is asymptotic...
- ▶ Good performance requires feature scaling & parameter tuning

Pitfalls (literally)



- ▶ Convexity required to guarantee *global minimum*
- ▶ Differentiable cost function required for *gradient descent*
- ▶ Convergence is asymptotic...
- ▶ Good performance requires feature scaling & parameter tuning

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

First $\frac{d}{dz} g^{-1}(z) = \frac{d}{dz} \frac{1}{1 + e^{-z}}$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

First

$$\begin{aligned}\frac{d}{dz} g^{-1}(z) &= \frac{d}{dz} \frac{1}{1 + e^{-z}} \\ &= \frac{d}{dz} (1 + e^{-z})^{-1}\end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

First

$$\begin{aligned}\frac{d}{dz} g^{-1}(z) &= \frac{d}{dz} \frac{1}{1 + e^{-z}} \\ &= \frac{d}{dz} (1 + e^{-z})^{-1} \\ &= -(1 + e^{-z})^{-2}(-e^{-z})\end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

First

$$\begin{aligned}\frac{d}{dz} g^{-1}(z) &= \frac{d}{dz} \frac{1}{1 + e^{-z}} \\ &= \frac{d}{dz} (1 + e^{-z})^{-1} \\ &= -(1 + e^{-z})^{-2}(-e^{-z}) \\ &= \frac{1}{1 + e^{-z}} \frac{e^{-z}}{1 + e^{-z}}\end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

First

$$\begin{aligned}\frac{d}{dz} g^{-1}(z) &= \frac{d}{dz} \frac{1}{1 + e^{-z}} \\&= \frac{d}{dz} (1 + e^{-z})^{-1} \\&= -(1 + e^{-z})^{-2}(-e^{-z}) \\&= \frac{1}{1 + e^{-z}} \frac{e^{-z}}{1 + e^{-z}} \\&= g^{-1}(z) (1 - g^{-1}(z))\end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} L(\beta) \\ &= \frac{\partial}{\partial \beta_j} \left(\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) \right) \end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} L(\beta) \\ &= \frac{\partial}{\partial \beta_j} \left(\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) \right) \\ &= \sum \frac{\partial}{\partial \beta_j} g^{-1}(x_i^T \beta) \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} L(\beta) \\ &= \frac{\partial}{\partial \beta_j} \left(\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) \right) \\ &= \sum \frac{\partial}{\partial \beta_j} g^{-1}(x_i^T \beta) \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum g^{-1}(x_i^T \beta) \left(1 - g^{-1}(x_i^T \beta) \right) \textcolor{red}{x_{ij}} \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} L(\beta) \\ &= \frac{\partial}{\partial \beta_j} \left(\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) \right) \\ &= \sum \frac{\partial}{\partial \beta_j} g^{-1}(x_i^T \beta) \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum g^{-1}(x_i^T \beta) \left(1 - g^{-1}(x_i^T \beta) \right) \cancel{x_{ij}} \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum x_{ij} \left(Y_i \left(1 - g^{-1}(x_i^T \beta) \right) - (1 - Y_i) g^{-1}(x_i^T \beta) \right) \end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = ?$$

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} L(\beta) \\ &= \frac{\partial}{\partial \beta_j} \left(\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) \right) \\ &= \sum \frac{\partial}{\partial \beta_j} g^{-1}(x_i^T \beta) \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum g^{-1}(x_i^T \beta) \left(1 - g^{-1}(x_i^T \beta) \right) \textcolor{red}{x_{ij}} \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum x_{ij} \left(Y_i \left(1 - g^{-1}(x_i^T \beta) \right) - (1 - Y_i) g^{-1}(x_i^T \beta) \right) \\ &= \sum x_{ij} \left(Y_i - g^{-1}(x_i^T \beta) \right) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right) = \sum_{i=1}^n x_{ij} (Y_i - \hat{Y}_i) \end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) - \lambda \frac{1}{2} \|\beta\|^2$$

$$\nabla L(\beta) = ?$$

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} L(\beta) \\ &= \frac{\partial}{\partial \beta_j} \left(\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) \right) \\ &= \sum \frac{\partial}{\partial \beta_j} g^{-1}(x_i^T \beta) \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum g^{-1}(x_i^T \beta) \left(1 - g^{-1}(x_i^T \beta) \right) \textcolor{red}{x_{ij}} \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum x_{ij} \left(Y_i \left(1 - g^{-1}(x_i^T \beta) \right) - (1 - Y_i) g^{-1}(x_i^T \beta) \right) \\ &= \sum x_{ij} \left(Y_i - g^{-1}(x_i^T \beta) \right) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right) = \sum_{i=1}^n x_{ij} (Y_i - \hat{Y}_i) \end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) - \frac{\lambda}{2} \|\beta\|^2$$

$$\nabla L(\beta) = ?$$

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} L(\beta) && -\lambda \beta \\ &= \frac{\partial}{\partial \beta_j} \left(\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) \right) \\ &= \sum \frac{\partial}{\partial \beta_j} g^{-1}(x_i^T \beta) \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum g^{-1}(x_i^T \beta) \left(1 - g^{-1}(x_i^T \beta) \right) \cancel{x_{ij}} \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum x_{ij} \left(Y_i \left(1 - g^{-1}(x_i^T \beta) \right) - (1 - Y_i) g^{-1}(x_i^T \beta) \right) \\ &= \sum x_{ij} \left(Y_i - g^{-1}(x_i^T \beta) \right) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right) = \sum_{i=1}^n x_{ij} (Y_i - \hat{Y}_i) \end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) - \lambda |\beta|$$

$$\nabla L(\beta) = ?$$

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} L(\beta) \\ &= \frac{\partial}{\partial \beta_j} \left(\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) \right) \\ &= \sum \frac{\partial}{\partial \beta_j} g^{-1}(x_i^T \beta) \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum g^{-1}(x_i^T \beta) \left(1 - g^{-1}(x_i^T \beta) \right) \textcolor{red}{x_{ij}} \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum x_{ij} \left(Y_i \left(1 - g^{-1}(x_i^T \beta) \right) - (1 - Y_i) g^{-1}(x_i^T \beta) \right) \\ &= \sum x_{ij} \left(Y_i - g^{-1}(x_i^T \beta) \right) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right) = \sum_{i=1}^n x_{ij} (Y_i - \hat{Y}_i) \end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) - \lambda |\beta|$$

$$\nabla L(\beta) = ?$$

$$\begin{aligned} & \frac{\partial}{\partial \beta_j} L(\beta) \quad \pm \lambda \\ &= \frac{\partial}{\partial \beta_j} \left(\sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta)) \right) \\ &= \sum \frac{\partial}{\partial \beta_j} g^{-1}(x_i^T \beta) \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum g^{-1}(x_i^T \beta) \left(1 - g^{-1}(x_i^T \beta) \right) \textcolor{red}{x_{ij}} \left(\frac{Y_i}{g^{-1}(x_i^T \beta)} - \frac{(1 - Y_i)}{(1 - g^{-1}(x_i^T \beta))} \right) \\ &= \sum x_{ij} \left(Y_i \left(1 - g^{-1}(x_i^T \beta) \right) - (1 - Y_i) g^{-1}(x_i^T \beta) \right) \\ &= \sum x_{ij} \left(Y_i - g^{-1}(x_i^T \beta) \right) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right) = \sum_{i=1}^n x_{ij} (Y_i - \hat{Y}_i) \end{aligned}$$

Logistic regression

$$L(\beta) = \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla L(\beta) = \begin{bmatrix} \sum x_{i1} \left(Y_i - \frac{1}{1+e^{-x_i^T \beta}} \right) \\ \sum x_{i2} \left(Y_i - \frac{1}{1+e^{-x_i^T \beta}} \right) \\ \vdots \\ \sum x_{ip} \left(Y_i - \frac{1}{1+e^{-x_i^T \beta}} \right) \end{bmatrix}$$

$$\beta^{(t)} = \beta^{(k-1)} + \alpha \nabla L(\beta^{(k-1)})$$

Logistic regression

$$C(\beta) = - \sum Y_i \log g^{-1}(x_i^T \beta) + (1 - Y_i) \log (1 - g^{-1}(x_i^T \beta))$$

$$\nabla C(\beta) = - \begin{bmatrix} \sum x_{i1} \left(Y_i - \frac{1}{1+e^{-x_i^T \beta}} \right) \\ \sum x_{i2} \left(Y_i - \frac{1}{1+e^{-x_i^T \beta}} \right) \\ \vdots \\ \sum x_{ip} \left(Y_i - \frac{1}{1+e^{-x_i^T \beta}} \right) \end{bmatrix}$$

$$\beta^{(t)} = \beta^{(k-1)} - \alpha \nabla C(\beta^{(k-1)})$$

Potential Gradient Descent Drawbacks

1. Memory (data needs to fit)
2. Processor (cost function over all rows is expensive)

Potential Gradient Descent Drawbacks *Solutions*

- ▶ Observations contribute equal weight to the gradient

$$\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\boldsymbol{\beta}) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \boldsymbol{\beta}}} \right)$$

Potential Gradient Descent Drawbacks *Solutions*

- ▶ Observations contribute equal weight to the gradient

$$\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\boldsymbol{\beta}) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \boldsymbol{\beta}}} \right)$$

- ▶ We could just use only a single data point at each iteration?

$$\frac{\partial L_i}{\partial \beta_j}(\boldsymbol{\beta}) = x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \boldsymbol{\beta}}} \right)$$

Potential Gradient Descent Drawbacks *Solutions*

- ▶ Observations contribute equal weight to the gradient

$$\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\beta) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right)$$

- ▶ We could just use only a single data point at each iteration?

$$\frac{\partial L_i}{\partial \beta_j}(\beta) = x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right)$$

- ▶ The expected direction of the gradient would stay the same

$$\frac{1}{n} E \left[\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\beta) \right] = E \left[\frac{\partial L_i}{\partial \beta_j}(\beta) \right]$$

Potential Gradient Descent Drawbacks *Solutions*

- ▶ Observations contribute equal weight to the gradient

$$\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\beta) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right)$$

- ▶ We could just use only a single data point at each iteration?

$$\frac{\partial L_i}{\partial \beta_j}(\beta) = x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right)$$

- ▶ The expected direction of the gradient would stay the same

$$\frac{1}{n} \mathbb{E} \left[\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\beta) \right] = \mathbb{E} \left[\frac{\partial L_i}{\partial \beta_j}(\beta) \right]$$

- ▶ We could also use a batch of data points at each iteration

Potential Gradient Descent Drawbacks *Solutions*

- ▶ Observations contribute equal weight to the gradient

$$\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\beta) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right)$$

- ▶ We could just use only a single data point at each iteration?

$$\frac{\partial L_i}{\partial \beta_j}(\beta) = x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right)$$

- ▶ The expected direction of the gradient would stay the same

$$\frac{1}{n} E \left[\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\beta) \right] = E \left[\frac{\partial L_i}{\partial \beta_j}(\beta) \right]$$

- ▶ We could also use a batch of data points at each iteration

This would address memory/processing limitations

Potential Gradient Descent Drawbacks *Solutions*

- ▶ Observations contribute equal weight to the gradient

$$\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\beta) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right)$$

- ▶ We could just use only a single data point at each iteration?

$$\frac{\partial L_i}{\partial \beta_j}(\beta) = x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right)$$

- ▶ The expected direction of the gradient would stay the same

$$\frac{1}{n} E \left[\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\beta) \right] = E \left[\frac{\partial L_i}{\partial \beta_j}(\beta) \right]$$

- ▶ We could also use a batch of data points at each iteration

This would address memory/processing limitations

It has been empirically proven to also often converge faster!

Potential Gradient Descent Drawbacks *Solutions*

- ▶ Observations contribute equal weight to the gradient

$$\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\beta) = \sum_{i=1}^n x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right)$$

- ▶ We could just use only a single data point at each iteration?

$$\frac{\partial L_i}{\partial \beta_j}(\beta) = x_{ij} \left(Y_i - \frac{1}{1 + e^{-x_i^T \beta}} \right)$$

- ▶ The expected direction of the gradient would stay the same

$$\frac{1}{n} E \left[\frac{\partial L_{1,\dots,n}}{\partial \beta_j}(\beta) \right] = E \left[\frac{\partial L_i}{\partial \beta_j}(\beta) \right]$$

- ▶ We could also use a batch of data points at each iteration

This would address memory/processing limitations

It has been empirically proven to also often converge faster!

It does tend to oscillate around as it nears the minimum...