

# Parallel Programming for Data Scientists

Benjamin S. Skrainka

July 11, 2016

# Objectives

Today's objectives:

- Describe basic components of a computer
- Describe basic components of an operating system (OS)
- List components of a process
- State difference between a task & a process
- List issues involved in parallelizing computation

# References

A couple references:

- The Design & Implementation of the FreeBSD Operating System
- Python documentation

# Basic computer architecture

Draw picture

# Operating system basics

An operating system is a program which manages a computer's resources:

- Synchronization to coordinate work
- Mutual exclusion to protect shared resources
- Scheduling of work
  - ▶ Usually 'fair'
  - ▶ Can change priority with `nice`

Work runs inside processes:

- Runs program *text*
  - ▶ Only one copy is in memory on the entire computer
  - ▶ Read-only
  - ▶ From OS, programs, and libraries
- Contains:
  - ▶ *data*: initialized global & static variables
  - ▶ *bss*: uninitialized global & static variables
  - ▶ One copy per process

# States of a process

Processes can be in one of the following states:

- Ready/executing
- Blocked
- Delayed
- Suspended

# Lifecycle of a process

The process life cycle is:

- Fork
- Exec
- Exit
- Reaped by parent



# Components of a process

A process has the following main components

- OS control information:
  - ▶ PID & PPID
  - ▶ File descriptor table
  - ▶ Mapping of standard input, output, and error
  - ▶ Error status
  - ▶ Signal handlers
- Thread of execution
- Stack: local storage for thread
- Heap: general memory for process to allocate dynamically

A process should be relatively insulated from other processes

# Process vs. thread

A thread is a lighter-weight concept:

- One or more run inside of a process
- Each thread has its own stack
- All threads uses the same heap/global memory
  - ▶ Easy to communicate
  - ▶ Easy to cause *race conditions* if threads do not coordinate access to shared memory

Use parallelization to speed up big jobs when:

- *Embarassingly parallel*: can break work up into independent chunks
- Operations can block or fail
- Application decomposes into different types of work or stages
- Have more data than fits in a single computer

# Tools for parallelization

OS and computer language provide support for parallel programming:

- *OpenMP* works within one node (multi-core) via shared memory
- *MPI* works between nodes, e.g., over a network

Python provides:

- Processes: use `multiprocessing` module
- Threads: use `threading` module

# Python Global Interpreter Lock

Python has *Global Interpreter Lock* (GIL):

- CPython only lets one thread in a process at a time run
- To avoid compromising shared/global data structures
- Makes parallelization difficult
- To get parallelization, must run multiple Python jobs as separate processes

# When to use a process

Use a process for longer running jobs:

- Length of the job must offset the cost of launching process
- Circumvent GIL
- Need extra fault tolerance
- Common on clusters using Condor, PBS, etc.
- Robust to errors

# When to use a thread

Use a thread for parallelization when:

- Quick creation/destruction vs. processes
- Easy communication via shared memory

Beware of trade-offs:

- Processes vs. threads: robustness vs. speed
- Cost of launching a process/thread vs. length of work

Other issues:

- Fork, then load data; not vice-versa
- Parallel programming is hard to debug



# Conclusion

- List components of an OS
- List components of a process
- Define states of a process
- When to use process vs. threads