

Using AWS

Jack Bennetto

March 30, 2016

Objectives

Morning objectives:

- Understand motivation for cloud computing and AWS
- Describe core AWS service and concepts
- Configure your laptop to use AWS CLI
- Use SSH key to access EC2 instances
- Access S3

Agenda

Morning:

- Introduction to AWS
- Install and configure AWS CLI
- Configure and use SSH
- Access EC2
- Access S2
- Advanced topics

Afternoon:

- High-performance python

References

Google it...

Amazon Web Services (AWS)

AWS provides on-demand use of computing resources in the cloud

- No need to build data centers
- Easy to create a new business
- Only pay for what you use
- Handle spikes in computational demand
- Secure, reliable, flexible, scalable, cost-effective

AWS skills are much in demand

AWS core services:

- Elastic compute cloud (EC2): computers for diverse problems
- Elastic block store (EBS): virtual hard disks to use with EC2
- Simple storage solution (S3): long-term bulk storage
- DynamoDB: NoSQL database
- And much, much more

Elastic compute cloud (EC2)

Spin up EC2 instances for on-demand computing power:

- Instance: a type of hardware you can rent, e.g., 'm3.xlarge' or 't2.micro'
- Amazon Machine Image (AMI), an OS you can run on an instance
- Region: a geographic region, e.g., Oregon a.k.a. 'us-west-2'
- Availability Zone (AZ): a specific subset of a region, often a data center, such as 'us-west-2a'

Elastic block store (EBS)

EBS provides disk-drive style storage:

- Create a virtual hard disk
- Then mount virtual hard disk on EC2 instances
- SSD or magnetic
- Can store data even when you aren't running any EC2 instances
- Built-in redundancy
- Lower latency than S3 but more expensive

Simple storage solution (S3)

S3 provides cheap, bulk storage:

- Create a bucket which serves as a container for files and directories
- Specify permissions using an access control list (ACL)
- Access via URL or AWS CLI or suitable API
- Higher latency than EBS but less expensive

Use the AWS command-line interface (CLI):

- To debug your configuration
- To manage AWS instances in EC2
- To access S3

Install AWS CLI

Brew is the best way to install the AWS CLI:

```
$ brew install awscli
```

Otherwise, see Amazon's [doc](#)

Install AWS CLI (OS/X or Linux)

On Linux, run:

```
$ curl "https://s3.amazonaws.com/aws-cli/awscli-bundle.zip" \
  -o "awscli-bundle.zip"
$ unzip awscli-bundle.zip
$ sudo ./awscli-bundle/install -i /usr/local/aws \
  -b /usr/local/bin/aws
```

Use this method to update AWS CLI on EC2 instances

Obtain your AWS credentials

Find your **AWS credentials**:

- 1 Login to AWS at <https://console.aws.amazon.com/console>
- 2 Click on your name in the upper right menu bar
- 3 Select **Security Credentials**
- 4 Click ****Users*** and select your name (or add it)
- 5 Click User Actions > Manage Access Keys
- 6 Create Access Key, Show User Security Credentials
- 7 Copy key into...

Make sure you choose Oregon (us-west-2) as your region

AWS credentials

The screenshot shows the AWS IAM console interface. The left sidebar contains navigation links: Dashboard, Details, Groups, Users, Roles, Policies, Identity Providers, Account Settings, Credential Report, and Encryption Keys. The main content area displays a table of users with columns: User Name, Groups, Password, Password Last Used, Access Keys, and Creation Time. A modal window titled 'Manage Access Keys' is open, showing details for the user 'alessandro'. The modal includes a table of access keys and a note about key rotation.

Manage Access Keys

Use access keys to make secure REST or Query protocol requests to any AWS service API.

Access Key ID	Created	Last Used	Last Used Service	Last Used Region	Status
[Redacted]	2015-08-24 17:55 PDT	N/A	N/A	N/A	Active (Make Inactive Delete)
[Redacted]	2015-08-25 09:58 PDT	2015-08-25 11:32 PDT	s3	N/A	Active (Make Inactive Delete)

Note: For your protection, you should never share your secret keys with anyone. In addition, industry best practice recommends frequent key rotation.
▶ [Learn more about Access Keys](#)

Cancel Create Access Key

Figure 1:
Using AWS

Configure AWS CLI (1/3)

Easiest way: run `aws configure`:

- Creates default profile in `~/.aws/config`
- Stores credentials in `~/.aws/credentials`
- Can create multiple profiles:

```
$ aws configure --profile fancy_profile
```

- Can also set credential on CLI or via environment variables

Configure AWS CLI (2/3)

Create AWS configuration info in ~/.aws:

```
$ aws configure
```

```
AWS Access Key ID [None]: AKIAIOSFODNN7EXAMPLE
```

```
AWS Secret Access Key [None]: wJalrXUtnFEMI/K7MDENG/bPxRfiCYEX
```

```
Default region name [None]: us-west-2
```

```
Default output format [None]: json
```

See Amazon's [doc](#) for details

Configure AWS CLI (3/3)

Some tools get AWS credentials via environment variables. Set the following in `~/.bash_profile` or equivalent:

```
export AWS_ACCESS_KEY_ID='your access key'  
export AWS_SECRET_ACCESS_KEY='your secret key'
```

Verify configuration

Check S3:

```
$ aws s3 ls
2015-08-25 10:42:43 dsci
2015-08-25 11:30:33 seattle-dsi
$ aws s3 ls s3://seattle-dsi
    PRE cohort1/
    PRE dsi_interviews/
    PRE skrainka/
```

Check EC2:

```
$ aws ec2 describe-instances --output table
$ aws ec2 describe-instances --output json
```

Help with AWS CLI

See the built-in help for more details:

```
$ aws help
```

```
$ aws ec2 help
```

. . . or Google

Configure and use SSH

To use AWS:

- Setup Key Pairs to access EC2
- Configure SSH on your laptop
- Can use SSH to access any remote machine running an SSH server

SSH

The secure shell protocol allows you to:

- Login to remote machines, such as EC2 using ssh
- Transfer files between remote machines using scp and sftp
- Execute commands on remote machines using ssh

Do not use telnet, rlogin, or FTP, which are older, insecure protocols!

Public-key encryption

Pairs of matched keys that can be used to encrypt or decrypt documents:

- Encrypt with public key \rightarrow decrypt with private key
- Encrypt with private key \rightarrow decrypt with public key

But:

- Trivial to calculate public key from private
- Very, very, very hard to calculate private key from public

Public-key encryption

SSH uses public-key encryption to protect access and encrypt communication:

- Generate a public & private key pair
- Keep private key safe
- Create a key pair for each resource you want to access (AWS, GitHub, etc.) \Rightarrow can revoke individual keys in case of a security breach

See AWS [doc](#) for details

Setup Key Pairs

Create and configure key pair to access EC2:

- Create and import key pair as described in *AWS* documentation
- Set permission on private key to 400:

```
$ chmod 400 ~/.ssh/aws-master.pem
```

```
$ chmod 444 ~/.ssh/aws-master.pub
```

- Can also generate key pair with `ssh-keygen`

Configuring SSH

Modify SSH config to:

- Create alias for long-running instance
- Forward X11 or security information
- Specify which key to use
- And, much much more...
- `man ssh_config` for details

Example ~/.ssh/config

```
Host github.com
  HostName github.com
  User git
  ForwardAgent yes
  IdentityFile /Users/jackbenn/.ssh/git-hub-id_rsa
```

Example ~/.ssh/config

If you master will run for a long time, setup an alias:

```
Host master
  HostName ec2-54-186-136-57.us-west-2.compute.amazonaws.com
  User ubuntu
  ForwardAgent yes
  ForwardX11Trusted yes
  TCPKeepAlive yes
  IdentityFile /Users/jackbenn/.ssh/aws-master.pem
```

Now, `ssh master` will connect to your EC2 instance

Accessing an EC2 instance with ssh

Connect to your machine via ssh:

- 1 Launch an EC2 instance from console
- 2 Use ssh from command line to connect to the instances **public DNS** (Shown in EC2 Dashboard):

```
$ ssh -i ~/.ssh/aws-master.pem \
    ubuntu@ec2-54-186-136-57.us-west-2.compute.amazonaws.co
```

or

```
$ ssh -i ~/.ssh/aws-master.pem \
    ec2user@ec2-54-186-136-57.us-west-2.compute.amazonaws
```

Example: ssh to EC2

```
$ ssh -i ~/.ssh/aws-master.pem ubuntu@ec2-54-186-136-57
The authenticity of host 'ec2-54-186-136-57.us-west-2.compu RS
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'ec2-54-186-136-57.us-west-2.com
Welcome to Ubuntu 14.04.2 LTS (GNU/Linux 3.13.0-48-generic
* Documentation:  https://help.ubuntu.com/
System information as of Tue Aug 25 21:58:16 UTC 2015
System load: 0.0                Memory usage: 5%    Processes
Usage of /:  9.8% of 7.74GB      Swap usage:  0%    Users log
```

Example: ssh to EC2

```
$ ssh -i ~/.ssh/aws-master.pem ubuntu@ec2-54-186-136-57
The authenticity of host 'ec2-54-186-136-57.us-west-2.compu RS
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'ec2-54-186-136-57.us-west-2.com
Welcome to Ubuntu 14.04.2 LTS (GNU/Linux 3.13.0-48-generic
* Documentation:  https://help.ubuntu.com/
System information as of Tue Aug 25 21:58:16 UTC 2015
System load: 0.0                Memory usage: 5%    Processes
Usage of /:  9.8% of 7.74GB      Swap usage:  0%    Users log
```

Transferring files with scp

To copy files between machines, use `scp`:

- Works just like regular copy
- Good for simple operations
- ... if you specify remote user and machine (IP,DNS) correctly
- Reference remote location as *user@host:path*

```
$ scp -i ~/.ssh/aws-master.pem ./toy_data.txt \  
ubuntu@54.186.136.57:/home/ubuntu/data  
toy_data.txt 100% 136 0.1KB/s 00:00
```

Transferring files with sftp

To copy files interactively, use sftp:

- Interactive shell for transferring files
- Use to transfer many files
- Use when you don't know the location of a file

```
$ sftp -i ~/.ssh/aws-master.pem ubuntu@54.186.136.57  
Connected to 54.186.136.57.  
sftp> help
```


Managing a session with `tmux`

Use `tmux` to persist jobs across multiple sessions:

- On logout, all child processes terminate
- Use `tmux` to safely disconnect from a session
- Reconnect on next login
- Install `tmux` via `brew` or Linux package manager
- See `tmux` exercise

Launching an EC2 instance

To launch an EC2 instance, follow this [tutorial](#)
... or the official [tutorial](#)

EC2 pro-tips

A few tips to make EC2 easier to deal with:

- Always create instances with tags so that you can find them easily
- Choose the appropriate hardware type for your problem
- If in doubt, use Ubuntu because it is a friendly flavor of Linux (for now, use `galvanize-dsi-ami`)
- Use `tmux` when you login in case you need to disconnect or your connection dies
- Be paranoid: sometimes Amazon will reboot or reclaim your instances
- Put data you need to persist in EBS or a database
- Never put AWS keys in GitHub because someone will steal them

To master the basics, see this [tutorial](#) or `lecture_notes.ipynb` in the repo

- Can find URL to access a file from S3 console
- Set properties (access) via S3 console
- Make sure names conform to S3 conventions:
- lowercase bucket names of at least four characters
- no leading or terminal ‘.

Boto config

To access S3 via Python, use the boto package

- Should be installed if you followed setup instructions
- Make sure boto is up to date:

```
$ conda update boto
```

- Uses credentials in `~/.aws/credentials` which you setup earlier
- Can also read directly from Pandas if you specify S3 URL

Advanced: accessing ipython notebook

- Use an ssh tunnel to run ipython notebook on a remote instance
- On remote host:

```
$ ipython notebook --no-browser --port=8889
```

- On local machine:

```
$ ssh -N -f -L localhost:8888:localhost:8889 \  
    remote_user@remote_host
```

- galvanize-dsi-ami already has a notebook server running, but you need to change security group to allow access
- Access notebook via browser at URL localhost:8888
- Run a notebook server:
- [Official documentation](#)
- Blog on [ipython notebook server](#)