

# High Performance Programming

# Goals

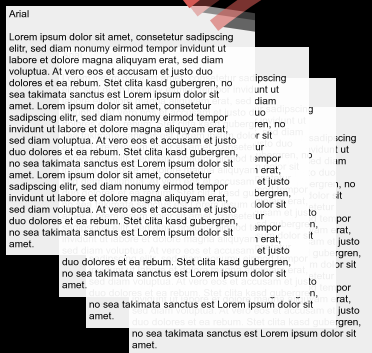
- Motivation
- Intro to computing resources
- Multi-core processing
  - ★ Parallelism
- Threading
  - ★ Concurrency

# Goals

- Motivation
- Intro to computing resources
- Multi-core processing
  - ★ Parallelism
- Threading
  - ★ Concurrency

# Counting words in documents

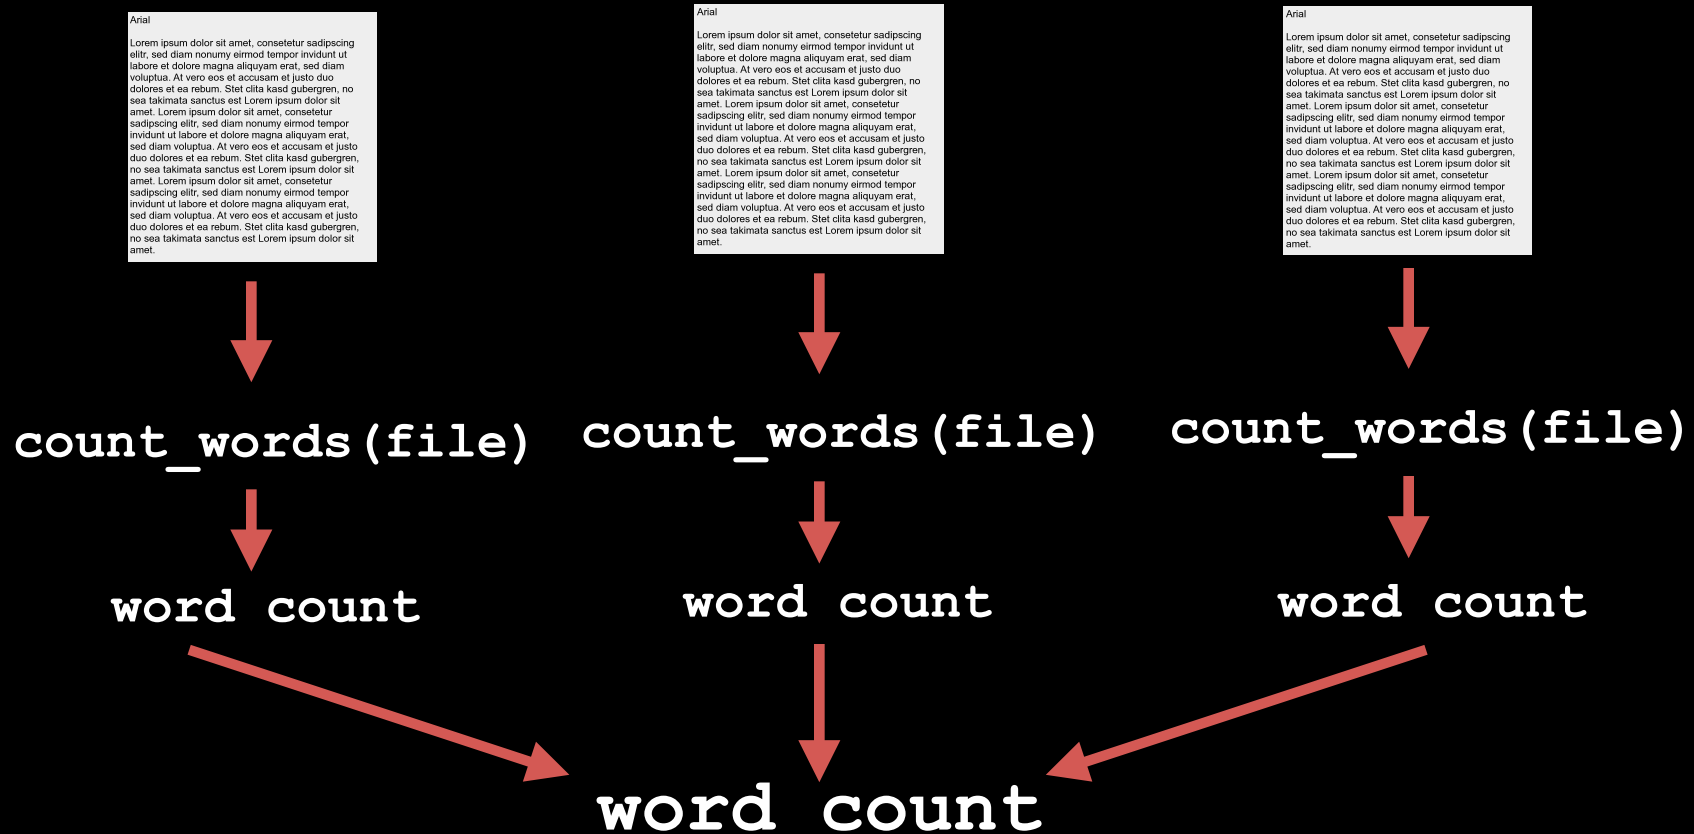
for loop



count\_words(file)

word count

# Counting words in documents (in parallel)



# Motivation

- Process biggish data ( $\geq 5\text{GB}$  depending on task)
- Saves time
- More efficient use of CPU resources

# Goals

- Motivation
- Intro to computing resources
- Multi-core processing
  - ★ Parallelism
- Threading
  - ★ Concurrency

# Types of Computing Resources

- **Central Processing Unit (CPU)**


- ★ Unit: GHz (instructions/second)

- **Random Access Memory (RAM / Memory)**

- ★ Unit: GB (Gigabytes)

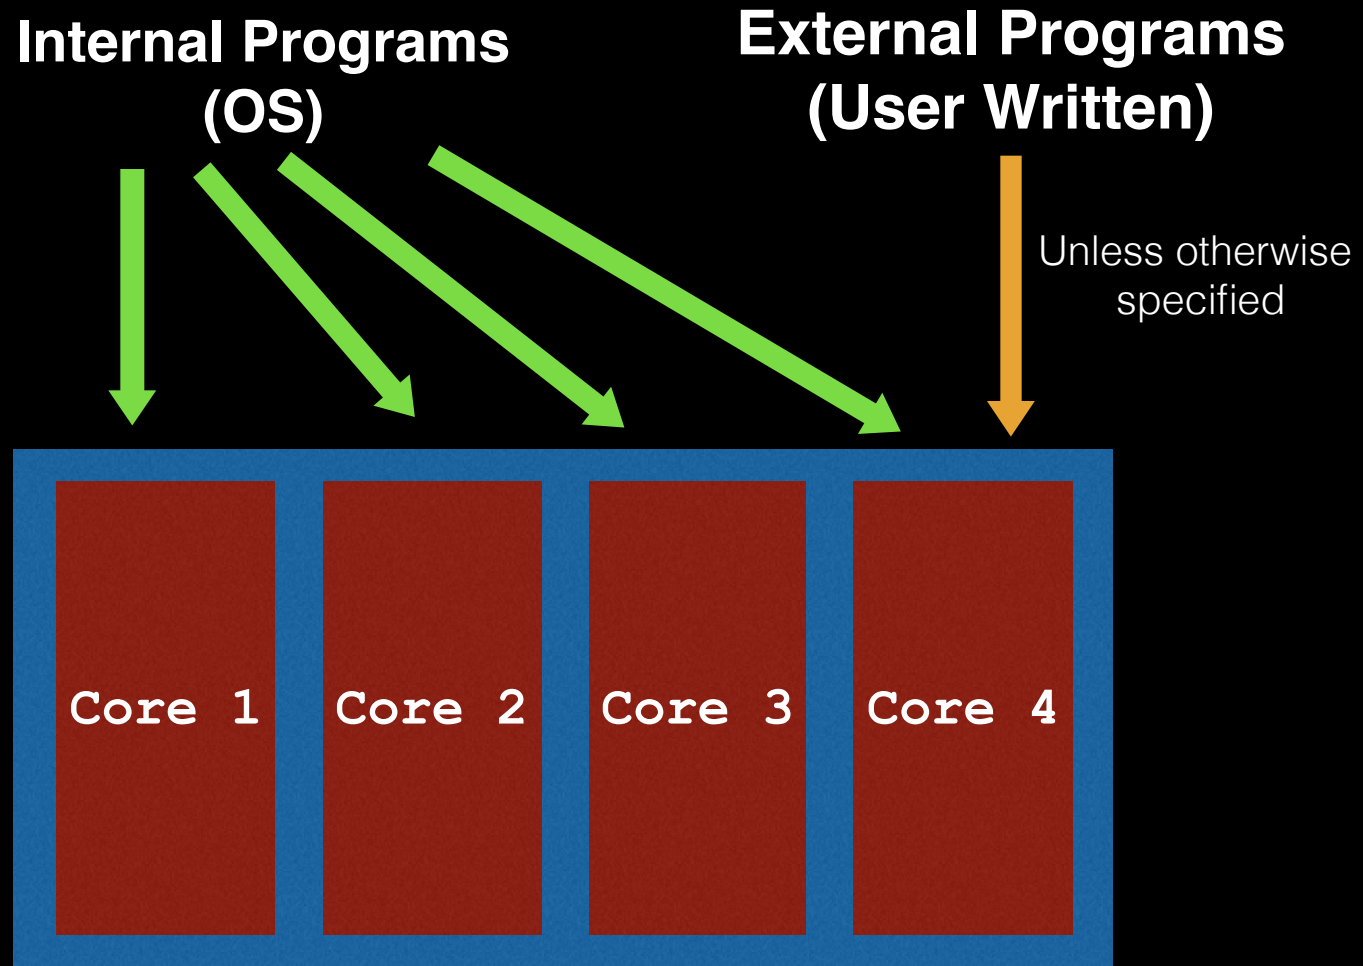


# CPU Cores



- A CPU can hold multiple cores
- Each core is a self-contained processor that can execute programs

# How are CPU Cores used



# What happens when a program runs

Program



Process

```
def edges_from_osmdb(osmdb, vertex_namespace, slogs, profiledb=None):
    """generates (vertex1_label, vertex2_label, edgepayload) from osmdb"""
    street_id_counter = 0
    street_names = {}

    # for each edge in the osmdb
    for i, (id, parent_id, node1, node2, distance, geom, tags) in enumerate(osmdb.edges()):

        # Find rise/fall of edge, if profiledb is given
        rise=0
        fall=0
        if profiledb:
            profile = profiledb.get(id)
            if profile:
                rise, fall = get_rise_and_fall(profile)

        # insert end vertices of edge to graph
        vertex1_label = "%s-%s"%(vertex_namespace,node1)
        vertex2_label = "%s-%s"%(vertex_namespace,node2)

        # create ID for the way's street
        street_name = tags.get("name")
        if street_name is None:
            street_id_counter += 1
            street_id = street_id_counter
        else:
            if street_name not in street_names:
                street_id_counter += 1
                street_names[street_name] = street_id_counter
                street_id = street_names[street_name]

        # Create edges to be inserted into graph
        s1 = Street(id, distance, rise, fall)
        s2 = Street(id, distance, fall, rise, reverse_of_source=True)
        s1.way = street_id
        s2.way = street_id

        # See if the way's highway tag is penalized with a 'slog' value; if so, set it in the edges
        slog = slogs.get(tags.get("highway"))
        if slog:
            s1.slog = s2.slog = slog

        # Add the forward edge and the return edge if the edge is not oneway
        yield vertex1_label, vertex2_label, s1

    oneway = tags.get("oneway")
    if oneway != "true" and oneway != "yes":
        yield vertex2_label, vertex1_label, s2
```



1. Process ID (pid)

2. Program Code

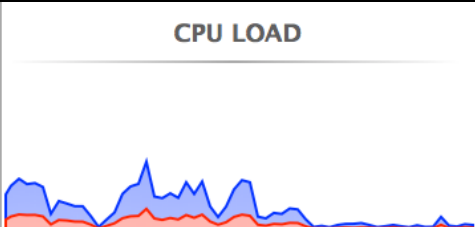


3. CPU Scheduling

4. Memory

# How much CPU / Memory

Activity Monitor (All Processes)						
CPU   Memory   Energy   Disk   Network						
Process Name	% CPU	CPU Time	Threads	Idle Wake Ups	PID	User
hidd	2.2	6:16.90	4	0	69	root
Activity Monitor	2.0	19.17	7	1	159	jeffreytar
kernel_task	1.9	15:50.58	82	168	0	root
Google Chrome Helper	1.2	2:11.83	11	1	281	jeffreytar
mongod	0.5	1:25.84	11	60	322	jeffreytar
WindowServer	0.3	20:12.10	5	3	92	_window
Google Chrome	0.1	15:58.10	38	1	154	jeffreytar
sysmond	0.1	2:05.99	3	0	41	root
Google Chrome Helper	0.1	48.90	13	1	2605	jeffreytar
mdworker	0.1	0.32	4	0	2666	jeffreytar
HyperDock	0.1	1:45.87	11	2	326	jeffreytar
Google Chrome Helper	0.0	9:46.70	4	1	209	jeffreytar
mds	0.0	7:06.23	7	1	60	root

System:	1.68 %		Threads:	706
User:	1.74 %		Processes:	139
Idle:	96.58 %			

# How much CPU / Memory

## Command line tool: top

```
Processes: 163 total, 921 running, 07 stuck, 1154 sleeping, 785 threads
Load Avg: 1.66, 1.71, 1.72 CPU usage: 0.67% user, 3.36% sys, 95.96% idle SharedLibs: 71M resident, 0B data, 12M linkedit.
MemRegions: 42943 total, 2229M resident, 63M private, 1895M shared. PhysMem: 7607M used (1073M wired), 283M unused.
VM: 425G vsize, 1313M framework vsize, 1159144(0) swapins, 2470739(0) swapouts. Networks: packets: 481571/399M in, 200872/77M out.
Disks: 1065467/56G read, 274229/20G written.
14:09:10
```

PID	COMMAND	%CPU	TIME	#TH	#WQ	#PORT	#MREG	MEM	RPRVT	PURG	CMPRS	VPRVT	VSIZE	PGRP	PPID	STATE	UID	FAULTS	COW
3274	screencaptur	0.0	00:00.12	2	0	47	91	1916K	888K	16K	0B	21M	2439M	162	162	sleeping	501	3753	243
3273	top	10.3	00:02.84	1/1	0	23	40	2524K+	2296K+	0B	0B	45M	2403M	3273	3252	running	0	38122+	91
3252	zsh	0.0	00:00.06	1	0	19	51	1960K	1816K	0B	0B	37M	2396M	3252	3251	sleeping	501	1874	501
3251	login	0.0	00:00.02	2	0	30	46	908K	580K	0B	0B	53M	2411M	3251	155	sleeping	0	811	143
3245	CVMCompiler	0.0	00:00.08	2	1	32	67	12M	11M	0B	0B	63M	2440M	3245	143	sleeping	501	5723	244
3224	zsh	0.0	00:00.06	1	0	19	52	1976K	1820K	0B	0B	45M	2404M	3224	3223	sleeping	501	1844	479
3223	login	0.0	00:00.02	2	0	30	46	996K	676K	0B	0B	53M	2411M	3223	155	sleeping	0	814	140

# Checkpoint

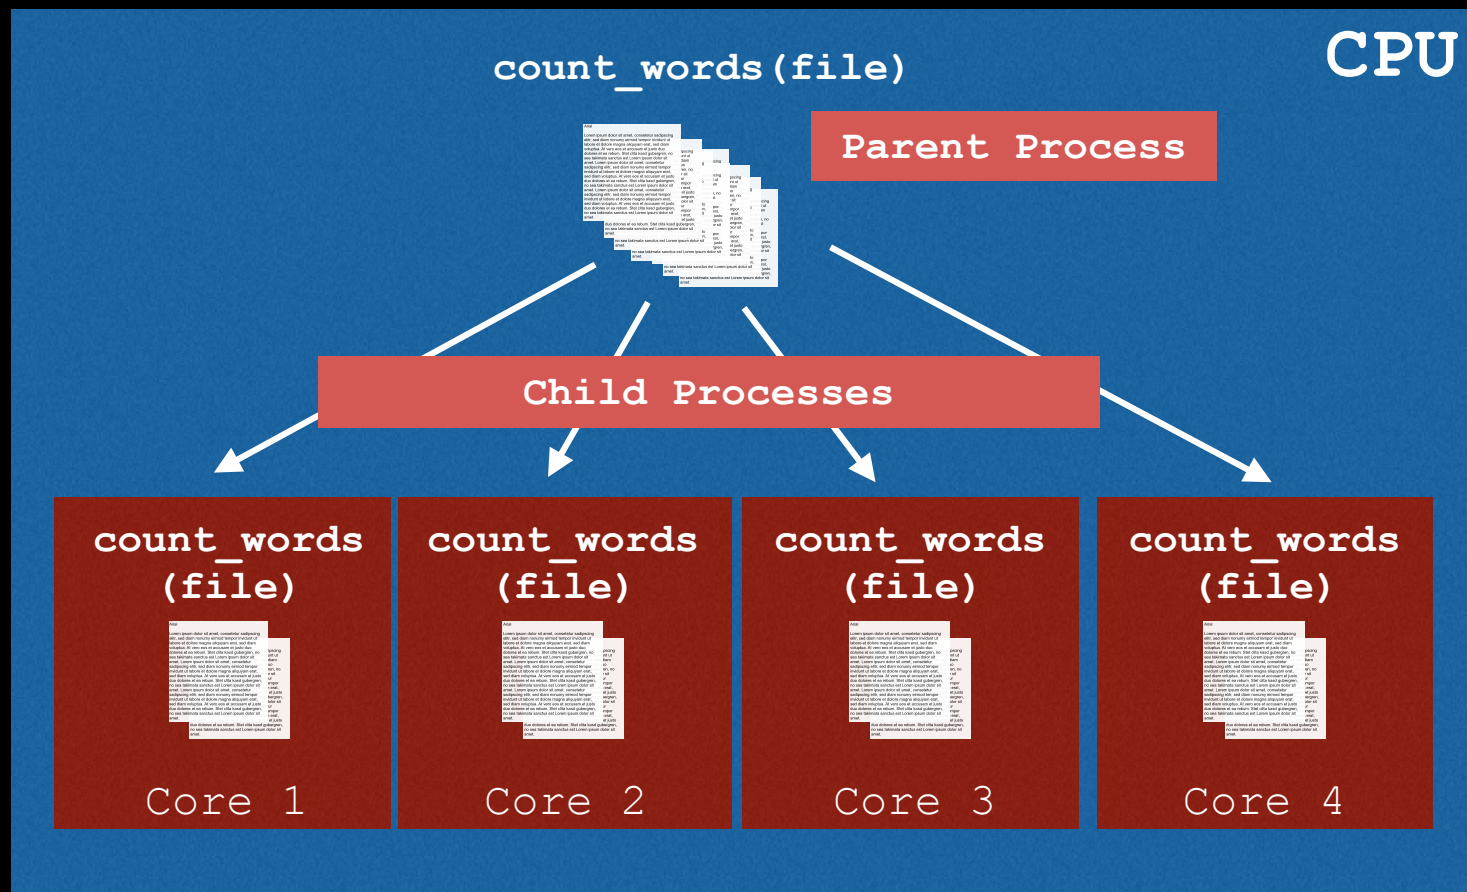
- **You should now understand:**
  - ★ CPU
  - ★ Memory
  - ★ Program and Process

# Goals

- Motivation
- Intro to computing resources
- Multi-core processing
  - ★ Parallelism
- Threading
  - ★ Concurrency

# Parallelism

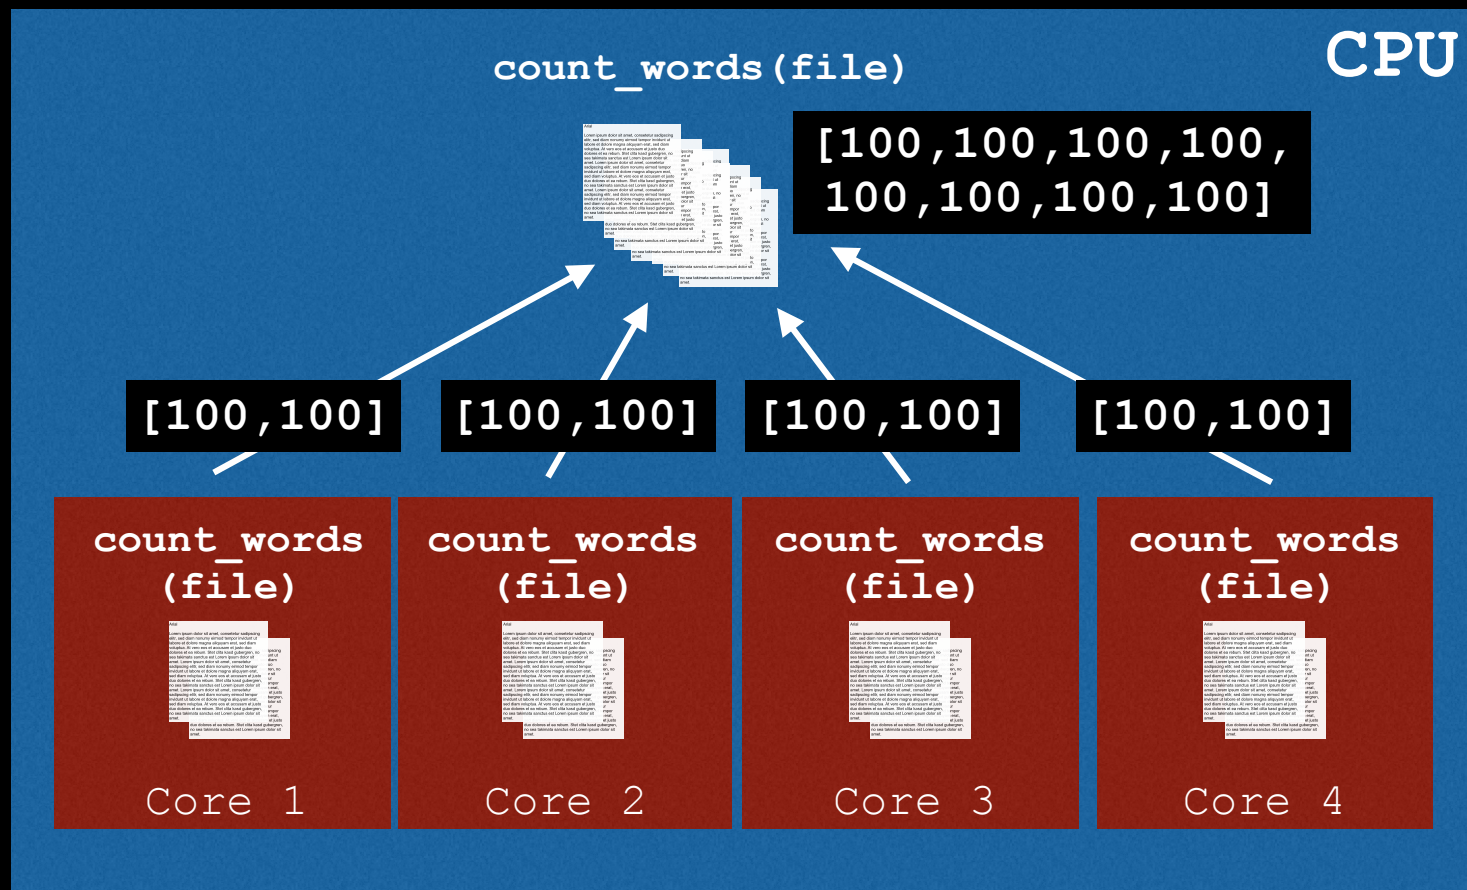
## Multicore Processing





# Parallelism

## Multicore Processing



```
from multiprocessing import Pool
import os

# Count the number of words in a file
def word_count(f):
    return len(open(f).read().split())

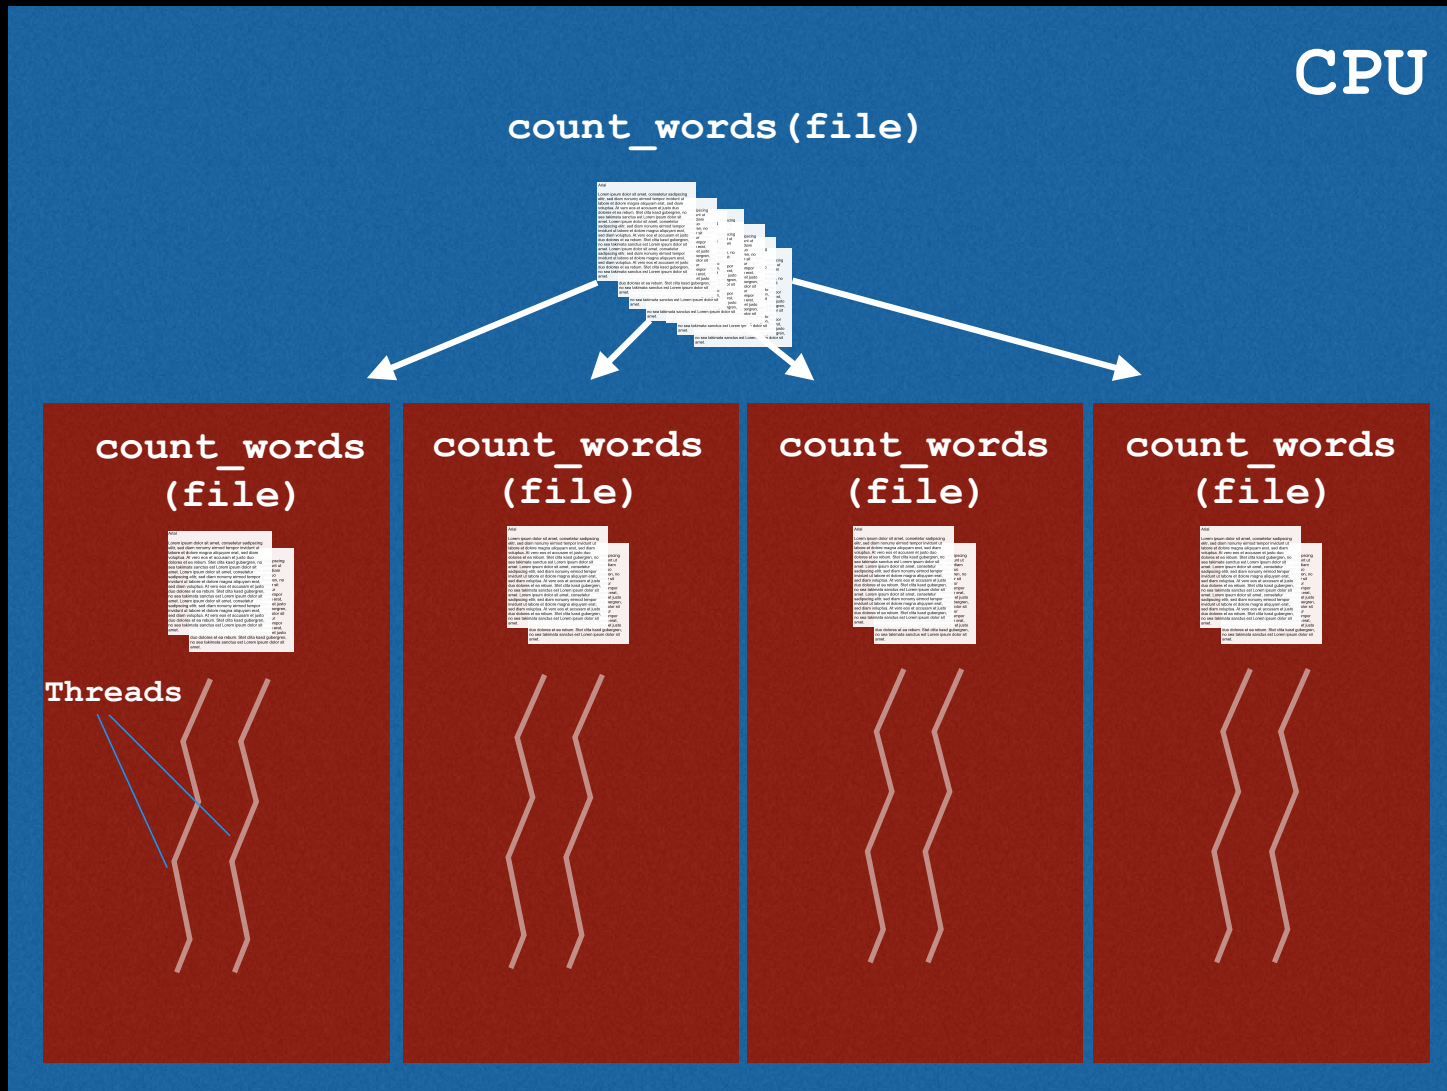
# Use a regular loop to count words in files
def sequential_word_count(lst_of_files):
    return sum([word_count(f) for f in lst_of_files])

# Use multiple cores to count words in files
def parallel_word_count(lst_of_files):
    pool = Pool(processes=4)
    results = pool.map(word_count, lst_of_files)
    print results
    # [100,100,100,100,100,100,100,100] # 8 files
    return sum(results)
```

# Goals

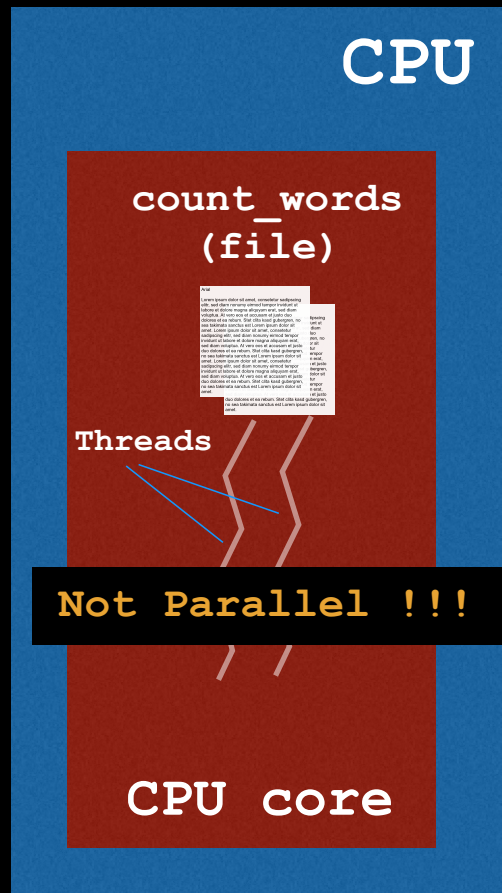
- Motivation
- Intro to computing resources
- Multi-core processing
  - ★ Parallelism
- Threading
  - ★ Concurrency

# Concurrency Threading




# Threads

- Threads within ***the same process*** are ***not parallel***



# Concurrency

- Threads **in the same process / core** are **concurrent**
- **1 process switching between multiple tasks**
- Not multi-core (multiple processes)




# Sharing between Threads

- **Share with other threads in the process:**
  - Memory allocated to the process
  - Code of the process
  - Context (i.e. variables)




# Parallelism vs Concurrency Analogy


Parallelism



Concurrency



# Parallelism Use Case




# Parallelism Use Case

- **Table 1** and **Table 2** both order dishes all the time and constantly need service
- CPU-bound problem

**Tasks that use heavy computations  
with no wait time in between**

# Concurrency Use Case



# Concurrency Use Case

- **Table 1** and **Table 2** order and can only order again when each finish the dish
- The wait time allow 1 waiter to service 2 tables
- I/O bound problem
  - **Read / Writing Files**
  - **Making Web Requests**

# Python Library

## Threading

```
import threading

jobs = []

# Initiate and Start Threads
for i in range(num_threads):
    t = threading.Thread(target=target_function, args=(arg1, arg2))
    jobs.append(t)
    t.start()

# "join" to make sure wait until the thread terminates
results = []
for t in jobs:
    t.join()
    # Access the result of the thread (if any)
    # Append to list
    results.append(t.result)
```

# Summary

- Parallelism on multiple cores
- Threading within a core
- **multiprocessing** for parallelism
- **threading** for concurrency
- Parallelism —> CPU-bound problems
- Concurrency —> I/O-bound problems