

Using AWS

Miles Erickson

October 24, 2016

Objectives

Morning objectives:

- Identify use cases for cloud computing and AWS
- Describe core AWS services
- Configure your computer to use AWS CLI
- Use SSH key to access Amazon EC2 instances
- Create and access an Amazon S3 bucket

Agenda

Morning:

- Introduction to AWS
- Install and configure AWS CLI
- Configure and use SSH
- Access EC2
- Access S2
- Advanced topics

Afternoon:

- High-performance python

References

- [AWS in Plain English](#)
- [Open Guide to AWS](#)
- [AWS Documentation](#)
- [AWS Support](#)

Amazon Web Services (AWS)

AWS provides on-demand use of computing resources in the cloud

- No need to build data centers
- Easy to create a new business
- Only pay for what you use
- Handle spikes in computational demand
- Secure, reliable, flexible, scalable, cost-effective

AWS skills are much in demand

AWS core services:

- Elastic compute cloud (EC2): computers for diverse problems
- Elastic block store (EBS): virtual hard disks to use with EC2
- Simple storage solution (S3): long-term bulk storage
- DynamoDB: NoSQL database
- And much, much more

Elastic compute cloud (EC2)

Spin up EC2 instances for on-demand computing power:

- Instance: a type of hardware you can rent, e.g., 'm3.xlarge' or 't2.micro'
- Amazon Machine Image (AMI), an OS you can run on an instance
- Region: a geographic region, e.g., Oregon a.k.a. 'us-west-2'
- Availability Zone (AZ): a specific subset of a region, often a data center, such as 'us-west-2a'

Elastic block store (EBS)

EBS provides disk-drive style storage:

- Create a virtual hard disk
- Then mount virtual hard disk on EC2 instances
- SSD or magnetic
- Can store data even when you aren't running any EC2 instances
- Built-in redundancy
- Lower latency than S3 but more expensive

Simple storage solution (S3)

S3 provides cheap, bulk storage:

- Create a bucket which serves as a container for files and directories
- Specify permissions using an access control list (ACL)
- Access via URL or AWS CLI or suitable API
- Higher latency than EBS but less expensive

Use the AWS command-line interface (CLI):

- To debug your configuration
- To manage AWS instances in EC2
- To access S3

Install AWS CLI

If you're using Anaconda:

- `pip install awscli`

If you're using system Python (you're not): `* sudo pip install awscli`

Obtain your AWS credentials

- 1 Login to AWS at <https://console.aws.amazon.com/>
- 2 Click on your name in the upper right menu bar
- 3 Select **Security Credentials** and go to **IAM Users**
- 4 Click **Users** and select your name (or add it)
- 5 Click User Actions > Manage Access Keys
- 6 Create Access Key, Show User Security Credentials
- 7 Copy key into...

Make sure you choose Oregon (us-west-2) as your region.

Configure AWS CLI (1/3)

Easiest way: run `aws configure`:

- Creates default profile in `~/.aws/config`
- Stores credentials in `~/.aws/credentials`
- Can create multiple profiles:

```
$ aws configure --profile fancy_profile
```

- Can also set credential on CLI or via environment variables

Configure AWS CLI (2/3)

Create AWS configuration info in ~/.aws:

```
$ aws configure
```

```
AWS Access Key ID [None]: AKIAIOSFODNN7EXAMPLE
```

```
AWS Secret Access Key [None]: wJalrXUtnFEMI/K7MDENG/bPxRfiCYEX
```

```
Default region name [None]: us-west-2
```

```
Default output format [None]: json
```

See Amazon's [doc](#) for details

Configure AWS CLI (3/3)

Some tools get AWS credentials via environment variables. Set the following in `~/.bash_profile` or equivalent:

```
export AWS_ACCESS_KEY_ID='your access key'  
export AWS_SECRET_ACCESS_KEY='your secret key'
```

Verify configuration

Check S3:

```
$ aws s3 ls
2015-08-25 10:42:43 dsci
2015-08-25 11:30:33 seattle-dsi
$ aws s3 ls s3://seattle-dsi
    PRE cohort1/
    PRE dsi_interviews/
```

Check EC2:

```
$ aws ec2 describe-instances --output table
$ aws ec2 describe-instances --output json
```


Help with AWS CLI

See the built-in help for more details:

```
$ aws help
```

```
$ aws ec2 help
```

. . . or Google

Configure and use SSH

To use AWS:

- Setup Key Pairs to access EC2
- Configure SSH on your laptop
- Can use SSH to access any remote machine running an SSH server

The secure shell protocol allows you to:

- Login to remote machines, such as EC2 using ssh
- Transfer files between remote machines using scp and sftp
- Execute commands on remote machines using ssh

Do not use telnet, rlogin, or FTP, which are older, insecure protocols!

Public-key encryption

Pairs of matched keys that can be used to encrypt or decrypt documents:

- Encrypt with public key → decrypt with private key
- Encrypt with private key → decrypt with public key

But:

- Trivial to calculate public key from private
- Very, very, very hard to calculate private key from public

Public-key encryption

SSH uses public-key encryption to protect access and encrypt communication:

- Generate a public & private key pair
- Keep private key safe
- Create a key pair for each resource you want to access (AWS, GitHub, etc.) \Rightarrow can revoke individual keys in case of a security breach

See AWS [doc](#) for details

Setup Key Pairs

Create and configure key pair to access EC2:

- Create and import key pair as described in *AWS* documentation
- Set permission on private key to 400:

```
$ chmod 400 ~/.ssh/aws-master.pem
```

```
$ chmod 444 ~/.ssh/aws-master.pub
```

- Can also generate key pair with `ssh-keygen`

Configuring SSH

Modify SSH config to:

- Create alias for long-running instance
- Forward X11 or security information
- Specify which key to use
- And, much much more...
- `man ssh_config` for details

Example ~/.ssh/config

```
Host github.com
  HostName github.com
  User git
  IdentityFile ~/.ssh/git-hub-id_rsa
```


Example ~/.ssh/config

Setup an alias:

```
Host master
  HostName ec2-54-186-136-57.us-west-2.compute.amazonaws.com
  User ubuntu
  ForwardAgent yes
  TCPKeepAlive yes
  IdentityFile /Users/jackbenn/.ssh/aws-master.pem
```

Now, `ssh master` will connect to your EC2 instance

Accessing an EC2 instance with ssh

Connect to your machine via ssh:

- 1 Launch an EC2 instance from console
- 2 Use ssh from command line to connect to the instances **public DNS** (Shown in EC2 Dashboard):

```
$ ssh -i ~/.ssh/aws-master.pem \  
ubuntu@ec2-54-186-136-57.us-west-2.compute.amazonaws.co
```

or

```
$ ssh -i ~/.ssh/aws-master.pem \  
ec2user@ec2-54-186-136-57.us-west-2.compute.amazonaws
```

Example: ssh to EC2

```
$ ssh -i ~/.ssh/aws-master.pem ubuntu@ec2-54-186-136-57
The authenticity of host 'ec2-54-186-136-57.us-west-2.compu RS
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'ec2-54-186-136-57.us-west-2.com
Welcome to Ubuntu 14.04.2 LTS (GNU/Linux 3.13.0-48-generic
* Documentation:  https://help.ubuntu.com/
System information as of Tue Aug 25 21:58:16 UTC 2015
System load: 0.0                Memory usage: 5%    Processes
Usage of /:  9.8% of 7.74GB      Swap usage:  0%    Users log
```

Example: ssh to EC2

```
$ ssh -i ~/.ssh/aws-master.pem ubuntu@ec2-54-186-136-57
The authenticity of host 'ec2-54-186-136-57.us-west-2.compu RS
Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'ec2-54-186-136-57.us-west-2.com
Welcome to Ubuntu 14.04.2 LTS (GNU/Linux 3.13.0-48-generic
* Documentation:  https://help.ubuntu.com/
System information as of Tue Aug 25 21:58:16 UTC 2015
System load: 0.0                Memory usage: 5%    Processes
Usage of /:  9.8% of 7.74GB     Swap usage:   0%    Users log
```

Transferring files with scp

To copy files between machines, use `scp`:

- Works just like regular copy
- Good for simple operations
- ... if you specify remote user and machine (IP,DNS) correctly
- Reference remote location as *user@host:path*

```
$ scp -i ~/.ssh/aws-master.pem ./toy_data.txt \  
ubuntu@54.186.136.57:/home/ubuntu/data  
toy_data.txt 100% 136 0.1KB/s 00:00
```

Transferring files with sftp

To copy files interactively, use sftp:

- Interactive shell for transferring files
- Use to transfer many files
- Use when you don't know the location of a file

```
$ sftp -i ~/.ssh/aws-master.pem ubuntu@54.186.136.57  
Connected to 54.186.136.57.  
sftp> help
```

Managing a session with `tmux`

Use `tmux` to persist jobs across multiple sessions:

- On logout, all child processes terminate
- Use `tmux` to safely disconnect from a session
- Reconnect on next login
- Install `tmux` via `brew` or Linux package manager
- See `tmux` exercise

EC2 pro-tips

A few tips to make EC2 easier to deal with:

- Always create instances with tags so that you can find them easily
- Choose the appropriate hardware type for your problem
- If in doubt, use Ubuntu because it is a friendly flavor of Linux (optionally, use `galvanize-dsi-ami`)
- Use `tmux` when you login in case you need to disconnect or your connection dies
- Be paranoid: sometimes Amazon will reboot or reclaim your instances
- Put data you need to persist in EBS or a database
- Never put AWS keys in GitHub because someone will steal them

Launching an EC2 instance

- Class exercise: Launch a t2.micro instance running Ubuntu

To master the basics, see this [tutorial](#) or `lecture_notes.ipynb` in the repo

- Can find URL to access a file from S3 console
- Set properties (access) via S3 console
- Make sure names conform to S3 conventions:
- lowercase bucket names of at least four characters
- no leading or terminal ‘.

Boto config

To access S3 via Python, use the boto package

- Should be installed if you followed setup instructions
- Make sure boto is up to date:

```
$ conda update boto
```

- Uses credentials in `~/.aws/credentials` which you setup earlier
- Can also read directly from Pandas if you specify S3 URL

Advanced: accessing ipython notebook

- Use an ssh tunnel to run ipython notebook on a remote instance
- On remote host:

```
$ jupyter notebook --no-browser --port=8889
```

- On local machine:

```
$ ssh -N -f -L localhost:8888:localhost:8889 \  
  remote_user@remote_host
```

- Access notebook via browser at URL localhost:8888