

# Image Featurization

## CovNet style

# Objectives

- Be able to explain what a convolution is, and how it works
- Understand the basic structure of a convolutional neural network
- Comprehend the three basic ideas behind convolutional networks :
  - (1) Local receptive field
  - (2) Shared weights
  - (3) Pooling
- Be aware of general strategies for building convolutional neural networks

# Convolutions

In image processing, a kernel, convolution matrix, or mask is a small matrix useful for blurring, sharpening, embossing, edge-detection, and more. This is accomplished by means of **convolution** between a kernel and an image.


-Wikipedia

# Convolutions

A Kernel

1	0	1
0	1	0
1	0	1

Applied Over



A Simple Image

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

1 <sub>x1</sub>	1 <sub>x0</sub>	1 <sub>x1</sub>	0	0
0 <sub>x0</sub>	1 <sub>x1</sub>	1 <sub>x0</sub>	1	0
0 <sub>x1</sub>	0 <sub>x0</sub>	1 <sub>x1</sub>	1	1
0	0	1	1	0
0	1	1	0	0

Image

4		

Convolved  
Feature

# Example Kernels - Edge Detectors

Vertical Edge Detector

-1	0	+1
-2	0	+2
-1	0	+1

Horizontal Edge Detector

-1	-2	-1
0	0	0
+1	+2	+1


**Sobel filter**



How do we tell that  
this is a door?




What if we applied our edge detectors to the image?



# Convolutional Neural Networks (CNNs)

What if we could get a computer to build it's own kernels, apply those to images, and then interpret those results to perform object recognition?



Enter convolutional neural networks....






# General Structure


Input Layer




Convolutional  
Layers




Pooling  
Layers



Fully  
Connected  
Layers




Output  
Layer




# Three Main Ideas...

- Local Receptive Fields
- Shared Weights
- Pooling


# General Structure



# Input Image




# Local Receptive Fields




- This group of pixels is a **\*\*local receptive field\*\***, and its size is defined by the size of your kernel
- We then slide this kernel across the entire image

# Convolutional Layers

We apply multiple kernels to the image, which results in multiple learned kernels per hidden layer



# General Structure

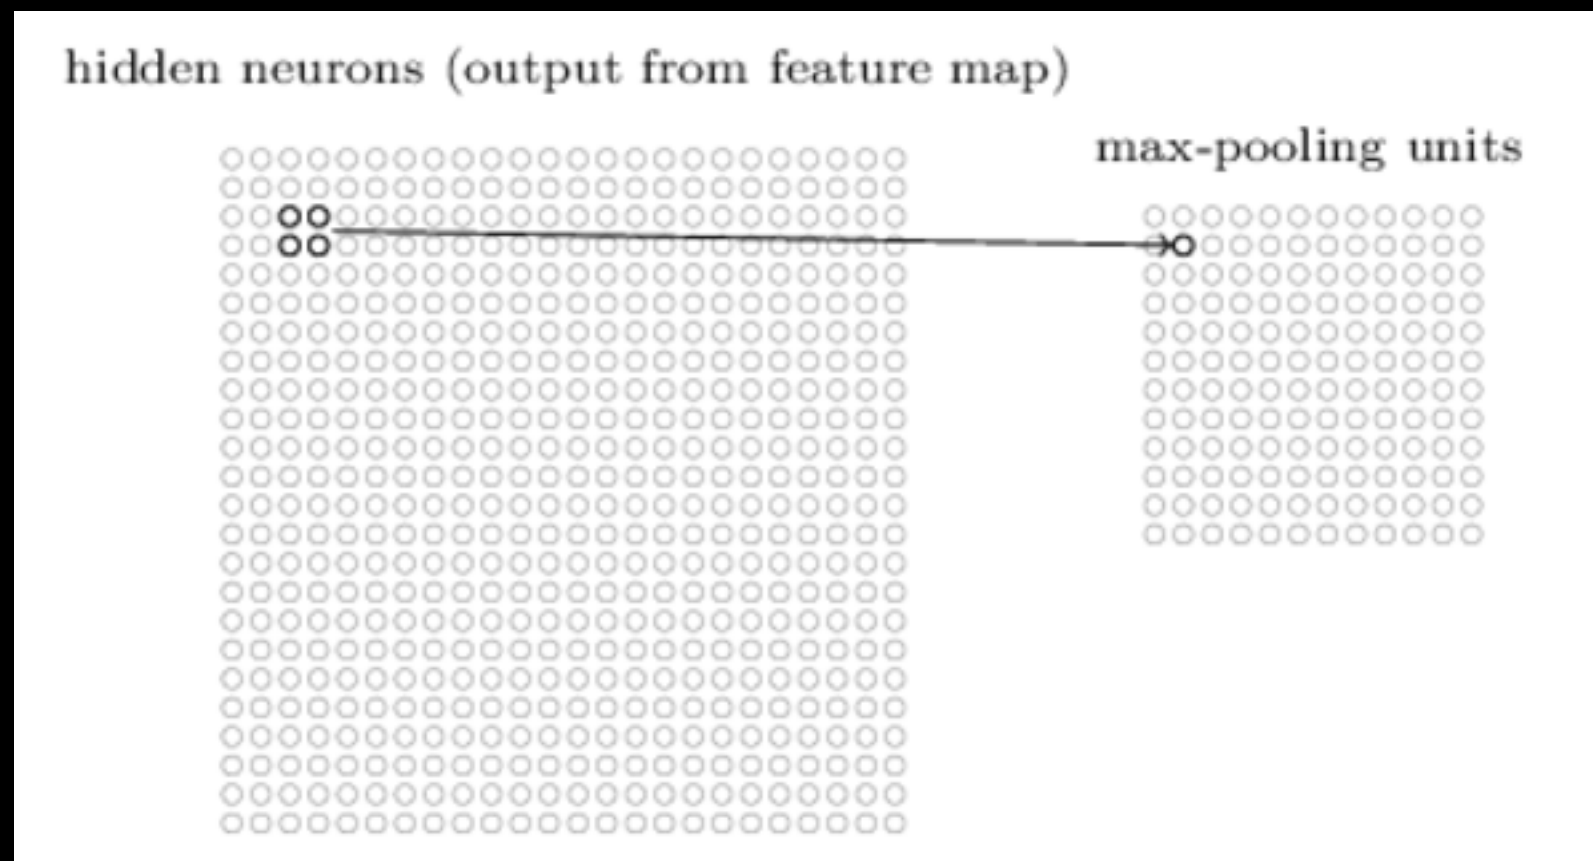


# Pooling Layers

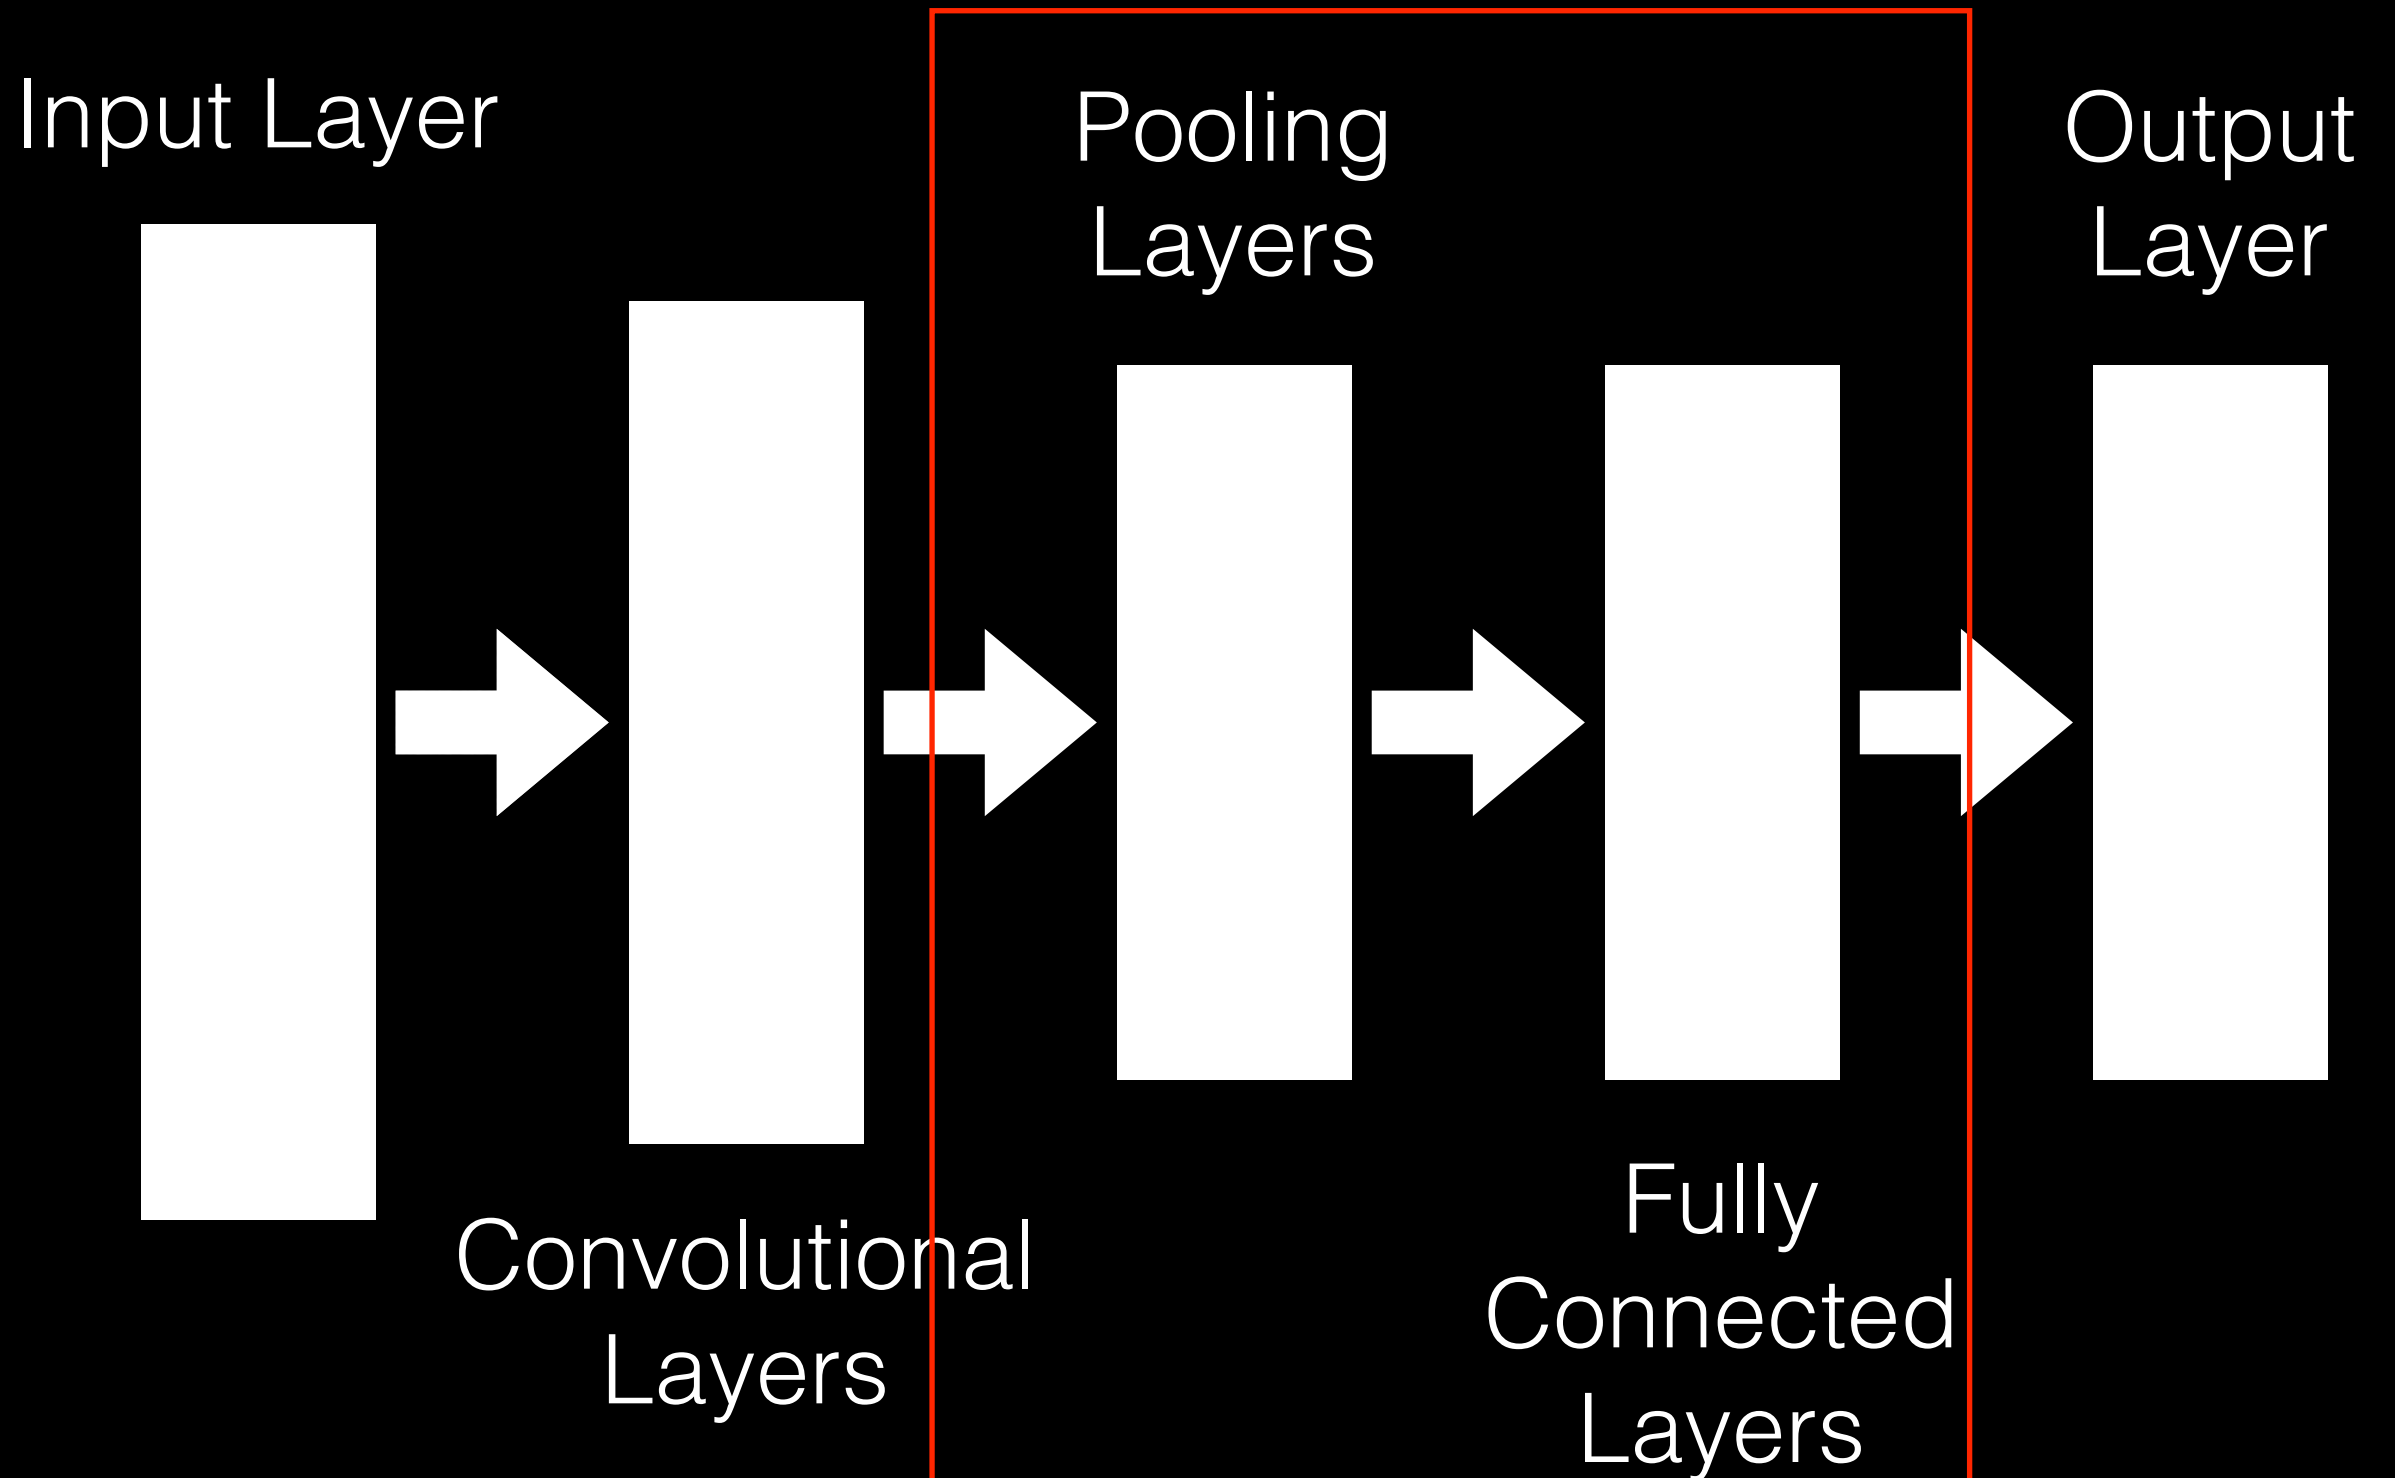
Pooling layers are used immediately after convolutional layers, and simplify the information in the output from the convolutional layer.



# Max Pooling



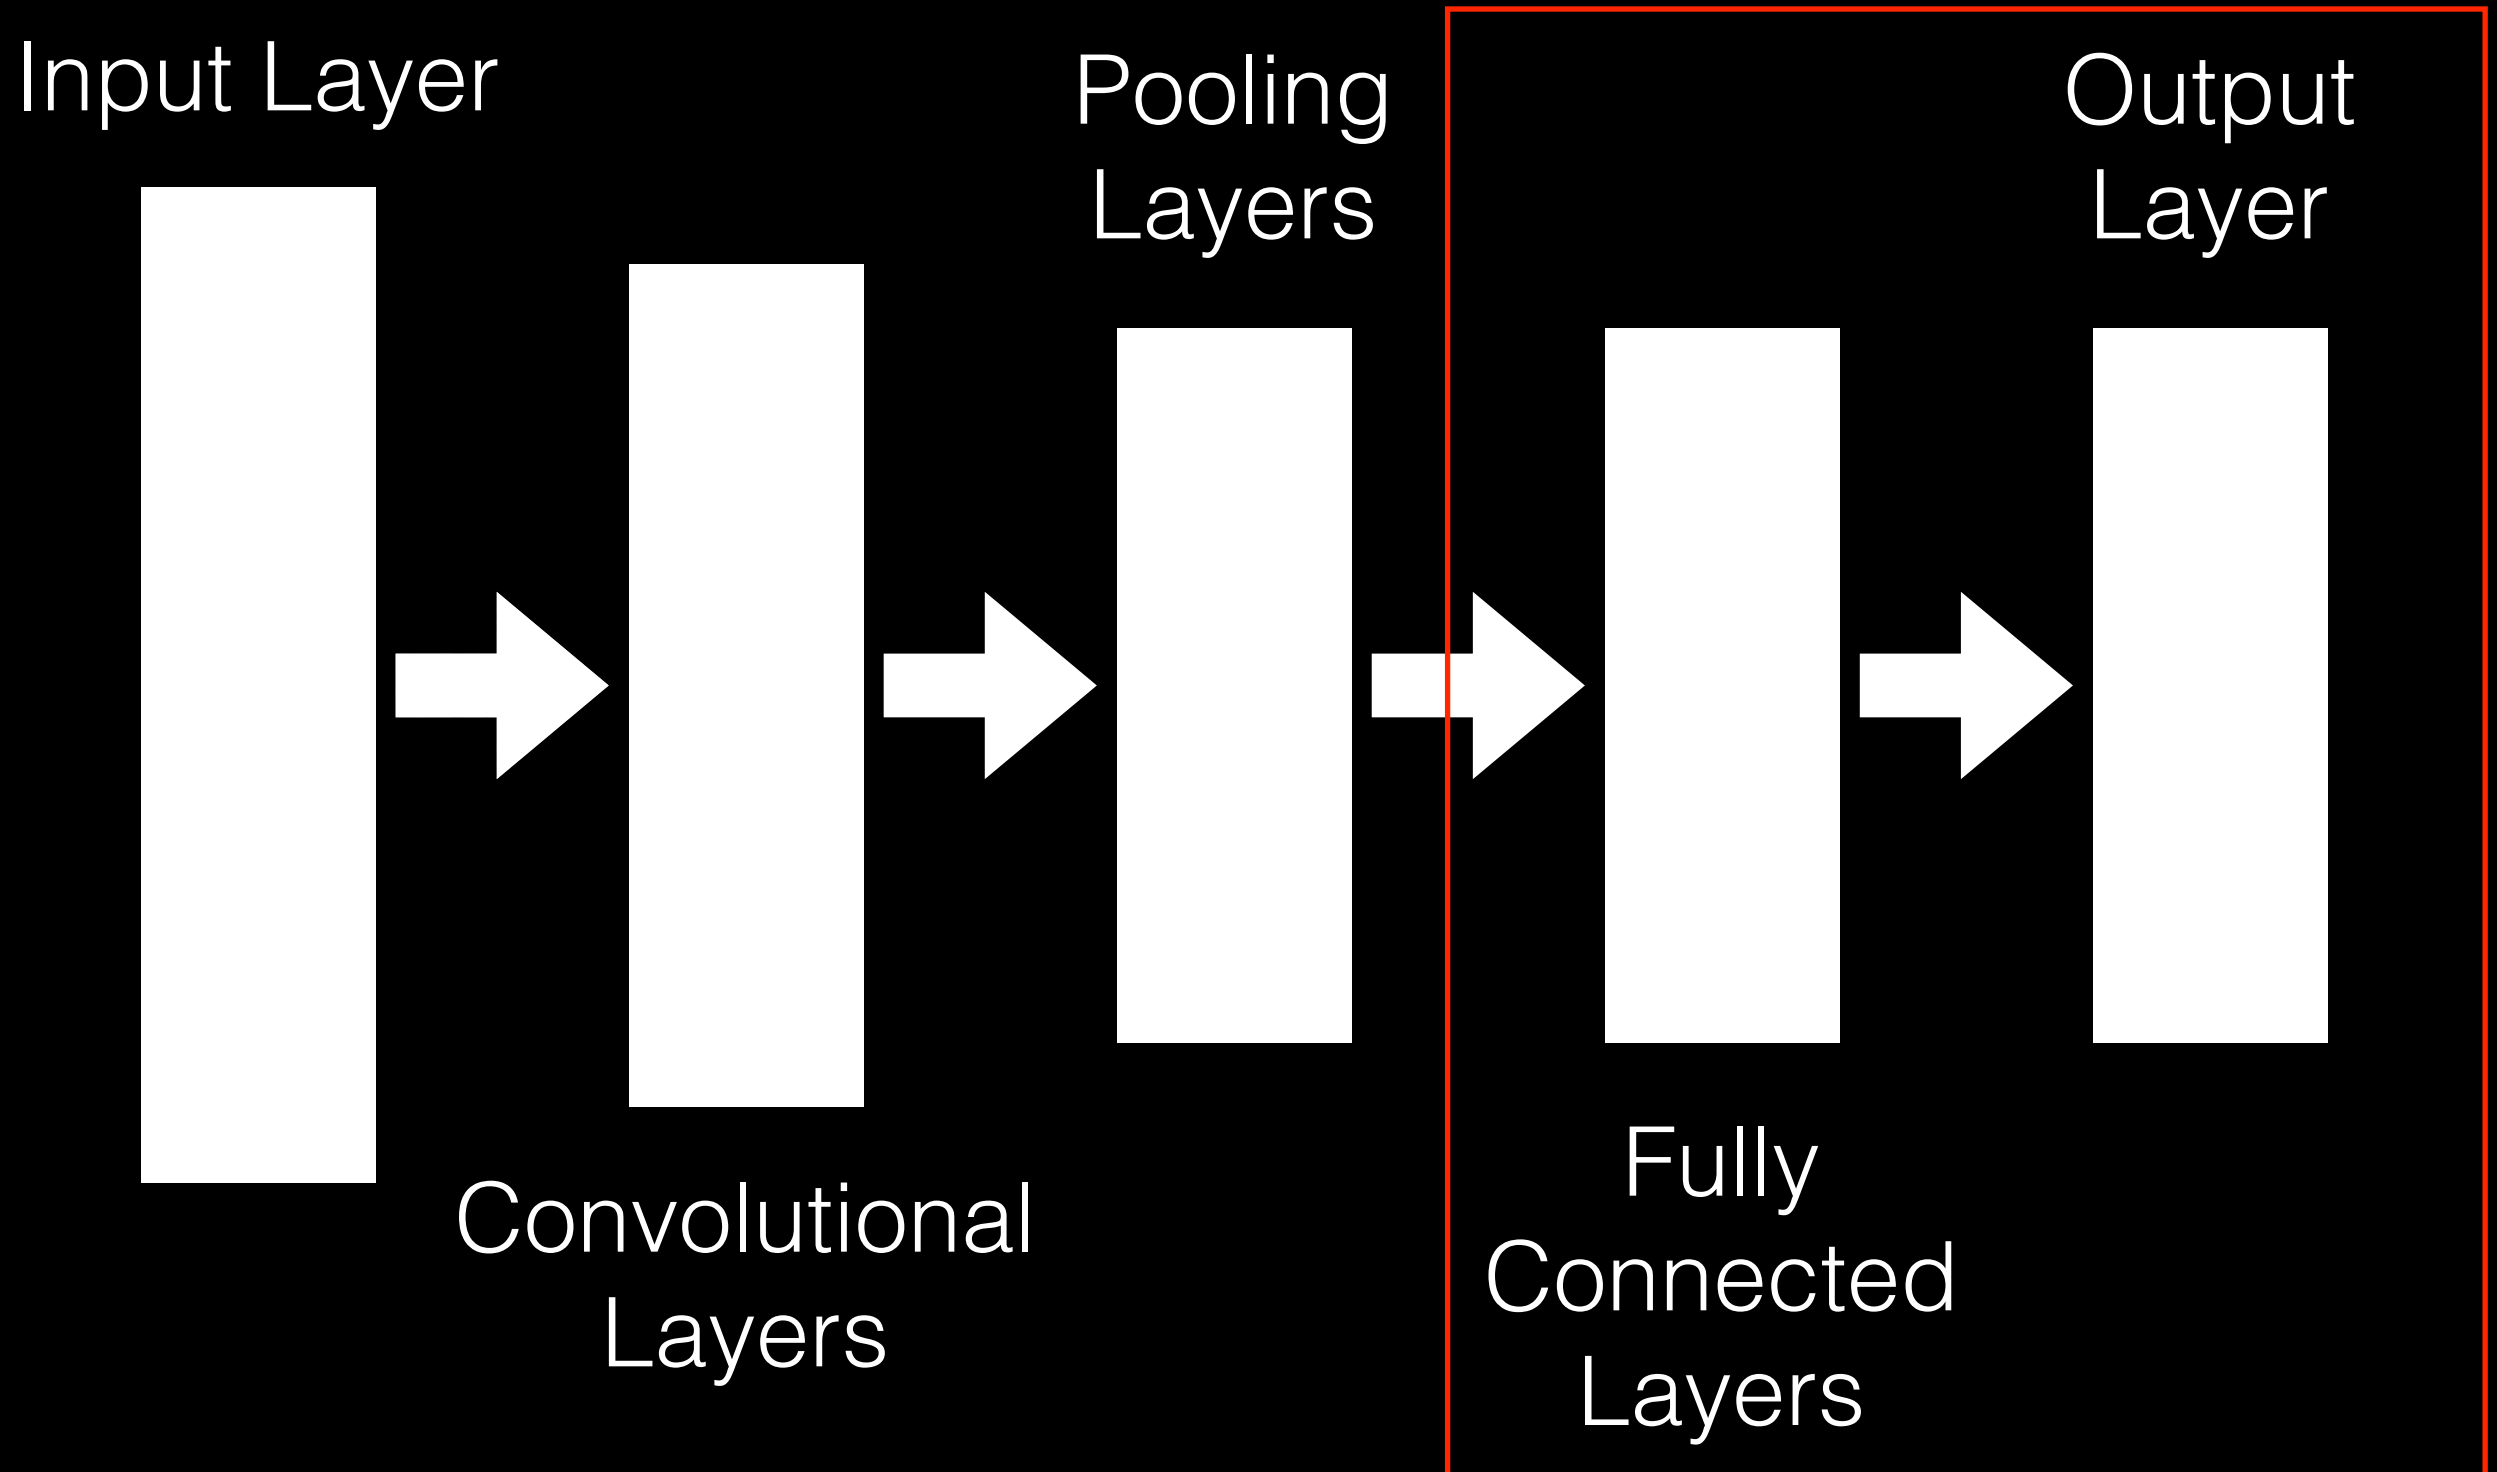
# General Structure



# Fully Connected Layers

Used to aggregate all of the information that has been learned in the convolutional and pooling layers.

# General Structure



# Output Layer

- For object recognition, always composed of softmax activation units
- Outputs the probability that your image is of a certain class

# Tips and Tricks

- We of course have all the image processing techniques that we learned about - resizing, denoising, etc.  
(denoising is not actually common to use with CNN's, but available)
- New set of image processing techniques for getting more images - rotate images, flip images, etc.

# General Structure

## Normalities

- It is not too common to use dropout after convolutional layers (but it is common to use it after you're fully connected layers)
- It's common to have multiple convolutional layers in between pooling layers
- ReLU activation units are incredibly popular with CNN's

# Don't know where to start?

If you're confused about the general CNN structure that you should start off, you should find a research paper in your domain space that uses CNN's. Start off trying to get something working that uses the same structure they did, and go from there.