


Logistic Regression 1/2

Classification, metrics
and ROC curves

DSI, jf.omhover, Dec 6, 2016



Logistic Regression 1/2

Classification, metrics
and ROC curves

DSI, jf.omhover, Dec 6, 2016

OBJECTIVES (morning)

- **Relate** Regression to Classification in the context of supervised learning
- **Compare** Logistic Regression to Linear Regression
- **Define** and **compute** metrics for evaluating classifiers

OBJECTIVES (afternoon)

- **Describe** the process for computing parameter values in LogReg
- **Use** the parameters of a LogReg model to **compute** the class of an observation





Supervised Learning

Learning / Estimating FUNCTIONS based on examples

Reality VS Model : assumptions and learning



REALITY

	type	income	education	prestige
accountant	prof	62	86	82
pilot	prof	72	76	83
architect	prof	75	92	90
author	prof	55	90	76
chemist	prof	64	86	90
minister	prof	21	84	87
professor	prof	64	93	93
dentist	prof	80	100	90
reporter	wc	67	87	52
engineer	prof	72	86	88
undertaker	prof	42	74	57
lawyer	prof	76	98	89

data

(x_1, y_1)

...

(x_n, y_n)

$x \ y$

OBJECTIVE:
descriptive
predictive
normative

...

$$\sum (y_i - \hat{f}(x_i))^2$$

COST FUNCTION


MODEL

$$y = f(x) + \epsilon$$

take a function as
an assumption

$$\hat{y} = \hat{f}(x)$$

Estimator
of the function




Linear Regression - General Process



REALITY


- 1) Having a data sample
Observing an underlying behavior



MODEL

- 2) Make an assumption
on the model underlying the data

$$y \approx \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$



linear relation
(+ assumptions)

- 3) Find the instance of the model
that fits with data sample

Multi-Linear Regression



COST FUNCTION (Residual Sum of Squares)


$$RSS(\beta) = (y - X\beta)^T(y - X\beta)$$

O.L.S.

REALITY

DATA

$$X = \begin{bmatrix} 1 & x_{1,1} & \cdots & x_{1,p-1} \\ 1 & x_{2,1} & \cdots & x_{2,p-1} \\ \vdots & \vdots & & \vdots \\ 1 & x_{n,1} & \cdots & x_{n,p-1} \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = y$$



MODEL

model class

$$y \approx X\beta$$

PROBLEM

$$\hat{y} = X\hat{\beta}$$

**model instance
estimator
parameters**

SOLUTION

$$\hat{\beta} = (X^T X)^{-1} X^T y$$



Classification

Learning / Estimating “models of classes” based on examples

Reality vs Model : assumptions and learning



REALITY

$$X = \begin{bmatrix} 1 & x_{1,1} & \cdots & x_{1,p-1} \\ 1 & x_{2,1} & \cdots & x_{2,p-1} \\ \vdots & \vdots & & \vdots \\ 1 & x_{n,1} & \cdots & x_{n,p-1} \end{bmatrix} \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix}$$

$= y$

OBJECTIVE:
descriptive
predictive
normative
...

$$\sum (y_i - \hat{f}(x_i))^2$$

COST FUNCTION

MODEL

$$y = f(x) + \epsilon$$

take a function as
an assumption

$$\hat{y} = \hat{f}(x)$$

Estimator
of the function

Mapping // Classification algorithms



Logistic Regression

k-NN


Decision Trees

Random Forest, Boosting

Support Vector Machines (SVM)

Neural Networks


...




Regression vs Classification



Quantitative response y in R



Categorical response y




Assigning $y = 0$ to neg, $y = 1$ to pos


Regression vs Classification



Quantitative response y in \mathbb{R}




Categorical response y in $\{0,1\}$




Regression vs Classification



Quantitative response y in \mathbb{R}




Categorical response y in $\{0,1\}$



Trying to apply LinReg to y



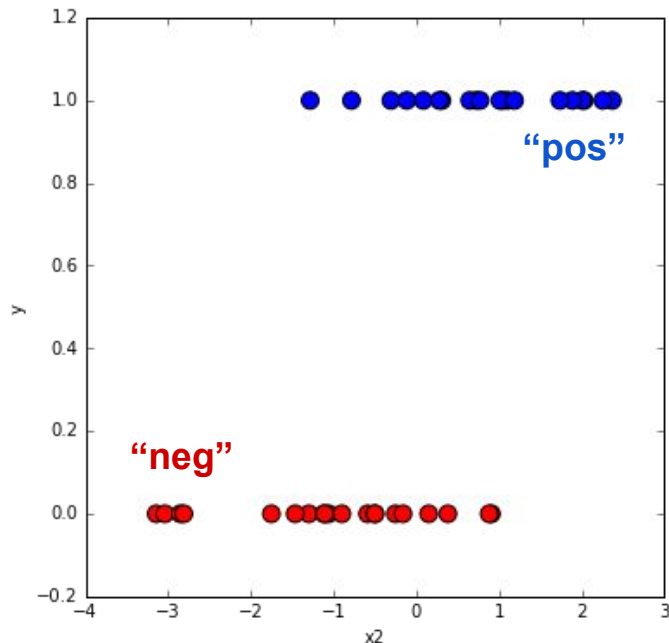
Quantitative response y in R



$$y \approx \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$

Negative probabilities ?
How to cut-off ?


Categorical response y in {0,1}



LogReg as model of probability




Quantitative response y in \mathbb{R}



$$y \approx \beta_0 + \beta_1 x_1 + \beta_2 x_2 + \dots + \beta_p x_p$$

Negative probabilities ?
How to cut-off ?

Categorical response y in $\{0,1\}$




$$p(X) = h(\beta_0 + \beta_1 \cdot x_1 + \dots + \beta_p \cdot x_p)$$

Idea : model probability of being positive
as a function of a linear model

LogReg in a nutshell



REALITY




*It (badly) translates as :
computes the probability
of being in one of the two
classes
depending on of the side
and distance of the plan*

MODEL

$$p(X) = h(\beta_0 + \beta_1 \cdot x_1 + \dots + \beta_p \cdot x_p)$$

$$h : \mathbb{R} \rightarrow [0, 1]$$

$$h(t) = \frac{1}{1+e^{-t}}$$





How to evaluate a classifier ?

Conforming a classifier to the actual response



x_1, x_2, x_3, x_4				y	$\hat{y} = \hat{f}(x)$	
...	1	1	0.95
...	0	0	0.21
...	0	1	0.55
...	1	0	0.43
...	1	1	0.77
...	1	0	0.44
...	0	0	0.15
...	1	1	0.81

		$\hat{y} = \hat{f}(x)$		
		Pred P	Pred N	
y	Actual P	3	2	P = 5
	Actual N	1	2	N = 3
		P*	N*	

Confusion Matrix



$$\hat{y} = \hat{f}(x)$$

		$\hat{y} = \hat{f}(x)$	
		Pred P	Pred N
y	Actual P	True Positive	False Negative
	Actual N	False Positive	True Negative
		P*	N*
		P = 5	
		N = 3	

Confusion Matrix - Metrics



The proportion of observations that are correctly classified ?

Accuracy :

The proportion of positives that are correctly identified as such ?

True Pos Rate :

(aka recall, sensitivity)

The proportion of negatives that are correctly identified as such

True Neg Rate :

(aka specificity)

$\hat{y} = \hat{f}(x)$

		$\hat{y} = \hat{f}(x)$	
		Pred P	Pred N
y	Actual P	True Positive	False Negative
	Actual N	False Positive	True Negative
		P*	N*

P = 5

N = 3

Confusion Matrix - Metrics



The proportion of observations that are correctly classified ?

$$\text{Accuracy} : (TN + TP) / (N + P)$$

The proportion of positives that are correctly identified as such ?

$$\text{True Pos Rate} : TP / P$$

(aka recall, sensitivity)

The proportion of negatives that are correctly identified as such

$$\text{True Neg Rate} : TN / N$$

(aka specificity)

$\hat{y} = \hat{f}(x)$

		$\hat{y} = \hat{f}(x)$	
		Pred P	Pred N
y	Actual P	True Positive	False Negative
	Actual N	False Positive	True Negative
		P*	N*

P = 5

N = 3

Confusion Matrix - Metrics



The proportion of observations that are
NOT correctly classified ?

Error rate :

The proportion of positives that are
NOT correctly identified as such ?

False Neg Rate :

(aka fall-out)

The proportion of negatives that are
NOT correctly identified as such

False Pos Rate :

(aka 1-specificity)

$\hat{y} = \hat{f}(x)$

		$\hat{y} = \hat{f}(x)$		
		Pred P	Pred N	
y	Actual P	True Positive	False Negative	P = 5
	Actual N	False Positive	True Negative	N = 3
		P*	N*	

Confusion Matrix - Metrics



The proportion of observations that are
NOT correctly classified ?

$$\text{Error rate : } (FN + FP) / (N + P)$$

The proportion of positives that are
NOT correctly identified as such ?

$$\text{False Neg Rate : } FN / P$$

(aka fall-out)

The proportion of negatives that are
NOT correctly identified as such

$$\text{False Pos Rate : } FP / N$$

(aka 1-specificity)

$\hat{y} = \hat{f}(x)$

		$\hat{y} = \hat{f}(x)$		
		Pred P	Pred N	
y	Actual P	True Positive	False Negative	P = 5
	Actual N	False Positive	True Negative	N = 3
		P*	N*	

Confusion Matrix - Metrics



The proportion of actual positives
in those identified as such ?

Precision : $TP / (FP + TP)$

The proportion of positives that are
correctly identified as such ?


Recall : TP / P
(aka TPR, sensitivity)

$\hat{y} = \hat{f}(x)$

	Pred P	Pred N	
Actual P	True Positive	False Negative	P = 5
Actual N	False Positive	True Negative	N = 3
	P*	N*	

y

Confusion Matrix - type I and type II error



y

$$\hat{y} = \hat{f}(x)$$

	Pred P	Pred N	
Actual P	good	Type II error	P = 5
Actual N	Type I error	good	N = 3
	P*	N*	

Using response probabilities



x_1, x_2, x_3, x_4	y	$P > 0.5$
...	1	1 0.95
...	0	0 0.21
...	0	1 0.55
...	1	0 0.43
...	1	1 0.77
...	1	0 0.44
...	0	0 0.15
...	1	1 0.81

$\hat{y} = \hat{f}(x)$

	Pred P	Pred N	
y	Actual P	True Positive	False Negative
	Actual N	False Positive	True Negative
	P^*	N^*	

$P = 5$
 $N = 3$

Cut-offs on probabilities



x_1, x_2, x_3, x_4 y					$P > 0.5$		$P > 0.6$		$P > 0.7$		$P > 0.8$		$P > 0.9$	
...	1	1	0.95	1	0.95	1	0.95	1	0.95	1	0.95
...	0	0	0.21	0	0.21	0	0.21	0	0.21	0	0.21
...	0	1	0.55	0	0.55	0	0.55	0	0.55	0	0.55
...	1	0	0.43	0	0.43	0	0.43	0	0.43	0	0.43
...	1	1	0.77	1	0.77	1	0.77	0	0.77	0	0.77
...	1	0	0.44	0	0.44	0	0.44	0	0.44	0	0.44
...	0	0	0.15	0	0.15	0	0.15	0	0.15	0	0.15
...	1	1	0.81	1	0.81	1	0.81	1	0.81	0	0.81

Those are sure ones ! \Rightarrow low FPR ! (high precision)
But we miss so many ones ! \Rightarrow low TPR ! (low recall)


Cut-offs on probabilities



x_1, x_2, x_3, x_4 y					$P > 0.5$	$P > 0.4$	$P > 0.3$	$P > 0.2$	$P > 0.1$
...	1	1	1	1	1	1
...	0	0	0	0	1	1
...	0	1	1	1	1	1
...	1	0	1	1	1	1
...	1	1	1	1	1	1
...	1	0	1	1	1	1
...	0	0	0	0	0	1
...	1	1	1	1	1	1

We have so many FP ! \Rightarrow high FPR ! (low precision)
But we capture all the ones ! \Rightarrow high TPR ! (high recall)


ROC curve (receiver operating characteristic)



For LogReg, think of it as sliding the purple/red line along the sigmoid function



Comparing classifiers based on their ROC curve



Possible metric : AUC
Area-under-curve

What is the “ideal” / “worst” classifier ?

