

# Logistic Regression & ROC

Miles Erickson

September 23, 2016

# DSI Standards: Logistic Regression

- Place logistic regression in the taxonomy of ML algorithms
- Explain the key similarities and differences between logistic and linear regression
- Implement and interpret a logistic regression model in scikit-learn
- Use the odds ratio to interpret the coefficients of linear regression
- Thoroughly explain ROC curves

# Binary Classification - Problem Motivation

Common examples of classification problems:

- detecting spam emails
- identifying potentially fraudulent transactions
- predicting the probability of default on a loan
- determining whether someone has a disease

# Binary Classification - Mathematical Description

- Classifier model: a mapping between your feature space and a finite set
- Binary classifier: maps  $X$  to  $y$  in  $\{0, 1\}$
- Example
  - ▶ Features: GPA  $[1.3, 4.0]$ , SAT score  $[600, 2400]$
  - ▶ Target: Not admitted  $\{0\}$ , Admitted  $\{1\}$
- Binary classifiers can generalize to multiple classes

# Your turn: plot a dataset

income	car
11	0
17	0
32	1
83	0
147	1
55	0
67	1
52	1
15	1

## Plot: income, car

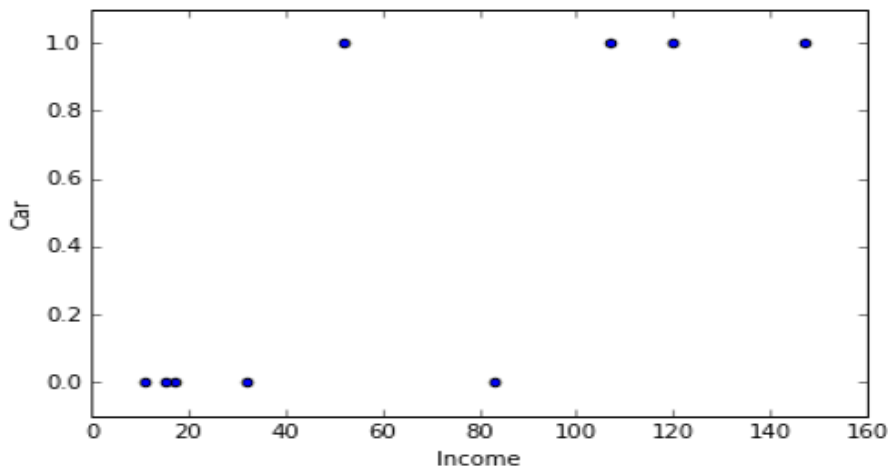


Figure 1: Income vs. Car

# Logistic Regression - Introduction

- Baseline model for classification (predicting probability)
- Estimates probability that an observation is in a given category based on the observation's features
- Regression step estimates the probability
- Classification step rounds the probability to 0 or 1