

NoSQL with MongoDB

Objectives

- Compare and contrast SQL and NoSQL
 - Explain use cases for NoSQL
 - Explain use cases for SQL
 - Discuss why we need both options
- Develop basic familiarity with MongoDB

SQL Review

- What is SQL?
 - Structured Query Language
- SQL allows us to interact with Relational Database Management Systems (RDBMS)
 - Model **relations** in the data
 - Stores data about one object across multiple tables
 - {student_id, course_id}, {student_id, student_name} -> join on student_id and filter on course_id for names of all students in a course
 - Query data and relations efficiently
 - Maintain data consistency and integrity

SQL Review

- Tables, Columns (fields), Rows (records)
 - Each column is of a certain data type
 - Each row is an entry in the table
 - It holds values for each one of the columns
- Tables are specified by a **schema** that defines the structure of the data
 - We specify the schema ahead of time

NoSQL

- NoSQL
 - “Non SQL”, “Non relational”, “Not Only SQL”
 - “*Not Only SQL*” because stacks often use both NoSQL and SQL for different purposes
- Many NoSQL databases (though not all) are **document-oriented**
 - Each object (row/document) is stored in one place
 - Each object can be completely different from all others
- There is **no schema** (actually, there is, but it's *implicit*)
 - Each document can have or not have whatever fields are appropriate for that particular document

Roughly Analogous:

- | | | |
|--------------|--------------|---------------|
| • NoSQL | • SQL | • Spreadsheet |
| • MongoDB | • PostgreSQL | • Excel |
| • Database | • Database | • File |
| • Collection | • Table | • Sheet |
| • Field | • Column | • Column |
| • Document | • Row | • Row |

<code>SELECT id, user_id, status FROM people</code>	<code>db.people.find({ }, { user_id: 1, status: 1 })</code>
<code>SELECT user_id, status FROM people</code>	<code>db.people.find({ }, { user_id: 1, status: 1, _id: 0 })</code>
<code>SELECT * FROM people WHERE status = "A"</code>	<code>db.people.find({ status: "A" })</code>
<code>SELECT user_id, status FROM people WHERE status = "A"</code>	<code>db.people.find({ status: "A" }, { user_id: 1, status: 1, _id: 0 })</code>
<code>SELECT * FROM people WHERE status != "A"</code>	<code>db.people.find({ status: { \$ne: "A" } })</code>
<code>SELECT * FROM people WHERE status = "A" AND age = 50</code>	<code>db.people.find({ status: "A", age: 50 })</code>
<code>SELECT * FROM people WHERE status = "A" OR age = 50</code>	<code>db.people.find({ \$or: [{ status: "A" }, { age: 50 }] })</code>
<code>SELECT * FROM people WHERE age > 25</code>	<code>db.people.find({ age: { \$gt: 25 } })</code>
<code>SELECT * FROM people WHERE age < 25</code>	<code>db.people.find({ age: { \$lt: 25 } })</code>
<code>SELECT * FROM people WHERE age > 25 AND age <= 50</code>	<code>db.people.find({ age: { \$gt: 25, \$lte: 50 } })</code>
<code>SELECT * FROM people WHERE user_id like "%bc%"</code>	<code>db.people.find({ user_id: /bc/ })</code>

COMMAND HELPERS

The following table lists some common help commands which are available in the mongo shell:

<code>help</code>	Show help.
<code>db.help()</code>	Show help for database methods.
<code>db.<collection>.help()</code>	Show help on collection methods. The <collection> can be the name of an existing collection or a non-existing collection.
<code>show dbs</code>	Print a list of all databases on the server.
<code>use <db></code>	Switch current database to <db> . The mongo shell variable <code>db</code> is set to the current database.
<code>show collections</code>	Print a list of all collections for the current database.
<code>show users</code>	Print a list of users for the current database.
<code>show roles</code>	Print a list of all roles, both user-defined and built-in, for the current database.
<code>show profile</code>	Print the five most recent operations that took 1 millisecond or more.
<code>show databases</code>	Print a list of all available databases.

 FOR MORE INFORMATION

<https://docs.mongodb.com/manual/reference/mongo-shell/>

CRUD METHODS

Queries typically take the following form:

```
db.<collection>.<method>( <filter>, <options>)
```

db refers to the current database. **<collection>** is the name of the target collection for your query. For **<method>**, substitute the desired query method, (examples below). Each method has its own **<options>** for what it will do with the matching document(s).

<code>db.collection.insertOne()</code>	Inserts a document into a collection.
<code>db.collection.insertMany()</code>	Inserts multiple documents into a collection.
<code>db.collection.find()</code>	Selects documents in a collection based on the filter and returns a cursor to the selected documents.
<code>db.collection.updateOne()</code>	Updates a single document within the collection based on the filter.
<code>db.collection.updateMany()</code>	Updates all documents within the collection that match the filter.
<code>db.collection.replaceOne()</code>	Replaces a single document within the collection based on the filter.
<code>db.collection.deleteOne()</code>	Removes a single document from a collection based on the filter.
<code>db.collection.deleteMany()</code>	Removes all documents that match the filter from a collection.

 **FOR MORE INFORMATION**

<https://docs.mongodb.com/manual/crud/>

Mongo Commands - Creating a database and collection

- To create a database and collection, all we have to do is switch into the database (using `use`), and then insert a record into it.

```
use class_db # Creates class_db if it doesn't exist, and  
             # otherwise logs into it.  
db.disney_chars.insert{name: 'Jasmine', age: 22}
```

Note: Mongo will create any databases that don't exist, as well as any collections when you try to insert into them.

QUERY FILTER PARAMETERS AND WHAT THEY MATCH

MongoDB uses a key-value structure to create query filter parameters, which you can use in the `mongo` shell or with a driver in a client application. For example, the following query finds all documents in the collection named `inventory` in which the `qty` field contains a value greater than 10:

```
db.inventory.find({ "qty" : { $gt: 10 } })
```

Queries take documents as query filter parameters, shown as examples below. Multiple filter parameters can be included in one document, separated by commas.

<code>{a: 10}</code>	Docs where a is 10 or an array containing the value 10 .
<code>{a: 10, b: "hello"}</code>	Docs where a is 10 and b is "hello" .
<code>{a: {\$gt: 10}}</code>	Docs where a is greater than 10 . <i>Also available:</i> \$lt (<), \$gte (>=), \$lte (<=), and \$ne (!=).
<code>{a: {\$in: [10, "hello"]}}</code>	Docs where a is either 10 or "hello" .
<code>{a: {\$all: [10, "hello"]}}</code>	Docs where a is an array containing both 10 and "hello" .

Mongo Commands - The Mongo “SELECT” and “FROM”

- Using `.find` is the Mongo equivalent of SQL's SELECT (kind of - below we are selecting all **fields**)

```
db.disney_chars.find() # Returns all documents.
```

```
db.disney_chars.find().limit(5) # Returns first 5 documents.
```

- Note that the collection comes after the `db`, and this is the equivalent of using SQL's FROM and putting a table name after it

Mongo Commands - The Mongo “WHERE”

- We can specify many ways of finding observations in Mongo:

```
db.disney_chars.find({name: 'Mulan'}) # Find by single field
db.disney_chars.find({age: 26}) # Find by single field
db.disney_chars.find({age: {$exists : true }}) # Find by
                                                # presence
                                                # of field
db.disney_chars.find({friends: 'BoBo'}) # Find by whether
                                           # value in array
```

Mongo Commands - The Mongo “SELECT” II

- To specify only certain fields to keep, we have to pass those in as a second argument:

```
# Select all documents with name 'Mulan', and only return  
# back the name field (_id is returned by default, so  
# it'll be returned as well)
```

```
db.disney_chars.find({name: 'Mulan'}, {name: true})
```

```
# Select all documents, returning their friends field  
# (_id is returned by default, so it'll be returned  
# as well)
```

```
db.disney_chars.find({}, {friends: true})
```

Mongo Commands - The Mongo “SELECT” III

*# Select all documents, returning ****only**** their friends
field, and ****not**** _id with it.*

```
db.disney_chars.find({}, {friends: true, _id: false})
```

*# Select all documents, returning ****only**** their friends
field, and ****not**** _id with it, but ****only**** if they
have a friends field.*

```
db.disney_chars.find({friends: {$exists: true}},  
                    {friends: true, _id: false})
```

Mongo Commands - The Mongo “SELECT” IV

```
# Find those documents without name 'Mulan'  
db.disney_chars.find({name : {$ne: 'Mulan'}})
```

- We have other operators for the other equality-like comparisons:

Operator Syntax	Meaning
\$eq	Equals
\$gt	Greater than
\$gte	Greater than or equal to
\$lt	Less than
\$lte	Less than or equal to
\$ne	Not equal to
\$in	In (for arrays)
\$nin	Not in (for arrays)

FIELD UPDATE OPERATORS

<code>{\$inc: {a: 2}}</code>	Increment a by 2 .
<code>{\$set: {a: 5}}</code>	Set a to the value 5 .
<code>{\$unset: {a: 1}}</code>	Delete the a key.
<code>{\$max: {a: 10}}</code>	Set a to the greater value, either current or 10 . If a does not exist, set a to 10 .
<code>{\$min: {a: -10}}</code>	Set a to the lowest value, either current or -10 . If a does not exist, set a to -10 .
<code>{\$mul: {a: 2}}</code>	Set a to the product of the current value of a and 2 . If a does not exist set a to 0 .
<code>{\$rename: {a: "b"}}</code>	Rename field a to b .
<code>{\$setOnInsert: {a: 1}}, {upsert: true}</code>	Set field a to 1 in case of upsert operation.

<code>{ \$push: { a: { \$each: [50, 60, 70], \$position: 0 } } }</code>	Insert 50 , 60 , and 70 starting at position 0 of the array a . <i>\$position can only be used with the \$each modifier.</i>
<code>{ \$addToSet: { a: 1 } }</code>	Append the value 1 to the array a (if the value doesn't already exist).
<code>{ \$addToSet: { a: { \$each: [1, 2] } } }</code>	Append both 1 and 2 to the array a (if they don't already exist).
<code>{ \$pop: { a: 1 } }</code>	Remove the last element from the array a .
<code>{ \$pop: { a: -1 } }</code>	Remove the first element from the array a .
<code>{ \$pull: { a: { \$gt: 5 } } }</code>	Remove all values greater than 5 from the array a .
<code>{ \$pullAll: { a: [5, 6] } }</code>	Remove multiple occurrences of 5 or 6 from the array a .

 **FOR MORE INFORMATION**

<http://docs.mongodb.org/manual/reference/operator/update/>

Mongo Updates

- To update a record, we can do the following:

```
# Update the **first instance** with name Mulan,  
# setting their age to 29.
```

```
db.disney_chars.update({name: "Mulan"}, {$set : {age: 29}})
```

```
# Update all instances.
```

```
db.disney_chars.update({name: "Mulan"}, {$set : {age :29}},  
                        {multi: true})
```

```
# Update all instances found, or insert a document  
# if not found.
```

```
db.disney_chars.update({name: "Mulan"}, {$set : {age :29}},  
                        {multi: true, upsert: true})
```

Mongo Updates Breakdown

**What to update
(like WHERE)**

**HOW TO
update**

- Let's break down that last result a bit...

*# Update all instances found or insert a document
if not found.*

```
db.disney_chars.update({name: "Mulan"}, {$set : {age :29}},  
                      {multi: true, upsert: true})
```

**Update all documents that
match (as opposed to
just the first one)**

**Insert a document like
this if none exist**

AGGREGATION FRAMEWORK:

The aggregation pipeline, part of the MongoDB query language, is a framework for data aggregation modeled on the concept of data processing pipelines. Documents enter a multi-stage pipeline that transforms the documents into aggregated results. Pipeline stages appear in an array. Documents pass through the stages in sequence. Structure an aggregation pipeline using the following syntax:

```
db.<collection>.aggregate( [ { <stage1> }, { <stage2> } ... ] )
```

COMMON AGGREGATION FRAMEWORK STAGES

<code>{ \$match: { a: 10 } }</code>	Passes only documents where a is 10 .	Similar to find()
<code>{ \$project: { a: 1, _id: 0 } }</code>	Reshapes each document to include only field a , removing others.	Similar to find() projection
<code>{ \$project: { new_a: "\$a" } }</code>	Reshapes each document to include only _id and the new field new_a with the value of a .	<code>{a:1} => {new_a:1}</code>
<code>{ \$project: { a: { \$add: ["\$a", "\$b"] } } }</code>	Reshapes each document to include only _id and field a , set to the sum of a and b .	<code>{a:1, b:10} => {a: 11}</code>

Mongo Aggregations I

- Aggregations in Mongo are a little more involved (and much less pretty) than in SQL:

```
db.disney_chars.aggregate( [ { $group :  
  {  
    _id: "$name",  
    total: { $sum: "$age"}  
  }  
} ] )
```

```
db.disney_chars.aggregate([  
  { $match: { name: { $in : ["Mulan", "Jasmine"] } } },  
  { $group: { _id: "$name", total: { $sum: "$age" } } },  
  { $sort: { total: 1 } }  
])
```

Mongo Aggregations I Breakdown

what to grab
(SQL
WHERE)

GROUP BY and
AGGREGATION
function

- Let's break down that last result a bit...

```
db.disney_chars.aggregate([
  { $match: { name: { $in : ["Mulan", "Jasmine"] } } },
  { $group: { _id: "$name", total: { $sum: "$age" } } },
  { $sort: { total: 1 } }
])
```

whether to
SORT

Mongo Aggregations II

- Other aggregation functions we have:

Aggregation command	Purpose
aggregate	Performs aggregation tasks
count	Counts the number of items meeting some criteria
distinct	Displays distinct values for a specified field(s)
group	Groups observations in some way

Connection with PyMongo I

- To connect from within Python, we still have to **start a server**, and then in code we can do the following:

```
from pymongo import MongoClient

client = MongoClient() # Instantiate a client that will
                        # be connected to Mongo.
database = client['class_db'] # Create a variable holding
                               # a reference to the db you
                               # want to connect to.
collection = database['disney_chars'] # Create a variable
                                       # to hold the
                                       # collection you
                                       # want to connect to.
```

- Note that we access the database and/or collection by effectively using the *name* as a key (it's like indexing into a dictionary)