

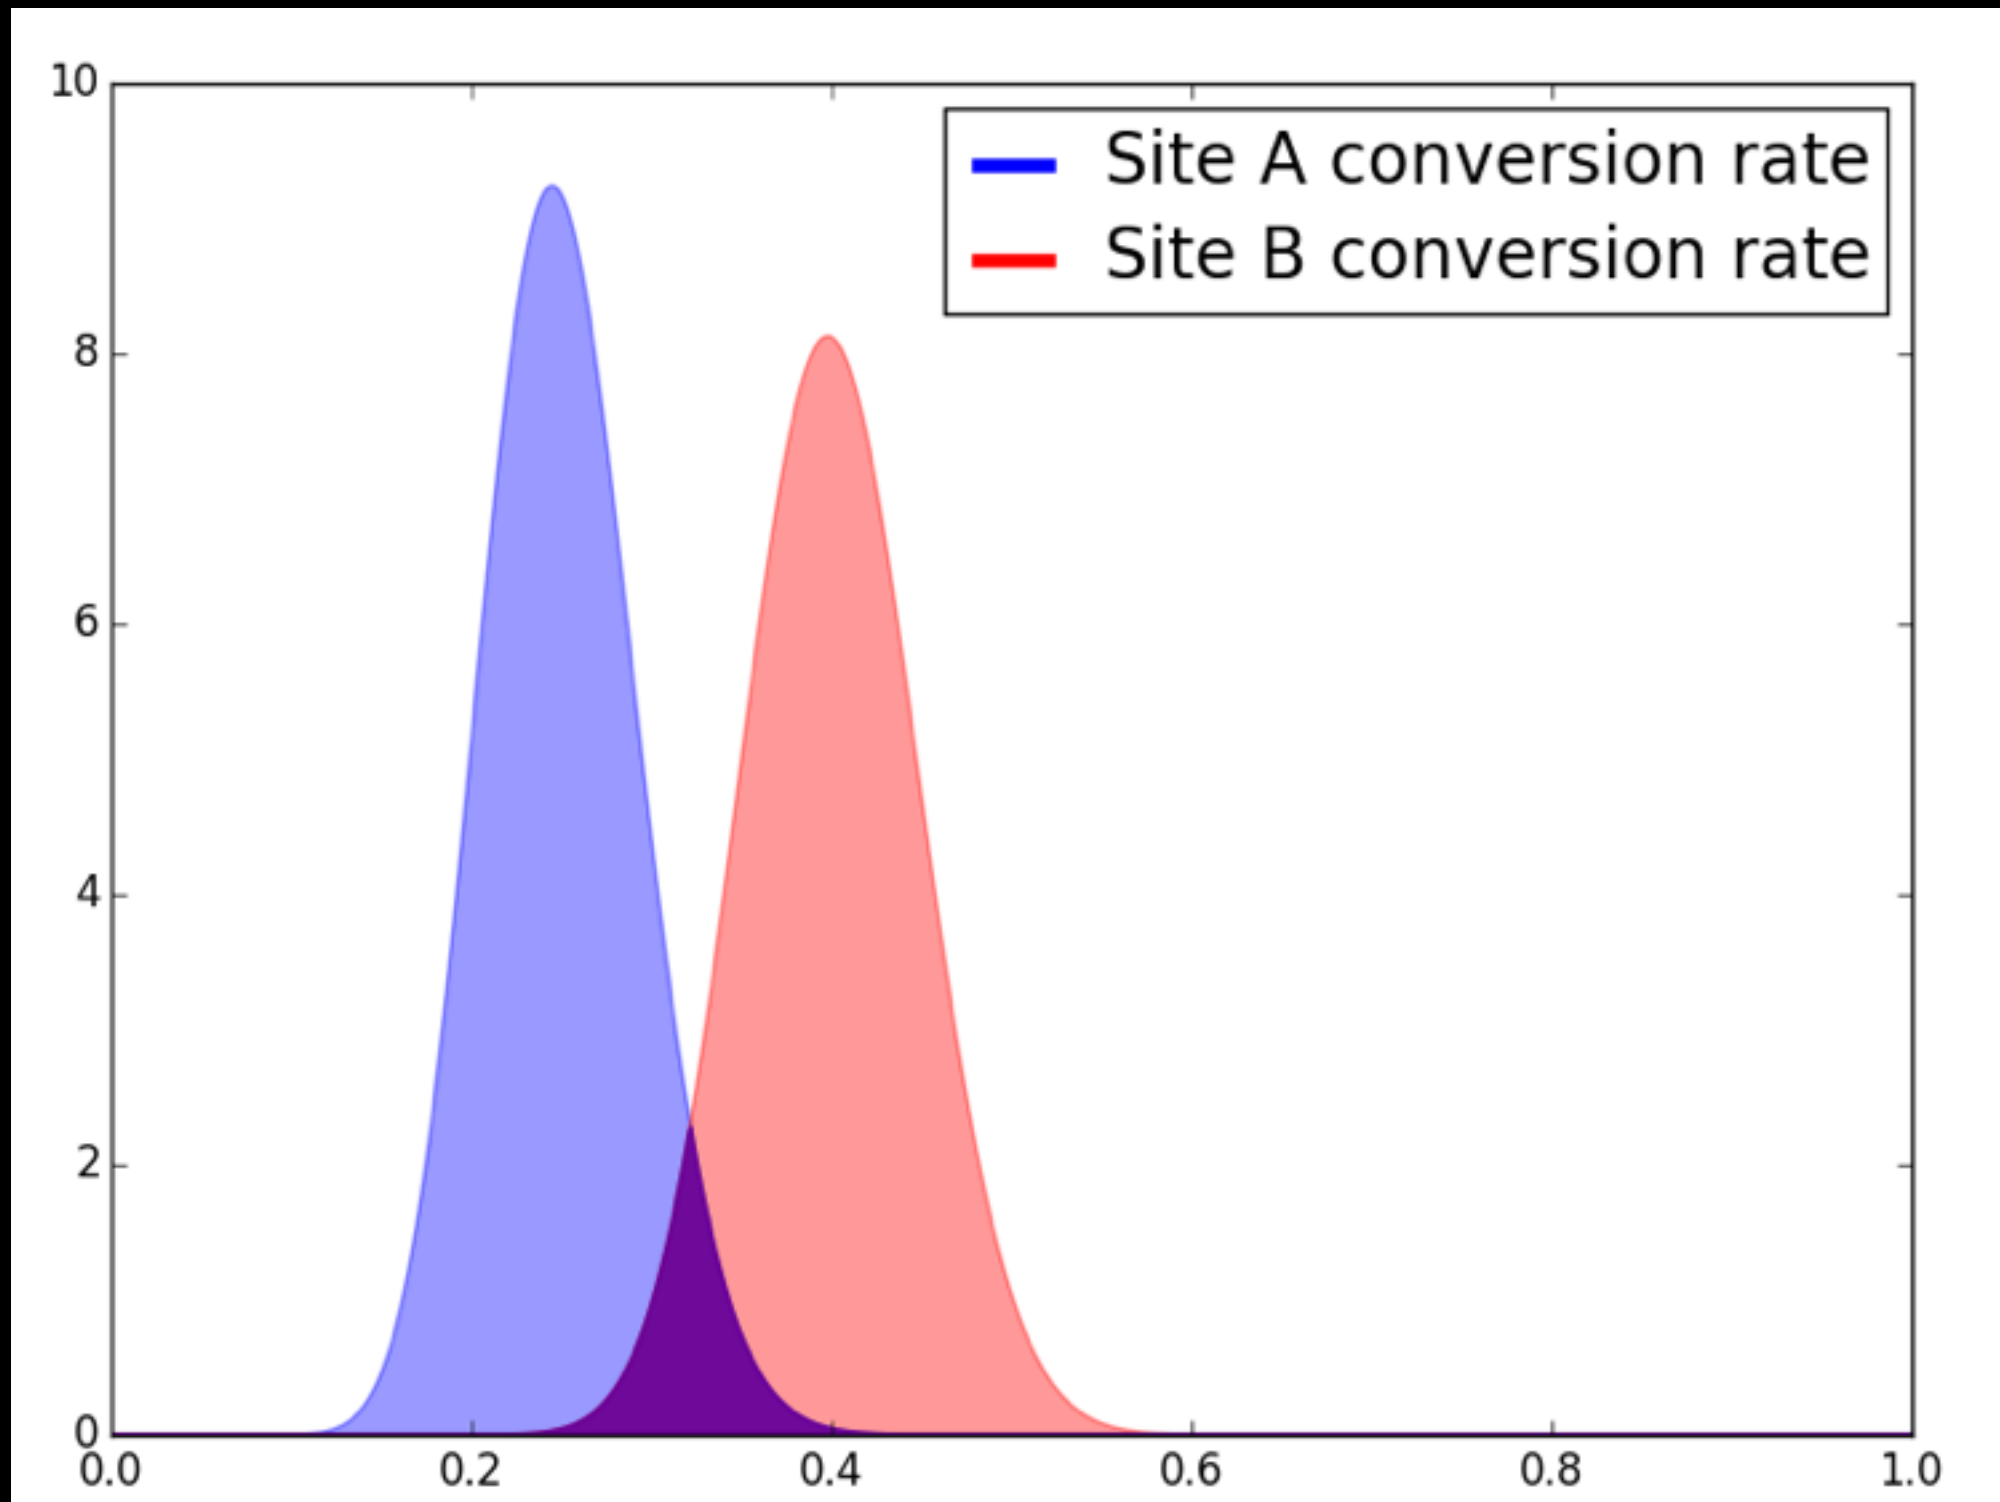
Bayesian A/B Testing

Beta Distribution and Multi-Arm Bandit

A/B Testing

- **Conversion Rate:** percent of users who view a site who perform the desired action.
 - e.g. Click on an ad, called *Click-through rate (CTR)*; or sign up for an account
 - value between 0 and 1
- We use the data to determine the distribution of the conversion rate for each version of the site.
- Determine probability that site A is better than site B

Bayesian A/B Testing



Bayes Theorem

Diagram illustrating Bayes Theorem with labels for its components:

- Posterior Probability** (labeled above $\Pr(\theta|y)$)
- Likelihood of Observations** (labeled above $\Pr(y|\theta)$)
- Prior Probability** (labeled above $\Pr(\theta)$)
- Normalizing Constant** (labeled below $\Pr(y)$)



$$\Pr(\theta|y) = \frac{\Pr(y|\theta)\Pr(\theta)}{\Pr(y)}$$

- **prior:** initial belief
- **likelihood:** likelihood of data given outcome
- **posterior:** updated belief

Bayes Theorem

$$\text{posterior} \propto \text{prior} \times \text{likelihood}$$

The Distribution



1 = conversion

0 = non conversion

Beta Distribution

$$\frac{p^{\alpha-1} (1-p)^{\beta-1}}{B(\alpha, \beta)}$$

- p : conversion rate (between 0 and 1)
- α, β : shape parameters
 - $\alpha = 1 + \text{number of conversions}$
 - $\beta = 1 + \text{number of non conversions}$
- Beta Function (B) is a normalizing constant
- $\alpha = \beta = 1$ gives the *uniform distribution*

Binomial (Likelihood)

$$\binom{n}{k} p^k (1 - p)^{n-k}$$

- p : conversion rate (between 0 and 1)
- n : number of visitors
- k : number of conversions

Conjugate Priors

posterior \propto prior \times likelihood

beta \propto beta \times binomial

THE MATH:


$$\begin{aligned}\text{posterior} &\propto \text{prior} \times \text{likelihood} \\ &= \frac{p^{\alpha-1}(1-p)^{\beta-1}}{B(a, b)} \times \binom{n}{k} p^k (1-p)^{n-k} \\ &\propto p^{\alpha-1}(1-p)^{\beta-1} \times p^k (1-p)^{n-k} \\ &\propto p^{\alpha+k-1} (1-p)^{\beta+n-k-1}\end{aligned}$$

The result is a Beta Distribution with these shape parameters:

$$\alpha + k \text{ and } \beta + n - k$$

A/B Testing

- We want to know if this is true:
conversion rate of site A > conversion rate of site B
- We can also answer if this is true:
conversion rate of site A > conversion rate of site B + 5%



Method:

- Sample a large number from both distributions
- Count the percent of times site A wins

The code

```
num_samples = 10000

A = np.random.beta(1 + num_clicks_A,
                   1 + num_views_A - num_clicks_A,
                   size=num_samples)

B = np.random.beta(1 + num_clicks_B,
                   1 + num_views_B - num_clicks_B,
                   size=num_samples)

### The probability that A wins:
print np.sum(A > B) / float(num_samples)

### The probability that A > B + 0.5%:
print np.sum(A > (B + 0.05)) / float(num_samples)
```

Multi-Arm Bandit


Smarter A/B Testing

Traditional A/B Testing

- First: pure *exploration*, in which you assign equal numbers of users to Group A and Group B
- Second: pure *exploitation*, in which you stop the experiment and send all your users to the more successful version of your site.

Multi-Arm Bandit

- Start *exploiting* the likely best solution *before* you're done *exploring*
- versions of the site
= **bandits** (slot machines)
- How to pick which one to play....



Epsilon-Greedy Algorithm

- Explore with probability epsilon (often 10%)
- Exploit the rest of the time (play the bandit with the best performance so far)

Regret

- To evaluate a multi-arm bandit algorithm, *minimize regret*
- **Regret** is a measure of how often you chose a suboptimal bandit

$$\begin{aligned}\text{regret} &= \sum_{i=1}^k (p_{\text{opt}} - p_i) \\ &= k \cdot p_{\text{opt}} - \sum_{i=1}^k p_i\end{aligned}$$

Softmax Algorithm

- Choose the bandit randomly in proportion to its estimated value:

$$\frac{e^{p_A/\tau}}{e^{p_A/\tau} + e^{p_B/\tau} + e^{p_C/\tau}}$$

UCB1 Algorithm (Upper Confidence Bound)

- Choose the bandit where this value is the largest

$$p_A + \sqrt{\frac{2 \log N}{n_A}}$$

where n_A = number of times bandit A has been played
and N = total number of times any bandit has been played

Bayesian Bandit Algorithm

- Model each of the bandits with a beta distribution with shape parameters:
 $\alpha = 1 + \text{number of times bandit has won}$
 $\beta = 1 + \text{number of times bandit has lost}$
- Take a random sample from each bandit's distribution and choose the bandit with the highest value.