

# Natural Language Processing

(NLP)

# Objectives for the morning

- Explain the steps of turning raw text data into useful features
- Explain the difference between stemming and lemmatization
- Compute the TF-IDF of documents

# Motivation

- You run Google News and you want to group news articles by topic
- You run a legal tech firm and you need to sift through 1,000,000 pages of legal documents to find the relevant ones

# Text Featurization Pipeline

- Tokenization
- Lower case conversion
- Punctuation removal
- Stop words removal
- Stemming / Lemmatization
- Bag of words / N-grams

# Terminology

- Corpus:
  - A dataset of text
  - E.g. newspaper articles, tweets
- Document:
  - A single entry from our corpus
  - E.g. article, sentence, tweet
- Vocabulary:
  - All the words that appear in the corpus
- Token:
  - An entity
  - E.g. a word

# Sample Sentence

“Students are learning from other students”

# Tokenization

Take the document and split it into a list of tokens

“Students are learning from other students”



[“Students”, “are”, “learning”, “from”, “other”, “students”]

# Lower case conversion

["Students", "are", "learning", "from", "other", "students"]



["students", "are", "learning", "from", "other", "students"]



# Stop word removal

- Remove the words we ignore in our analysis that are too common to be useful (since they are too common, they do not provide predictive power)
- Sklearn and nltk have standard list of stop words

["students", "are", "learning", "from", "other", "students"]



["students", "learning", "other", "students"]

# Stemming/Lemmatization

- Stemming:

- Removes morphical suffixes
  - speaking -> speak
  - actually -> actual
- Does it without context
- Sometimes weird
- Pretty fast

- Lemmatization:

- Replace words with canonical form
  - better -> good
  - speaking -> speak
- Uses hash tables
- Slower

# Stemming/Lemmatization

["students", "learning", "other", "students"]



["student", "learn", "other", "student"]

# Corpus with 3 documents

Doc 1: “Students are learning from other students”

- [“student”, “learn”, “other”, “student”]

Doc 2: “I am teaching at Galvanize”

- [“teach”, “galvanize”]

Doc 3: “There are students learning at Galvanize”

- [“student”, “learn”, “galvanize”]

# Bag of words

A document represented as a vector of word counts is called “bag of words”

Vector for our corpus: (galvanize, learn, other, student, teach)

document	galvanize	learn	other	student	teach
Doc 1					
Doc 2					
Doc 3					

# Bag of words

A document represented as a vector of word counts is called “bag of words”

Vector for our corpus: (galvanize, learn, other, student, teach)

document	galvanize	learn	other	student	teach
Doc 1	0	1	1	2	0
Doc 2	1	0	0	0	1
Doc 3	1	1	0	1	0

# Bag of words

- Issues with bag of words:
  - Word counts
    - Counts emphasize results from longer documents
  - Every word has equal weighting
    - “other” and “student” have different predictive power.

# Term Frequency - Inverse Document Frequency

- Measures the relevance of a word
- Adjusts word count by document length and how often it shows up in the corpus



# Term Frequency

Normalize counts within a document to frequency

$$tf(t, d) = \frac{\text{total count of term } t \text{ in document } d}{\text{total count of all terms in document } d}$$

document	galvanize	learn	other	student	teach
Doc 1					
Doc 2					
Doc 3					

# Term Frequency

Normalize counts within a document to frequency

$$tf(t, d) = \frac{\text{total count of term } t \text{ in document } d}{\text{total count of all terms in document } d}$$

document	galvanize	learn	other	student	teach
Doc 1	0	$\frac{1}{4} = 0.25$	$\frac{1}{4} = 0.25$	$\frac{2}{4} = 0.5$	0
Doc 2	$\frac{1}{2} = 0.5$	0	0	0	$\frac{1}{2} = 0.5$
Doc 3	$\frac{1}{3} = 0.33$	$\frac{1}{3} = 0.33$	0	$\frac{1}{3} = 0.33$	0

# Term Frequency

Issues with term frequency:

- Words found only in one document should have highest weighting
- Words found in every document should have lowest weighting

# Inverse Document Frequency

$$idf(t, D) = \log \frac{\text{total number of document in corpus } D}{\text{count of document containing term } t}$$

document	galvanize	learn	other	student	teach
Doc 1		X	X	X	
Doc 2	X				X
Doc 3	X	X		X	
$idf(t, D)$	$\log(3/2)$	$\log(3/2)$	$\log(3/1)$	$\log(3/2)$	$\log(3/1)$

# TF-IDF

$$tfidf(t, d, D) = tf(t, d) \cdot idf(t, D)$$

document	galvanize	learn	other	student	teach
Doc 1	0	$0.25 \times \log(3/2)$ = 0.101	$0.25 \times \log(3/1)$ = 0.275	$0.5 \times \log(3/2)$ = 0.203	0
Doc 2	$0.5 \times \log(3/2)$ = 0.203	0	0	0	$0.5 \times \log(3/1)$ = .549
Doc 3	$0.33 \times \log(3/2)$ = 0.135	$0.33 \times \log(3/2)$ = 0.135	0	$0.33 \times \log(3/2)$ = 0.135	0

# Comparing TF-IDF vectors of documents

Cosine similarity:  $similarity = \cos\theta = \frac{A \cdot B}{\|A\| \|B\|}$

- Doc 1 vs. Doc 2:

- ["student", "learn", "other", "student"] vs ["teach", "galvanize"]

$(0, 0.101, 0.275, 0.203, 0)$  vs.  $(0.203, 0, 0, 0, 0.275)$

$$similarity = \frac{0}{0.36 \times 0.34} = 0$$

- Doc 1 vs. Doc 3:

- ["student", "learn", "other", "student"] vs ["student", "learn", "galvanize"]

$(0, 0.101, 0.275, 0.203, 0)$  vs.  $(0.135, 0.135, 0, 0.135, 0)$

$$similarity = \frac{0.041}{0.36 \times 0.23} = 0.34$$

# N-Grams & Skip Grams

- Sliding window on text
- Handle phrases or words that go together (to capture the relationship between the consecutive words)

“The rain in Spain falls mainly on the plain”

Bi-gram: [“the rain”, “rain in”, “in spain”, “spain falls”, ... ]

1-skip-bi-gram: [“the in”, “rain spain”, “in falls”, ... ]

# Morning Assignment