

Recurrent Neural Networks

Objectives

- Introduce RNNs
 - Why are RNNs good for sequential analysis?
- Tangent: Time Series (common application for RNNs)
- State important differences between MLPs and RNNs
- State a common shortcoming of “vanilla” RNNs
- Explain, at a high-level, what an LSTM is (assignment)
- Predict stock price using an LSTM (assignment)

Recurrent Neural Networks

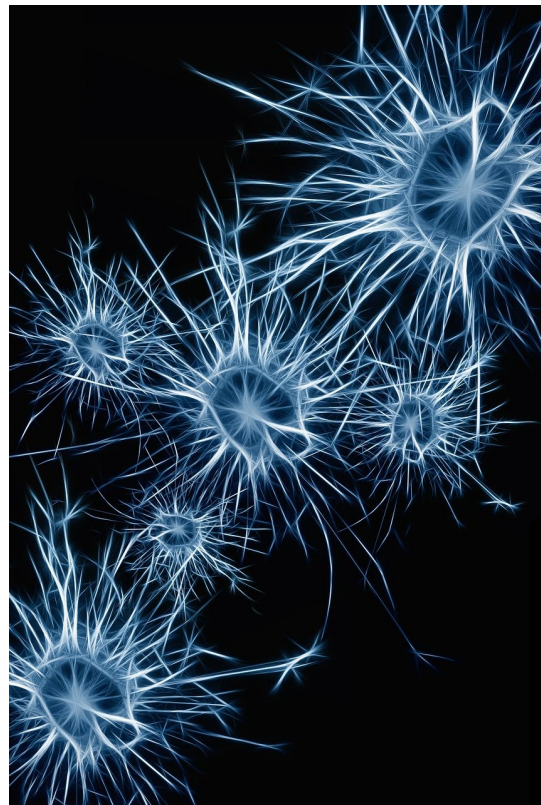
Models based on the connection of simple computational units, loosely analogous to neurons in the human brain.

What distinguishes neural networks is that connections between neurons can form a [directed cycle](#). This gives a network the ability to maintain a state based on previous inputs. So it can model temporal, sequential behavior.

RNN use cases:

Pattern recognition: [handwriting](#), [captioning images](#)

Sequential data: [speech recognition](#), [stock price prediction](#), and [generating text](#) and [news stories](#)



Time Series Tangent

Time Series

A **time series** is a series of data points indexed (or listed or graphed) in time order. Most commonly, a time series is a sequence taken at successive *equally spaced* points in time. The data are ordered (are not independent of each other).

examples: heights of ocean tides, the daily closing value of the DJIA

Examples of time series:

`time-series-trends-and-seasonality.ipynb`

For your reference:

Classical time series analysis: [ARIMA](#)

`time-series-lecture.ipynb`

DSI repo: [dsi-time-series](#)

Time Series - using models you know

X					y
Value @ t-4	Value @ t-3	Value @ t-2	Value @ t-1	Value @ t	Target Value at t+1
1	2	3	4	5	6
2	3	4	5	6	7
3	4	5	6	7	8

But be careful about the train-test split. It's not fair to train on values in the future to predict the past.

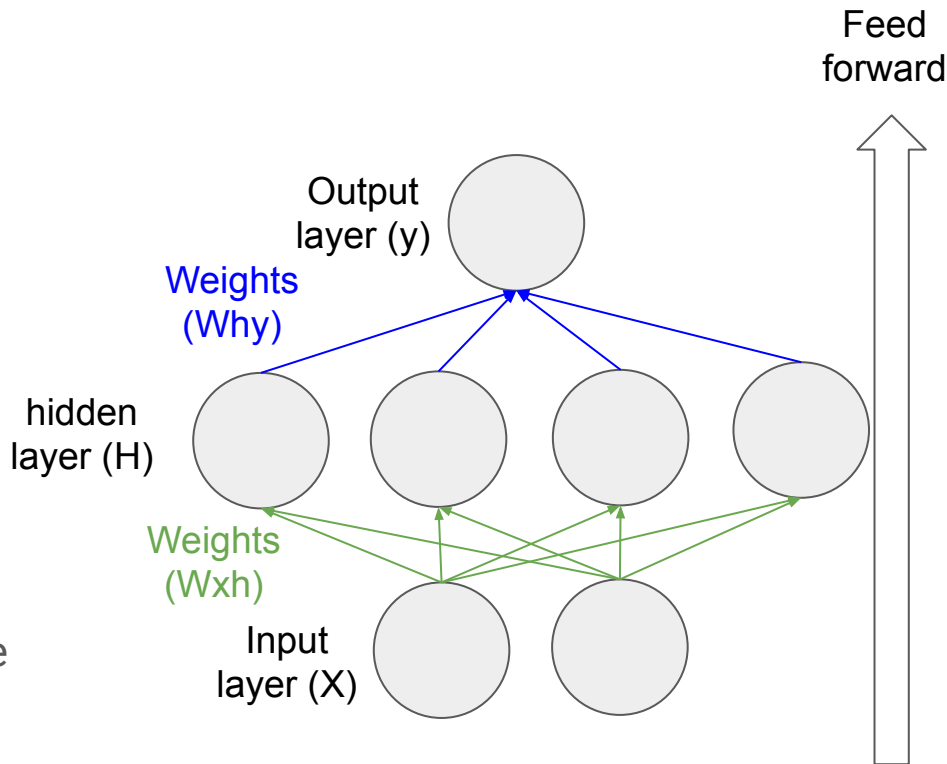
Sklearn [TimeSeriesSplit](#)

Demo: `time-series-with-conventional-model.ipynb`

RNNs can be very good at Time Series
so back to RNNs

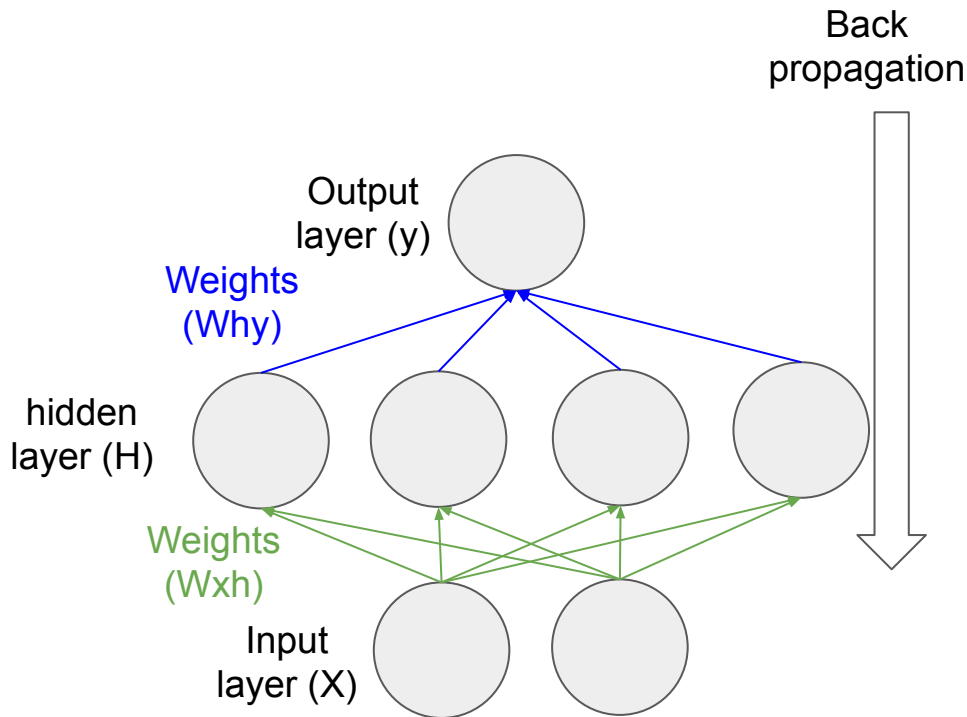
Multi-layer perceptron

- Maps input data to corresponding outputs. Calculation *feeds forward* through the network.
- Nodes are arranged in layers.
- Nodes have values for any input given the sum of inputs to the node and an *activation function* that transforms the sum to a non-linear output.
- The “learning” in the network is held by the trained values of the connections (the weights). These weights start with random values.



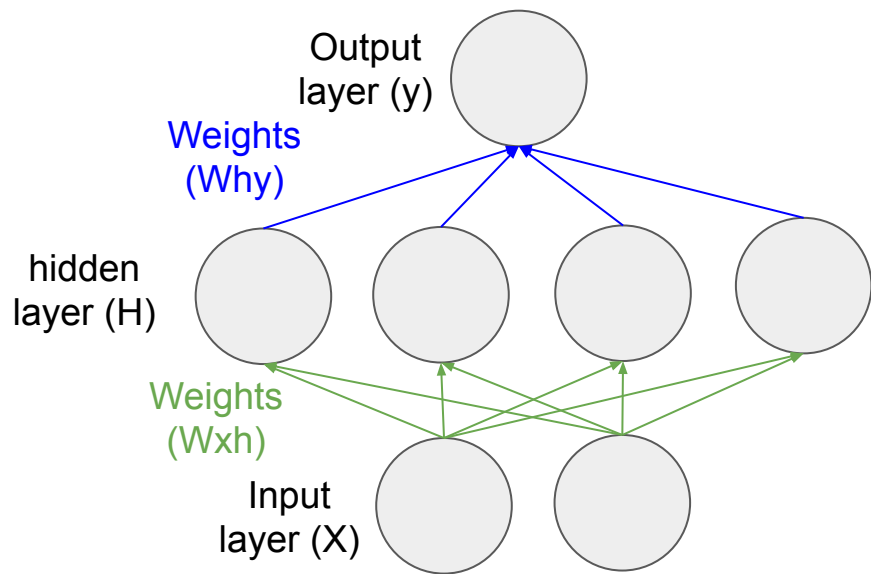
Review: Multilayer perceptron - learning the weights

- After feed forward yields a predicted output (\mathbf{y}_p), its difference (loss) is calculated from the real output (\mathbf{y}).
- Back propagation estimates how much the loss varies due to each weight in the network (the gradient).
- Gradient descent uses the gradient and learning rate to tweak each of the weights so that the network predicts a little better next time.
- When the total loss reaches an acceptable level, stop. Now it's trained and ready to predict.

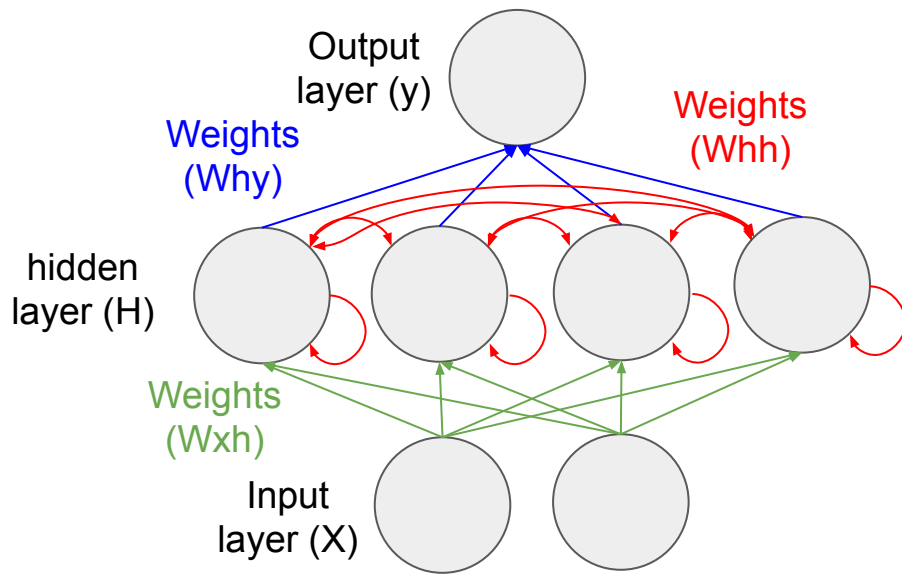


Comparing the simplest version of these neural nets

Vanilla MLP



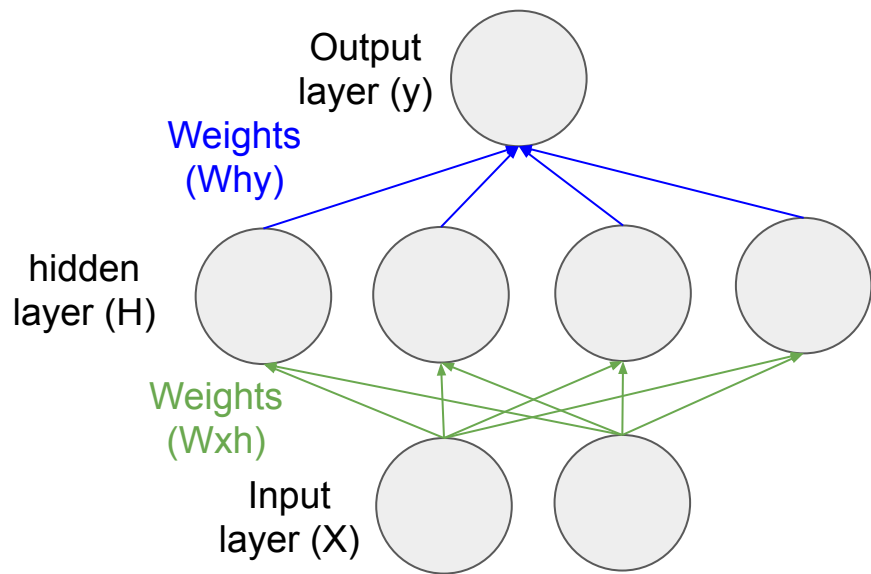
Vanilla RNN



A double arrow indicates a weight in each direction (2 weights).

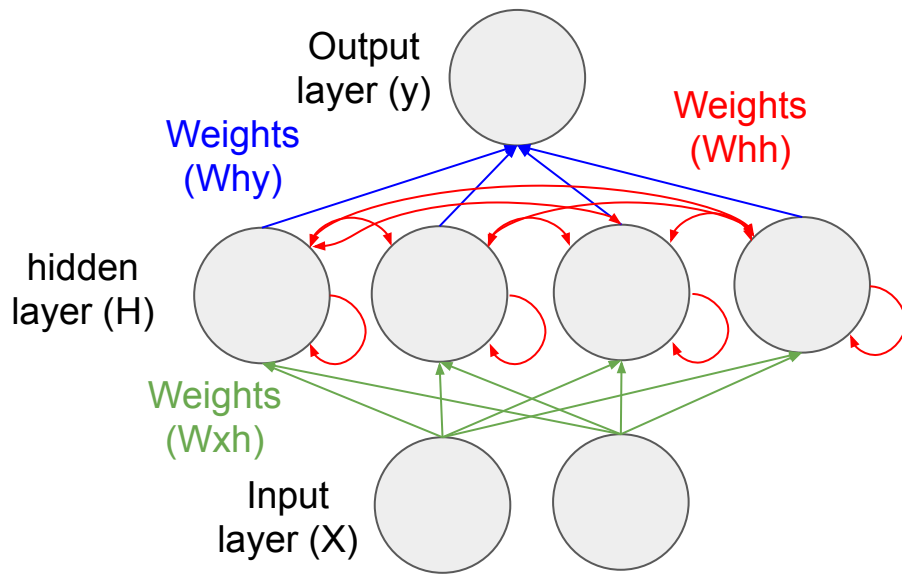
So how many weights in each architecture?

Vanilla MLP



$W_{hy} =$
 $W_{xh} =$

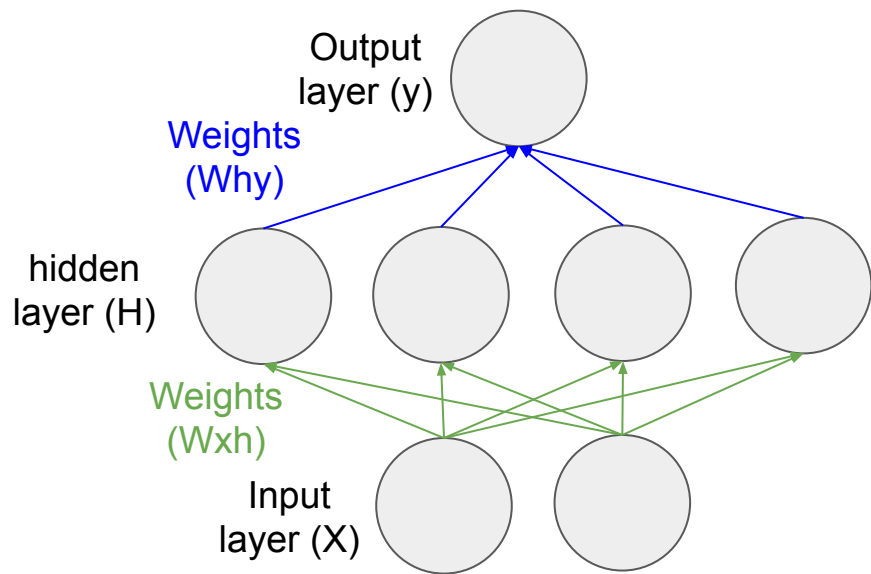
Vanilla RNN



$W_{hy} =$
 $W_{hh} =$
 $W_{xh} =$

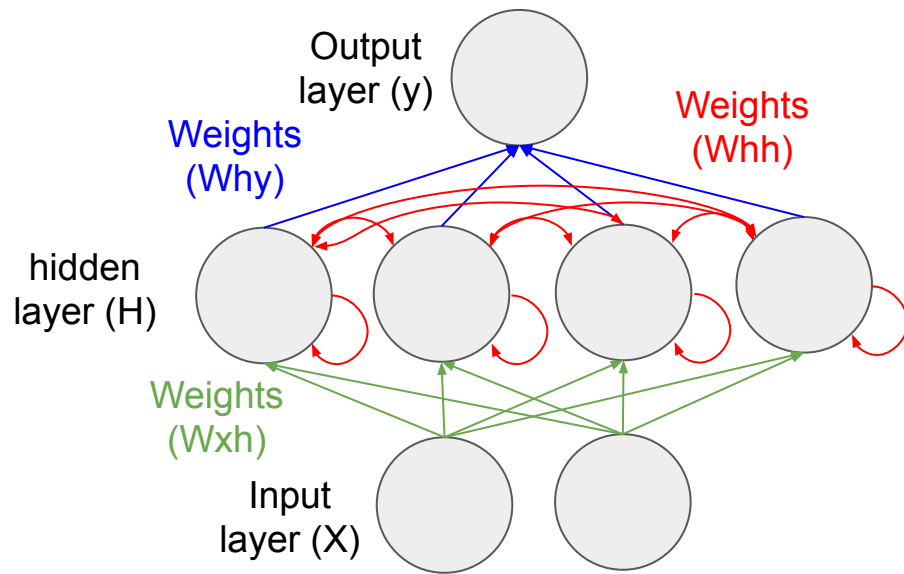
So how many weights in each architecture?

Vanilla MLP



$$W_{hy} = 4$$
$$W_{xh} = 8$$

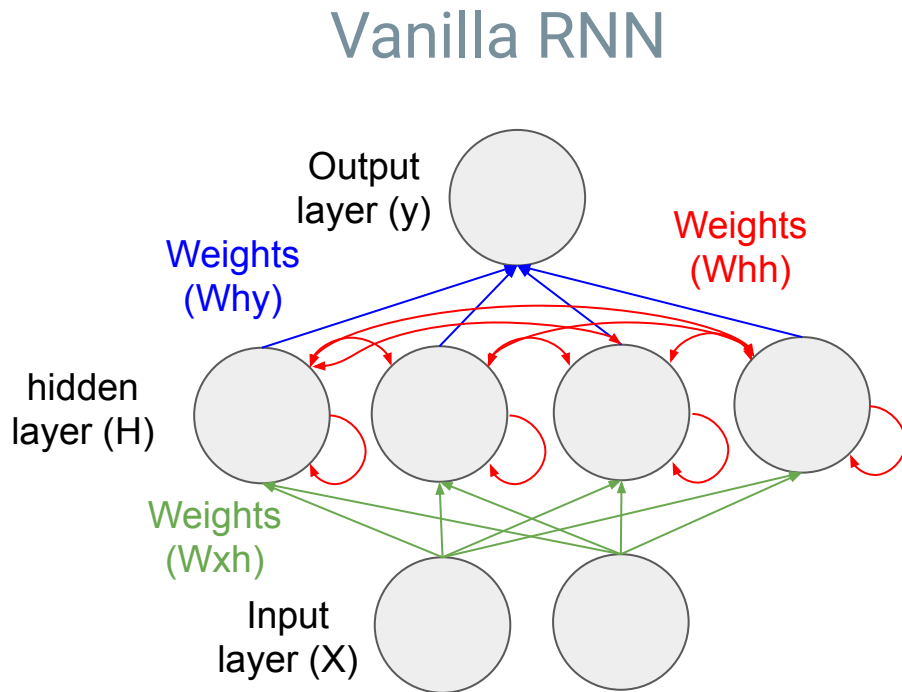
Vanilla RNN



$$W_{hy} = 4$$
$$W_{hh} = 16$$
$$W_{xh} = 8$$


Benefit of the intra - layer recurrent connections

- The previous state of a node in a recurrent hidden layer (**H_{prev}** for coding purposes) can affect the value of itself or other nodes in the layer in the present time (it's a directed cycle).
- This gives the net the ability to model sequential data.
- Feedforward and backpropagation work the same way.
- Learn **W_{hh}** like all the other weights. In a trained model all the weights are fixed. It's the activations of the nodes that changes with changes in sequence.



Exercise: RNN - text is sequential data

Use the `min-char-rnn.py` code to learn Dr. Seuss. As the model trains it will eventually write some new books!



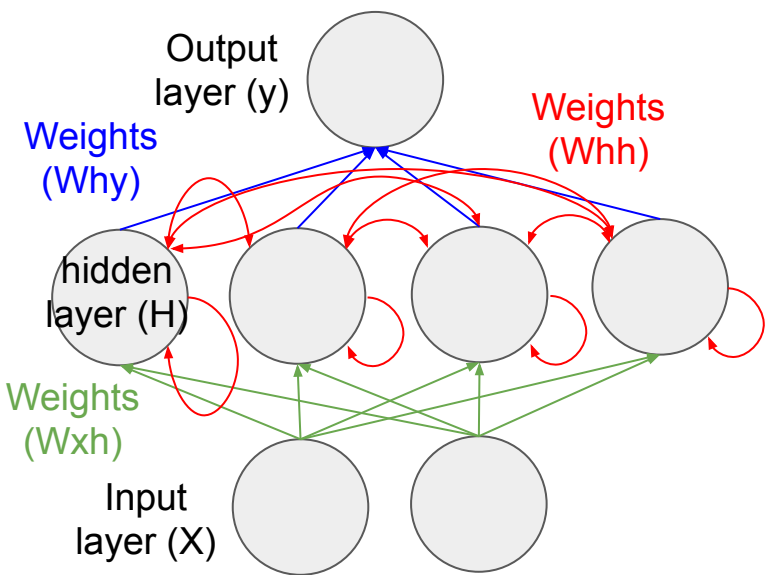


As you play with the model, try to answer the following questions:

- 1) What is the model's architecture?
(#inputs, #layers, activation functions, #outputs)
- 2) How is the model predicting new characters?

Moving into multilayer RNNs

Multiple layers (and more nodes in each layer) allow more difficult sequences to be learned. They are also harder to train. Exploding and vanishing gradients cause convergence problems, too. Let's stop building things from scratch.



Keras

[Keras](#) is a high-level neural networks API, written in Python and capable of running on top of either TensorFlow, CNTK or Theano.

We use it in the DSI for capstone projects. MLPs, CNNs, RNNs, some Reinforcement Learning too.

[Will become TensorFlow's default API.](#)


Available recurrent layers:

- Recurrent

- SimpleRNN

- Long Short-Term Memory (LSTM)

- Gated Recurrent Unit (GRU)

 Keras Documentation

Search docs

[Home](#)
[Getting started](#)

- [Guide to the Sequential model](#)
- [Guide to the Functional API](#)
- [FAQ](#)

[Models](#)

- [About Keras models](#)
- [Sequential](#)
- [Model \(functional API\)](#)

[Layers](#)

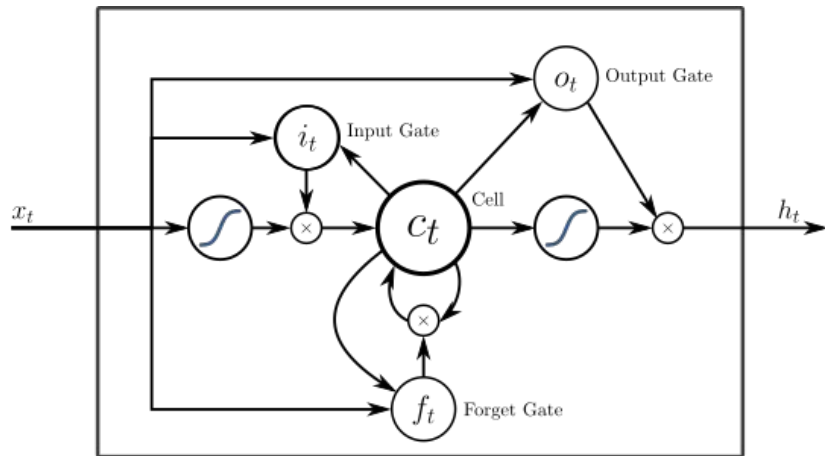
- [About Keras layers](#)
- [Core Layers](#)
- [Convolutional Layers](#)
- [Pooling Layers](#)
- [Locally-connected Layers](#)

[Recurrent Layers](#)

- [Recurrent](#)
- [SimpleRNN](#)
- [GRU](#)
- [LSTM](#)

LSTM

- Long short-term memory (LSTM) is an architecture (an artificial neural network) proposed in 1997.
- LSTM network is well-suited ... when there are time lags of unknown size and bound between important events.
- LSTM practical applications: natural language text compression, handwriting recognition, speech recognition, translation. ([See Wikipedia](#))
- Your assignment will ask you to describe how one works, and then you'll use them in Keras to predict stock price!



[Attribute](#)

RNN resources

[Andrej Karpathy's NN course](#), [Karpathy Github](#)

[Ian Goodfellow's Deep Learning book](#)

[Iamtrask's blog](#)

[Christopher Olah's blog](#)

[Jakob Aungiers's blog](#), [Aungiers Github](#)

[Github page for my Meetup!](#)

Past capstones:

[Solving LSAT Puzzles with Parsey McParseface and Seq2Seq](#)

[DJ ABC LSTM](#) (making music influenced by Irish folk songs, Enya, Michael Jackson, and Bach)

Objectives

- Introduce RNNs
 - Why are RNNs good for sequential analysis?
- Tangent: Time Series (common application for RNNs)
- State important differences between MLPs and RNNs
- State a common shortcoming of “vanilla” RNNs
- Explain, at a high-level, what an LSTM is (assignment)
- Predict stock price using an LSTM (assignment)