

Profit Curves and Imbalanced Classes

Ryan Henning
Frank Burkholder



“I just made the most amazing predictive model.”

“Do tell.”

“99.99% accuracy on fraud/not fraud test data.”

“What ratio of instances in the training data are fraud?”

“Fraud is super rare! Only 1 in 10,000!”

What would you ask next?

Problem Motivation

- Classification datasets can be “imbalanced”.
 - i.e. many observations of one class, few of another
 - Will give concrete examples later, but even a minority class of comprising 33% of the data can be considered imbalanced.
- Costs (in time, money, or life!) of a false positive is often different from cost of a false negative. Need to consider external (e.g. business) costs.
 - e.g. missing fraud can be more costly than screening legitimate activity
 - False negative in disease screening vs False negative in email spam classification
- Accuracy-driven models will over-predict the majority class.

Solutions

Practical steps (help your model fit better):

- Stratifying train_test_split
- Change weighting of training data for poorly represented class

Cost-sensitive learning (use outside costs & benefits to set prob. thresh):

- thresholding (aka “profit curves”)

Sampling (reduce imbalance with more/less data):

- Oversampling
- Undersampling
- SMOTE - Synthetic Minority Oversampling Technique

Dealing with imbalanced classes: Practical steps

- Stratifying train_test_split
- Change weighting of training data for poorly represented class

If you have a minority class, are you sure it's represented in the same proportion in your `y_train` and `y_test` datasets?

```
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

If you have a minority class, are you sure it's represented in the same proportion in your `y_train` and `y_test` datasets?

```
X_train, X_test, y_train, y_test = train_test_split(X, y)
```

Maybe we'll get lucky!?

But this is better:

```
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify = y)
```

In objective function minimization, all classes are weighted equally by default:

```
class sklearn.linear_model. LogisticRegression (penalty='l2', dual=False, tol=0.0001, C=1.0, fit_intercept=True,
intercept_scaling=1, class_weight=None, random_state=None, solver='liblinear', max_iter=100, multi_class='ovr',
verbose=0, warm_start=False, n_jobs=1) ¶
```

[\[source\]](#)

Option 1)

class_weight : dict or 'balanced', optional

Weights associated with classes in the form `{class_label: weight}`. If not given, all classes are supposed to have weight one.

The “balanced” mode uses the values of y to automatically adjust weights inversely proportional to class frequencies in the input data as

```
n_samples / (n_classes * np.bincount(y))
```

Note that these weights will be multiplied with sample_weight (passed through the fit method) if sample_weight is specified.

Option 2)

fit (X, y[, sample_weight]) Fit the model according to the given training data.

Jupyter notebook

Dealing with imbalanced classes:

Cost sensitive learning

- Quantify relative costs of TP, FP, TN, FN
- Construct a confusion matrix for each probability threshold, and use a cost-benefit matrix to calculate a “profit” for each threshold. Pick the threshold that give the highest profit.


Recall the ROC Curve:


- ROC shows $FPR = (1 - TNR)$ vs TPR (aka Recall)
- doesn't give preference to one over the other

Q: How to handle unequal error costs?

Q: Which threshold to pick? (assessment)


A: Plot expected profit (4 steps)!





Computing Expected Profit

Step 1 - Estimate error probabilities based on a given threshold (will review thresholding in a bit)



Computing Expected Profit

Step 2 - Define the cost-benefit matrix (based on your out-of-model knowledge)

Example
(maybe filling out a survey)

		Actual	
		p	n
Predicted	Y	49	-1
	N	0	0

\$50 - \$1

		Actual	
		p	n
Predicted	Y	TP $b(Y,p)$	FP $c(Y,n)$
	N	FN $c(N,p)$	TN $b(N,n)$

To close this section on estimated profit, we emphasize two pitfalls that are common when formulating cost-benefit matrices:

- It is important to make sure the signs of quantities in the cost-benefit matrix are consistent. In this book we take benefits to be positive and costs to be negative. In many data mining studies, the focus is on minimizing cost rather than maximizing profit, so the signs are reversed. Mathematically, there is no difference. However, it is important to pick one view and be consistent.
- An easy mistake in formulating cost-benefit matrices is to “double count” by putting a benefit in one cell and a negative cost for the same thing in another cell (or vice versa). A useful practical test is to compute the *benefit improvement* for changing the decision on an example test instance.

For example, say you’ve built a model to predict which accounts have been defrauded. You’ve determined that a fraud case costs \$1,000 on average. If you decide that the benefit of catching fraud is therefore +\$1,000/case on average, *and* the cost of missing fraud is -\$1,000/case, then what would be the *improvement in benefit* for catching a case of fraud? You would calculate:

$$b(\mathbf{Y}, \mathbf{p}) - b(\mathbf{N}, \mathbf{p}) = \$1000 - (-\$1000) = \$2000$$

But intuitively you know that this improvement should only be about \$1,000, so this error indicates double counting. The solution is to specify either that the benefit of catching fraud is \$1,000 or that the cost of missing fraud is -\$1,000, but not both. One should be zero.

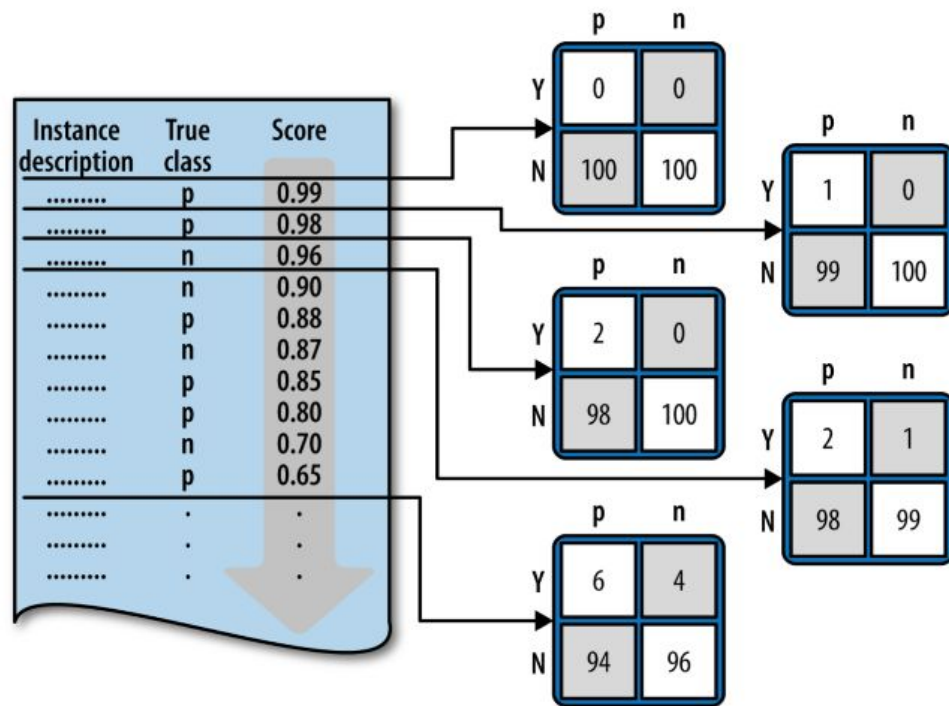
Computing Expected Profit

Step 3 - Combine probabilities and payoffs.

$$E[Profit] = \begin{matrix} & \text{TP} & & \text{FP} \\ & P(Y, p) \cdot b(Y, p) & + & P(Y, n) \cdot c(Y, n) \\ & P(N, p) \cdot c(N, p) & + & P(N, n) \cdot b(N, n) \\ & \text{FN} & & \text{TN} \end{matrix}$$

Find the profit-maximizing threshold

- Starting with the highest threshold (most probable) and working down compute expected profit.
- Then select threshold with highest expected profit (except if a budget is coming into play - next slide)
- Benefits of ranking (high prob. to low) clear when we have a budget and want to spend money on the most probable cases




Step 4 - Plot profit

Cost - benefit matrix

		Actual	
		p	n
Predicted Y	Y	49	-1
	N	0	0

Note no profit / cost on this matrix for FN, TN

		p	n
Y	Y	0	0
	N	10	990



		p	n
Y	Y	10	990
	N	0	0

Threshold = 1.0

<- Sorted by threshold ->


Threshold = 0.0

Profit Curve - budget could influence number of users you target

Cost - benefit matrix

		Actual	
		p	n
Predicted	Y	49	-1
	N	0	0

Note no profit / cost on this matrix for FN, TN



Say there are 1000 instances. It costs \$1 to check an instance.



Dealing with imbalanced classes: Sampling techniques

- Undersampling
- Oversampling
- SMOTE - Synthetic Minority Oversampling Technique

Sampling Techniques: Undersampling

- Undersampling randomly discards majority class observations to balance training sample.
- **PRO:** Reduces runtime on very large datasets.
- **CON:** Discards potentially important observations.

Mention imbalanced-learn package
(<https://github.com/scikit-learn-contrib/imbalanced-learn>)

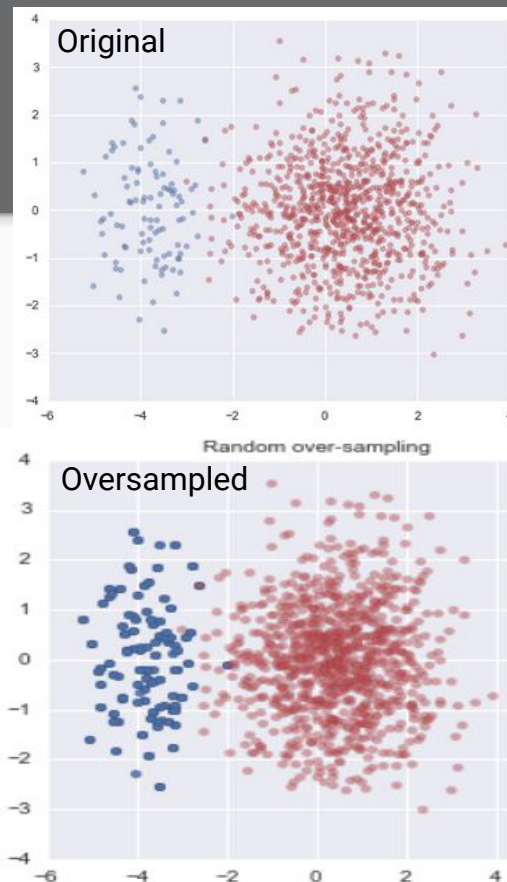


<https://github.com/scikit-learn-contrib/imbalanced-learn>

Sampling Techniques - Oversampling


- Oversampling replicates observations from minority class to balance training sample.
- **PRO:** Doesn't discard information.
- **CON:** Likely to overfit.

(Often better to use SMOTE)



Sampling Techniques: SMOTE

- SMOTE - Synthetic Minority Oversampling Technique
- Generates new observations from minority class.
- For each minority class observation and for each feature, randomly generate between it and one of its k-nearest neighbors.



SMOTE pseudocode

```
synthetic_observations = []  
while len(synthetic_observations) + len(minority_observations) < target:  
    obs = random.choice(minority_observations):  
    neighbor = random.choice(kNN(obs, k)) # randomly selected neighbor  
    new_observation = {}  
    for feature in obs:  
        weight = random() # random float between 0 and 1  
        new_feature_value = weight*obs[feature] \  
                               + (1-weight)*neighbor[feature]  
        new_observation[feature] = new_feature_value  
    synthetic_observations.append(new_observation)
```

of desired minority observations
↙

Sampling Techniques - Distribution

What's the right amount of over-/under-sampling?

- If you know the cost-benefit matrix:
 - Maximize profit curve over target proportion
- If you don't know the cost-benefit matrix:
 - No clear answer...
 - ROC's AUC might be more useful...

Cost Sensitivity vs Sampling

- Neither is strictly superior.
- Oversampling tends to work better than undersampling on small datasets.
- Some algorithms don't have an obvious cost-sensitive adaptation, requiring sampling.

See also "Cost-Sensitive Learning vs. Sampling: Which is Best for Handling Unbalanced Classes with Unequal Error Costs?" <http://storm.cis.fordham.edu/gweiss/papers/dmin07-weiss.pdf>

Review

Practical steps:

- Stratifying train_test_split
- Change weighting of training data for poorly represented class

Cost-sensitive learning:

- thresholding (aka “profit curves”)

Sampling:

- Oversampling
- Undersampling
- SMOTE

Best: Can you get more data?!?