

Bagging and Random Forests

Jack Bennetto

September 28, 2016

Objectives

Morning Objectives

- Explain & construct a random forest (classification or regression).
- Explain the relationship and difference between random forest and bagging.
- Explain why random forests are more accurate than a single decision tree.

Afternoon Objectives

- Get feature importances from a random forest.
- Explain how OOB error is calculated and what is it an estimate of.

Agenda

2

Morning Agenda:

- Discuss ensemble methods
- Review bias/variance tradeoff
- Review decision trees
- Discuss bagging (bootstrap aggregation)
- Discuss random forests

Afternoon Agenda:

- Discuss out-of-bag error
- Discuss feature importance

What is an Ensemble Method?

Ensemble: many weak learners combined to form a strong learner

3

Train multiple different models on the data. The overall prediction is

- the average prediction, for a regressor, or
- the plurality choice, for a classifier (fraction of models is probability).

4

Ensembles: Intuition

Supposed we have 5 *independent* binary classifiers each 70% accurate.

- What's the overall accuracy?

$$\binom{5}{5}0.7^5 + \binom{5}{4}0.7^40.3 + \binom{5}{3}0.7^30.3^2 \approx 0.83$$

With 101 such classifiers we can achieve 99.9% accuracy.

What's the limitation?

How to Make Them Independent?

If the learners are all the same, ensembles don't help.

5

Train each learner on different subset of data.

- Why is this better than a single good model?

6

Bias and Variance

Bias: Error from failure to match training set

Variance: Error from sampling training set

- What is the bias of an unpruned decision tree?

Decision trees are easy to overfit.

Review: Classification Trees

Training:

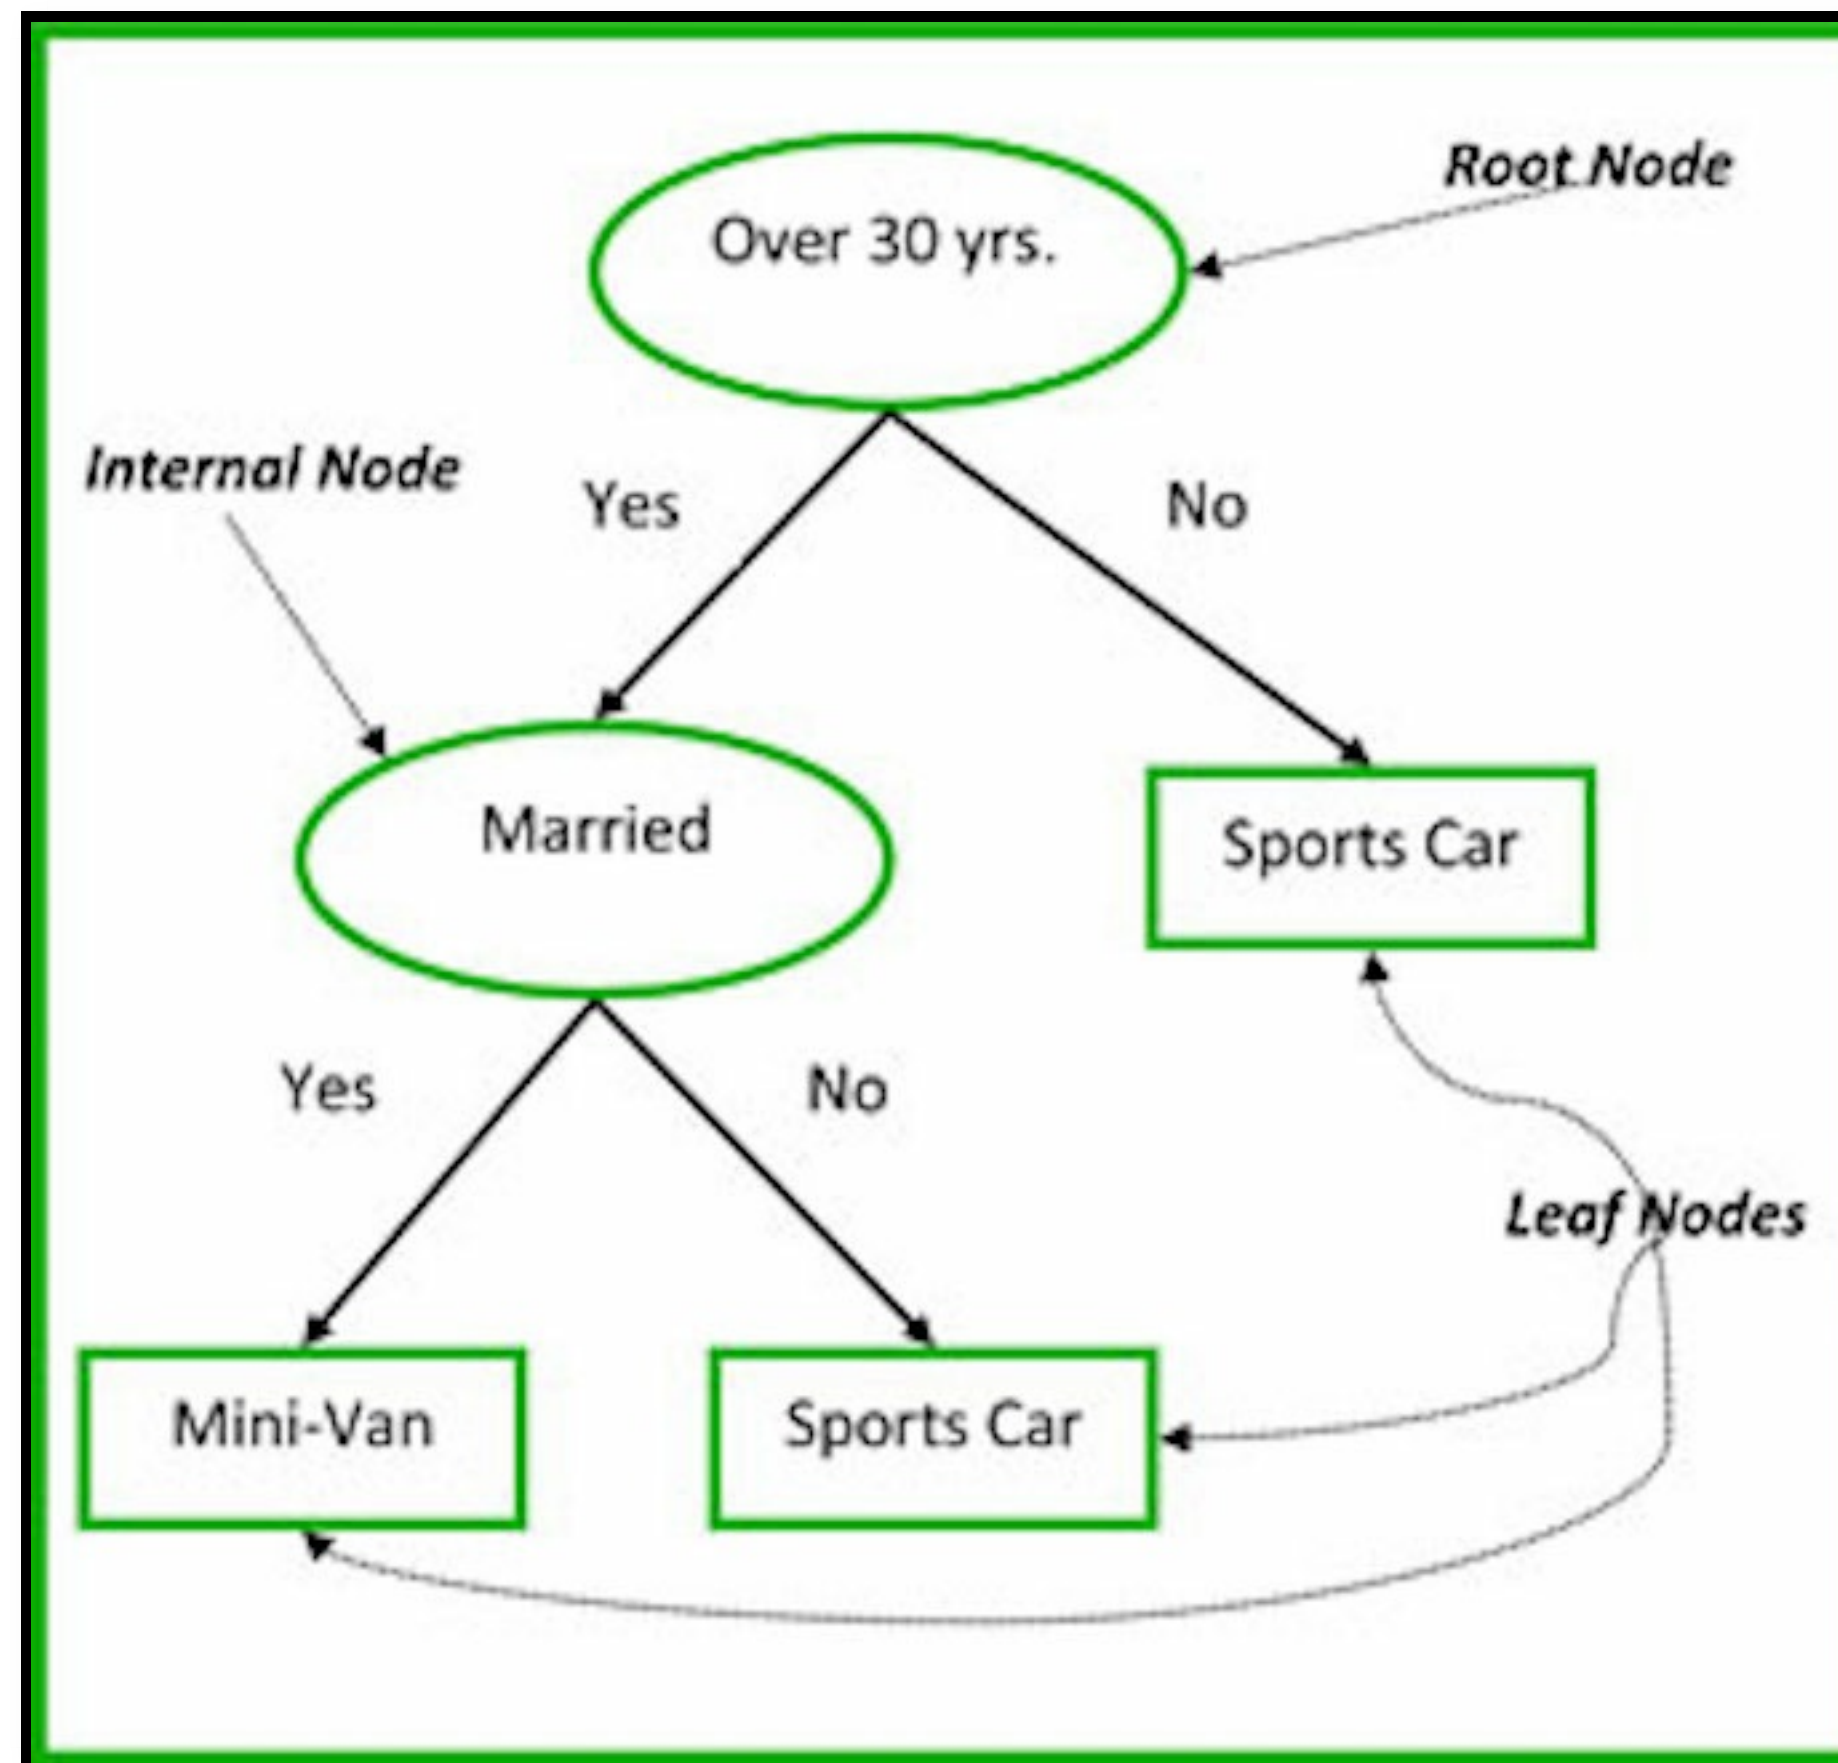
- Iteratively divide the nodes such that (entropy/gini impurity) is minimized
- Various stopping conditions like a depth limit
- Prune trees by merging nodes

Inference:

- Take the most common class in the leaf node

Review: Classification Trees

8



Review: Regression Trees

Predicting a number, not a class

Training:

- Iteratively divide the nodes such that *total squared error* is minimized

$$\sum_{i \in L} (y_i - m_L)^2 + \sum_{i \in R} (y_i - m_R)^2$$

- Various stopping conditions like a depth limit
- Prune trees by merging nodes

Inference:

- Take the most average value of samples the leaf node

Regression Trees: Example

10

x_1	x_2	y
1	1	1
0	0	2
1	0	3
0	1	4

Prior to the split we guess the mean, 2.5, for everything, giving total squared error:

$$E = (1 - 2.5)^2 + (2 - 2.5)^2 + (3 - 2.5)^2 + (4 - 2.5)^2 = 5$$

After we split on x_1 we guess 2 for rows 1 & 3 and 3 for rows 2 & 4:

$$E = (1 - 2)^2 + (3 - 2)^2 + (2 - 3)^2 + (4 - 3)^2 = 4$$

Decision Tree Summary

11

What are the pros and cons?

Pros

- No feature scaling needed
- Model nonlinear relationships
- Can do both classification and regression
- Robust
- Highly interpretable

Cons

- Can be expensive to train
- Often poor predictors - high variance

Review: Bootstrapping

12

Questions:

- What is a bootstrap sample?

Given n data points we select a sample of n points with replacement

- What have we learned that bootstrap samples are good for so far?

We use bootstrap samples to construct confidence intervals around sample statistics.

Example: get a confidence interval around sample median

- Take 1000 bootstrap samples.
- Take the median of each sample.
- The 95% confidence interval for the median is between the 25th and 975th largest samples.

Bagging

We could repeatedly sample from the population, building decision tree models and averaging the results. But we only have one sample.

13

- Simulate multiple draws from the data by using multiple bootstrap samples
- Bagging stands for Bootstrap Aggregation

14

Bagging

In a bit more detail:

- Take a bunch of bootstrap samples - say n
- Train a high variance, low bias model on each of them
- Average the results - this can reduce your variance by up to \sqrt{n}
 - Question: Why is the reduction in variance less than \sqrt{n} ?

Correlation Between the Trees

Why is the reduction in variance less than \sqrt{n} ?

- We are thinking about the population of all possible decision tree models on our data.
- If I take n samples *iid* from this distribution and average them the variance goes down by \sqrt{n}
- There is some correlation between models because they are all trained on bootstrap samples from the same draw.

Random Forests

Random Forests improve on Bagging by de-correlating the trees using a technique called Subspace Sampling.

16

- At each decision tree split only m (often $m = \sqrt{n}$) features are considered.

17

Random Forest Parameters

Random Forest Parameters

- Total number of trees
- Number of features to use at each split
- Individual decision tree Parameters
 - e.g., tree depth, pruning, split criterion

In general, RF are fairly robust to the choice of parameters and overfitting.

Pros and Cons of Random Forest

18

Pros

- Often give near state-of-the-art performance
- Good out-of-the-box performance
- No feature scaling needed
- Model nonlinear relationships

Cons

- Can be expensive to train (though can be done in parallel)
- Not interpretable

Afternoon Lecture

Interpreting Random Forests

19

20

Objectives

Morning Objectives

- Explain & construct a random forest (classification or regression).
- Explain the relationship and difference between random forest and bagging.
- Explain why random forests are more accurate than a single decision tree.

Afternoon Objectives

- Get feature importances from a random forest.
- Explain how OOB error is calculated and what is it an estimate of.

Agenda

21

Morning Agenda

- Discuss ensemble methods
- Review bias/variance tradeoff
- Review decision trees
- Discuss bagging (bootstrap aggregation)
- Discuss random forests

Afternoon Agenda

- Discuss out-of-bag error
- Discuss feature importance

Out-Of-Bag Error

22

Measuring error of a bagged model.

- Out-of-bag (OOB) error is a method of estimating the error of ensemble methods that use Bagging.
- About 37% of the estimators will not have been trained on each data point.
- Test each data point only on the estimators that didn't see that data point during training.
- Often use cross validation anyway because we're comparing with other models and want to measure the accuracy the same way.

Feature Importances

23

Bringing interpretability to random forests

- Determining which features are important in predicting the target variable is often a critically important business question.
- Example: Churn analysis - it's generally more important to understand *why* customers are churning than to predict which customers are going to churn.

How should we measure it?

Feature Importances: Mean Decrease Impurity

24

How much does each feature decrease the impurity?

To compute the importance of the j^{th} feature:

- For each tree, each split is made in order to reduce the total impurity of the tree (Gini/entropy/MSE); we can record the magnitude of the reduction.
- Then the importance of a feature is the average decrease in impurity across trees in the forest, as a result of splits defined by that feature.
- This is implemented in sklearn.

Feature Importances: Mean Decrease Accuracy

25

How much does randomly mixing values of a feature affect accuracy?

To compute the importance of the j^{th} feature:

- When the b^{th} tree is grown, use it to predict the OOB samples and record accuracy.
- Scramble the values of the j^{th} feature in the OOB samples and do the prediction again. Compute the new (lower) accuracy.
- Average the decrease in accuracy across all trees.

Feature Importances: ipython

See example in ipython notebook.

26

27