

Bagging and Random Forests

Jack Bennetto

June 15, 2016

Objectives

Morning Objectives:

- Thoroughly explain the construction of a random forest (classification or regression) algorithm.
- Explain the relationship and difference between random forest and bagging.
- Explain why random forests are more accurate than a single decision tree.

Afternoon Objectives:

- Explain how to get feature importances from a random forest using an algorithm.
- Explain how OOB error is calculated and what it is an estimate of.

Agenda

Morning Agenda

- Discuss ensemble methods
- Review bias/variance tradeoff
- Review decision trees
- Discuss bagging (bootstrap aggregation)
- Discuss random forests

Afternoon Agenda

- Discuss feature importance
- Discuss out-of-bag error

What is an Ensemble Method?

Ensemble: many weak learners combined to form a strong learner

Train multiple different models on the data. To predict:

- For regressor, average the predictions of the models
- For classifier, take the plurality choice

Ensembles: Intuition

Supposed we have 5 *independent* binary classifiers each 70% accurate.

Overall accuracy:

$$\binom{5}{3} 0.7^3 0.3^2 \times \binom{5}{4} 0.7^4 0.3 \times \binom{5}{5} 0.7^5 \approx 0.83$$

With 101 such classifiers we can achieve 99.9% accuracy.

- Why isn't this easy?

How to Make Them Independent?

If the learners are all the same, ensembles don't help.

Train each learner on different subset of data.

- Why is this better than a single good model?

Bias and Variance

Bias: Error from failure to match training set

Variance: Error from sampling training set

- What accuracy can you expect from a decision tree?

Bias and Variance

Bias: Error from failure to match training set

Variance: Error from sampling training set

- What accuracy can you expect from a decision tree?

Decision trees are easy to overfit.

Review: Classification Trees

Training:

- Iteratively divide the nodes into subnodes such that (entropy/gini impurity) is minimized
- Various stopping conditions like a depth limit
- Prune trees by merging nodes

Inference:

- Take the most common class in the leaf node

Review: Classification Trees

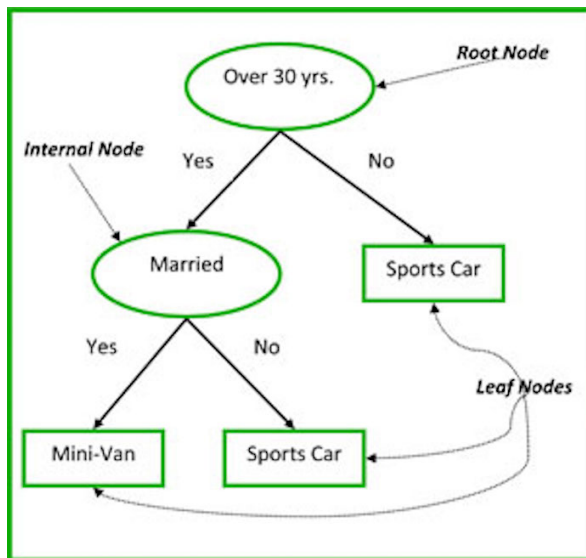


Figure 1: (<http://www.hypertextbookshop.com/>)

Regression Trees

Similar to Classification Trees but:

- Instead of predicting a class label we're not trying to predict a number
- We minimize *total squared error* instead of entropy or impurity

$$\sum_{i \in R} (y_i - m_R)^2 + \sum_{i \in S} (y_i - m_S)^2$$

- For inference take the mean of the leaf node

Regression Trees: Example

x_1	x_2	y
1	1	1
0	0	2
1	0	3
0	1	4

Prior to the split we guess the mean, 2.5, for everything, giving total squared error:

$$E = (1 - 2.5)^2 + (2 - 2.5)^2 + (3 - 2.5)^2 + (4 - 2.5)^2 = 5$$

After we split on x_1 we guess 2 for rows 1 & 3 and 3 for rows 2 & 4:

$$E = (1 - 2)^2 + (3 - 2)^2 + (2 - 3)^2 + (4 - 3)^2 = 4$$

Decision Tree Summary

Pros

- No feature scaling needed
- Model nonlinear relationships
- Highly interpretable
- Can do both classification and regression

Cons

- Can be expensive to train
- Often poor predictors - high variance

Review: Bootstrapping

Questions:

What is a bootstrap sample?

Review: Bootstrapping

Questions:

What is a bootstrap sample?

- Given n data points we select a sample of n points with replacement

What have we learned that bootstrap samples are good for so far?

Review: Bootstrapping

Questions:

What is a bootstrap sample?

- Given n data points we select a sample of n points with replacement

What have we learned that bootstrap samples are good for so far?

- We use bootstrap samples to construct confidence intervals around sample statistics
- Ex: If I want a confidence interval around my sample median I could take 200 bootstrap samples and take the median of all 200 bootstrap samples.

I am about 95% confident that the true population median is between the 5th from smallest and the 5th from largest.

Bagging

We've seen that repeatedly sampling from the population, building decision tree models and averaging the results gives a good estimate. But we only have one sample.

- Simulate multiple draws from the data by using multiple bootstrap samples
- Bagging stands for Bootstrap Aggregation

Second slide saying the same thing again to emphasize that Bagging is important.

- Take a bunch of bootstrap samples - say n
- Train a high variance, low bias model on each of them
- Average the results - this can reduce your variance by up to \sqrt{n}
 - ▶ Question: Why is the reduction in variance less than \sqrt{n} ?

Correlation Between the Trees

Why is the reduction in variance less than \sqrt{n} ?

- We are thinking about the population of all possible decision tree models on our data.
- If I take n samples *iid* from this distribution and average them the variance goes down by \sqrt{n}
- There is some correlation between my models because they are all trained on bootstrap samples from the same draw.

Random Forests

Random Forests improve on Bagging by de-correlating the trees using a technique called Subspace Sampling.

- At each decision tree split only m (often $m = \sqrt{n}$) features are considered for the split.

Random Forest Parameters

Random Forest Parameters

- Total number of trees
- Number of features to use at each split
- Number of points for each bootstrap sample
- Individual decision tree Parameters
 - ▶ e.g., tree depth, pruning

In general, RF are fairly robust to the choice of parameters and overfitting.

Pros and Cons of Random Forest

Pros

- Often give near state-of-the-art performance
- Good out-of-the-box performance.
- No feature scaling needed.
- Model nonlinear relationships

Cons

- Can be expensive to train (though can be done in parallel)
- Not interpretable

Afternoon Lecture: Interpreting Random Forests

Objectives

Morning Objectives:

- Thoroughly explain the construction of a random forest (classification or regression) algorithm
- Explain the relationship and difference between random forest and bagging.
- Explain why random forests are more accurate than a single decision tree.

Afternoon Objectives:

- Explain how to get feature importances from a random forest using an algorithm.
- Explain how OOB error is calculated and what is it an estimate of.

Agenda

Morning Agenda

- Discuss ensemble methods
- Review bias/variance tradeoff
- Review decision trees
- Discuss bagging (bootstrap aggregation)
- Discuss random forests

Afternoon Agenda

- Discuss feature importance
- Discuss out-of-bag error

Out Of Bag Error

- Out Of Bag error is a method of estimating the error of ensemble methods that use Bagging.
- About $1/3$ of the estimators will not have been trained on each data point. (Why?)
- Test each data point only on the estimators that didn't see that data point during training.

Out Of Bag Error

- Often we'll use cross validation anyway because we're comparing random forest to other models and we want to measure the accuracy the same way.

Feature Importances

- Determining which features are important in predicting the target variable is often a critically important business question.
- Example: Churn analysis - it's generally more important to understand *why* customers are churning than to predict which customers are going to churn.

Feature Importances: Mean Decrease Impurity

- For each tree, each split is made in order to reduce the total impurity of the tree (Gini Impurity for classification, mean squared error for regression); we can record the magnitude of the reduction.
- Then the importance of a feature is the average decrease in impurity across trees in the forest, as a result of splits defined by that feature.
- This is implemented in sklearn.

Feature Importances: Mean Decrease Accuracy

To compute the importance of the j th variable:

- When the b th tree is grown, use it to predict the OOB samples and record accuracy.
- Scramble the values of the j th variable in the OOB samples and do the prediction again. Compute the new (lower) accuracy.
- Average the decrease in entropy across all trees.

Feature Importances: ipython

See example in ipython notebook.