

Recommender Systems

Benjamin S. Skrainka

November 30, 2015

Objectives

Today's objectives:

- Describe primary approaches to recommender systems
- Build a recommender using collaborative filtering and similarity
- Build a recommender using collaborative filtering and matrix factorization

Agenda

Today's plan:

- ➊ Overview of types of recommender systems
- ➋ Collaborative filtering with similarity
- ➌ Collaborative filtering with matrix factorization
- ➍ Best practices
- ➎ Appendix

References

A couple references, from the machine learning perspective:

- Mining of Massive Datasets
- Recommender Systems: An Introduction
- Music Recommendation and Discovery: The Long Tail, Long Fail, and Long Play in the Digital Music Space
- Matrix Factorization Techniques for Recommender Systems
- Amazon.com recommendations: Item-to-Item Collaborative Filtering
- Dato/GraphLab documentation & blog

Introduction

Recommendation business problem

Recommendation problem takes several forms:

- Goal of recommender:
 - ▶ predict missing ratings
 - ▶ May be sufficient to just predict a subset of items with high expected rankings
 - ▶ May be sufficient to just predict general trends, such as *trending* news
- Long-tail:
 - ▶ Scarcity \Rightarrow brick & mortar stocks items based on average user
 - ▶ Online \Rightarrow cater to individual, not average user \Rightarrow stock everything, both popular and long tail
- Often described as *personalization*
- Examples: Movies (Netflix), Products (Amazon), Music (Pandora), and News articles (CNN)

Approaches to recommender systems

There are several approaches to building a recommender:

- Content-based: recommend based on properties/characteristics
- Collaborative filtering (CF): recommend based on similarity
- Hybrid: Content-based + Collaborative filtering
- Applications:
 - ▶ Product recommendations
 - ▶ Movie recommendations
 - ▶ News articles

Typically, data is a *utility (rating)* matrix, which captures user preferences/well-being:

- User rating of items
- User purchase decisions for items
- Unrated are coded as 0 or missing
- Most items are unrated \Rightarrow matrix is sparse
- Use recommender:
 - ▶ Determine which attributes users think are important
 - ▶ Predict ratings for unrated items
 - ▶ Better than trusting 'expert' opinion

Types of data

Data can be:

- *Explicit:*
 - ▶ User provided ratings (1 to 5 stars)
 - ▶ User like/non like
- *Implicit:*
 - ▶ Infer user-item relationships from behavior
 - ▶ More common
 - ▶ Example: buy/not buy; view/not-view
- To convert implicit to explicit, create a matrix of 1s (yes) and 0s (no)

Example: explicit utility matrix

Example 9.1 in [Mining of Massive Datasets](#):

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3		5	1		3

Example: implicit utility matrix

Based on example 9.1 in [Mining of Massive Datasets](#):

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	1			1	1		
B	1	1	1				
C				1	1	1	
D		1		1	1		1

Collaborative filtering using similarity

Overview of CF using similarity

Use similarity to recommend items:

- Make recommendations based on similarity:
 - ▶ Between users
 - ▶ Between items
- Similarity measures:
 - ▶ Pearson
 - ▶ Cosine
 - ▶ Jaccard

Types of collaborative filtering

Two types of similarity-based CF:

- *User-based*: predict based on similarities between users
 - ▶ Performs well, but slow if many users
 - ▶ Use item-based CF if $|Users| \gg |Items|$
- *Item-based*: predict based on similarities between items
 - ▶ Faster if you precompute item-item similarity
 - ▶ Usually $|Items| \ll |Users| \Rightarrow$ item-based CF is most popular
 - ▶ Items tend to be more stable:
 - ★ Items often only in one category (e.g., action films)
 - ★ Stable over time
 - ★ Users may like variety or change preferences over time

Collaborative filtering recipe

Compute predictions by similarity:

- 1 Normalize (demean) utility matrix
- 2 Reduce dimensionality: SVD, NMF, or UV (optional)
- 3 Compute similarity of users or items
- 4 Predict ratings for unrated items
- 5 Add prediction to average rating of user/item

Note:

- Precompute utility matrix for each user – it is relatively stable
- Only compute predictions at runtime

Review: measuring similarity

Example 9.1 in [Mining of Massive Datasets](#):

	HP1	HP2	HP3	TW	SW1	SW2	SW3
A	4			5	1		
B	5	5	4				
C				2	4	5	
D		3		5	1		3

- What is the Jaccard distance between A & B? A & C?
- What is the Cosine distance between A & B? A & C?
- See text for examples with normalization and rounding

Choosing a similarity measure

Chose the appropriate similarity measure for your data:

- Cosine:
 - ▶ Use for ratings (non-Boolean) data
 - ▶ Treat missing ratings as 0
 - ▶ Cosine + de-meaned data is the same as Pearson
- Jaccard:
 - ▶ Use only Boolean (e.g., buy/not buy) data
 - ▶ Loses information with ratings data

Then compute *similarity matrix* of pair-wise similarities between items (users)

Predict ratings from similarity

Predict using a similarity-weighted average of ratings:

$$\hat{r}_{ui} = \frac{\sum_{j \in I_u} \text{similarity}(i, j) \cdot R_{uj}}{\sum_{j \in I_u} \text{similarity}(i, j)}$$

where

- \hat{r}_{ui} is u 's predicted rating for item i
- $I_u \equiv$ set of items rated by u
- R_{uj} is utility matrix, i.e., $R_{uj} \equiv$ user u 's rating of item j

Check for mastery

How would you modify the prediction formula below for a user-based recommender?

$$\hat{r}_{ui} = \frac{\sum_{j \in I_u} \text{similarity}(i, j) \cdot R_{uj}}{\sum_{j \in I_u} \text{similarity}(i, j)}$$

Recommend best items

Recommend items with highest predicted rating:

- Sort predicted ratings \hat{r}_{ui}
- Optimize by only searching a neighborhood which contains the n items most similar to i
- Beware of 'cyberbalkanization':
 - ▶ Consumers like variety
 - ▶ Don't recommend every Star Trek film to someone who liked first film
 - ▶ Best to offer several different types of item

Dimensionality reduction (optional)

May use SVD or similar method to reduce dimension:

```
U, Sigma, VT = np.linalg(m_ratings)
# Set n_top_eig to capture most of the variance
m_sigma = np.mat(np.eye(n_top_eig) * Sigma[:n_top_eig])
m_new_ratings = m_ratings.T * U[:, :n_top_eig] * m_sigma.I
```

Collaborative filtering using matrix factorization

Collaborative filtering using matrix factorization

Predict ratings from *latent factors*:

- Compute latent factors q_i and p_u via matrix factorization
- *Latent factors* are unobserved user or item attributes:
 - ▶ Describe some user or item concept
 - ▶ Affect behavior
 - ▶ Example: escapist vs. serious, male vs. female films
- Predict rating: $\hat{r}_{ui} = q_i^T p_u$
- Assumes:
 - ▶ Utility matrix is product of two simpler matrices (long, thin):
 - ▶ \exists small set of users & items which characterize behavior
 - ▶ Small set of features determines behavior of most users
- Can use NMF, UV , or SVD

Review: SVD

Q: What is SVD?

Q: How do you compute it? (optional)

Q: How do you compute the variance in the data that a factor explains?

Q: What do the different matrices in decomposition represent?

Q: How can you use it to reduce dimensions?

Review: NMF

Q: What is NMF?

Q: How do you compute it?

Q: What do the different matrices in decomposition represent?

SVD vs. NMF

SVD:

- Must know all ratings – i.e., no unrated items
- Assumes can minimize squared Frobenius norm
- Very slow if matrix is large & dense

NMF:

- Can estimate via alternating least squares (ALS) or stochastic gradient descent (SGD)
- Must regularize
- Can handle big data, biases, interactions, and time dynamics

Using NMF in recommendation systems

NMF is a 'best in class' option for many recommendation problems:

- Includes overall, user & item bias as well as latent factor interactions
- Can fit via SGD or ALS
- No need to impute missing ratings
- Use regularization to avoid overfitting
- Can handle time dynamics, e.g., changes in user preferences
- Used by winning entry in Netflix challenge

NMF problem formulation

To factor the utility matrix:

$$\operatorname{argmin}_{\{q_i, p_u\}} \sum_{(u,i) \in \mathcal{K}} (r_{ui} - q_i^T p_u)^2 + \lambda(\|q_i\|^2 + \|p_u\|^2)$$

where

- $\mathcal{K} \equiv$ all (u, i) in the training set with known ratings
- λ is amount of regularization
- r_{ui} is user u 's rating of item i
- p_u is latent factor for user u
- q_i is latent factor for item i

NMF problem formulation with bias

Should account for bias:

$$\operatorname{argmin}_{\{q_i, p_u, \mu, b_u, b_i\}} \sum_{(u,i) \in \mathcal{K}} (r_{ui} - \mu - b_u - b_i - q_i^T p_u)^2 + \lambda(\|q_i\|^2 + \|p_u\|^2 + b_u^2 + b_i^2)$$

where

- μ : overall bias (average rating)
- b_u : user bias
- b_i : item bias

Estimating NMF

Two methods to estimate NMF factors:

- Stochastic gradient descent (SGD):
 - ▶ Easier and faster than ALS
 - ▶ Must tune learning rate
 - ▶ Sometimes called 'Funk SGD' after originator
- Alternating least squares (ALS):
 - ▶ Use least squares, alternate between fixing q_i and p_u
 - ▶ Available in Spark/MLib
 - ▶ Fast if you can parallelize
 - ▶ Better for implicit (non-sparse) data
- Beware of local optima!

To get best performance with NMF:

- Model bias (overall, user, and item)
- Model time dynamics, such as changes in user preferences
- Add side or implicit information to handle cold-start
- See [Matrix Factorization Techniques for Recommender Systems](#)

Building a recommender with NMF

Use **GraphLab**:

- Supports many types of recommenders
- Provides (near) best in class performance
- Reasonable licensing terms
- To improve performance, focus on:
 - ▶ Data collection and quality
 - ▶ Cold-start problem
 - ▶ Feature engineering

Best practices

Overview:

Will discuss:

- Cold-start problem
- Evaluation
- GraphLab ProTips
- (GraphLab) model selection

The cold-start problem

Difficult to build a recommender without ratings:

- *Cold-start* problem:
 - ▶ Need utility matrix to recommend
 - ▶ Can ask users to rate items
 - ▶ Infer ratings from behavior, e.g., viewing an item
- Must also handle new users and new items
- Approaches:
 - ▶ Use ensemble of (bad) recommenders until you have enough ratings
 - ▶ Use content-based recommender
 - ▶ Exploit implicit, tag, and other side data
 - ▶ Use `ItemSimilarityModel` until you have enough rating data

Choose right evaluation criteria:

- Historically, used RMSE or MAE
- But, only care about predicting top n items
 - ▶ Should you compute metric over all missing ratings in test set?
 - ▶ No need to predict items undesirable items well
- *Precision at n* : percentage of top n predicted ratings that are 'relevant'
- *Recall at n* : percentage of relevant items in top n predictions
- Lift or hit rate are more relevant to business

Evaluation issues

Evaluation is difficult:

- Performance of recommender should be viewed in context of *user experience* (UX)
- \Rightarrow run A/B test on entire system
- Cross validation is hard:
 - ▶ What do you use for labels because of missing data?
 - ▶ Users choose to rate only some items \Rightarrow selection bias
 - ▶ Not clear how to fix this bias, which is always present
- Beware of local optima \Rightarrow use multiple starts

Cross-validation (for item-based recommender):

- Randomly sample ratings to use in training set
- Split on users
- Be careful if you split temporally
- Do not split on items

Building a production recommender is also challenging:

- Part of entire UX
- Should consider:
 - ▶ Diversity of recommendations
 - ▶ Privacy of personal information
 - ▶ Security against attacks on recommender
 - ▶ Social effects
 - ▶ Provide explanations
- See [Recommender systems: from algorithms to user experience](#)

GraphLab provides best in class performance:

- Start with `MatrixFactorizationModel`:
 - ▶ Switch to `LinearRegressionModel` if too slow
 - ▶ Switch to `FactorizationModel` if need interactions
- Focus on cold-start and side information to obtain best performance
- Tune settings with `graphlab.toolkits.model_params_search()`
- Compare models with `graphlab.recommender.util.compare_models()`

Select model based on data and business metric:

- For best ranking performance:
 - ▶ Use `ItemSimilarityModel`, `MatrixFactorizationModel`, or `FactorizationModel`
 - ▶ Set `ranking_regularization` $\in (0,1)$
 - ▶ With implicit data, add rating column of 1s and set `unobserved_rating_value=0`
- For best ratings prediction with real ratings:
 - ▶ Use `MatrixFactorizationModel`, `FactorizationModel`, or `LinearRegressionModel`
 - ▶ `LinearRegressionModel` uses user & item features and user & item popularity bias
 - ▶ Matrix models add user & item latent factors
 - ▶ `FactorizationModel` adds interaction between latent and side features

Dato documentation

Dato's documentation is excellent:

- Documentation
- Basic example
- Million song example:

Computation tips:

- Compute offline:
 - ▶ Matrix factorization
 - ▶ Similarity matrix
 - ▶ User/item neighborhoods (via clustering)
- Compute predicted ratings/rankings live

Summary

You should now be able to explain:

- Content-based vs. collaborative filtering recommenders?
- Item-based vs. user-based CF?
- Compute measures of similarity (Jaccard, Pearson, cosine)?
- State which GraphLab recommender model is right for which problem?
- Describe how to tune and evaluate a recommender?
- Explain how to overcome the cold-start problem

Appendix: similarity measures

Similarity measures

Recommenders use distance to quantify similarity:

- Cosine similarity:

- ▶ $\text{cosine}(x, y) = \cos \theta = \frac{\mathbf{x} \cdot \mathbf{y}}{\|\mathbf{x}\| \cdot \|\mathbf{y}\|}$
- ▶ $\text{similarity}(x, y) = \frac{1}{2} + \frac{1}{2} \cdot \text{cosine}(x, y)$
- ▶ Same as Pearson if you de-mean data
- ▶ Treat blanks as 0

- Jaccard distance:

- ▶ Jaccard index: $J(A, B) = \frac{|A \cap B|}{|A \cup B|} = \frac{|A \cap B|}{|A| + |B| - |A \cap B|}$
- ▶ Jaccard distance: $d_J(A, B) = 1 - J(A, B)$
- ▶ Use for binary data
- ▶ Loses information with non-Boolean data
- ▶ Example:
 - ★ Let $U_k \equiv \{i \in \text{Users} \mid R_{ik} \neq 0\}$, i.e. user i rated item k
 - ★ $\text{similarity}(a, b) = J(U_a, U_b) = \frac{|U_a \cap U_b|}{|U_a \cup U_b|}$

Other distance measures

Two other measures of similarity:

- Similarity:

- ▶ Constructed from Euclidean distance so $similarity(x, y) \in (0, 1)$

- ▶ $similarity(x, y) = \frac{1}{1 + \|\mathbf{x} - \mathbf{y}\|}$

- Pearson correlation: $pearson(x, y) = \frac{cov(x, y)}{\sigma(x) \cdot \sigma(y)}$

- ▶ Renormalize to be in (0,1): $similarity(x, y) = \frac{1}{2} + \frac{1}{2} \cdot pearson(x, y)$
- ▶ Use Numpy `corrcoef()`

Appendix: matrix factorization

Review: matrix factorization (1/4)

Use matrix factorization to predict ratings:

- Discover *latent factors*, unobserved characteristics which determine behavior
- Reduce dimension
- Consider: SVD, UV, or NMF
- Avoid PCA (why?)

Review: SVD (2/4)

Decompose rating matrix, M , into $U \cdot \Sigma \cdot V^T$

- U : $m \times d$ unitary matrix, represents user latent factors
- Σ :
 - ▶ $d \times d$ diagonal matrix of singular values
 - ▶ Σ^2 is the variance of each factors
- V^T :
 - ▶ $d \times n$ matrix
 - ▶ Transpose of item latent factors
- Keep only factors which explain the top ~90% of variance
- Caveat: doesn't work with missing values

Review: UV (3/4)

Decompose rating matrix, M , into $U \cdot V$

- U : $m \times d$ unitary matrix, represents user latent factors
- V : $d \times n$ matrix
- Can fit via stochastic gradient descent or alternating least squares (ALS)
- Use regularization to avoid overfitting

$U \cdot V$ decomposition

M is an m by n matrix

- $M \approx U \cdot V$, U is m by d and V is d by n
- Use entries from $U \cdot V$ to predict missing ratings
- Fit by minimizing RMSE of $M - U \cdot V$:
 - ▶ Has multiple local optima
 - ▶ Use multiple starts & algorithms
 - ★ Start from $\sqrt{\frac{\text{ave}(\{m_{ij} \in M | m_{ij} \neq 0\})}{d}}$
 - ★ Perturb for other starts
 - ★ Vary path for visiting elements during optimization
 - ▶ Compute via ALS or update rule
 - ★ Minimize RMSE of $\sum (m_{ij} - (U \cdot V)_{ij})^2$
 - ★ Overfitting
 - ★ Use (stochastic) gradient descent to optimize

Non-negative matrix factorization :

- Includes overall, user & item bias as well as latent factor interactions
- Can fit via stochastic gradient descent or alternating least squares (ALS)
- Use regularization to avoid overfitting
- Used by winning entry in Netflix challenge

Appendix: content-based recommenders

Overview of content-based recommenders

Use features to determine similarity:

- Recommend based on item properties/characteristics
- ① Construct item *profile* of characteristics
- ② Construct item features:
 - ▶ Text: use TF-IDF and use top N features or features over a cutoff
 - ▶ Images: use tags – only works if tags are frequent & accurate
- ③ Compute document similarity: Jaccard, Cosine
- ④ Construct user profile

Item profile

- Consists of (feature, value) pairs
- Consider setting feature to 0 or 1
- Consider how to scale non-Boolean features

User profile

- Describes user preferences (utility matrix)
- Consider how to aggregate item features per user:
 - ▶ Compute “weight” a user puts on each feature
 - ▶ E.g., “Julia Roberts” feature = average rating for films with “Julia Roberts”
- Normalize: subtract average utility per user
 - ▶ E.g., “Julia Roberts” feature = average rating for films with “Julia Roberts” - average rating

Content-based recommendations

- Compute (cosine) distance between user profile and item profiles
- May want to bucket items first using random-hyperplane and locality-sensitivity-hashing (LSH)
- ML approach:
 - ▶ Use random forest or equivalent to predict on a per-user basis
 - ▶ Computationally intensive – usually only feasible for small problems