An Otherwise Blank Slide

Recommender Systems

Overview

- What are Recommender Systems?
- Three common approaches
- A product development story

What

Approaches

Story

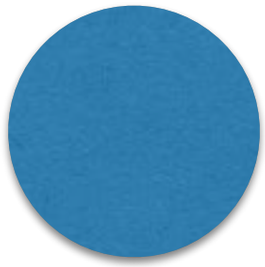

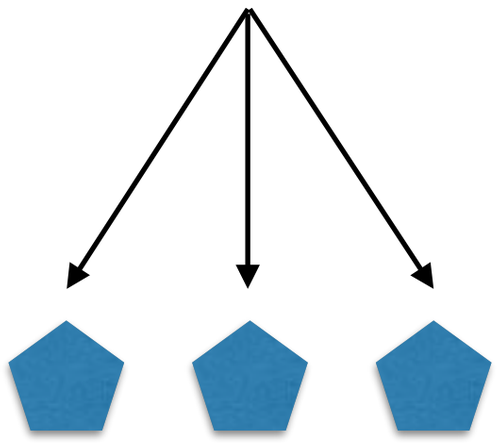

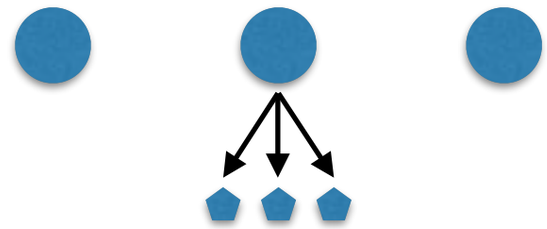

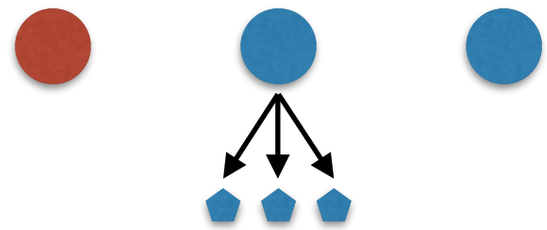

What are Recommender Systems?

What are Recommender Systems?

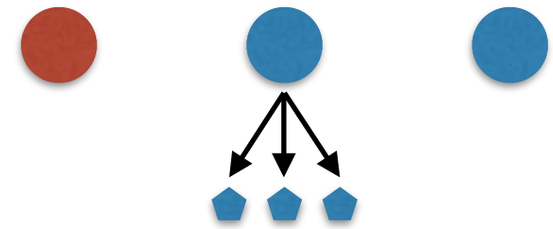

- Goal: Offer a user something they might be interested in
  - What is “interest”?
    - Binary yes/no
      - Here are songs you’ll like
    - Ranked list
    - Collection of items
      - You might want a memory card with that camera
- What can we recommend?
  - another person, place, thing, experience, company, employee
  - dating, travel, shopping, media, HR

What are Recommender Systems?

- Goal: Offer something the user cares about

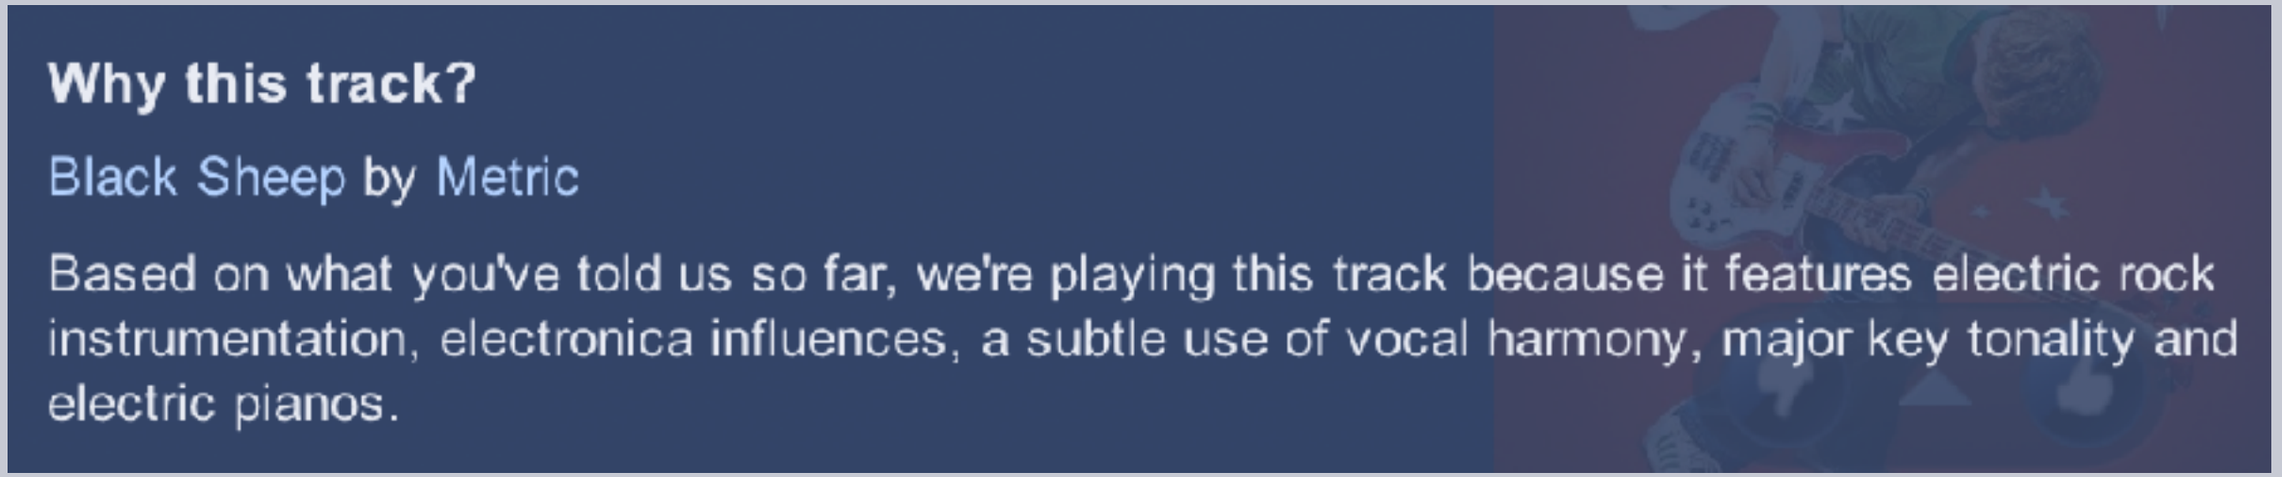

- Field: Blend of AI, HCI, CogPsy
    - personalization, serendipity, motivation, trust, confidence

What are Recommender Systems?

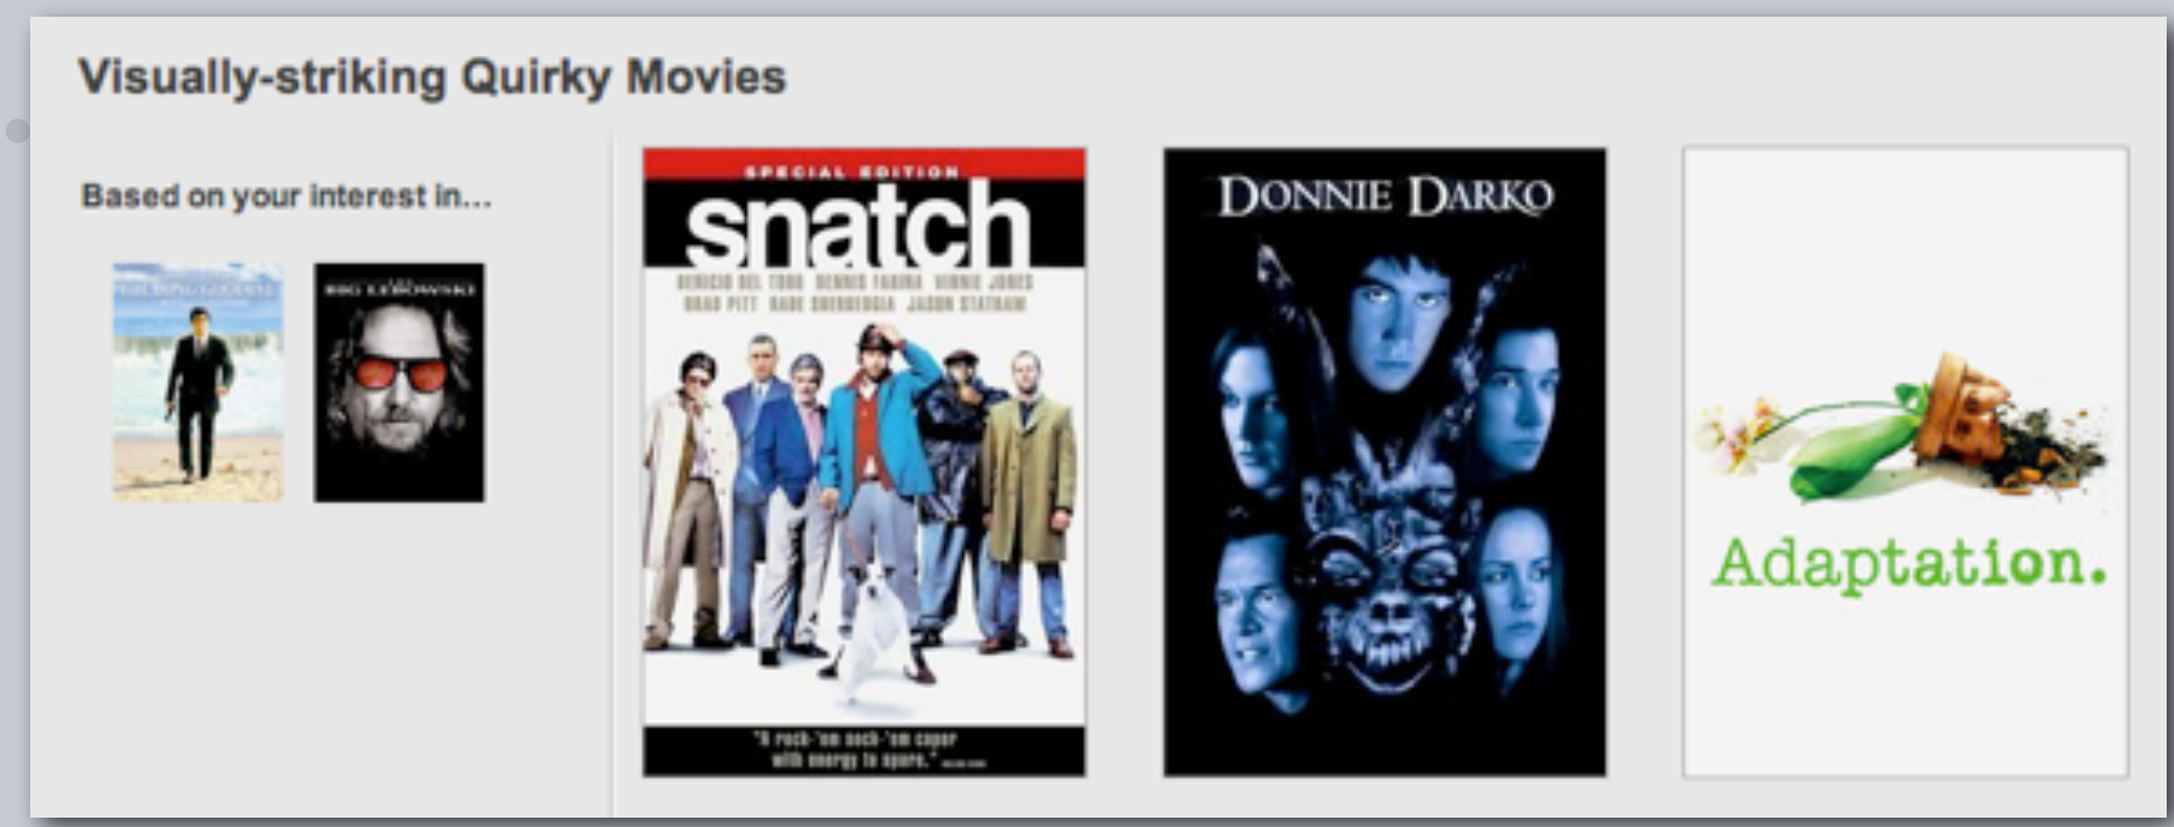

- Field: Blend of AI, HCI, CogPsy
    - personalization, serendipity, motivation, trust, confidence

What are Recommender Systems?

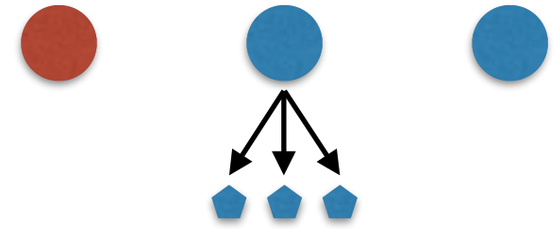

- Field: Blend of CS, AI, HCI, CogPsy
    - Serendipity
    - Personalization
    - Diversity
    - Persistence
    - Modality
    - Privacy
    - Motivation
    - Trust

Kinds of Data

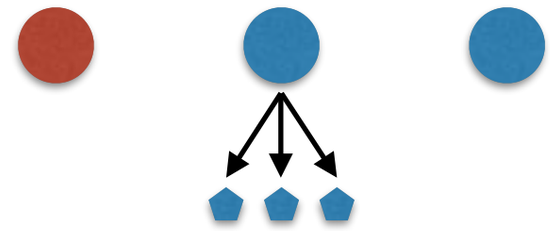

- Target person (recommendee/active user)
  - Are they middle-aged? What's their income? How often do they log in?
- Target item
  - Is it expensive? Does it feature Nicholas Cage?
- User-item interactions
  - Did many users buy/star/like/save/wishlist/rate this item? What was the typical rating?
  - Did this user buy/star/etc many items? What was their typical rating behavior?

Encoding Information

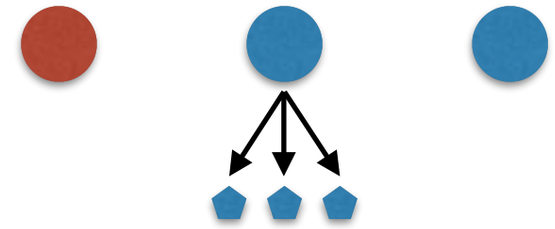

- What does a “like” mean? How about a “view”? How about 10 views? How about a “saved but kept not buying it”? (Did they think it was funny? Are they actually using ‘like’ as a ‘bookmark’ functionality?)
- Explicit v Implicit
    - Did the user take a clear and explicit action?
        - Thumbs up/down on Pandora
            - Did the user, UI/UX designer, and data scientist all interpret this action to be the same thing?
    - Did the user take some actions that might indicate something, but which weren’t explicit?
        - Did they view a product? Did they share it? Did they spend a long time on the page? (Was their mouse active while they were on the page?)
- Incentive pollution
    - Did you ask your friend to review your recommender system/app?

User Utility

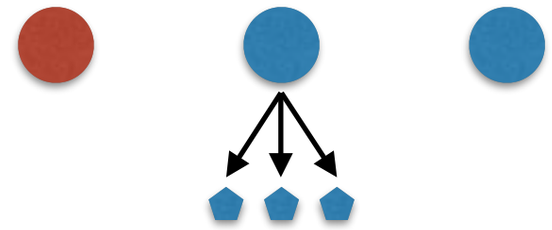

- How and when do the recommendations become useful?
  - As a function of data, some algorithms…
    - rely on a single user rating many items before becoming useful
    - rely on many users rating many items, collaboratively
    - require lots of info about the users themselves
    - have a higher or lower ceiling of learning
  - As a function of time, some algorithms…
    - are easy to update/recalculate, others are very hard
- The approach used when your company/data is small isn't necessarily what you'll use later

Three Common Approaches

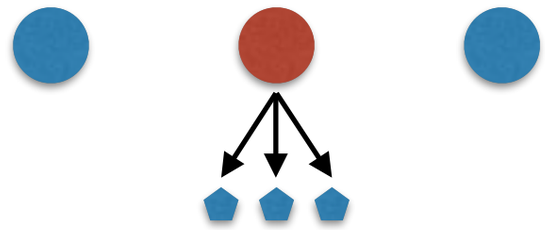

- Many ways to frame the challenge
- Different data requirements, runtime requirements, types of success
- Lots of moving parts; ‘core’ algorithms at the heart of various approaches

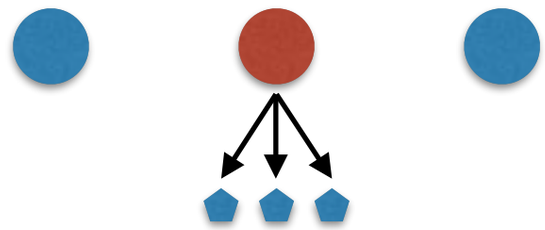

Classification/Regression

Classification/Regression

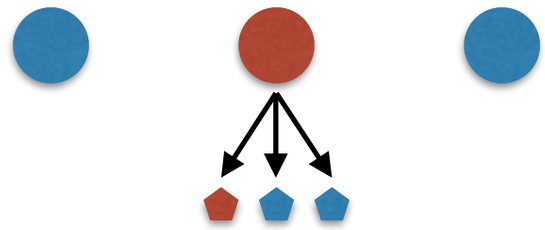

- Classifier to predict interaction/rating, using past interactions/ratings as training data
  - Using features of the people I've liked so far on <dating site> as training data, predict/rank who I'll like
  - Using features of the people who have bought product X, predict/rank who else is likely to buy it

Classification/Regression

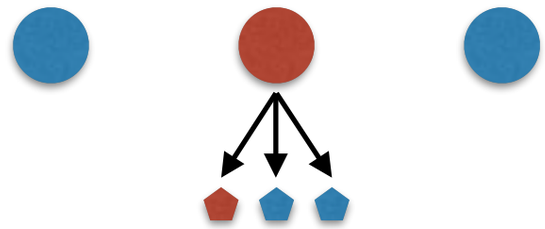

6’2	Blond	5 o’clock shadow	Y
5’11	Brown	Clean-shaven	N
6’01	Black	Gandalf	N
6’0	Blond	Moustache	?

Classification/Regression

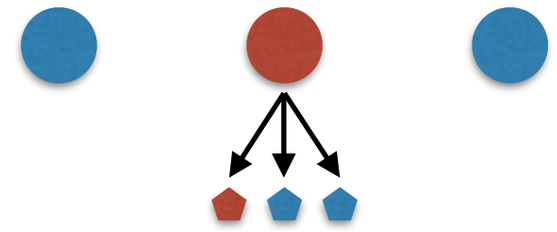

- Advantages
    - Familiar tools/algorithms
    - Well studied success metrics (for one classifier)
        - Diagnose bias/variance, swap algorithms, add features
    - Well studied confidence metrics
    - Potential for very personalized results, learning overall trends
    - Potential for good explanations (confidence&trust)
    - Single-user case is OK
    - Information privacy maintained

Classification/Regression

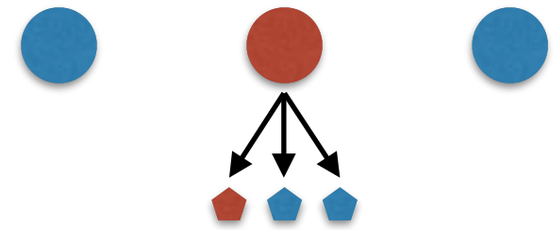

- Limitations
    - Training data required: interactions/ratings (Y)
    - Feature engineering required: meaningful descriptive data (X)
    - Tuning one classifier is tricky; tuning n classifiers?
        - Devops challenge (training, caching)
        - Data Science challenge (success metric for experiments)
- Time to first utility (from User's POV): interactions, compute time
- Learning overall trends may be an oversimplification, compressing modalities
    - Complex decision boundaries can be learned w enough data, still a UX issue
- Rating/Interaction gathering&interpretation creates product restrictions
- Solipsistic (not inherently social)

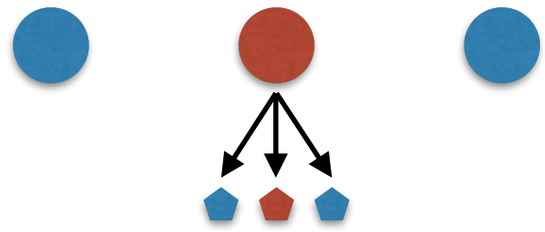

Collaborative Filtering

Collaborative Filtering

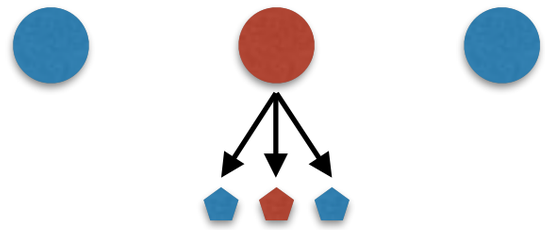

- Use similarity to other users/products to predict interaction/rating, using past interactions/ratings as training data

- To the degree to which I'm similar to Bob, I may feel like he did about this item. Likewise for Carly, Danielle, Edgar…

- I'm similar to Bob to the degree to which we've agreed on other items

Collaborative Filtering

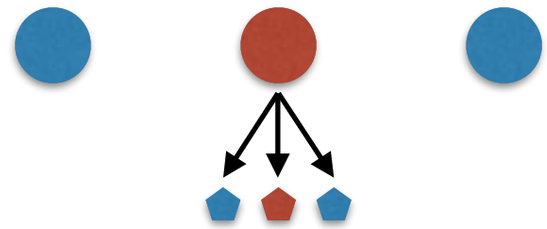

Bob	0	1	1
Carly	1	0	0
Danielle	1	1	1
Dan	0	1	?

Collaborative Filtering

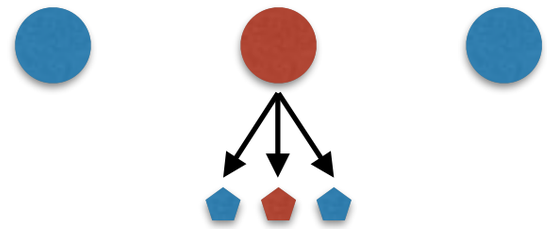

$$prcd(p_i, t_i) = \bar{r}_{p_i} + \frac{\sum_{p_j \in P, i \neq j} sim(p_i, p_j) \times (r_{p_j,t_i} - \bar{r}_{p_j})}{\sum_{p_j \in P, i \neq j} sim(p_i, p_j)}$$

Collaborative Filtering

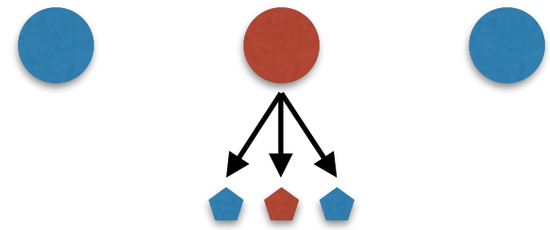

- Variations
  - Neighborhoods
    - Similarity amongst n most similar people/items
      - Faster (can also cache similarity)
  - Flip the matrix: items “choose” people rather than vice versa
  - Use a different similarity function: “content-boosted” collaborative filtering
    - Dan and Bob are similar because they have similar Facebook profiles, rather than because they rated things similarly

Collaborative Filtering

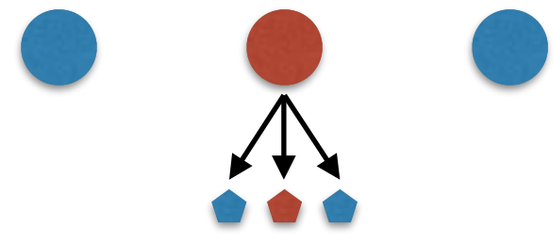

- Advantages
  - Content agnostic
    - Feature engineering not required
  - Inherently social
    - Potential for social explanations
    - Popularity naturally built-in, in a way
- One rating helps many people (“collaborative”)
  - Time to first utility (interactions/ratings)

Collaborative Filtering

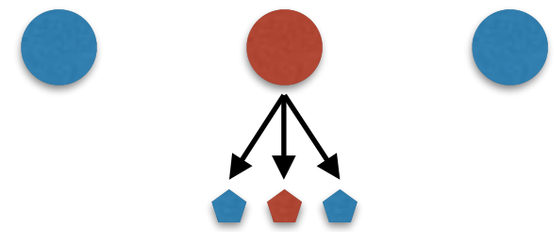

- Limitations
  - Content agnostic
    - May not take advantage of all information known about people/ items
  - Inherently social
    - Potentially violates user privacy expectations
    - Less-rated areas can lead to lower-quality recommendations
- Sparsity … it's really tough
  - Low ratings or overlaps (similarity) can lead to relatively little information going into each prediction
  - High person/item churn can lead to low overlaps
  - Single-user case is a fail case

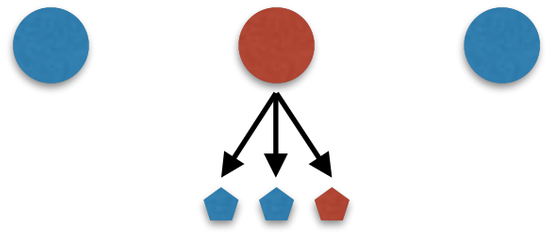

Similarity

Similarity

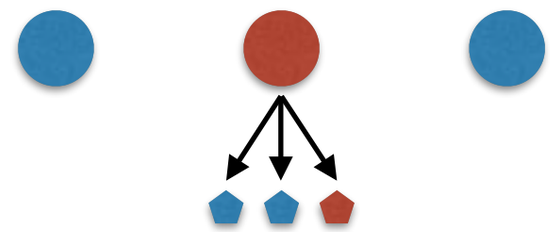

- Go all-in on similarity: do a really good job of figuring out which things are similar to one another
  - Use that similarity to make connections: because I liked A and A is similar to B (in reduced dimensionality subspace), I may like B
  - "Smooth out" similarity space, and make distance more meaningful (anti-Curse of Dimensionality): Dimensionality Reduction
    - Loss of information, nuances may be lost
    - More feature 'overlaps', more ways to compare
    - Depending on alg, can be very explainable (eg LDA topic modeling is relatively easy to interpret)
- Runtime: run a potentially-computationally-expensive offline dimensionality reduction (eg, LDA/topic modeling/NMF), then index similarity in reduced-dimensionality subspace for fast queries

Similarity

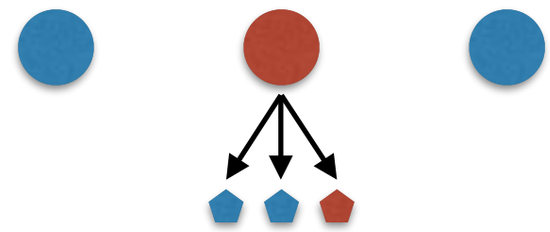

- Advantages
  - User information requirements are very low
    - Interactions/ratings not needed to precompute similarities
    - Time to first utility is 1 rating
  - Potential for high interpretability of similarities and recommendations
  - Modality preserved*
  - Runtime requirements are low
    - Heavy computation can be done offline
  - Single-user use case OK
  - Potentially useful side-effects: similarities, topics, topic distributions

Similarity

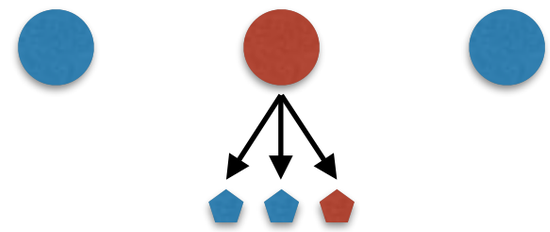

- Limitations
  - Item information requirements fairly high
    - Feature engineering required: meaningful descriptive data about items
  - Information about person not taken into account
  - High compute requirements for dimensionality reduction, indexing
  - Can be difficult to evaluate quality of dimensionality reduction
  - Solipsistic (not inherently social)
  - Overall trends ignored; pointwise recommendations only

Three Common Approaches

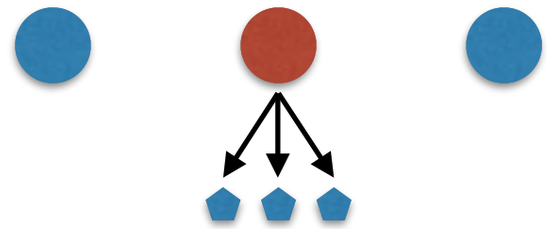

- Classification/Regression
- Collaborative Filtering
- Similarity (Dimensionality Reduction)

Three Common Approaches

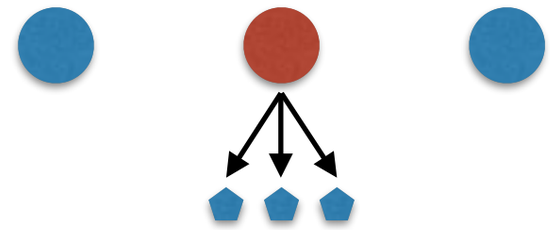

- Recommendation can be framed in many ways
  - These approaches can be stacked & combined & inverted, etc.
    - Reduced dimensions as input to classification?
    - Collaborative Filtering matrix blanks filled in with regression output?
    - Similarity function based on dimensionality reduction?
    - Inversion of person<>item framing?

Three Common Approaches

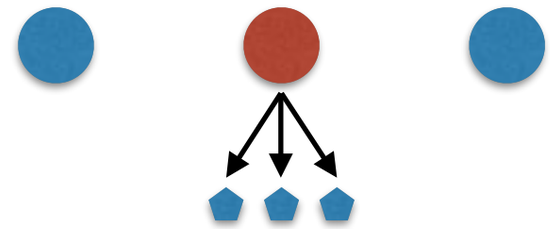

- Additional considerations
    - User state (already owns X, needs Y?)
        - basket analysis
- User preference shift (used to like X; decay?)
- Overall success evaluation (A/B test bad recs?)

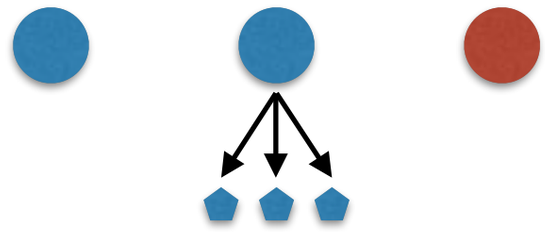

A Story

Thanks!