

Bias/Variance and Cross-Validation

Review: Linear Regression

Overfitting and Underfitting

The Bias/Variance Tradeoff

Cross-Validation

K-fold Cross-Validation

Subset Selection of Predictors

Bias and Variance

Bias and Variance

One Goal of Data Science: Make Future Predictions

One goal is to make accurate *predictions* on future (unseen) data.

1. Define a business goal.

e.g. make Tesla cars the most dependable vehicles on the market

2. Collect training data.

e.g. Tesla cars' event logs + historical record of parts replaced

3. Train a model.

e.g. **features:** event statistics, **target:** time until failure

4. Deploy the model.

e.g. monitor cars' events in real time, send mechanics to replace parts that will soon fail

Review: Linear Regression

We assume the world is built on linear relationships. Under that assumption, we can model the relationship between *features* and a *target* like this:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p + \epsilon$$

Review: Linear Regression

We assume the world is built on linear relationships. Under that assumption, we can model the relationship between *features* and a *target* like this:


$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p + \epsilon$$

The diagram illustrates the components of the linear regression equation $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p + \epsilon$. Arrows of different colors point from descriptive labels to specific terms in the equation:

- A yellow arrow points from "target (aka, outcome)" to Y .
- Four blue arrows point from "model parameters (aka, coefficients)" to $\beta_0, \beta_1, \beta_2,$ and β_p .
- Three green arrows point from "features (aka, predictors)" to $X_1, X_2,$ and X_p .
- A red arrow points from "sampling error" to ϵ .

Review: Linear Regression

We can make linear regression non-linear by inserting extra “interaction” features or higher-order features.




Review: Linear Regression

We can make linear regression non-linear by inserting extra “interaction” features or higher-order features.

Examples:

$$\text{Sickliness} = \beta_0 + \beta_1 * \text{age}$$

$$\text{Sickliness} = \beta_0 + \beta_1 * \text{age} + \beta_2 * \text{age}^2$$



Is R^2 all that matters?


We *could* just keep inserting interaction features until $R^2 = 1$.

Boom. I solved data science. Here's my idea:

```
def train_super_awesome_perfect_model(X, y):  
    while True:  
        model = LinearRegression()  
        model.fit(X, y)  
        if calculate_r2(model, X, y) >= 0.999:  
            return model  
        else:  
            X = insert_random_interaction_feature(X)
```

Why is this a
bad idea?

Oh the woes of overfitting...



Underfitting and Overfitting

Underfitting: The model doesn't fully capture the relationship between predictors and the target. The model has *not* learned the data's signal. The model is not flexible enough.

→ What should we do if our model underfits the data?

Overfitting: The model has tried to capture the sampling error. The model has learned the data's signal *and* the noise. I.e., the model attributes to signal that which is truly noise. The model is too flexible.

→ What should we do if our model overfits the data?

The Bias/Variance Tradeoff

We assume the true predictor/target relationship is given by an unknown function plus some sampling error:

$$Y = f(X) + \epsilon$$

We estimate the true (unknown) function by fitting a model over the training set.

$$\hat{Y} = \hat{f}(X)$$

Let's evaluate this model using a test observation $(\mathbf{x}_0, \mathbf{y}_0)$ drawn from the population. What is the model's expected squared prediction error on this test observation?

$$E[(y_o - \hat{f}(x_0))^2] = \dots$$

The Bias/Variance Tradeoff

The Bias/Variance Tradeoff

The Bias/Variance Tradeoff

The Bias/Variance Tradeoff

The Bias/Variance Tradeoff

Our model's expected squared prediction error will depend on (1) the variability of \mathbf{y}_0 and (2) the variability of the training set used to train our model. We can break this into three pieces:

$$E[(y_o - \hat{f}(x_0))^2] = \dots = \text{Var}(\hat{f}(x_0)) + \text{Bias}^2(\hat{f}(x_0)) + \text{Var}(\epsilon)$$


The variance of our model's prediction of \mathbf{x}_0 over all possible training sets

The difference between the true prediction and our model's average prediction over all possible training sets

The variance of the irreducible error.

$$\text{Bias}(\hat{f}(x_0)) = E[\hat{f}(x_0)] - f(x_0)$$

The Bias/Variance Tradeoff




How is the **bias/variance tradeoff** related to **underfitting** and **overfitting**?

How can we find the best tradeoff point?
I.e. The optimum model complexity

Cross-Validation

Main idea: **Don't** use all your data for training.

Instead: **Split** your data into a “training set” and a “validation set”.



Cross-Validation

1. Split your data into training/validation sets.

70/30 or 90/10 splits are commonly used

2. Use the training set to train several models of varying complexity.

e.g. linear regression (w/ and w/out interaction features), neural nets, decision trees, etc.


(we'll talk about hyperparameter tuning, grid search, and feature engineering later)

3. Evaluate each model using the validation set.


calculate R^2 , MSE, accuracy, or whatever you think is best

4. Keep the model that performs best over the validation set.

Let's predict MPG from horsepower



Cross-Validation Example



You will see
this shape all
the time!

You will wrestle
with the
bias/variance
tradeoff
constantly...

E.g. linear regression w/
varying degree of
polynomial


Recall our goal: Making accurate future predictions

Fitting the training set perfectly is *easy*.

How?

Fitting future (unseen) data is *not easy*.

Cross validation helps us choose a model that performs well on unseen data.




k-Fold Cross-Validation

1. Split the dataset into k “folds”.
2. Train using $(k-1)$ folds.
Validate using the one left-out fold. Record a validation metric such as RSS or accuracy.
3. Train k models, leaving out a different fold for each one.
4. Average the validation results.

Commonly, $k=5$ or $k=10$

randomize!




Leave-one-out Cross-Validation

Assume we have n training examples.

A special case of k -fold CV is when $k=n$. This is called *leave-one-out cross-validation*.


Useful (only) if you have a tiny dataset where you can't afford a large validation set.



Overfitting in high dimensions is easy, even with simple models.

If our data has high dimensionality (many many predictors), then it becomes easy to overfit the data.

Even linear regression might be too complex of a model for high dimensional data (and the smaller the dataset, the worse this problem is).



Linear regression in high dimensions

“HELP, my model is overfitting!”

You have a few options.

1. **Get more data:** not always possible/practical
2. **Subset Selection:** keep only a subset of your predictors (i.e, dimensions)
3. **Regularization:** restrict your model's parameter space
4. **Dimensionality Reduction:** project the data into a lower dimensional space

Subset Selection

Best subset: Try every model. Every possible combination of p predictors

Computationally intensive. 2^p possible subsets of p predictors

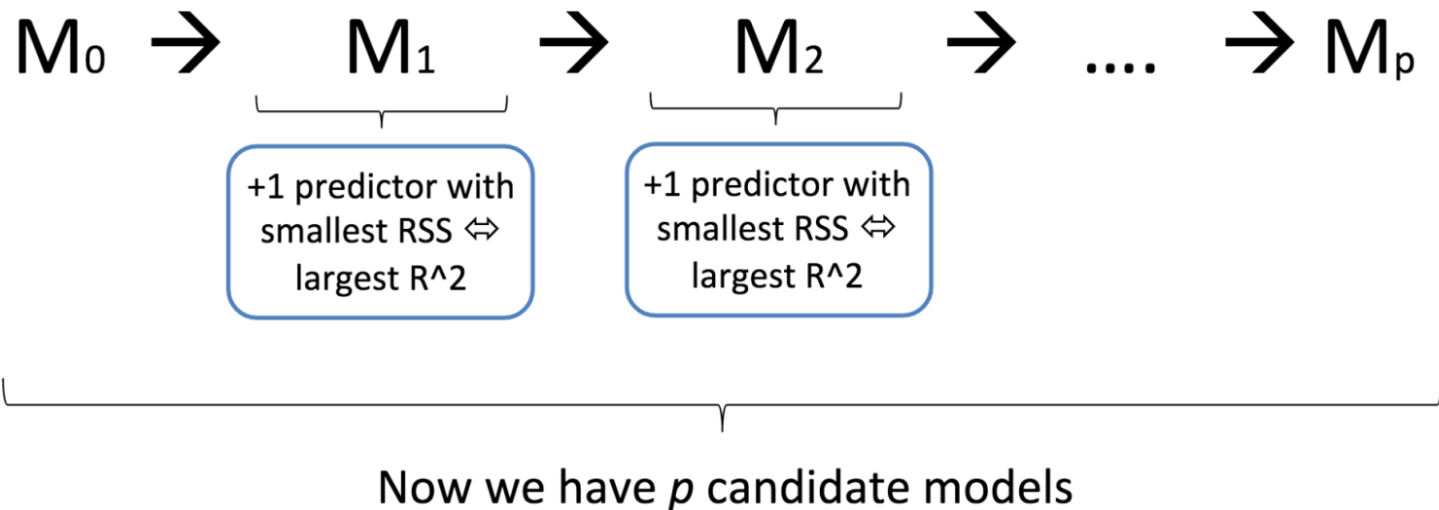
High chance of finding a “good” model by random chance.

... A sort-of monkeys-Shakespeare situation ...

Stepwise: Iteratively pick predictors to be in/out of the final model.

Forward, backward, forward-backward strategies

Forward Stepwise Selection



Are RSS and R^2 good ways to decide amongst the resulting p candidates?

Answer: Don't use RSS or R^2 for this part.
Use Mallows's C_p , or AIC, or BIC, or Adjusted R^2 .

... or just use cross-validation with any error measurement.

Subset Selection: Comparing models of varying number of predictors...

$$C_p = \frac{1}{n}(RSS + \underline{2p}\hat{\sigma}^2)$$

Mallow's C_p

p is the total # of parameters

$\hat{\sigma}^2$ is an estimate of the variance of the error, ε

$$AIC = -2\log L + 2 \cdot \underline{p}$$

L is the maximized value of the likelihood function for the model estimated

$$BIC = \frac{1}{n}(RSS + \log(n)\underline{p}\hat{\sigma}^2)$$

This is C_p , except 2 is replaced by $\log(n)$.
 $\log(n) > 2$ for $n > 7$, so BIC generally exacts a heavier penalty for more variables

$$\text{Adjusted } R^2 = 1 - \frac{RSS/(n - \underline{p} - 1)}{TSS/(n - 1)}$$

Similar to R^2 , but pays price for more variables

Side Note: Can show AIC and Mallow's C_p are equivalent for linear case

Subset Selection: Comparing models of varying number of predictors...

Dep. Variable:	y	R-squared:	0.933			
Model:	OLS	Adj. R-squared:	0.928			
Method:	Least Squares	F-statistic:	211.8			
Date:	Mon, 03 Nov 2014	Prob (F-statistic):	6.30e-27			
Time:	14:45:06	Log-Likelihood:	-34.438			
No. Observations:	50	AIC:	76.88			
Df Residuals:	46	BIC:	84.52			
Df Model:	3					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[95.0% Conf. Int.]	
x1	0.4687	0.026	17.751	0.000	0.416	0.522
x2	0.4836	0.104	4.659	0.000	0.275	0.693
x3	-0.0174	0.002	-7.507	0.000	-0.022	-0.013
const	5.2058	0.171	30.405	0.000	4.861	5.550
Omnibus:	0.655		Durbin-Watson:		2.896	
Prob(Omnibus):	0.721		Jarque-Bera (JB):		0.360	
Skew:	0.207		Prob(JB):		0.835	
Kurtosis:	3.026		Cond. No.		221.	