



Bias/Variance and Cross-Validation

J.F. Omhover

slides of Ryan Henning

- Explain the bias/variance tradeoff
- Define overfitting and underfitting
- Relate flexibility and complexity of a model to bias / variance, to overfitting / underfitting, and to training / testing error
- Define and implement cross-validation on a dataset



Quick Review: Regression vs. Classification

(in machine learning)

What is regression?

Use features to predict real valued targets. E.g. predict future sales/revenue

What is classification?

Use features to predict categorical targets. E.g. predict yes/no, male/female, 0-9



One goal is to make accurate **predictions** on future (unseen) data.

1. Define a business goal.

e.g. make Tesla cars the most dependable vehicles on the market

2. Collect training data.

e.g. Tesla cars' event logs + historical record of parts replaced

3. Train a model.

e.g. **features**: event statistics, **target**: time until failure

4. Deploy the model.

e.g. monitor cars' events in real time, send mechanics to replace parts that will soon fail

Questions!

Review: Linear Regression

We assume the world is built on linear relationships. Under that assumption, we can model the relationship between *features* and a *target* like this:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p + \epsilon$$

The diagram illustrates the components of the linear regression equation $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p + \epsilon$. Arrows of different colors point from descriptive text to specific terms in the equation:

- A yellow arrow points from "target (aka, outcome)" to Y .
- Four blue arrows point from "model parameters (aka, coefficients)" to $\beta_0, \beta_1, \beta_2,$ and β_p .
- Three green arrows point from "features (aka, predictors)" to $X_1, X_2,$ and X_p .
- A red arrow points from "sampling error" to ϵ .

Review: Linear Regression




We can make linear regression non-linear by inserting extra “interaction” features or higher-order features.

Example:

$$Y = \beta_0 + \beta_1 * \text{age}$$

$$Y = \beta_0 + \beta_1 * \text{age} + \beta_2 * \text{age}^2$$






We *could* just keep inserting interaction features until $R^2 = 1$.

Boom. I solved data science. Here's my idea:

```
def train_super_awesome_perfect_model (X, y):  
    while True:  
        model = LinearRegression()  
        model.fit(X, y)  
        if calculate_r2(model, X, y) >= 0.999:  
            return model  
        else:  
            X = insert_random_interaction_feature(X)
```

Why is this a
bad idea?

Oh the woes of overfitting...



What's bad about the first model?

What's bad about the second model?

What's bad about the third model?



Underfitting and Overfitting

Underfitting: The model doesn't fully capture the relationship between predictors and the target. The model has *not* learned the data's signal.

→ What should we do if our model underfits the data? (assume using lin. reg.)

Overfitting: The model has tried to capture the sampling error. The model has learned the data's signal *and* the noise.


→ What should we do if our model overfits the data? (harder... any guesses?)

The Bias/Variance Tradeoff

Boardwork... build intuition...

Let's get an intuitive feel for the bias and the variance of a model... we'll see more math on the next slide.

Note: **Bias** and **Variance** are terms you will use A TON as a data scientist! Exciting times!





We assume the true predictor/target relationship is given by an unknown function plus some sampling error:

$$Y = f(X) + \epsilon$$

We estimate the true (unknown) function by fitting a model over the training set.

$$\hat{Y} = \hat{f}(X)$$

Let's evaluate this model using a test observation $(\mathbf{x}_o, \mathbf{y}_o)$ drawn from the population. What is the model's expected squared prediction error on this test observation?

$$E[(y_o - \hat{f}(x_o))^2] = \dots$$

Our model's expected squared prediction error will depend on (1) the variability of \mathbf{y}_0 and (2) the variability of the training set used to train our model. We can break this into three pieces:

$$E[(y_o - \hat{f}(x_0))^2] = \dots = \text{Var}(\hat{f}(x_0)) + \text{Bias}^2(\hat{f}(x_0)) + \text{Var}(\epsilon)$$


The variance of our model's prediction of \mathbf{x}_0 over all possible training sets

The difference between the true target and our model's average prediction over all possible training sets

The variance of the irreducible error.

$$\text{Bias}(\hat{f}(x_0)) = E[\hat{f}(x_0)] - f(x_0)$$

The Bias/Variance Tradeoff



How is the **bias/variance tradeoff** related to **underfitting** and **overfitting**?

How can we find the best tradeoff point? I.e. The optimum model complexity




QUESTION TIME !

1. What are the ways we can increase complexity of a model?
2. Why does variance increase and bias decrease as model complexity increases?
3. How do you determine the optimal complexity to use?

Cross-Validation

Main idea: **Don't use all your data for training.**

Instead: **Split your data into a “training set” and a “validation set”.**






Cross-Validation

1. Split your data into training/validation sets.
70/30 or 90/10 splits are commonly used
2. Use the training set to train several models of varying complexity.
e.g. linear regression (w/ and w/out interaction features), neural nets, decision trees, etc.
(we'll talk about hyperparameter tuning, grid search, and feature engineering later)
3. Evaluate each model using the validation set.
calculate R^2 , MSE, accuracy, or whatever you think is best
4. Keep the model that performs best over the **validation** set.

Let's predict MPG from horsepower



Cross-Validation Example



You will see this shape all the time!

You will wrestle with the bias/variance tradeoff constantly...

E.g. linear regression w/ varying degree of polynomial

Recall our goal: Making accurate future predictions




Fitting the training set perfectly is *easy*.

How?

Fitting future (unseen) data is *not easy*.


Cross validation helps us choose a model that performs well on unseen data.



1. Split the dataset into k “folds”.
2. Train using $(k-1)$ folds. Validate using the one “leave out” fold. Record a validation metric such as RSS or accuracy.
3. Train k models, leaving out a different fold for each one.
4. Average the validation results.

Commonly, $k=5$ or $k=10$


randomize!



Assume we have n training examples.

A special case of k -fold CV is when $k=n$. This is called *leave-one-out cross-validation*.

Useful (only) if you have a tiny dataset where you can't afford a large validation set.




Overfitting in high dimensions is easy, even with simple models.



If our data has high dimensionality (many many predictors), then it becomes easy to overfit the data.

This is one result of the so-called **Curse of Dimensionality** (look it up).

Even linear regression might be too complex of a model for high dimensional data (and the smaller the dataset, the worse this problem is).



Linear regression in high dimensions



QUESTION TIME !

1. Why do we often prefer to use K-Fold Cross Validation instead of a simple train test split?
2. How many models are created in a 5-fold cross validation to compare two models (model A and model B)?
3. What do we do with all the models we created and how do we determine what single model to keep?



You have a few options.

1. Get more data... (not usually possible/practical)
2. **Subset Selection:** keep only a subset of your predictors (i.e, dimensions)
3. **Regularization:** restrict your model's parameter space
4. **Dimensionality Reduction:** project the data into a lower dimensional space

Subset Selection

Best subset: Try every model. Every possible combination of p predictors

- Computationally intensive. 2^p possible subsets of p predictors
- High chance of finding a “good” model by random chance.
... A sort-of monkeys-Shakespeare situation ...

Stepwise: Iteratively pick predictors to be in/out of the final model.

- Forward, backward, forward-backward strategies

$M_0 \rightarrow M_1 \rightarrow M_2 \rightarrow \dots \rightarrow M_p$

+1 predictor with
smallest RSS \Leftrightarrow
largest R^2

+1 predictor with
smallest RSS \Leftrightarrow
largest R^2

Now we have p candidate models

Are RSS and R^2 good ways to decide amongst the resulting $(p+1)$ candidates?

Answer: Don't use RSS or R^2 for this part.
Use Mallows's C_p , or AIC, or BIC, or Adjusted R^2 .

... or just use cross-validation with any error measurement.

$$C_p = \frac{1}{n}(RSS + \underline{2p}\hat{\sigma}^2)$$

Mallow's C_p

p is the total # of parameters

$\hat{\sigma}^2$ is an estimate of the variance of the error, ε

$$AIC = -2\log L + 2 \cdot \underline{p}$$

L is the maximized value of the likelihood function for the model estimated

$$BIC = \frac{1}{n}(RSS + \log(n)\underline{p}\hat{\sigma}^2)$$

This is C_p , except 2 is replaced by $\log(n)$.
 $\log(n) > 2$ for $n > 7$, so BIC generally exacts a heavier penalty for more variables

$$\text{Adjusted } R^2 = 1 - \frac{RSS/(n - \underline{p} - 1)}{TSS/(n - 1)}$$

Similar to R^2 , but pays price for more variables

Side Note: Can show AIC and Mallow's C_p are equivalent for linear case



You have a few options.

1. Get more data... (not usually possible/practical)
2. **Subset Selection:** keep only a subset of your predictors (i.e, dimensions)
3. **Regularization:** restrict your model's parameter space
4. **Dimensionality Reduction:** project the data into a lower dimensional space

TO BE CONTINUED...