

Introduction to Spark II

Galvanize



Introduction to Spark II



OBJECTIVES

- **Explain** the differences between RDDs and DataFrames, and where Spark development is moving
- **Explain** what persisting/caching an RDD means, and situations where this is useful
- **Define** an out of memory error and why it happens
- **Describe** the difference between narrow and wide transformations
- **Discuss** next steps to expanding Spark to new domains

Spark is moving from RDDs to DataFrames

- Since Spark 2.0, the emphasis has shifted
- MLlib (RDD-based) -> ML (DF-based)
- GraphX (RDD-based) -> GraphFrames (DF-based)
- Focus on DataFrames for future proofing



SparkSQL



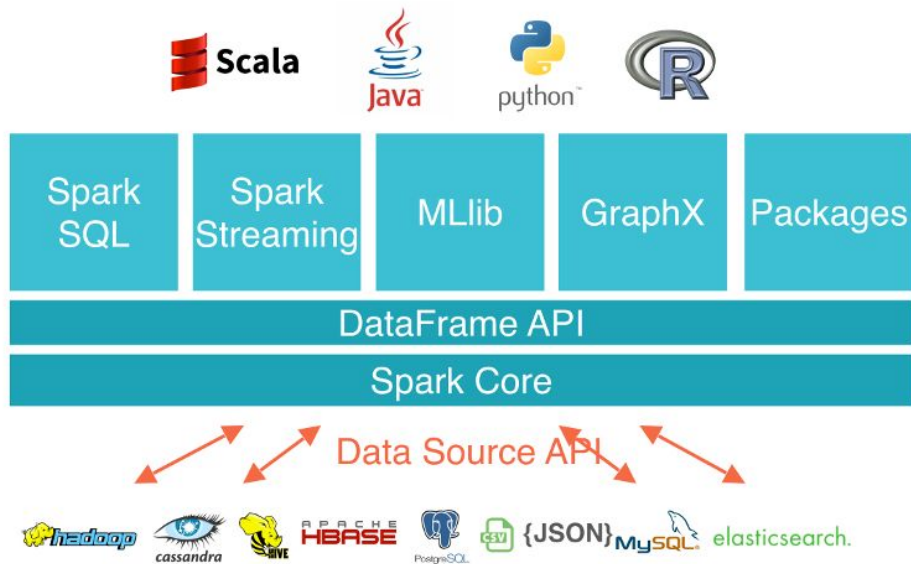
Introducing SparkSQL

- An in-memory database
- Access it using `SparkSession`
- The **primary abstraction is a DataFrame**
- Allows for the execution of SQL commands

```
import pyspark as ps

conf = ps.SparkConf() \
    .setAppName("My App") \
    .setMaster("local[*]")
sc = ps.SparkContext(conf=conf)

spark = ps.sql.SparkSession(sc)
```



What is a DataFrame?



DataFrames are...

- Immutable collections of data (like RDDs)
- Distributed across nodes in a cluster
- **Organized into named columns**
 - not schema-less rows, like RDDs
 - Schema = Table Names + Column Names + Column Types
- Think of them like **an RDD with a schema**

Why are they useful?

- They make large data processing easier
- Allow developers to formalize the structure of the data
- Performance parity
 - Unlike RDDs, which are slower on Python than Scala or Java

col count	
iHeartAwards	27299
BestFanArmy	19892
BestFans2017	12534
OneDBestFans	10829
GagaBestFans	9164
YOU_NEVER_WALK_ALONE	8229
BTS	7566
CamilaBestFans	7224
Lovatics	6491
HappyBirthdayHarry	5391
NOW2016	5345
BlackHistoryMonth	5042
RTした人全員フォローする	4855
ALDUB81stWeeksary	4567
ALDUBLoveMonth	4513
BestMusicVideo	4435
PBBPADALUCKMAYMAY	4060
ツインテールの日	4029
사설토토사이트추천	3933
gameinsight	3796

DataFrame Speed Comparison I



DataFrames are **fast**

- Imposing a schema allows for optimization
- **Catalyst Optimizer** is at the heart of Spark SQL
 - Eases the addition of new optimizations
- **Project Tungsten** improves memory/CPU use




Figure 3-1. Relative performance for RDD versus DataFrames based on SimplePerfTest computing aggregate average fuzziness of pandas


Source: *High Performance Spark*

DataFrame Speed Comparison II



RDDs/Python are **slow**

- There's an overhead between Python and the JVM
- RDDs (especially older versions) suffered from this
- Python is built to be slow (single threaded, not type safe)



Source: [Databricks](#)

Creating DataFrames



```
# read JSON
df = spark.read.json('data/sales.json')

# prints the schema
df.printSchema()

# some functions are still valid
print("line count: {}".format(df.count()))

# show the table in a oh-so-nice format
df.show()
```


Creating and Querying DataFrames



```
# read CSV
df_sales = spark.read.csv('data/sales.csv',
                           header=True,      # use headers or not
                           quote='\"',      # char for quotes
                           sep=";",        # char for separation
                           inferSchema=True) # do we infer schema or
                                           # not ?

# Now create an SQL table and issue SQL queries against it without
# using the sqlContext but through the SparkSession object.
# Creates a temporary view of the DataFrame
df_sales.createOrReplaceTempView("sales")

result = spark.sql('''
    SELECT state, AVG(amount) as avg_amount
    FROM sales
    GROUP BY state
''')
result.show()
```

```
# import the many data types
from pyspark.sql.types import *

# create a schema of your own
schema = StructType( [
    StructField('id',IntegerType(),True),
    StructField('date',StringType(),True),
    StructField('store',IntegerType(),True),
    StructField('state',StringType(),True),
    StructField('product',IntegerType(),True),
    StructField('amount',FloatType(),True) ] )

# feed that into a DataFrame
df = spark.createDataFrame(rdd_sales,schema)

# show the result
df.show()

# print the schema
df.printSchema()
```

Caching/Persisting DFs and RDDs

Use the `.cache()` method to cache

- This keeps a copy of your data at that point in time
- When an action is called Spark figures out the answer and then throws away all the data
- For multiple passes across your data (e.g. iterative algorithms), cache your results

Use the `.persist()` method to cache to disk

Level	Meaning
MEMORY_ONLY	Same as cache()
MEMORY_AND_DISK	Cache in memory then overflow to disk
MEMORY_AND_DISK_SER	Like above; in cache keep objects serialized instead of live
DISK_ONLY	Cache to disk not to memory

```
import random
num_count = 500*1000
num_list = [random.random() for i in range(num_count)]
rdd1 = sc.parallelize(num_list)
rdd2 = rdd1.sortBy(lambda num: num)
```

```
%time rdd2.count()
%time rdd2.count()
%time rdd2.count()
%time rdd2.count()
%time rdd2.count()
```

```
CPU times: user 8.34 ms, sys: 1.25 ms, total: 9.59 ms
Wall time: 1.25 s
CPU times: user 9.04 ms, sys: 1.79 ms, total: 10.8 ms
Wall time: 515 ms
CPU times: user 8.93 ms, sys: 1.5 ms, total: 10.4 ms
Wall time: 608 ms
CPU times: user 6.48 ms, sys: 1.37 ms, total: 7.85 ms
Wall time: 616 ms
CPU times: user 7.06 ms, sys: 1.54 ms, total: 8.6 ms
Wall time: 563 ms
```

500000

```
rdd2.cache()
%time rdd2.count()
%time rdd2.count()
%time rdd2.count()
%time rdd2.count()
%time rdd2.count()
```

```
CPU times: user 5.13 ms, sys: 1.1 ms, total: 6.23 ms
Wall time: 581 ms
CPU times: user 6.12 ms, sys: 1.5 ms, total: 7.62 ms
Wall time: 79.8 ms
CPU times: user 5.34 ms, sys: 1.63 ms, total: 6.97 ms
Wall time: 110 ms
CPU times: user 8.43 ms, sys: 2.13 ms, total: 10.6 ms
Wall time: 135 ms
CPU times: user 5.51 ms, sys: 1.28 ms, total: 6.79 ms
Wall time: 97.1 ms
```


500000

Narrow vs Wide Dependencies I




Narrow dependencies:

Each partition of the parent RDD is used by at most one partition of the child RDD.



Wide dependencies:

Each partition of the parent RDD may be depended on by multiple child partitions.



Narrow vs Wide Dependencies II




A ShuffleHashJoin is the Spark default

- Wide dependencies
- Maps data in two DFs -> shuffles based on field in join condition
-> reduces to join two datasets
- Best when evenly distributed data across cluster
- **Joins are expensive**

A BroadcastHashJoin can be more performant

- Narrow dependencies
- Broadcasts a small DF to each of the larger partitions
- Use when the smaller DF is small enough to fit on a single machine
- Minimizes data transfer

Shuffle Hash Join



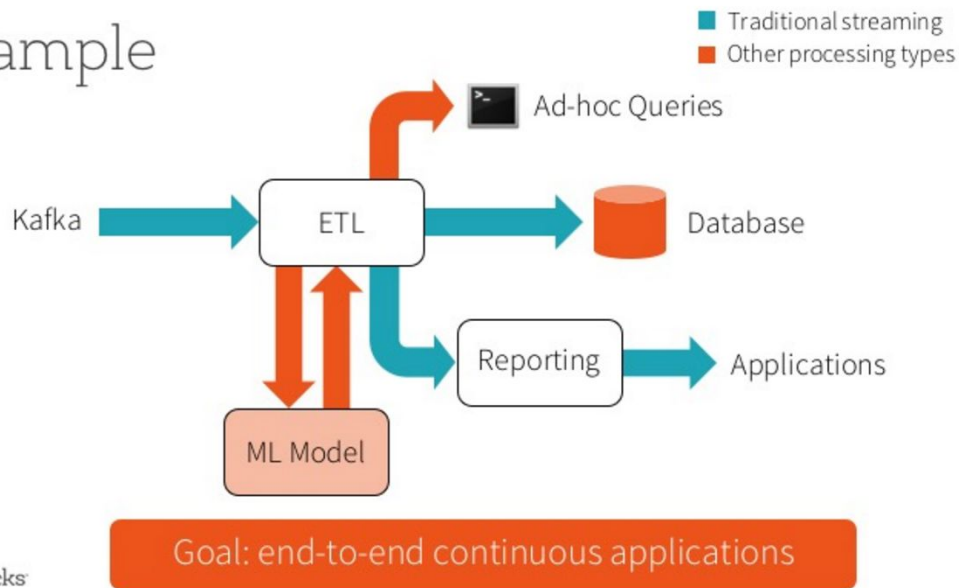
The Spark Ecosystem II



Spark can be the **center of your data product**

- Extract Transform Load (ETL)
- Ad-hoc queries (OLAP)
- Transactions (OLTP)
- Stream processing
- Check out the SMACK stack

Example



Source: [Databricks](https://databricks.com)



Possible **future avenues** for understanding Spark

- Probabilistic data structures
- Difference in local and distributed ML algorithms
- Look into the [SMACK stack](#)
- Learn Scala to maximize your use of Spark
 - Allows for Datasets API
 - Python can't use Datasets because it's not type safe
- Contribute to Spark core, or [community projects](#)



- [Spark Documents](#): Best stop for the most up to date (and versioned) information
- [JF's Walkthroughs](#): Galvanize instructor who made some helpful walkthroughs
- [Learning PySpark](#): An up-to-date intro to the python API to Spark. Good treatment of machine learning
- [Learning Spark](#): A good introduction to Spark written on Spark 1.X. Parts of it are still relevant for a high level overview, but most of it is outdated
- [High Performance Spark](#): More recent publication by the author of *Learning Spark*
- [Databrick's explanation of key terms](#)

Questions?

