

# SQL

Storing and querying structured data

# By the end of this lecture, you should be able to:

- Given some data, say whether or not a relational database would be appropriate.
- Explain what structured data is
- Relate relational database concepts to OOP concepts
- Connect to a database through a client
- Use a SELECT query to get data from a single entity
  - Limit the set of rows using WHERE
  - Perform aggregate functions using GROUP BY
- Get information from multiple entities
  - Simple in-line combinations using subclauses
  - Complicated combinations using JOIN
- Compose SELECT, WHERE, GROUP BY, JOIN into single query

# How to store data?

- The contents of Tolstoy's War and Peace? In Russian? 紅樓夢 by 曹雪芹? In Chinese?
- The temperature at SFO every minute over a 100 year span?
- The location and size of every window currently open on your computer?
- The current inventory of the Trader Joe's at 4th and Market?
- The bus and train (and cable car!) schedules for SF MUNI?

# Data stores

- Text files. (ASCII-encoded? UTF-8 encoded?)
- Arrays of floats. (Persisted? How?)
- In-memory data structure. (Persisted? Pickled?)
- CSV file? JSON file? How many CSV files?
- Database servers? What kind? **Relational??**

# Relational Database Management System (RDBMS)

- Bigish
- Persistent
- Structured

# Structured data

- Entities

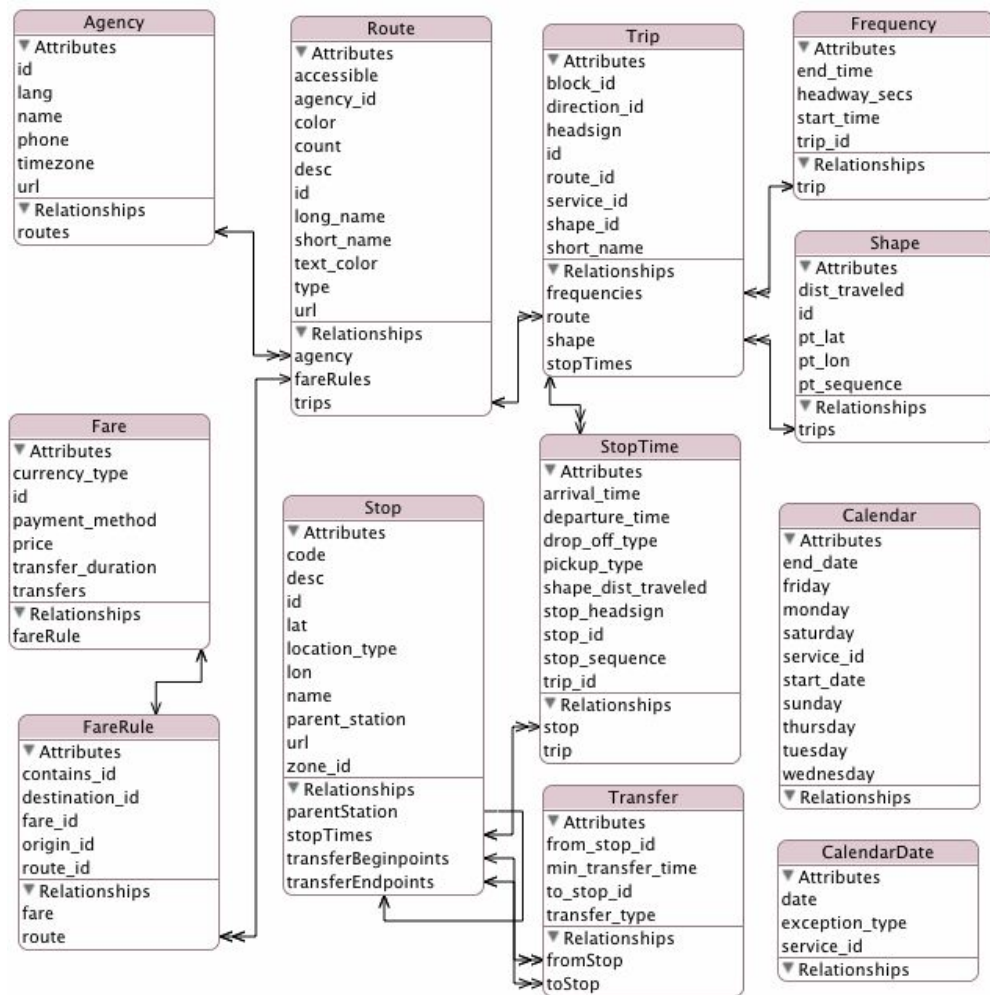
- Similar to classes
- -methods
- +flat namespace
- +ids

- External references

- Reference instances of other entities
  - E.g., “Product Count” references “Store”, “Product”
- Can reference own entity
  - Person has attribute ‘mother’, which is a Person

- Data validity

- Everyone has a mother.



# Connection to OOP concepts


- Class  $\leftrightarrow$  Table
- Instance  $\leftrightarrow$  Row
- Attribute  $\leftrightarrow$  Column
- Data Model  $\leftrightarrow$  Schema
- ???  $\leftrightarrow$  Methods
- Joins  $\leftrightarrow$  ???
- Indexing  $\leftrightarrow$  ???

**Warning:** Object-Relational (ORM) Systems. Tempting, but irritating.

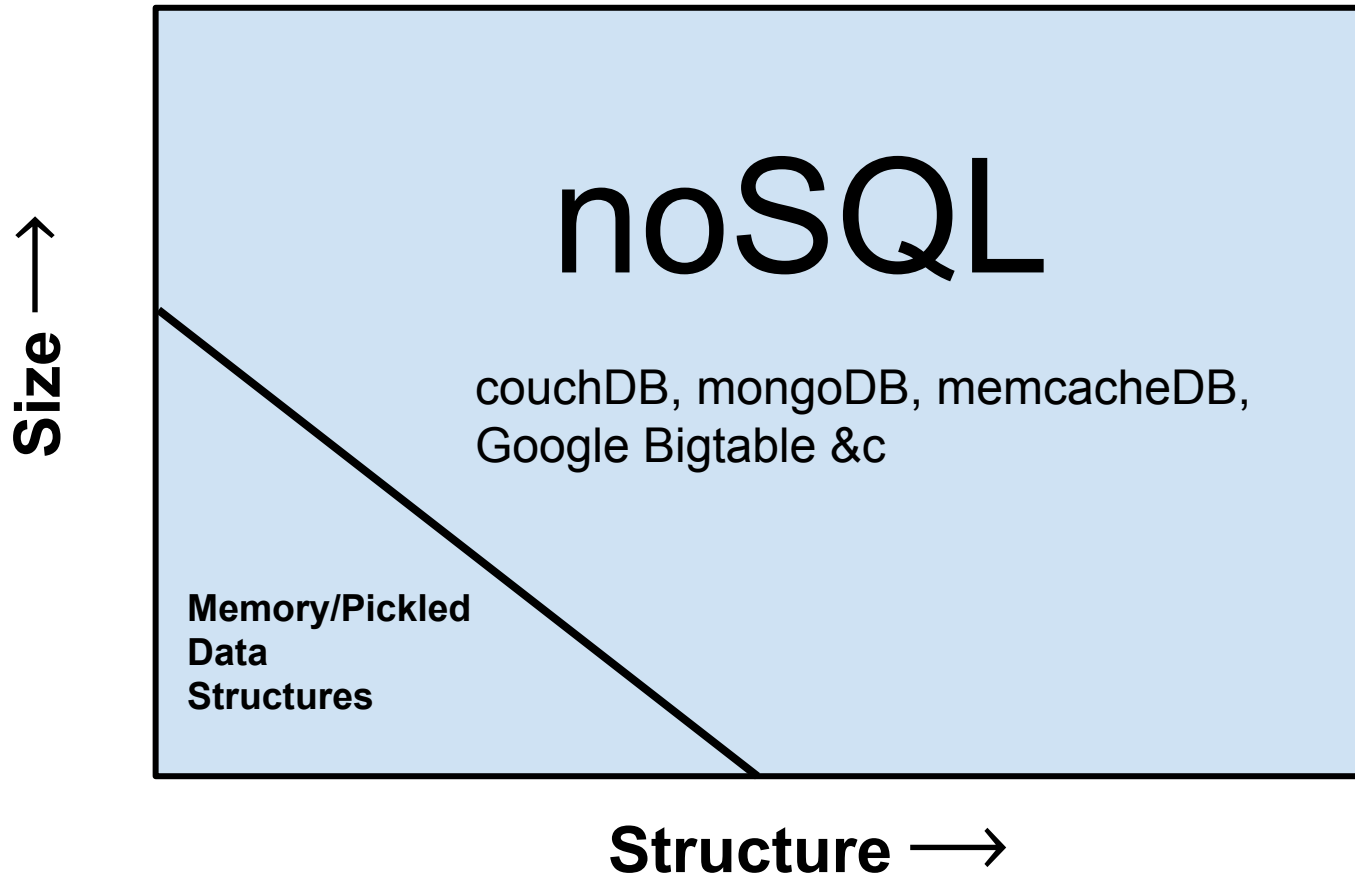


How do we all feel?


**1985-2008    Memory expensive, CPUs slow,  
programming languages inconvenient**



## 2008ish: The Dream



2008 - ?? : Memory cheap, CPUs fast,  
programming languages convenient



## Other DBMS/SQL use cases:


- Legacy system support
- Tightly bound to web frameworks
- Concurrent access
- Referential integrity guarantees
- Data security
- **Querying & report generation**
  - “Remixable” data

# Basic RDBMS/SQL Concepts



- **Schema** defines the structure of the data
- The **database** is composed of a number of user-defined **tables**
- Each **table** will have **columns** (aka fields) and **rows** (aka records)
- A column is of a given **data type**
- A row is an entry in a table with data for each column of that table

# RDBMS and SQL (Structured Query Language)



- SQL is the language used to query relational databases
- **All RDBMS use SQL** and the syntax and keywords are the same for the most part, across systems
- **SQL is used to interact** with RDBMS, allowing you to create tables, alter tables, insert records, update records, delete records, and query records within and across tables.
- Even non-relational databases like **Hadoop** usually have a SQL-like interface available.



# Client-Server Architecture

## Relational Database Server



```
$ postgres -D /usr/local/pgsql/data
```

postgres protocol

over TCP, sockets &c

Client

```
$ psql
```

Data files



Ways to use psql in the shell/term:

- |                                      |   |
|--------------------------------------|---|
| <code>\$ psql</code>                 | connects to postgres server                     |
| <code>\$ psql -U [USERNAME]</code>   | connects with given username                    |
| <code>\$ psql [DBNAME]</code>        | connects to a given database                    |
| <code>\$ psql &lt; script.sql</code> | reads file script.sql and send commands to psql |

*Try it live:*

- *Open file `sql/lecture_create.sql` with a text editor*
- *Use it to create a “dsilecture” database on your psql server*




Useful psql commands at the prompt [[link](#)]:

# \h	SQL help
# \?	psql commands help
# \l	List all the tables in the database
# \d	Describe the table schema
# \d db_name	Describe tables for a specific db
# \connect db_name	Connects to a database

*Try it live: Connect to “dsilecture” and describe schema of table “customer”*

# All together, now

```
$ cd ~/galvanize
$ mkdir sql-lecture
$ cd sql-lecture
$ wget
https://raw.githubusercontent.com/gSchool/DSI\_Lectures/master/sql/moses\_marshall/sql/lecture\_create.sql?token=AADLERkktd7xo-j9nehgcPHqdS4HRyBqks5bocjfwA%3D%3D -O lecture_create.sql
$ psql < lecture_create.sql
$ psql dsilecture
```



All SQL queries have three main ingredient :

<b>SELECT</b>	*What* data do you want?
<b>FROM</b>	*Where* do you want to get the data from?
<b>WHERE</b>	*Under what* conditions?

SQL is Declarative rather than Imperative. That is, you tell the machine what you want and it (database optimizer) decides how to do it

Advanced: You can use **EXPLAIN** to look at the how

## Basic SELECT

```
SELECT * FROM customers;
```

## Select columns

```
SELECT name, age, gender FROM customers;
```

## Restrict rows with WHERE

```
SELECT * FROM customers  
WHERE gender='M' ;
```

## Limit number of responses

```
SELECT * FROM customers LIMIT 3
```

Basic aggregate functions

```
SELECT count(*) FROM customers;
```

Basic aggregate functions

```
SELECT min(age) AS min_age,  
       max(age) AS max_age FROM customers;
```



# Aggregating groups with GROUP BY

Find average customer age for each state

```
SELECT state, AVG(age) as avg_age  
FROM customers  
GROUP BY state;
```

How to query the “visits” table to get number of visits per customer?

What happens when we do this?

```
SELECT * FROM customers GROUP BY state;
```

# Distinct, ordering

Select a distinct set

```
SELECT DISTINCT state FROM customers;
```

Ordering with ORDER BY

```
SELECT * FROM customers ORDER BY age;
```

To order by age in descending order:


```
SELECT * FROM customers ORDER BY age DESC;
```

How would we write a query to get the oldest customer?

# JOINS


The peanut butter and jelly of SQL

# JOIN types



## INNER JOIN


discards any entries  
that do not have a match  
between the tables  
based on the given keys.



```
SELECT <auswahl>  
FROM tabelleA A  
INNER JOIN tabelleB B  
ON A.key = B.key
```

## LEFT OUTER JOIN


keeps all entries  
in the left table  
regardless of  
whether a match is found  
in the right table



```
SELECT <auswahl>  
FROM tabelleA A  
LEFT JOIN tabelleB B  
ON A.key = B.key
```

## FULL OUTER JOIN

will keep the rows  
of both tables  
no matter what




```
SELECT <auswahl>  
FROM tabelleA A  
FULL OUTER JOIN tabelleB B  
ON A.key = B.key
```

# Inner Join

id	name
1	Elliott
2	Mark
3	Moses
4	Brandon

id	name
2	Llorente
3	Marsh
5	Engard
6	Van

```
SELECT id, first.name,
last.name
FROM first
INNER JOIN last
ON first.id=last.id
```



```
SELECT <auswahl>
FROM tabelleA A
INNER JOIN tabelleB B
ON A.key = B.key
```

[illegible]

# Inner Join


**first**

id	name
1	Elliott
2	Mark
3	Moses
4	Brandon

**last**

id	name
2	Llorente
3	Marsh
5	Engard
6	Van

```
SELECT id, first.name,  
last.name  
FROM first  
INNER JOIN last  
ON first.id=last.id
```



```
SELECT <auswahl>  
FROM tabelleA A  
INNER JOIN tabelleB B  
ON A.key = B.key
```


id	first	last
2	Mark	Llorente
3	Moses	Marsh

# Left Outer Join

id	name
1	Elliott
2	Mark
3	Moses
4	Brandon

id	name
2	Llorente
3	Marsh
5	Engard
6	Van

```
SELECT id, first.name,
last.name
FROM first
LEFT JOIN last
ON first.id=last.id
```



```
SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
```

[illegible]

# Left Outer Join


**first**

id	name
1	Elliott
2	Mark
3	Moses
4	Brandon

**last**

id	name
2	Llorente
3	Marsh
5	Engard
6	Van

```
SELECT id, first.name,  
last.name  
FROM first  
LEFT JOIN last  
ON first.id=last.id
```



```
SELECT <auswahl>  
FROM tabelleA A  
LEFT JOIN tabelleB B  
ON A.key = B.key
```

id	first	last
1	Elliott	
2	Mark	Llorente
3	Moses	Marsh
4	Brandon	




# Full Outer Join

id	name
1	Elliott
2	Mark
3	Moses
4	Brandon

id	name
2	Llorente
3	Marsh
5	Engard
6	Van

```
SELECT id, first.name,
last.name
FROM first
FULL OUTER JOIN last
ON first.id=last.id
```



```
SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
```

[illegible]

# Full Outer Join

**first**


id	name
1	Elliott
2	Mark
3	Moses
4	Brandon

**last**

id	name
2	Llorente
3	Marsh
5	Engard
6	Van


```
SELECT id, first.name,  
last.name  
FROM first  
FULL OUTER JOIN last  
ON first.id=last.id
```

```
SELECT <auswahl>  
FROM tabelleA A  
FULL OUTER JOIN tabelleB B  
ON A.key = B.key
```




id	first	last
1	Elliott	
2	Mark	Llorente
3	Moses	Marsh
4	Brandon	
5		Engard
6		Van


# Other types of joins




```
SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
```




```
SELECT <auswahl>
FROM tabelleA A
INNER JOIN tabelleB B
ON A.key = B.key
```




```
SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
```




```
SELECT <auswahl>
FROM tabelleA A
LEFT JOIN tabelleB B
ON A.key = B.key
WHERE B.key IS NULL
```



```
SELECT <auswahl>
FROM tabelleA A
RIGHT JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
```



```
SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
```



```
SELECT <auswahl>
FROM tabelleA A
FULL OUTER JOIN tabelleB B
ON A.key = B.key
WHERE A.key IS NULL
OR B.key IS NULL
```

# Composing SQL Queries

--- Return the customer\_ids of all customers who visited in June ---

```
SELECT c.id, v.created_at
FROM customers as c
JOIN visits as v
ON c.id = v.customer_id
WHERE date_part('month', v.created_at) = 6;
```

--- LEFT JOIN: return all customers from the customers table regardless of presence in visits

```
SELECT
  c.id
, v.created_at
FROM customers as c
LEFT JOIN visits as v
ON c.id = v.customer_id
WHERE date_part('month', v.created_at) = 6;
```

# Subqueries



In general, you can replace any table name with a SELECT statement.

```
SELECT . . . . . FROM (SELECT . . . . .)
```

If a query returns a **single value**, you can treat it as such.

```
WHERE var1 = (SELECT . . . .)
```

If a query returns a **single column**, you can treat it sort of like a list/vector

```
WHERE var1 IN (SELECT . . . .)
```

# Afternoon Lecture

More about SQL and RDBMSs

Given the following query, number what order the commands are executed:

```
SELECT a.userid, COUNT(*) AS recent_visits
FROM users AS a
LEFT JOIN visits AS b
ON a.userid = b.userid
WHERE b.dt > '2012-01-01'
GROUP BY a.userid
HAVING count(0) < 10
ORDER BY recent_visits;
```


# Order of Evaluation of a SQL SELECT Statement



1. FROM + JOIN: first the product of all tables is formed
2. WHERE: the where clause filters rows that do not meet the search condition
3. GROUP BY + (COUNT, SUM, etc): the rows are grouped using the columns in the group by clause and the aggregation functions are applied on the grouping
4. HAVING: like the WHERE clause, but can be applied after aggregation
5. SELECT: the targeted list of columns are evaluated and returned
6. DISTINCT: duplicate rows are eliminated
7. ORDER BY: the resulting rows are sorted



# Creating a table with a schema



# Inserting values in a table




table name

records and their values

```
INSERT INTO products (id, name, price) VALUES
(1, 'soccer ball', 20.5),
(2, 'iPod', 200),
(3, 'headphones', 50);
```

# SQL Queries for table creation / maintenance



Creating a table from query:

```
CREATE [TEMPORARY] TABLE table AS <SQL query>;
```

Inserting records in a table:

```
INSERT INTO table [(c1,c2,c3,...)] VALUES (v1,v2,v3,...);
```

Updating records:

```
UPDATE table SET c1=v1,c2=v2,... WHERE cX=vX;
```

Delete records:

```
DELETE FROM table WHERE cX=vX;
```


Change model (add, drop, modify columns):

```
ALTER TABLE table [DROP/ADD/ALTER] column [datatype];
```

Delete a table:

```
DROP TABLE table;
```

# Designing a database with keys




## Primary Key

A primary key is a special column of a table that uniquely identifies that entry.

A primary key is not always an integer; it could be a combination of columns, hash, timestamp..etc.,


## Foreign Keys

Foreign Keys are columns that reference some other entry in the database.





- Minimizes Redundancy, for example:
  - Details about a user(address, age) are only stored once (in a users table)
  - Any other table (eg. purchases) where this data might be relevant, only references the user\_id
  - Choose Normalized or Denormalized Schemas based on the use case:
    - Heavy reporting (Data Warehouse)
    - Transactional Systems (Ordering System)



- “self join”: JOIN a table with itself:  
<http://www.w3resource.com/sql/joins/perform-a-self-join.php>
- CROSS JOIN: join each row in table A with every row in table B:  
<http://www.w3resource.com/sql/joins/cross-join.php>
- window functions:  
<https://www.postgresql.org/docs/9.1/static/tutorial-window.html>
- COALESCE: often used to turn NULL values into non-null values:  
<https://www.postgresql.org/docs/9.5/static/functions-conditional.html>
- EXCEPT: return rows from one SELECT statement that are NOT returned by a second SELECT statement:  
<https://www.tutorialspoint.com/sql/sql-except-clause.htm>