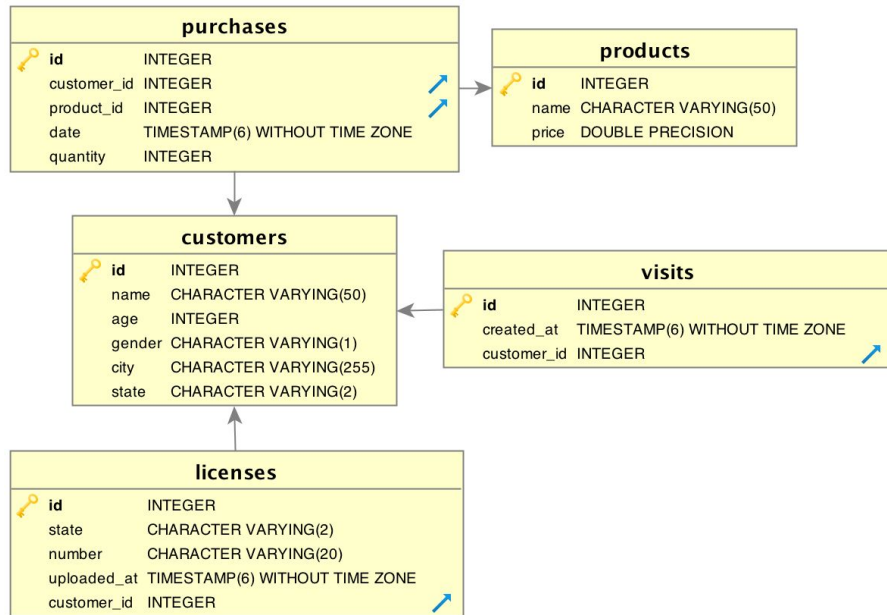




# SQL

DSI, Galvanize, Seattle

Jean-François Omhover  
Miles Erickson





# SQL

DSI, Galvanize, Seattle

Jean-François Omhover  
Miles Erickson

## OBJECTIVES

- **Connect** to a Postgres database
- **Design** and **write** queries to answer questions using an RDMBS (i.e. Postgres)



- Connect to a SQL database via command line (i.e. Postgres).
- Connect to a database from within a python program.
- State function of basic SQL commands.
- Write simple queries on a single table including SELECT, FROM, WHERE, CASE... clauses and aggregates.
- Write complex queries including JOINS and subqueries.
- Explain how indexing works in Postgres.
- Create and dump tables.
- Format a query to follow a standard style.
- Move data from SQL database to text file.



It is a **persistent data storage system**

- survives after the process in which it was created has ended
- is written to non-volatile storage
- is infrequently accessed and unlikely to be changed

RDBMS was the de facto standard for storing data

- Examples: Oracle, MySQL, SQLServer, Postgres
- With "Big Data", this is beginning to change

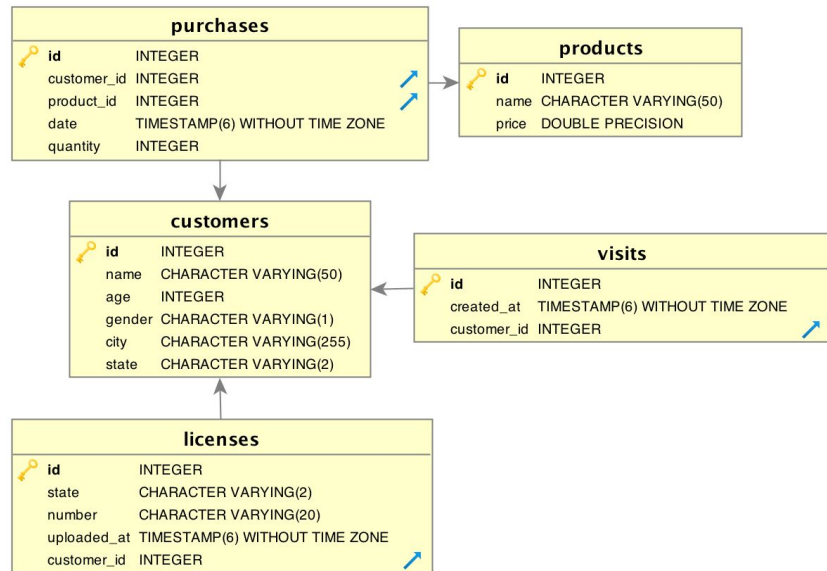
# Why RDBMS ?



An RDBMS provides the ability to

- **Model** relations in data
- **Query** data and their relations efficiently
- **Maintain** data consistency and integrity

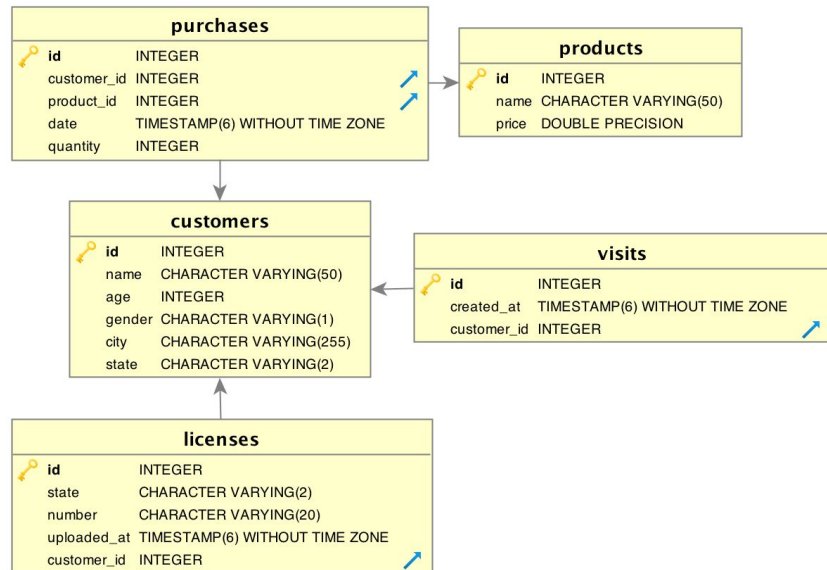
It will require a Data Model



# RDBMS Data Model



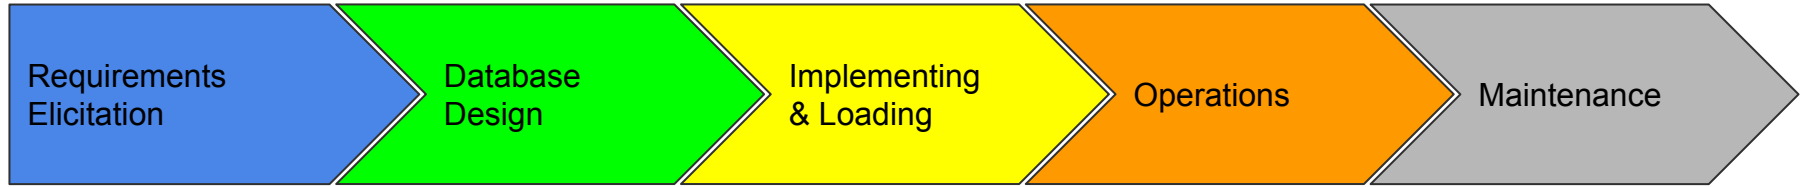
- **Schema** defines the structure of the data
- The **database** is composed of a number of user-defined **tables**
- Each **table** will have **columns** (aka fields) and **rows** (aka records)
- A column is of a given **data type**
- A row is an entry in a table with data for each column of that table



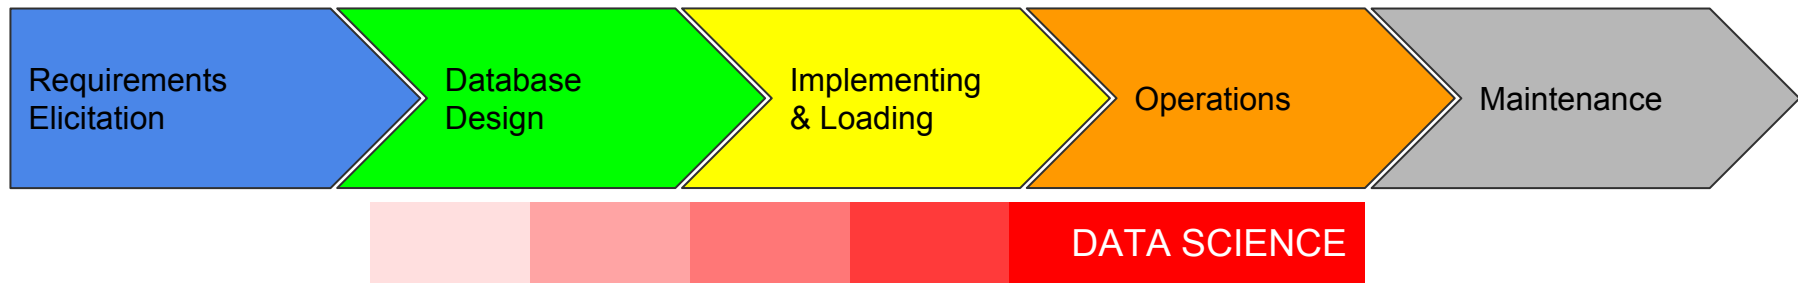


- SQL is the language used to query relational databases
- **All RDBMS use SQL** and the syntax and keywords are the same for the most part, across systems
- **SQL is used to interact** with RDBMS, allowing you to create tables, alter tables, insert records, update records, delete records, and query records within and across tables.
- Even non-relational databases like **Hadoop** usually have a SQL-like interface available.

# Database Life Cycle (DBLC)







DS / Operations: querying, aggregating

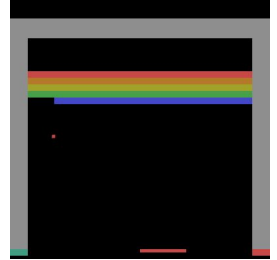
DS / Implementing: identifying, cleaning, pushing external data sources inside a RDBMS

DS / Design: recommendations on the model, specs on operations

# Question !

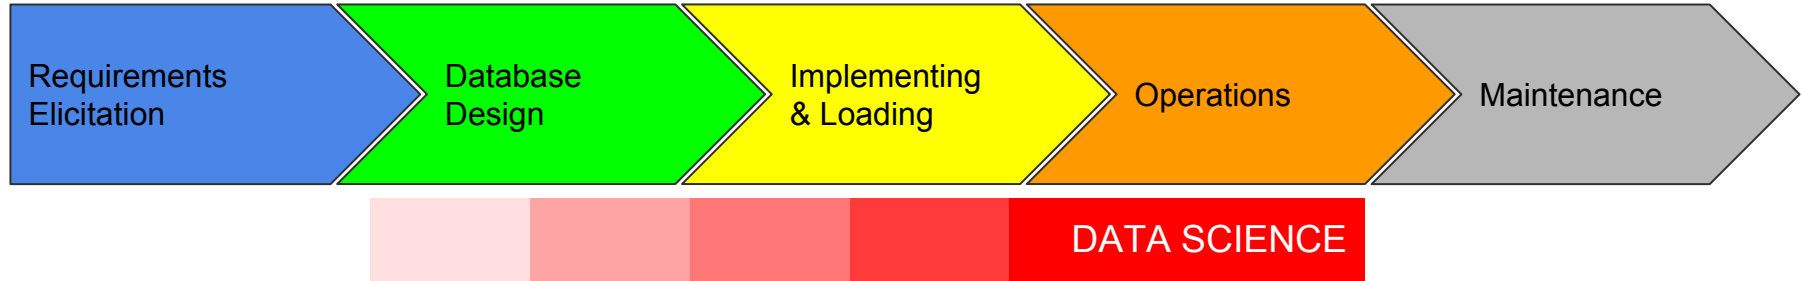


As a data scientist, what are the advantages of storing, querying, and maintaining data in a SQL database over curating your own flat files (e.g. csv files) ?



...

# Concepts of the day in the DBLC

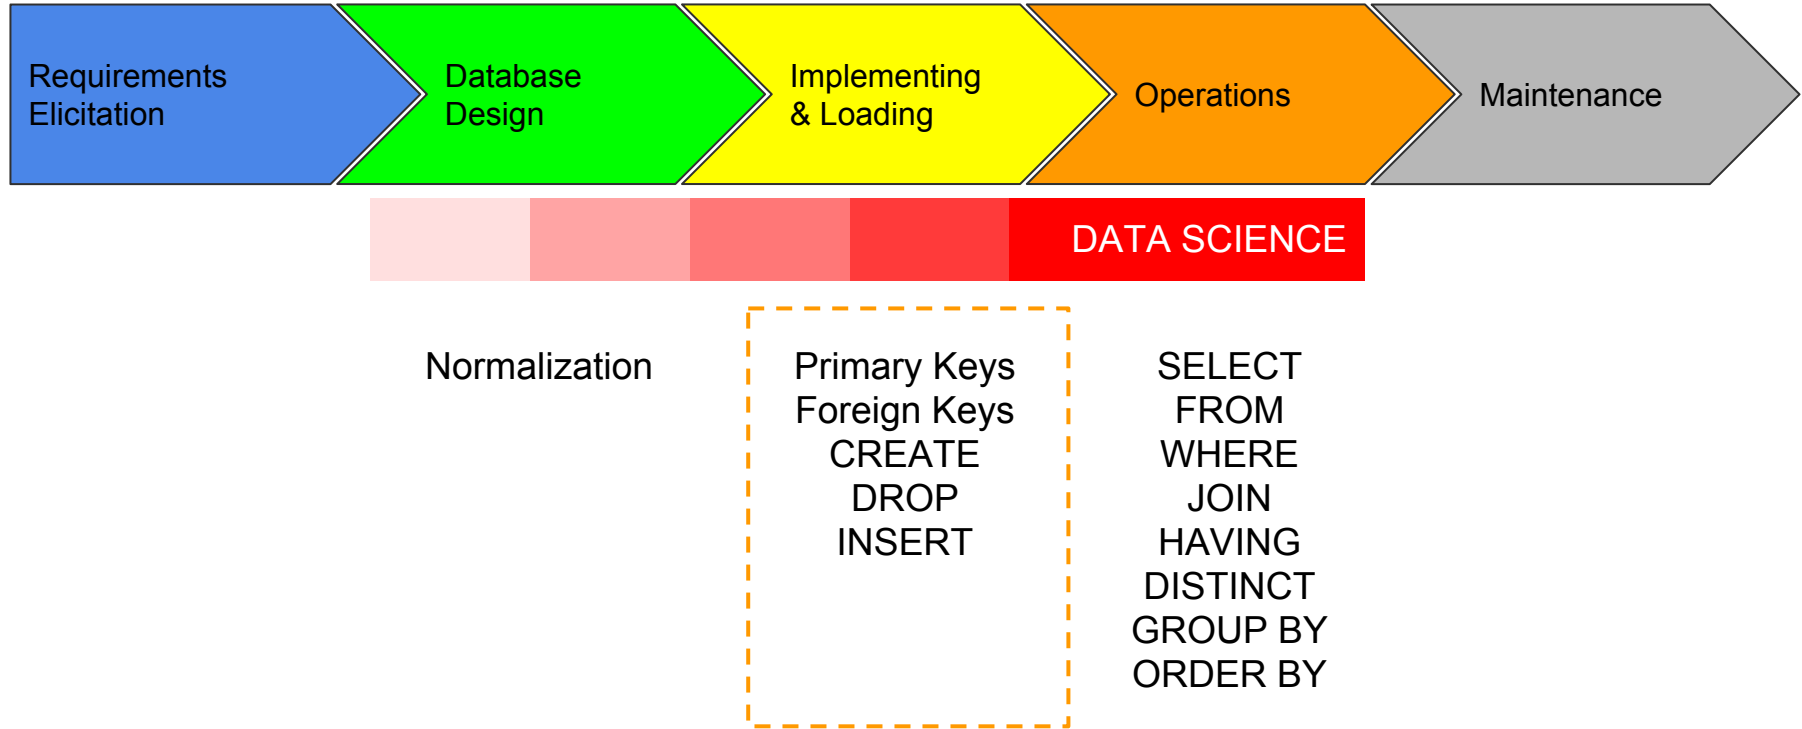


Normalization

Primary Keys  
Foreign Keys  
CREATE  
DROP  
INSERT

SELECT  
FROM  
WHERE  
JOIN  
HAVING  
DISTINCT  
GROUP BY  
ORDER BY

# Concepts of the day in the DBLC





Ways to use psql in the shell/term:

- |                                      |   |
|--------------------------------------|---|
| <code>\$ psql</code>                 | connects to postgres server                     |
| <code>\$ psql -U [USERNAME]</code>   | connects with given username                    |
| <code>\$ psql [DBNAME]</code>        | connects to a given database                    |
| <code>\$ psql &lt; script.sql</code> | reads file script.sql and send commands to psql |

*Try it live:*

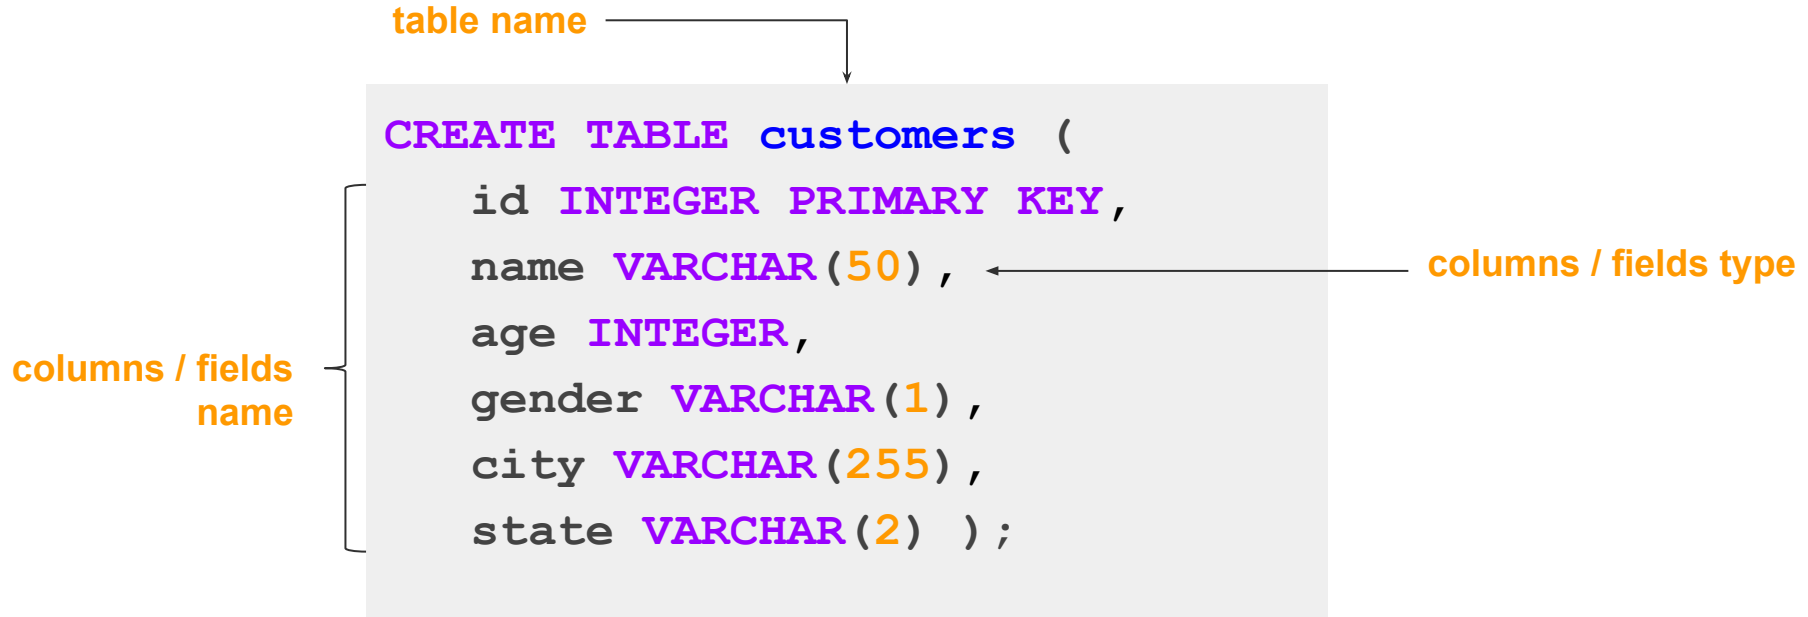
- *Open file `sql/lecture_create.sql` in atom*
- *Use it to create a “dsilecture” database on your psql server*

Useful psql commands at the prompt [[link](#)]:

# \h	SQL help
# \?	psql commands help
# \l	List all the tables in the database
# \d	Describe the table schema
# \d db_name	Describe tables for a specific db
# \connect db_name	Connects to a database

*Try it live: Connect to “dsilecture” and describe schema of table “customer”*

# Creating a table with a schema



# Inserting values in a table



table name

records and their values

```
INSERT INTO products (id, name, price) VALUES  
  (1, 'soccer ball', 20.5),  
  (2, 'iPod', 200),  
  (3, 'headphones', 50);
```



# SQL Queries for table creation / maintenance



Creating a table from query:

```
CREATE [TEMPORARY] TABLE table AS <SQL query>;
```

Inserting records in a table:

```
INSERT INTO table [(c1,c2,c3,...)] VALUES (v1,v2,v3,...);
```

Updating records:

```
UPDATE table SET c1=v1,c2=v2,... WHERE cX=vX;
```

Delete records:

```
DELETE FROM table WHERE cX=vX;
```

Change model (add, drop, modify columns):

```
ALTER TABLE table [DROP/ADD/ALTER] column [datatype];
```

Delete a table:

```
DROP TABLE table;
```

# Designing a database with keys



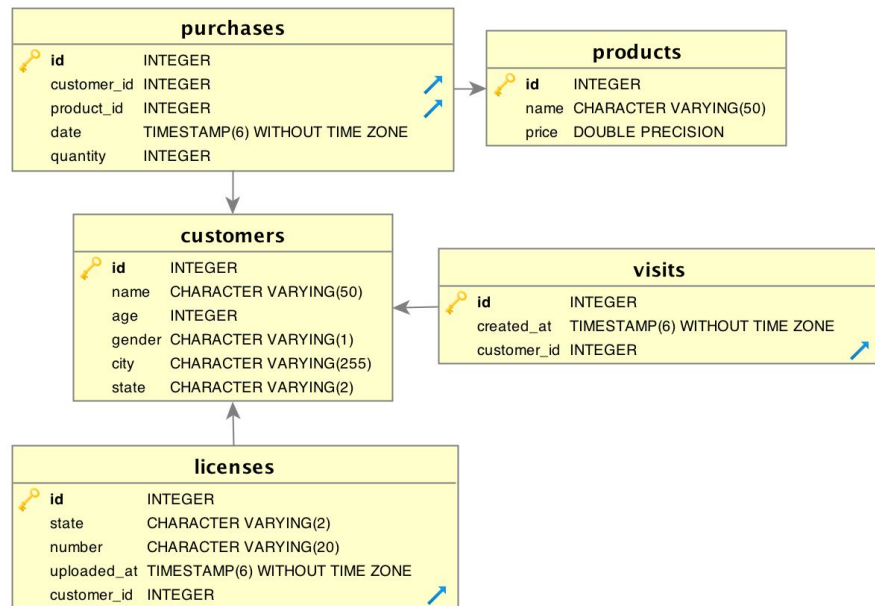
## Primary Key

A primary key is a special column of a table that uniquely identifies that entry.

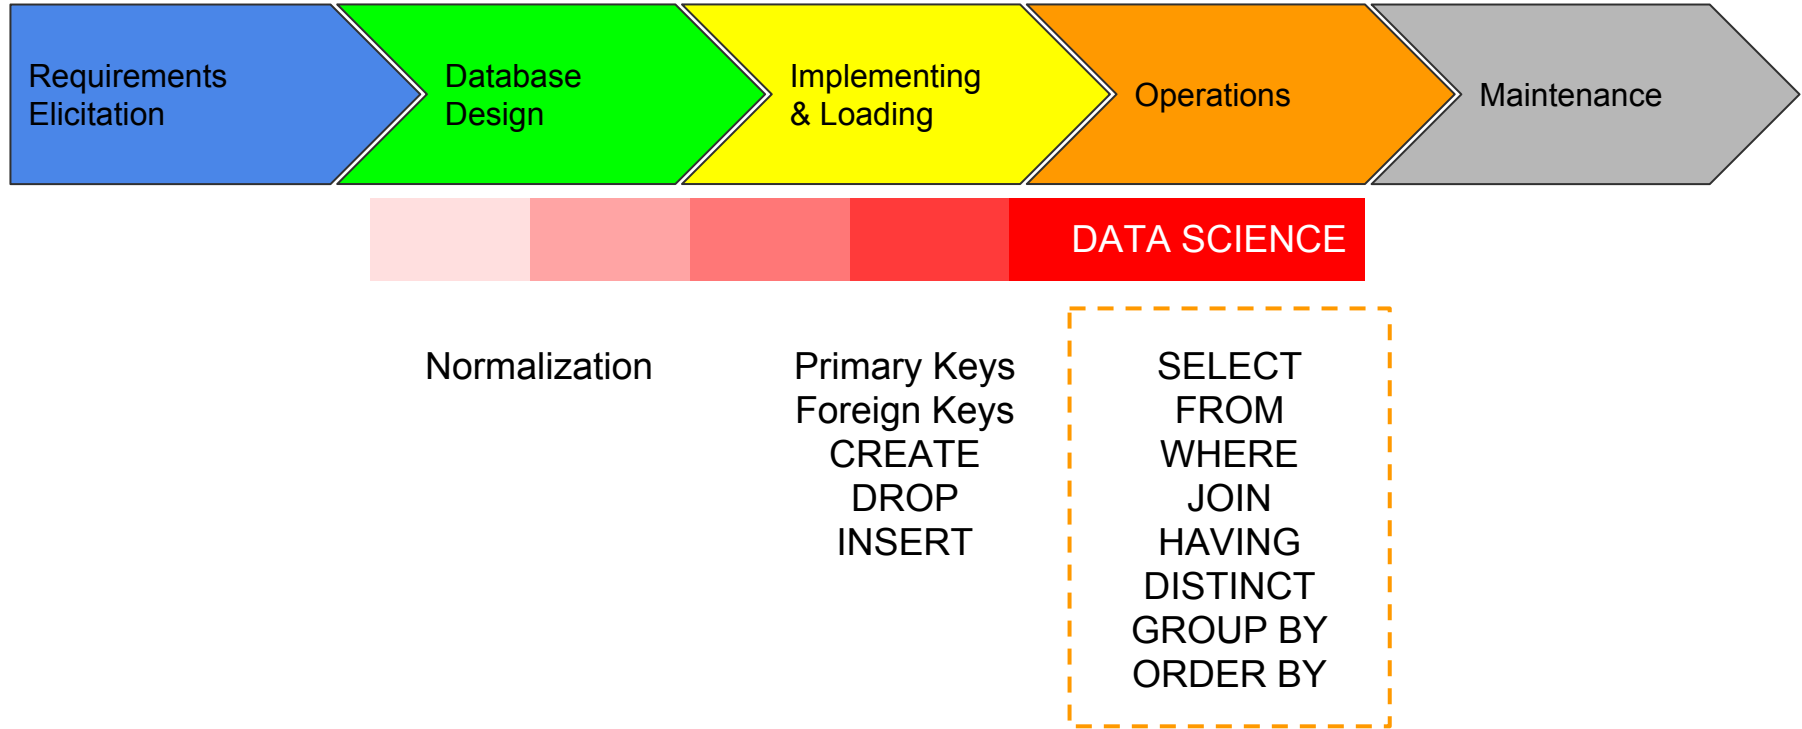
A primary key is not always an integer; it could be a combination of columns, hash, timestamp..etc.,

## Foreign Keys

Foreign Keys are columns that reference some other entry in the database.



# Concepts of the day in the DBLC





All SQL queries have three main ingredient :

<b>SELECT</b>	*What* data do you want?
<b>FROM</b>	*Where* do you want to get the data from?
<b>WHERE</b>	*Under what* conditions?

SQL is Declarative rather than Imperative. That is, you tell the machine what you want and it (database optimizer) decides how to do it

Advanced: You can use **EXPLAIN** to look at the how



Select the columns name, age from the table users.

```
SELECT name, age  
FROM customers
```

SQL always returns a table, so the output of the query above is a sub-table of users with 2 columns.

Select name and age for every user in users who live in CA.

```
SELECT name, age  
FROM customers  
WHERE state = 'CA'
```



*Open file sql/lecture\_examples.sql*

*Run them in psql*

JOIN clause used to query across multiple tables using foreign keys

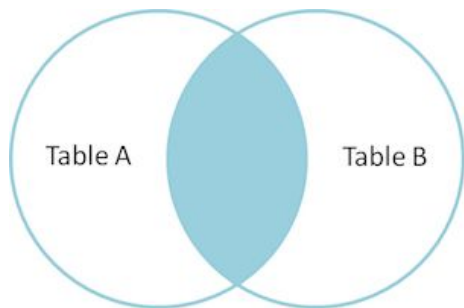
Every JOIN has two segments:

Specifying the tables to JOIN

Specifying the columns to match

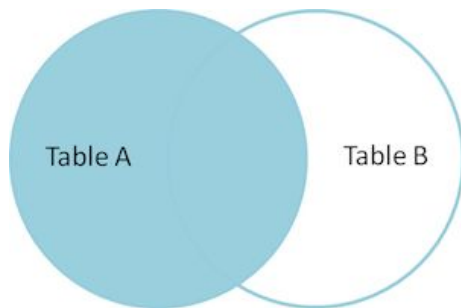
## INNER JOIN

discards any entries  
that do not have a match  
between the tables  
based on the given keys.




## LEFT OUTER JOIN

keeps all entries  
in the left table  
regardless of  
whether a match is found  
in the right table



## FULL OUTER JOIN

will keep the rows  
of both tables  
no matter what





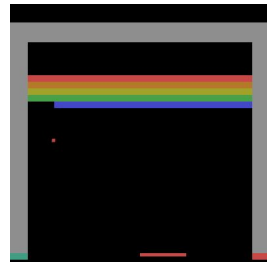
# Question ! (JOIN the fun)



How many rows would result in  
(inner join, left join, right join, full outer join) on `department_id` ?

employee_id	department_id	name	salary
2	1	Jon	40000
7	1	Linda	50000
12	2	Ashley	15000
1	0	Mike	80000

department_id	location
1	NY
2	SF
3	Austin



In general, you can replace any table name with a SELECT statement.

```
SELECT . . . . . FROM (SELECT . . . . .)
```

If a query returns a **single value**, you can treat it as such.

```
WHERE var1 = (SELECT . . . .)
```

If a query returns a **single column**, you can treat it sort of like a list/vector

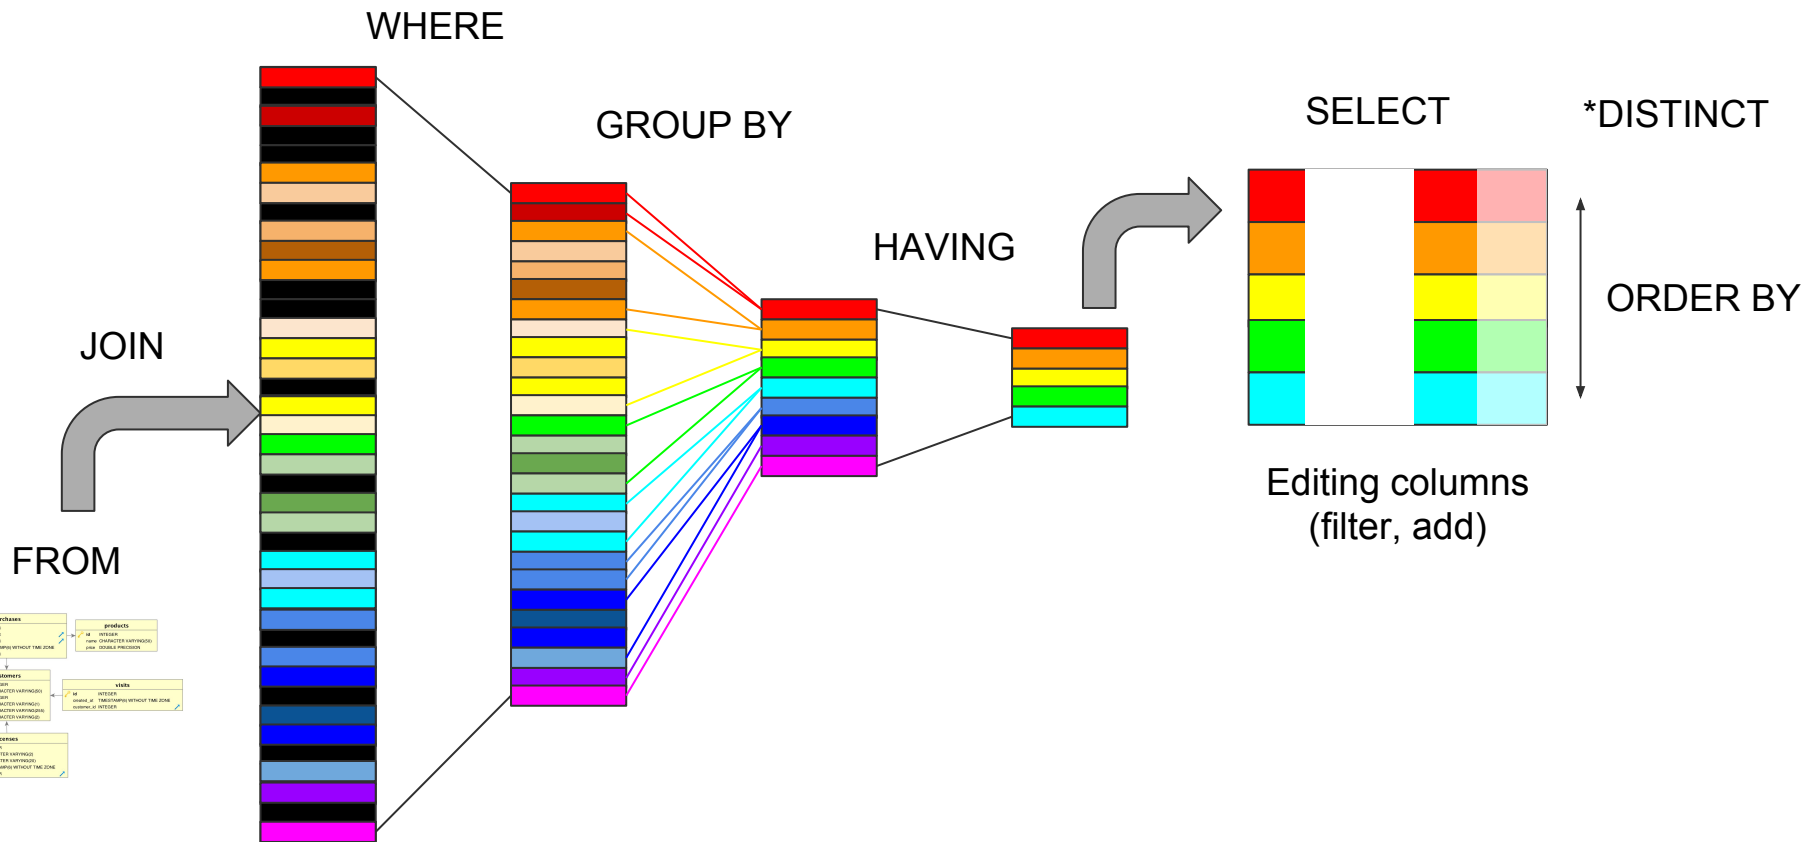
```
WHERE var1 IN (SELECT . . . .)
```

# Order of Evaluation of a SQL SELECT Statement

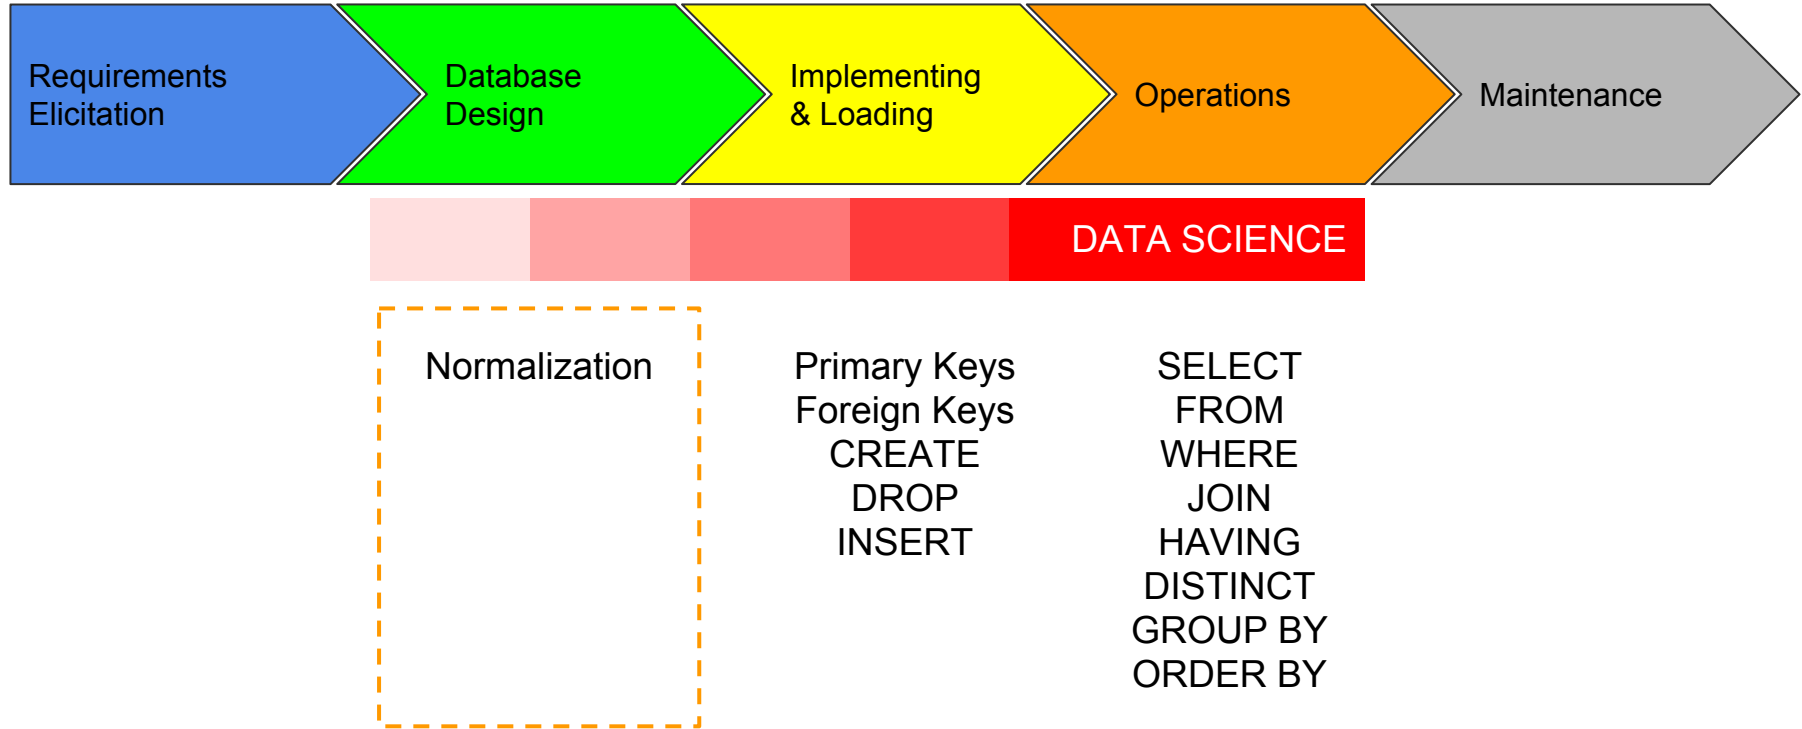


1. FROM + JOIN: first the product of all tables is formed
2. WHERE: the where clause filters rows that do not meet the search condition
3. GROUP BY + (COUNT, SUM, etc): the rows are grouped using the columns in the group by clause and the aggregation functions are applied on the grouping
4. HAVING: like the WHERE clause, but can be applied after aggregation
5. SELECT: the targeted list of columns are evaluated and returned
6. DISTINCT: duplicate rows are eliminated
7. ORDER BY: the resulting rows are sorted

# Order of Evaluation of a SQL SELECT Statement



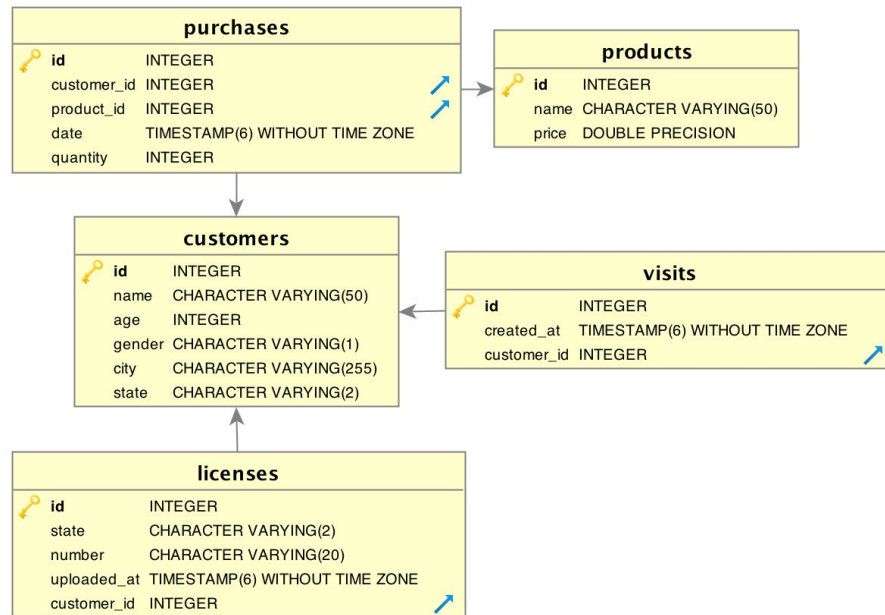
# Concepts of the day in the DBLC



# Database Normalization



- Minimizes Redundancy, for example:
  - Details about a user(address, age) are only stored once (in a users table)
  - Any other table (eg. purchases) where this data might be relevant, only references the user\_id
  - Choose Normalized or Denormalized Schemas based on the use case:
    - Heavy reporting (Data Warehouse)
    - Transactional Systems (Ordering System)





# SQL

DSI, Galvanize, Seattle

Jean-François Omhover  
Miles Erickson

