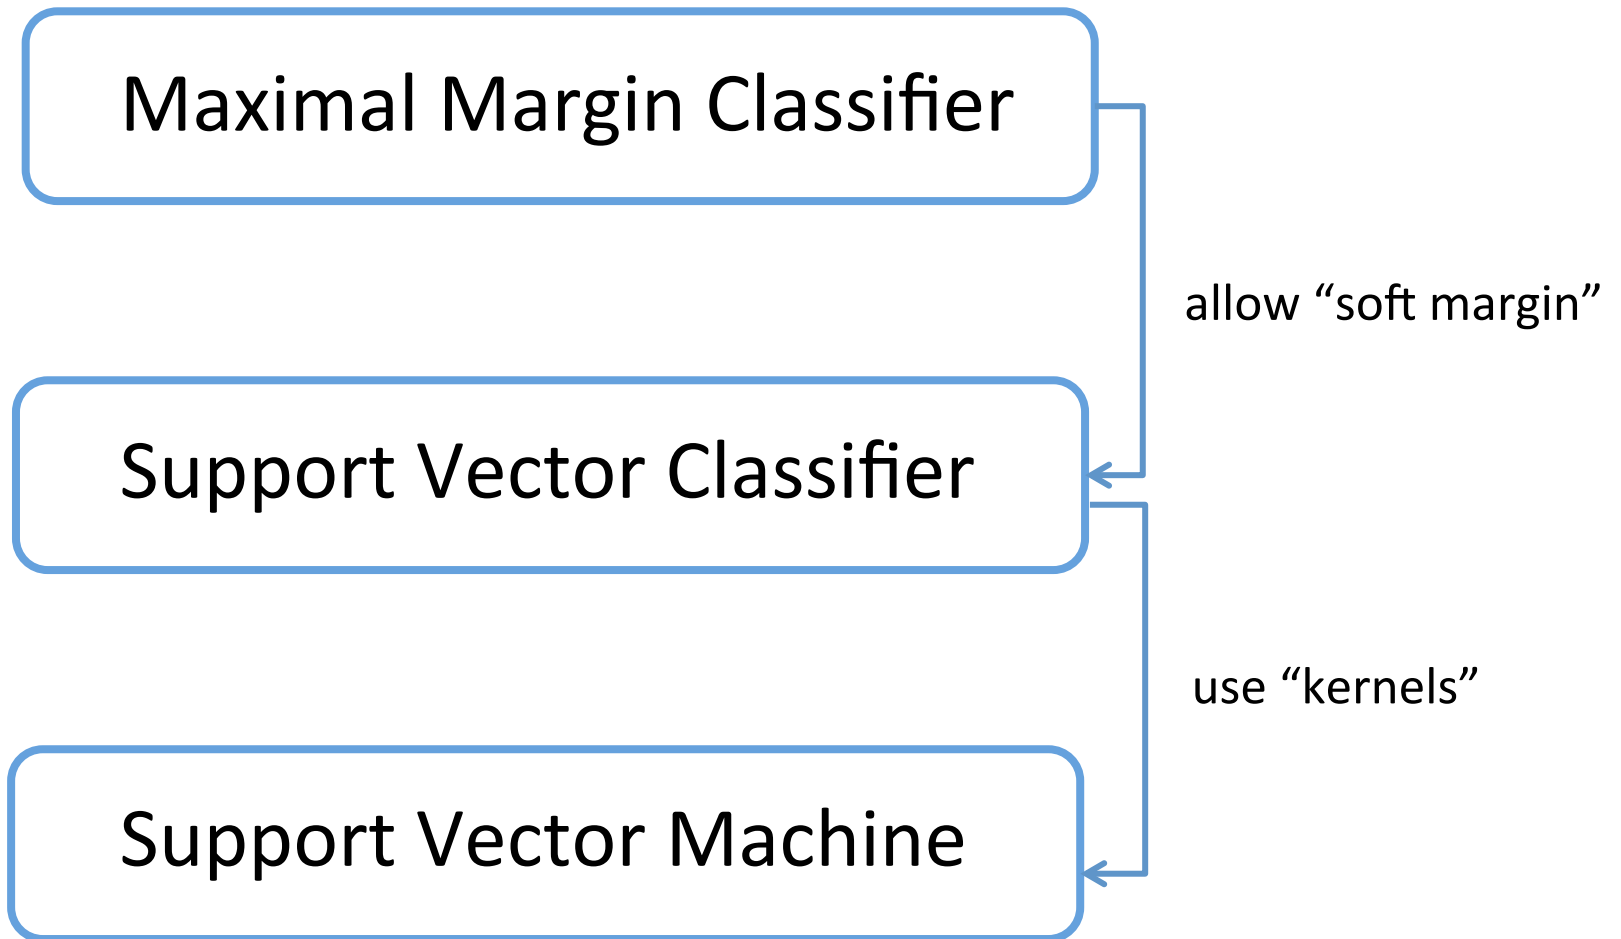



Support Vector Machines



Hyperplanes


In p -dimensional space, a flat affine subspace of dimension $p-1$

$p = 2$



$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 = 0$$

$p = 3$




$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 = 0$$

Generally, a hyperplane in p -dimensional space can be defined as

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p = 0$$

Hyperplanes




Can think of hyperplane as dividing p-dimensional space into two halves

Separating Hyperplane

Suppose we code...


If $y_i = \text{Blue} \implies y_i = +1$

If $y_i = \text{Red} \implies y_i = -1$




We have a *separating hyperplane*,
if for *all points*, we have...

$$\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} > 0 \text{ when } y_i = +1$$




$$\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} < 0 \text{ when } y_i = -1$$


In particular, we fit...



hmm....



need some sort of *budget*



$$\underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n}{\text{maximize}} \quad M$$

$$\text{subject to} \quad \sum_{j=1}^p \beta_j^2 = 1,$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i)$$

$$\epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C,$$


Budget that we can tune

Slack from
each point



- $\epsilon_i = 0$ for being on correct side of margin
- $\epsilon_i > 0$ for violating the margin
- $\epsilon_i > 1$ for being on wrong side of hyperplane

need some sort of *budget*



$\epsilon_i = 0$ for being on correct side of margin
 $\epsilon_i > 0$ for violating the margin
 $\epsilon_i > 1$ for being on wrong side of hyperplane

$$\begin{aligned}
 & \underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n}{\text{maximize}} && M \\
 & \text{subject to} && \sum_{j=1}^p \beta_j^2 = 1, \\
 & && y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \\
 & && \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C,
 \end{aligned}$$

Slack from each point
 \downarrow
 (points to ϵ_i in the constraint)


\uparrow
Budget that we can tune
 (points to C in the constraint)

Bias Variance Tradeoff

- C small \Leftrightarrow Low bias, High Variance
- C large \Leftrightarrow High bias, Low Variance (not quite as clear cut)




Afternoon

hmm....



[Cool video.](#)

<https://www.youtube.com/watch?v=3liCbRZPrZA>



Maximal Margin Classifier

Support Vector Classifier



Support Vector Machine

allow "soft margin"

use "kernels"

So what are SVMs again?

- Hyperplane that separates data as well as possible, while allowing some room for error (“soft margin”)
- Kernels are powerful way to accommodate non-linear class boundaries.



$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i K(x, x_i)$$

Kernels

Solution to SVC only involves inner product of observations

$$\langle x_i, x_{i'} \rangle = \sum_{j=1}^p x_{ij} x_{i'j}$$

$$f(x) = \beta_0 + \sum_{i=1}^n \alpha_i \langle x, x_i \rangle \leftarrow \text{SVC}$$

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i \langle x, x_i \rangle \quad \leftarrow \text{Only requires support vectors}$$

More generally, instead of just taking inner product, we can use *Kernels*

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i K(x, x_i) \quad \leftarrow \text{SVM, since using Kernels now}$$

$$K(x_i, x_{i'}) = \sum_{j=1}^p x_{ij} x_{i'j} \quad \text{Linear Kernel}$$



$$K(x_i, x_{i'}) = \left(1 + \sum_{j=1}^p x_{ij} x_{i'j}\right)^d \quad \text{Polynomial Kernel}$$

$$K(x_i, x_{i'}) = \exp\left(-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2\right) \quad \text{Radial Basis Function Kernel ("Gaussian")}$$

Radial Basis Kernel (Gaussian)

$$K(x_i, x_{i'}) = \exp\left(-\gamma \sum_{j=1}^p (x_{ij} - x_{i'j})^2\right)$$

$$f(x) = \beta_0 + \sum_{i \in \mathcal{S}} \hat{\alpha}_i K(x, x_i)$$



Polynomial Kernel

- Expand feature space by simply creating new features

$$(X_1, X_2) \longrightarrow (X_1, X_2, X_1^2, X_2^2, X_1X_2)$$

- Same idea of hyperplane decision boundary, but it is non-linear when projected down to X_1 vs. X_2 space

$$\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_1^2 + \beta_4 X_2^2 + \beta_5 X_1 X_2 = 0$$

- Cool Video: <https://www.youtube.com/watch?v=3liCbRZPrZA>

Turns out, we can rewrite the optimization


Loss + Penalty

$$\begin{aligned} & \underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n}{\text{maximize}} && M \\ & \text{subject to} && \sum_{j=1}^p \beta_j^2 = 1, \\ & && y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \\ & && \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq C, \end{aligned}$$

$$\underset{\beta_0, \beta_1, \dots, \beta_p}{\text{minimize}} \left\{ \sum_{i=1}^n \max[0, 1 - y_i f(x_i)] + \lambda \sum_{j=1}^p \beta_j^2 \right\}$$


Relation to Logistic Regression

$$\sum_{i=1}^n \max [0, 1 - y_i(\beta_0 + \beta_1 x_{i1} + \dots + \beta_p x_{ip})]$$



- Often results similar to logistic regression
- Robust to observations far from hyperplane

Dealing with Unbalanced Classes



- Suppose 10% Red
- Can adjust weights inversely proportional to class frequencies

Questions

- What's a hyperplane? Specify equation
- How does a Support Vector Classifier work?
 - Describe the “soft margin” aspect
 - Describe the “hinge loss” aspect – Slide 16
 - Contrast with Logistic Regression (as compared to hinge loss)
- How to handle unbalanced classes?
- How to tune?

Appendix

Rewriting the optimization


$$\begin{aligned} & \underset{\beta_0, \beta_1, \dots, \beta_p, \epsilon_1, \dots, \epsilon_n}{\text{maximize}} && M \\ & \text{subject to} && \sum_{j=1}^p \beta_j^2 = 1, \\ & && \underline{y_i(\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}) \geq M(1 - \epsilon_i)} \\ & && \epsilon_i \geq 0, \quad \sum_{i=1}^n \epsilon_i \leq \underline{C}, \end{aligned}$$

$$\underset{\beta_0, \beta_1, \dots, \beta_p}{\text{minimize}} \left\{ \sum_{i=1}^n \underline{\max[0, 1 - y_i f(x_i)]} + \underline{\lambda} \sum_{j=1}^p \beta_j^2 \right\}$$

The idea is to take a certain budget, C , and find the optimal β vector that achieves the maximum margin possible, M .

- $C \rightarrow \lambda$, both tuning parameters
- $\epsilon_i = 0$ for being on correct side of margin
 $\epsilon_i > 0$ for violating the margin
 $\epsilon_i > 1$ for being on wrong side of hyperplane

More details on SVM vs. Logistic Regression



For $Y = -1$ or $Y = +1$

Logistic Regression

- Loss Function, or “binomial Log Likelihood”*

$$L[Y, f(X)] = \log(1 + e^{-Yf(X)})$$

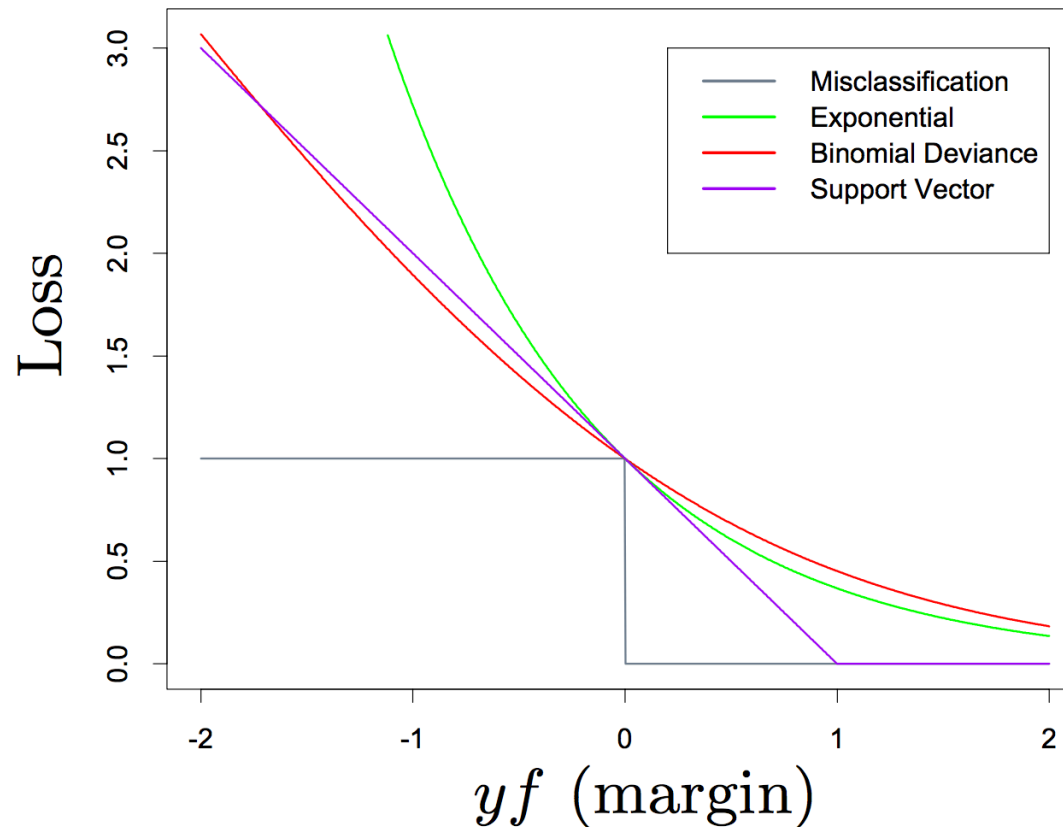
- Which estimates the logit

$$f(X) = \log \frac{\Pr(Y = 1|X)}{\Pr(Y = -1|X)}$$

SVMs vs. Logistic Regression

- When **classes nearly separable, SVM** tends to do better than Logistic Regression
- Otherwise, Logistic Regression (with Ridge) and SVMs are similar
- However, if you **want to estimate probabilities, Logistic Regression** is the choice.
- With kernels, SVMs work well. Logistic Regression works fine with kernels but can get computationally too expensive

SVM and other Loss functions



- Can consider training model using different loss functions and scoring each one
- In scikit-learn, SVC unfortunately does not have a loss function parameter but LinearSVC does!

SVM and Multiple Classes

Doesn't extend so nicely. Still, we can do...

- One vs. One classification
 - If $K > 2$ classes, simply compute $\binom{K}{2}$ classifiers
 - Take test observation and tally up times assigned to each of K classes
→ Most frequent class assigned
- One vs. All classification
 - Fit K classifiers, each comparing one class to $(K-1)$ classes
 - Take test observation and assign to each class we have highest confidence in \Leftrightarrow class for which $\beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \dots + \beta_p x_{ip}$ is highest

